



# Flink Positive

Caito Scherr, Nik Davis



# Safe Harbor


This presentation and the information herein (including any information that may be incorporated by reference) is provided for informational purposes only and should not be construed as an offer, commitment, promise or obligation on behalf of New Relic, Inc. ("New Relic") to sell securities or deliver any product, material, code, functionality, or other feature. Any information provided hereby is proprietary to New Relic and may not be replicated or disclosed without New Relic's express written permission.

Such information may contain forward-looking statements within the meaning of federal securities laws. Any statement that is not a historical fact or refers to expectations, projections, future plans, objectives, estimates, goals, or other characterizations of future events is a forward-looking statement. These forward-looking statements can often be identified as such because the context of the statement will include words such as "believes," "anticipates," "expects" or words of similar import.

Actual results may differ materially from those expressed in these forward-looking statements, which speak only as of the date hereof, and are subject to change at any time without notice. Existing and prospective investors, customers and other third parties transacting business with New Relic are cautioned not to place undue reliance on this forward-looking information. The achievement or success of the matters covered by such forward-looking statements are based on New Relic's current assumptions, expectations, and beliefs and are subject to substantial risks, uncertainties, assumptions, and changes in circumstances that may cause the actual results, performance, or achievements to differ materially from those expressed or implied in any forward-looking statement. Further information on factors that could affect such forward-looking statements is included in the filings New Relic makes with the SEC from time to time. Copies of these documents may be obtained by visiting New Relic's Investor Relations website at [ir.newrelic.com](http://ir.newrelic.com) or the SEC's website at [www.sec.gov](http://www.sec.gov).

New Relic assumes no obligation and does not intend to update these forward-looking statements, except as required by law. New Relic makes no warranties, expressed or implied, in this presentation or otherwise, with respect to the information provided.

# Who Are We?



# Who Are We?



Caito Scherr

Nik Davis







Portland Office




# New Relic



# Who Are We?



# Why Are We Here?





# Why Are We Here?



# Thesis

You can be successful in your Flink project using data-driven development

# Thesis

You can be successful in your Flink project using data-driven development

1. Proving viability early with a Proof of Concept

# Thesis

You can be successful in your Flink project using data-driven development

1. Proving viability early with a Proof of Concept
2. Letting data from our Flink monitoring drive our development



# Agenda



## The Challenge

Proof of Concept

Flink Monitoring

6000 Meter View

# Agenda

The Challenge

 **Proof of Concept**

Flink Monitoring

6000 Meter View

# Agenda

The Challenge

Proof of Concept

 **Flink-Monitoring**

6000 Meter View

# Agenda

The Challenge


Proof of Concept

Flink Monitoring


 **6000 Meter View**




# The Challenge



# Usage




# Usage



# Usage

## Agent records

```
{  
  "hostname": "Jane's host",  
  "timestamp": "2018-09-05T10:00:00+02:00"  
}  
...  
{  
  "hostname": "Jane's host",  
  "timestamp": "2018-09-05T11:00:00+02:00"  
}
```




## Summary *daily* record


```
{  
  "hostname": "Jane's host",  
  "daysSinceEpoch": 17779,  
  "hoursUsed": 2  
}
```




# Usage



# Usage





Usage



● Message ● Operations ● Data

Ingests **5M** messages


Runs **15M** operations

Processes **3GB** data

EVERY MINUTE



# Proof of Concept



# Proof of Concept: What

- Tiny working model



# Proof of Concept: What

- Tiny working model
- Data-driven development



# Proof of Concept: What

- Tiny working model
- Data-driven development
- Technically viable





# Proof of Concept: What

- Tiny working model
- Data-driven development
- Technically viable
  - Steel thread: start to finish testing



# Proof of Concept: What

- Tiny working model
- Data-driven development
- Technically viable
  - Steel thread: start to finish testing
  - Minimal, miniaturize



## Proof of Concept: Why

- Prove viability early



## Proof of Concept: Why

- Prove viability early
- Goal: miniaturize components to build quickly, keep scale and technical requirements

## Proof of Concept: Why

- Prove viability early
- Goal: miniaturize components to build quickly, keep scale and technical requirements
- **Make it quick, make it work**

# Proof of Concept: Pre-Work

- Rubric

# Proof of Concept: Pre-Work

- Rubric
  - Time windowing

# Proof of Concept: Pre-Work

- Rubric
  - Time windowing
  - Parallelizing

# Proof of Concept: Pre-Work

- Rubric
  - Time windowing
  - Parallelizing
  - Development velocity

# Proof of Concept: Pre-Work

- Rubric
  - Time windowing
  - Parallelizing
  - Development velocity
    - Maintainability



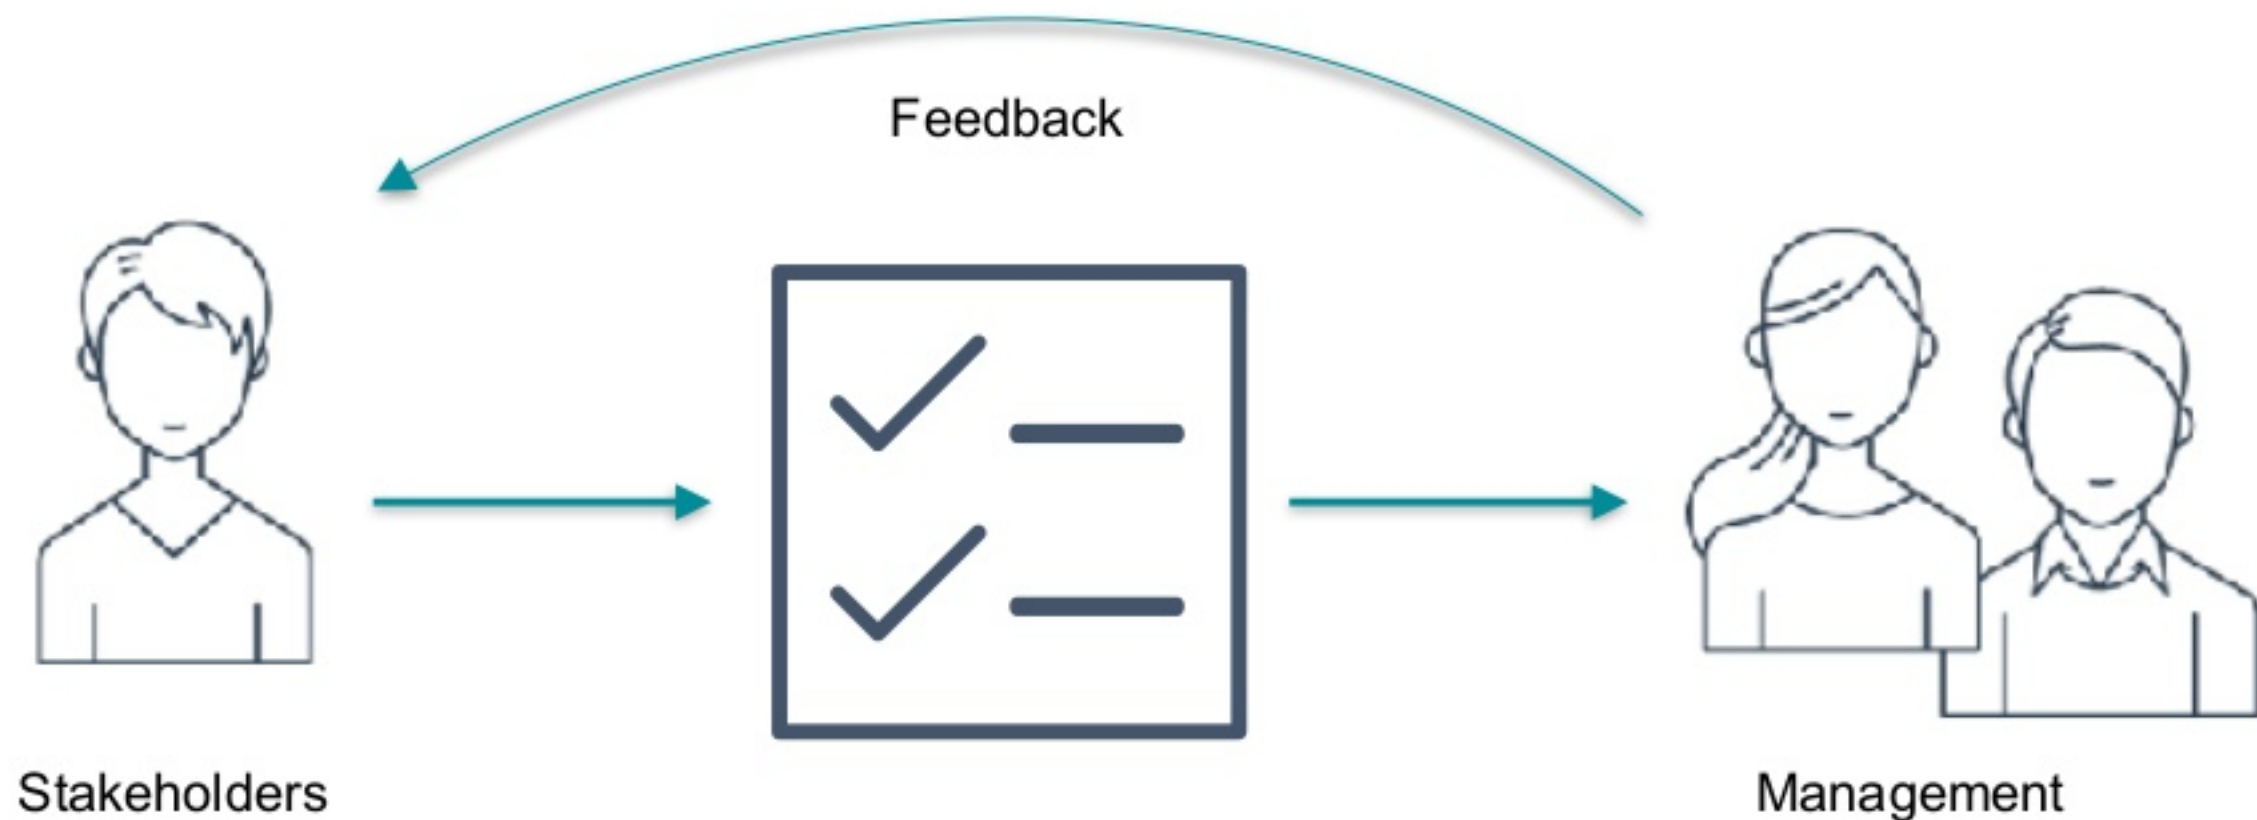
# Proof of Concept: Pre-Work

- Rubric
  - Time windowing
  - Parallelizing
  - Development velocity
    - Maintainability
  - Failure modes

# Proof of Concept: Pre-Work

- Rubric
  - Time windowing
  - Parallelizing
  - Development velocity
    - Maintainability
  - Failure modes
  - Testing

# Proof of Concept: Pre-Work



# Proof of Concept: How

Prove a steel thread

```
DataStream<DailyHostSummary> dailyHostStream = agentDataStream
    .keyBy("accountId", "hostId")
    .window(TumblingEventTimeWindows.of(Time.minutes(3))) // change to 24 hours (1440 minutes)
    .fold(new DailyHostSummary(), new DailyHostSummarizer());
```

# Proof of Concept: How

```
agentHeartbeats
    .assignTimestampsAndWatermarks(
        new BoundedOutOfOrdernessTimestampExtractor<AgentHeartbeat>{
            Time.seconds(env.getMaxAllowedLatenessApm())
        } {
            @Override
            public long extractTimestamp(AgentHeartbeat element) {
                return element.getHeartbeatTime() * 1000;
            }
        })
    .name(...).uid(...)
    .map(value -> new HostState(value))
    .keyBy["consumingAccountId", "hostId"]
    .window(TumblingEventTimeWindows.of(Time.minutes(env.getAggregationWindowDurationInMinutes())))
    .reduce(new ReduceHostUsageFunction())
    .name(...).uid(...)
    .map(new MapHostStateToEventFunction(env.isProduceApmProductionEvent()))
    .name(...).uid(...)
    .map(value -> (ApmDailyBillingRecord) value)
    .map(new MapOverrideInstanceSizeFunction()).name(...).uid(...)
    .map(new MapCloudInstanceSizeFunction()).name(...).uid(...)
    .map(new MapAgentCalculatedInstanceSizeFunction()).name(...).uid(...)
    .map(new MapDefaultInstanceSizeFunction()).name(...).uid(...)
    .map(new MapInstanceSizeCapFunction()).name(...).uid(...)
    .map(value -> (NrDailyHost) value);
```

# Proof of Concept: How

```
agentHeartbeats
    .assignTimestampsAndWatermarks(
        new BoundedOutOfOrdernessTimestampExtractor<AgentHeartbeat>{
            Time.seconds(env.getMaxAllowedLatenessApm())
        } {
            @Override
            public long extractTimestamp(AgentHeartbeat element) {
                return element.getHeartbeatTime() * 1000;
            }
        })
    .name(...).uid(...)
    .map(value -> new HostState(value))
    .keyBy("consumptionAccountid", "hostid")
    .window(TumblingEventTimeWindows.of(Time.minutes(env.getAggregationWindowDurationInMinutes())))
    .reduce(new ReduceHostUsageFunction())
    .name(...).uid(...)
    .map(new MapHostStateToEventFunction(env.isProduceApmProductionEvent()))
    .name(...).uid(...)
    .map(value -> (ApmDailyBillingRecord) value)
    .map(new MapOverrideInstanceSizeFunction()).name(...).uid(...)
    .map(new MapCloudInstanceSizeFunction()).name(...).uid(...)
    .map(new MapAgentCalculatedInstanceSizeFunction()).name(...).uid(...)
    .map(new MapDefaultInstanceSizeFunction()).name(...).uid(...)
    .map(new MapInstanceSizeCapFunction()).name(...).uid(...)
    .map(value -> (NrDailyHost) value);
```



# Proof of Concept: How

Make it quick

```
DataStream<DailyHostSummary> dailyHostStream = agentDataStream
    .keyBy("accountId", "hostId")
    .window(TumblingEventTimeWindows.of(Time.minutes(3))) // change to 24 hours (1440 minutes)
    .fold(new DailyHostSummary(), new DailyHostSummarizer());
```



# Proof of Concept: How

```
AGGREGATION_WINDOW_DURATION_IN_MINUTES =  
Integer.parseInt(getEnvWithDefault("AGGREGATION_WINDOW_DURATION_IN_MINUTES", "1440"));
```

# Proof of Concept: How

```
agentHeartbeats
    .assignTimestampsAndWatermarks(
        new BoundedOutOfOrdernessTimestampExtractor<AgentHeartbeat>{
            Time.seconds(env.getMaxAllowedLatenessApm())
        } {
            @Override
            public long extractTimestamp(AgentHeartbeat element) {
                return element.getHeartbeatTime() * 1000;
            }
        })
    .name(...).uid(...)
    .map(value -> new HostState(value))
    .keyBy("consumingAccountId", "hostId")
    .window(TumblingEventTimeWindows.of(Time.minutes(env.getAggregationWindowDurationInMinutes())))
    .reduce(new ReduceHostStateFunction())
    .name(...).uid(...)
    .map(new MapHostStateToEventFunction(env.isProduceApmProductionEvent()))
    .name(...).uid(...)
    .map(value -> (ApmDailyBillingRecord) value)
    .map(new MapOverrideInstanceSizeFunction()).name(...).uid(...)
    .map(new MapCloudInstanceSizeFunction()).name(...).uid(...)
    .map(new MapAgentCalculatedInstanceSizeFunction()).name(...).uid(...)
    .map(new MapDefaultInstanceSizeFunction()).name(...).uid(...)
    .map(new MapInstanceSizeCapFunction()).name(...).uid(...)
    .map(value -> (NrDailyHost) value);
```

# Proof of Concept: Overview

- Tiny, working model

## Proof of Concept: Overview

- Tiny, working model
- Having done the PoC meant that we had:

## Proof of Concept: Overview

- Tiny, working model
- Having done the PoC meant that we had:
  - Confidence in our choice and roadmap

## Proof of Concept: Overview

- Tiny, working model
- Having done the PoC meant that we had:
  - Confidence in our choice and roadmap
  - Understanding —> Velocity

## Proof of Concept: Overview


- Tiny, working model
- Having done the PoC meant that we had:
  - Confidence in our choice and roadmap
  - Understanding —> Velocity
  - Groundwork



## Proof of Concept: Overview

- Tiny, working model
- Having done the PoC meant that we had:
  - Confidence in our choice and roadmap
  - Understanding —> Velocity
  - Groundwork
  - Troubleshooting done early

# Flink Monitoring



# Flink Monitoring


Data-driven development leads to real gains in both code quality and developer velocity

# Flink Monitoring

It's not complicated!

# Flink Monitoring

Ship code




# Flink Monitoring

Collect data



# Flink Monitoring

Improve your product



# Flink Monitoring

What do we track?



# Flink Monitoring

What do we track? Examples

- State
  - size
  - backup duration



# Flink Monitoring

What do we track? Examples

- State
  - size
  - backup duration
- Throughput



# Flink Monitoring

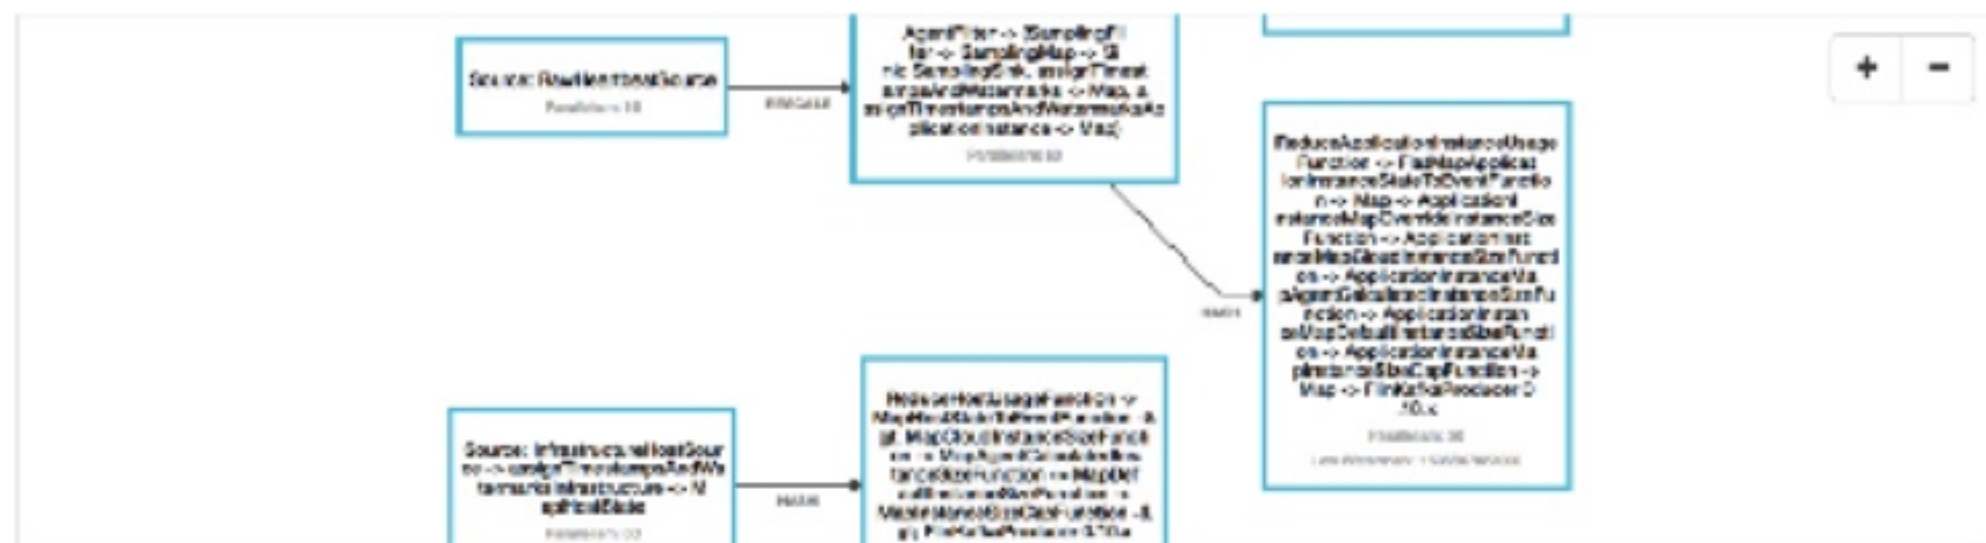
What do we track? Examples

- State
  - size
  - backup duration
- Throughput
- Watermark



# Flink Monitoring

Flink cluster UI is a start ...



☐ Aggregate task statistics by TaskManager

Start Time	End Time	Duration	Name	Bytes received	Records received	Bytes sent	Records sent	P...
2018-09-28, 14:03:35	2018-09-28, 22:44:03	1d 8h	Source: RawHeartbeatSource	0 B	0	2.70 TB	6,903,157,557	16

# Flink Monitoring

## Watermark Samples (sec past realtime)

Since 60 minutes ago



# Flink Monitoring

How?



# Flink Monitoring

Flink has a great metrics system

# Flink Monitoring

Flink has a great metrics system; use it!

# Flink Monitoring

You can ...

Use an existing framework

# Flink Monitoring

You can ...

Use an existing framework

Write your own

# Flink Monitoring—Flink's interfaces


```
1  public interface MetricReporter {
2
3      void open(MetricConfig config);
4
5      void close();
6
7      void notifyOfAddedMetric(Metric metric, String metricName, MetricGroup group);
8
9      void notifyOfRemovedMetric(Metric metric, String metricName, MetricGroup group);
10
11 }
12
13 public interface Scheduled {
14
15     void report();
16
17 }
```

# Flink Monitoring

## Checkpoint size

[Embed](#)

Since 26 Aug 15:00 PDT until 27 Aug 19:00 PDT




# Flink Monitoring

## Checkpoint size

Since 26 Aug 15:00 PDT until 27 Aug 19:00 PDT

 Embed





# Flink Monitoring

We have data, now what?

# Flink Monitoring

Use it!


# Flink Monitoring

Use it!




### APM lag

Since 60 minutes ago




### Watermark Samples (sec post realtime)

Since 60 minutes ago




### Synthetic Checks Lag

Since 60 minutes ago



### Kafka consumption

Since 60 minutes ago



### Infrastructure lag

Since 60 minutes ago




### Calculator Checkpoint Size

Since 60 minutes ago



### Throughput at Reducers

Since 60 minutes ago



# Flink Monitoring

Use it!



# Flink Monitoring

Now what? Keep consuming your monitoring

# Flink Monitoring

Information radiation






# Flink Monitoring

## Checkpoint size

Since 26 Aug 15:00 PDT until 27 Aug 19:00 PDT

 Embed




# Flink Monitoring


## Checkpoint size

Since 26 Aug 15:00 PDT until 27 Aug 19:00 PDT

 Embed



# Flink Monitoring



# Flink Monitoring

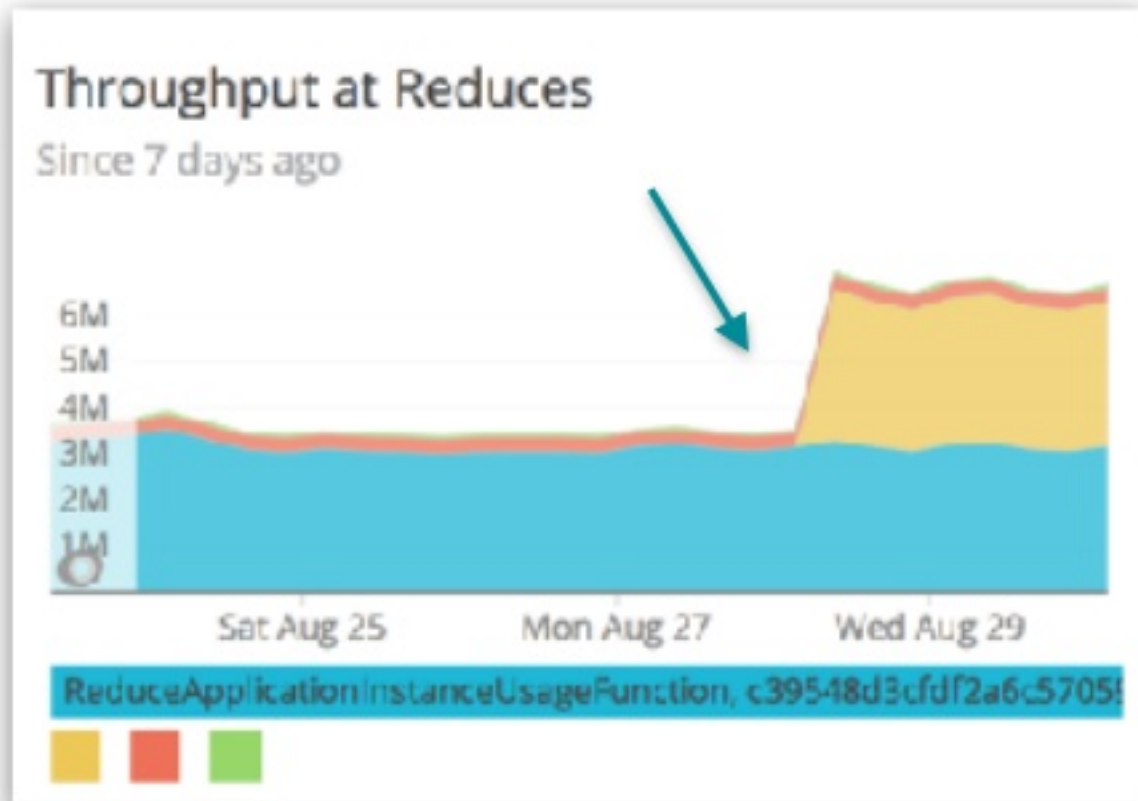
Metrics have some area for growth

- No distributed tracing, Flink-managed memory


# Flink Monitoring

Metrics have some area for growth

- No distributed tracing, Flink-managed memory
- Can be buggy



# 6000 Meter View



# 6000 Meter View

PoCs:

Flink Monitoring:

## 6000 Meter View

PoCs: **make it quick, make it work**

Flink Monitoring:



## 6000 Meter View

PoCs: **make it quick, make it work**

Flink Monitoring: **use it, everywhere!**

# Thank You!

Flink Forward

AX Team (past and present)

Alysa Wood

Jen Hammond

Ralph Bodenner

Jeff Ostrin

Maureen Dugan

Moof Mayeda

Nate Borrebach

Ruby Andrews

Ron Crocker

Bill Green

## Keep in touch!

Caito: [cscherr@newrelic.com](mailto:cscherr@newrelic.com)

Nik: [ndavis@newrelic.com](mailto:ndavis@newrelic.com)

New Relic Blog: <https://blog.newrelic.com>

New Relic Blog, new Flink article! <https://blog.newrelic.com/engineering/what-is-apache-flink/>

Join our team!! <https://newrelic.com/about/culture>

# Photo Credits

Profile Photo, Caito  
Matthew Scharr, 2017

Profile Photo, Nik  
New Relic, 2017

AX Team Photo:  
Moof Mayeda, 2018

Flink Logo:

<https://flink.apache.org/material.html>  
<https://www.apache.org/foundation/marks/>

Portland Skyline:

Michael Wiley/  
<https://tinyurl.com/yc2aubdt>/CC BY

Model Train:

Kevin Phillips/  
<https://tinyurl.com/yaoow6jo>/CC BY