

# Detecting Patterns in Event Streams with Flink SQL

Dawid Wysakowicz, SOFTWARE ENGINEER

# ABOUT DATA ARTISANS



Original Creators of  
Apache Flink®

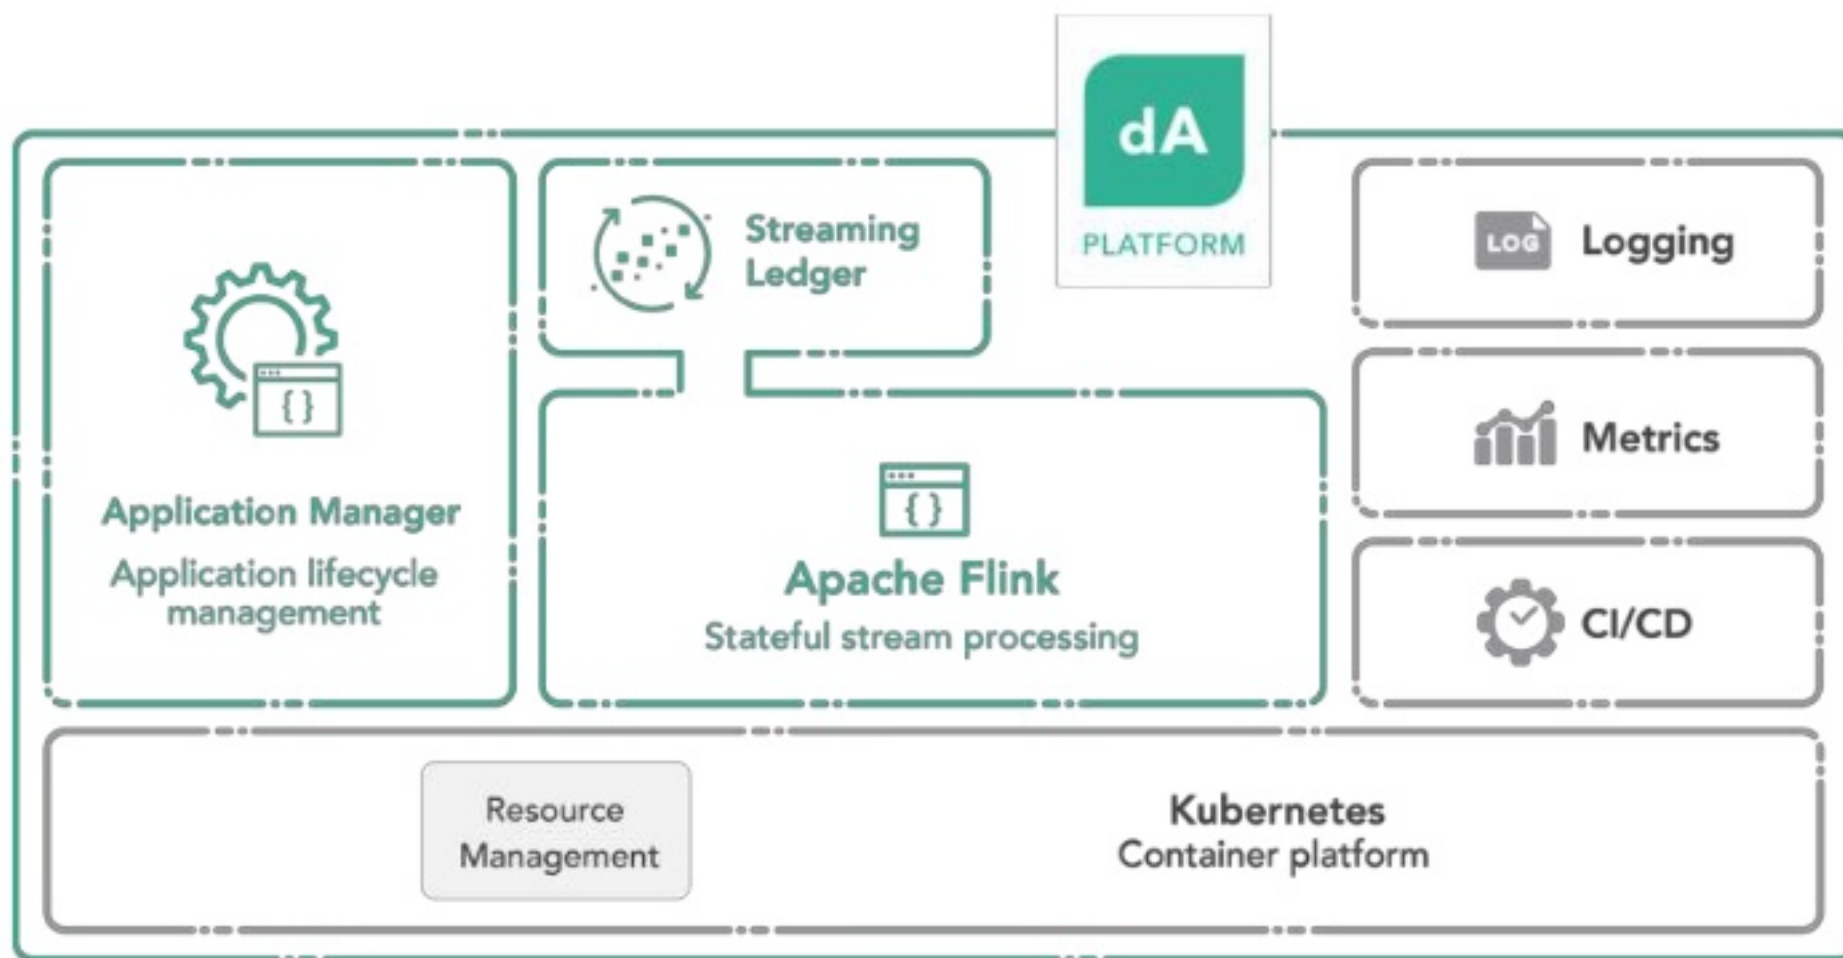


PLATFORM

Enterprise Ready  
Real Time Stream Processing



# FREE TRIAL DOWNLOAD



[data-artisans.com/download](https://data-artisans.com/download)



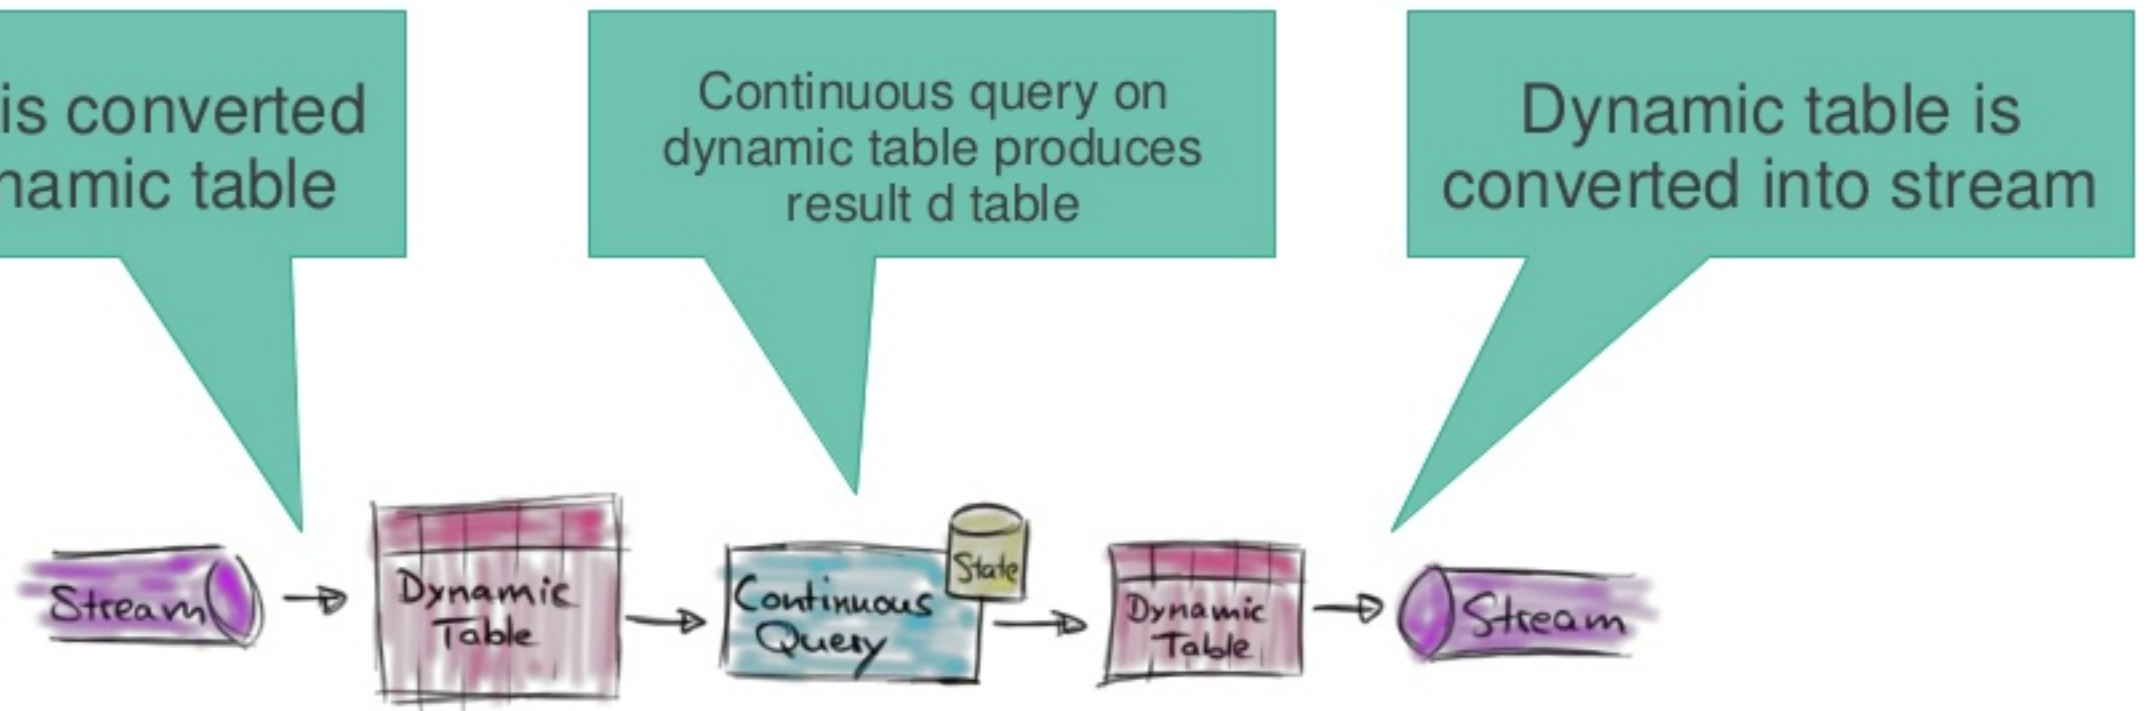
# Why people like SQL?

- Well known interface: "lingua franca of data processing"
- No programming required - easier to learn
- Declarative way to express your query, WHAT not HOW
- Out-of-the box optimization
- Reusability
  - built-in functions
  - user defined functions



# SQL on Streams

- Unified semantics for stream and batch input
  - Same query, same input, same result



# The New York Taxi Rides Data Set

- The New York City & Limousine Commission provides a public data set about past taxi rides in New York City
- We can derive a streaming table from the data

- Table: **TaxiRides**


rideId:	BIGINT	// ID of the taxi ride
taxiId:	BIGINT	// ID of the taxi cab
driverId:	BIGINT	// ID of the driver
isStart:	BOOLEAN	// flag for pick-up (true) or drop-off (false) event
lon:	DOUBLE	// longitude of pick-up or drop-off location
lat:	DOUBLE	// latitude of pick-up or drop-off location
rowTime:	TIMESTAMP	// time of pick-up or drop-off event





# Number of rides per 30 minutes per area

Number of rides



# SQL Example

SELECT

toCellId(lat, lon) as cellId,

COUNT(distinct rideId) as rideCount,

TUMBLE\_ROWTIME(rowTime, INTERVAL '30' minute) AS rowTime,

cast (TUMBLE\_START(rowTime, INTERVAL '30' minute) as TIMESTAMP) AS startTime,

cast(TUMBLE\_END(rowTime, INTERVAL '30' minute) as TIMESTAMP) AS endTime

FROM

TaxiRides

GROUP BY

toCellId(lat, lon),

TUMBLE(rowTime, INTERVAL '30' minute)





# SQL Support in Flink 1.6

- SELECT FROM WHERE
- GROUP BY / HAVING
  - Non-windowed, TUMBLE, HOP, SESSION windows
- JOIN
  - Windowed INNER, LEFT / RIGHT / FULL OUTER JOIN
  - Non-windowed INNER / LEFT / RIGHT / FULL OUTER JOIN
- Scalar, aggregation, table-valued UDFs
- Many built-in functions
  
- [streaming only] OVER / WINDOW
  - UNBOUNDED / BOUNDED PRECEDING



# SQL Support in Flink 1.6

- SELECT FROM WHERE
- GROUP BY / HAVING
  - Non-windowed, TUMBLE, HOP, SESSION windows
- JOIN
  - Windowed INNER, LEFT / RIGHT / FULL OUTER JOIN
  - Non-windowed INNER / LEFT / RIGHT / FULL OUTER JOIN
- Scalar, aggregation, table-valued UDFs
- Many built-in functions
- [streaming only] OVER / WINDOW
  - UNBOUNDED / BOUNDED PRECEDING

12:20-13:00 Kesselhaus  
Timo Walther, data Artisans  
Flink SQL in Action



# What is hard with SQL?

- Find rides with mid-stops



# Mid stops

```
Pattern.<Row>begin("S").where((row) -> {  
    return row.isStart == true;  
}).next("E").where((row) -> {  
    return row.isStart == true;  
});
```

```
CEP.pattern(input.keyBy("driverId"), pattern)  
    .flatMapSelect(  
        new PatternFlatSelectFunction<Row, Row>() {  
            @Override  
            public void flatSelect(  
                Map<String, List<Row>> pattern,  
                Collector<Row> out) throws Exception {  
                out.collect(  
                    pattern.get("S").get(0).getRideId  
                );  
            }  
        })  
    );
```



---

# MATCH\_RECOGNIZE

SQL:2016 extension



# Common use-cases

- stock market analysis
- customer behaviour
- tracking money laundering
- service quality
- network intrusion detection



# Syntax

tableReference:

tablePrimary

[ matchRecognize ]

[ [ AS ] alias [ '(' columnAlias [, columnAlias ]\* ')' ] ]

select:

SELECT [ STREAM ] [ ALL | DISTINCT ]

{ \* | projectItem [, projectItem ]\* }

FROM tableExpression

[ matchRecognize ]

[ WHERE booleanExpression ]

[ GROUP BY { groupItem [, groupItem ]\* } ]

....






# Rides with mid-stops

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.rideId as sRideId  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```



partition the data by  
given field = keyBy



# Rides with mid-stops


```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.rideId as sRideId  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```

- specify order
- primary order = Event or Processing time



# Rides with mid-stops

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.rideId as sRideId  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```



construct pattern




# Rides with mid-stops

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.rideId as sRideId  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```

extract measures from  
matched sequence



# Multi-Stop



# Rides with more than one mid-stop


```
SELECT *
FROM TaxiRides
MATCH_RECOGNIZE (
  PARTITION BY driverId
  ORDER BY rowTime
  MEASURES
    S.rideId as sRideId
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (S E)
  DEFINE
    S AS S.isStart = true,
    E AS E.isStart = true
)
```

```
SELECT *
FROM TaxiRides
MATCH_RECOGNIZE (
  PARTITION BY driverId
  ORDER BY rowTime
  MEASURES
    S.rideId as sRideId
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (S M{2,} E)
  DEFINE
    S AS S.isStart = true,
    M AS M.rideId <> S.rideId,
    E AS E.isStart = false
        AND E.rideId = S.rideId
)
```



# Rush (peak) hours - V shape

Number of rides





# Statistics per Area

```
CREATE VIEW RidesInArea AS

SELECT

    toCellId(lat, lon) as cellId,
    COUNT(distinct rideId) as rideCount,
    TUMBLE_ROWTIME(rowTime, INTERVAL '30' minute) AS rowTime,
    cast (TUMBLE_START(rowTime, INTERVAL '30' minute) as TIMESTAMP) AS startTime,
    cast(TUMBLE_END(rowTime, INTERVAL '30' minute) as TIMESTAMP) AS endTime

FROM

    TaxiRides

GROUP BY

    toCellId(lat, lon),

    TUMBLE(rowTime, INTERVAL '30' minute)
```




# Rush (peak) hours

Use previous table/view

```
SELECT * FROM RidesInArea
MATCH_RECOGNIZE(
  PARTITION BY cellId
  MEASURES
    FIRST(UP.rideCount) AS preCnt,
    LAST(UP.rideCount) AS rushCnt, FIRST(UP.startTime) as rushStart,
    LAST(UP.endTime) as rushEnd,
    LAST(DOWN.rideCount) AS postCnt
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (UP{4,} DOWN{2,} E)
  DEFINE UP AS UP.rideCount > PREV(UP.rideCount) or PREV(UP.rideCount) IS
NULL,
        DOWN AS DOWN.rideCount < PREV(DOWN.rideCount) AND
        DOWN.rideCount < LAST(UP.rideCount),
  E AS E.rideCount > LAST(DOWN.rideCount)
```

Number of rides



# Rush (peak) hours

```
SELECT * FROM RidesInArea
```

Use previous table/view

```
MATCH_RECOGNIZE(
```

```
  PARTITION BY cellId
```

```
  MEASURES
```

```
    FIRST(UP.rideCount) AS preCnt,
```

```
    LAST(UP.rideCount) AS rushCnt, FIRST(UP.startTime) as rushStart,
```

```
    LAST(UP.endTime) as rushEnd,
```

```
    LAST(DOWN.rideCount) AS postCnt
```

```
  AFTER MATCH SKIP PAST LAST ROW
```

```
  PATTERN (UP{4,} DOWN{2,} E)
```

```
  DEFINE UP AS UP.rideCount > PREV(UP.rideCount) or PREV(UP.rideCount) IS
```


```
NULL,
```

```
    DOWN AS DOWN.rideCount < PREV(DOWN.rideCount) AND
```

```
    DOWN.rideCount < LAST(UP.rideCount),
```

```
  E AS E.rideCount > LAST(DOWN.rideCount)
```

Number of rides




apply match to result of the inner query



# Rush (peak) hours

```
SELECT * FROM RidesInArea
MATCH_RECOGNIZE(
  PARTITION BY cellId
  MEASURES
    FIRST(UP.rideCount) AS preCnt,
    LAST(UP.rideCount) AS rushCnt, FIRST(UP.startTime) as rushStart,
    LAST(UP.endTime) as rushEnd,
    LAST(DOWN.rideCount) AS postCnt
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (UP{4,} DOWN{2,} E)
  DEFINE UP AS UP.rideCount > PREV(UP.rideCount) or PREV(UP.rideCount) IS
NULL,
    DOWN AS DOWN.rideCount < PREV(DOWN.rideCount) AND
    DOWN.rideCount < LAST(UP.rideCount),
  E AS E.rideCount > LAST(DOWN.rideCount)
```



access elements of  
looping pattern



# Rush (peak) hours

```
SELECT * FROM RidesInArea
```

```
MATCH_RECOGNIZE(
```

```
  PARTITION BY cellId
```

```
  MEASURES
```

```
    FIRST(UP.rideCount) AS preCnt,
```

```
    LAST(UP.rideCount) AS rushCnt, FIRST(UP.startTime) as rushStart,
```

```
    LAST(UP.endTime) as rushEnd,
```

```
    LAST(DOWN.rideCount) AS postCnt
```

```
AFTER MATCH SKIP PAST LAST ROW
```

```
PATTERN (UP{4,} DOWN{2,} E)
```


```
DEFINE UP AS UP.rideCount > PREV(UP.rideCount) or PREV(UP.rideCount) IS
```

```
NULL,
```

```
    DOWN AS DOWN.rideCount < PREV(DOWN.rideCount) AND
```

```
    DOWN.rideCount < LAST(UP.rideCount),
```

```
E AS E.rideCount > LAST(DOWN.rideCount)
```



# Feature set of MATCH\_RECOGNIZE (Experimental)

- Quantifiers support:
  - + (one or more), \* (zero or more), {x,y} (times)
  - greedy(default), ?(reluctant)
    - with some restrictions (not working for last pattern)
- After Match Skip
  - skip\_to\_first/last, skip\_past\_last, skip\_to\_next
- Not supported:
  - alter(|), permute, exclude '{- -}'
  - aggregates within MATCH\_RECOGNIZE





# Summary

- opens for new use cases, that assume order of events
- enables reuse of SQL goods
- still in experimental phase
  - [https://github.com/dawidwys/flink/SQL\\_MATCH\\_RECOGNIZE](https://github.com/dawidwys/flink/SQL_MATCH_RECOGNIZE)





---

# THANK YOU!

@dwysakowicz

@dataArtisans

@ApacheFlink

## WE ARE HIRING

[data-artisans.com/careers](https://data-artisans.com/careers)

**dataArtisans**

---

# Under the hood

Dissecting MATCH\_RECOGNIZE



# Under the hood

```
SELECT * FROM TaxiStatistics
MATCH_RECOGNIZE(
  PARTITION BY cellId
  MEASURES
    LAST(UP.rideCount) AS rushCnt,
    FIRST(UP.startTime) as rushStart,
    LAST(UP.endTime) as rushEnd
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (UP{4,} DOWN{2,} E)
  DEFINE
    UP AS UP.rideCount > PREV(UP.rideCount)
    or PREV(UP.rideCount) IS NULL,
    DOWN AS DOWN.rideCount <
      LAST(UP.rideCount),
    E AS E.rideCount >
      LAST(DOWN.rideCount)
)
```

```
Pattern.<Row>begin("UP", SKIP_PAST_LAST).timesOrMore(4).where((row, ctx) -> {
    Long prevRideCount = ctx.get("UP").tail().get("rideCount");
    return prevRideCount == null || row.get("rideCount") >
prevRideCount;
}).next("DOWN").timesOrMore(2).where( (row, ctx) -> {
    ...
}).next("E").where( (row, ctx) -> {
    ...
});

CEP.pattern(input.keyBy("cellId"), pattern).flatSelect(
  new PatternFlatSelectFunction<Row, Row>() {
    @Override
    public void flatSelect(Map<String, List<Row>> pattern,
      Collector<Row> out) throws Exception {
      Row last = pattern.get("UP").tail();
      Row first = pattern.get("UP").head();
      out.collect(Row.of(...));
    }
  }
);
```



# Under the hood

```
SELECT * FROM TaxiStatistics
MATCH_RECOGNIZE(
  PARTITION BY cellId
  MEASURES
    LAST(UP.rideCount) AS rushCnt,
    FIRST(UP.startTime) as rushStart,
    LAST(UP.endTime) as rushEnd
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (UP{4,} DOWN{2,} E)
  DEFINE
    UP AS UP.rideCount > PREV(UP.rideCount)
    or PREV(UP.rideCount) IS NULL,
    DOWN AS DOWN.rideCount <
      LAST(UP.rideCount),
    E AS E.rideCount >
      LAST(DOWN.rideCount)
)
```

```
Pattern.<Row>begin("UP", SKIP_PAST_LAST).timesOrMore(4).where((row, ctx) -> {
  Long prevRideCount = ctx.get("UP").tail().get("rideCount");
  return prevRideCount == null || row.get("rideCount") >
    prevRideCount;
}).next("DOWN").timesOrMore(2).where( (row, ctx) -> {
  ...
}).next("E").where( (row, ctx) -> {
  ...
});

CEP.pattern(input.keyBy("cellId"), pattern).flatMapSelect(
  new PatternFlatSelectFunction<Row, Row>() {
    @Override
    public void flatSelect(Map<String, List<Row>> pattern,
      Collector<Row> out) throws Exception {
      Row last = pattern.get("UP").tail();
      Row first = pattern.get("UP").head();
      out.collect(Row.of(...));
    }
  }
);
```



# Under the hood

```
SELECT * FROM TaxiStatistics
MATCH_RECOGNIZE(
  PARTITION BY cellId
  MEASURES
    LAST(UP.rideCount) AS rushCnt,
    FIRST(UP.startTime) as rushStart,
    LAST(UP.endTime) as rushEnd
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (UP{4,} DOWN{2,} E)
  DEFINE
    UP AS UP.rideCount > PREV(UP.rideCount)
      or PREV(UP.rideCount) IS NULL,
    DOWN AS DOWN.rideCount <
      LAST(UP.rideCount),
    E AS E.rideCount >
      LAST(DOWN.rideCount)
)
```

```
Pattern.<Row>begin("UP", SKIP_PAST_LAST).timesOrMore(4).where((row, ctx) -> {
    Long prevRideCount = ctx.get("UP").tail().get("rideCount");
    return prevRideCount == null || row.get("rideCount") >
prevRideCount;
}).next("DOWN").timesOrMore(2).where( (row, ctx) -> {
    ...
}).next("E").where( (row, ctx) -> {
    ...
});

CEP.pattern(input.keyBy("cellId"), pattern).flatSelect(
  new PatternFlatSelectFunction<Row, Row>() {
    @Override
    public void flatSelect(Map<String, List<Row>> pattern,
      Collector<Row> out) throws Exception {
      Row last = pattern.get("UP").tail();
      Row first = pattern.get("UP").head();
      out.collect(Row.of(...));
    }
  }
);
```





# Under the hood

```
SELECT * FROM TaxiStatistics
MATCH_RECOGNIZE(
  PARTITION BY cellId
  MEASURES
    LAST(UP.rideCount) AS rushCnt,
    FIRST(UP.startTime) as rushStart,
    LAST(UP.endTime) as rushEnd
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (UP{4,} DOWN{2,} E)
  DEFINE
    UP AS UP.rideCount > PREV(UP.rideCount)
    or PREV(UP.rideCount) IS NULL,
    DOWN AS DOWN.rideCount <
    LAST(UP.rideCount),
    E AS E.rideCount >
    LAST(DOWN.rideCount)
)
```

```
Pattern.<Row>begin("UP", SKIP_PAST_LAST).timesOrMore(4).where((row, ctx) ->
{
  Long prevRideCount = ctx.get("UP").tail().get("rideCount");
  return prevRideCount == null || row.get("rideCount") >
prevRideCount;
}).next("DOWN").timesOrMore(2).where( (row, ctx) -> {
  ...
}).next("E").where( (row, ctx) -> {
  ...
});

CEP.pattern(input.keyBy("cellId"), pattern).flatSelect(
new PatternFlatSelectFunction<Row, Row>() {
  @Override
  public void flatSelect(Map<String, List<Row>> pattern,
Collector<Row> out) throws Exception {
    Row last = pattern.get("UP").tail();
    Row first = pattern.get("UP").head();
    out.collect(Row.of(...));
  }
});
```



---

# THANK YOU!

@dwysakowicz

@dataArtisans

@ApacheFlink

## WE ARE HIRING

[data-artisans.com/careers](https://data-artisans.com/careers)

**dataArtisans**