
FLINK AS A LIBRARY (AND STILL AS A FRAMEWORK)

- GARY YAO

ABOUT DATA ARTISANS



Original Creators of
Apache Flink®

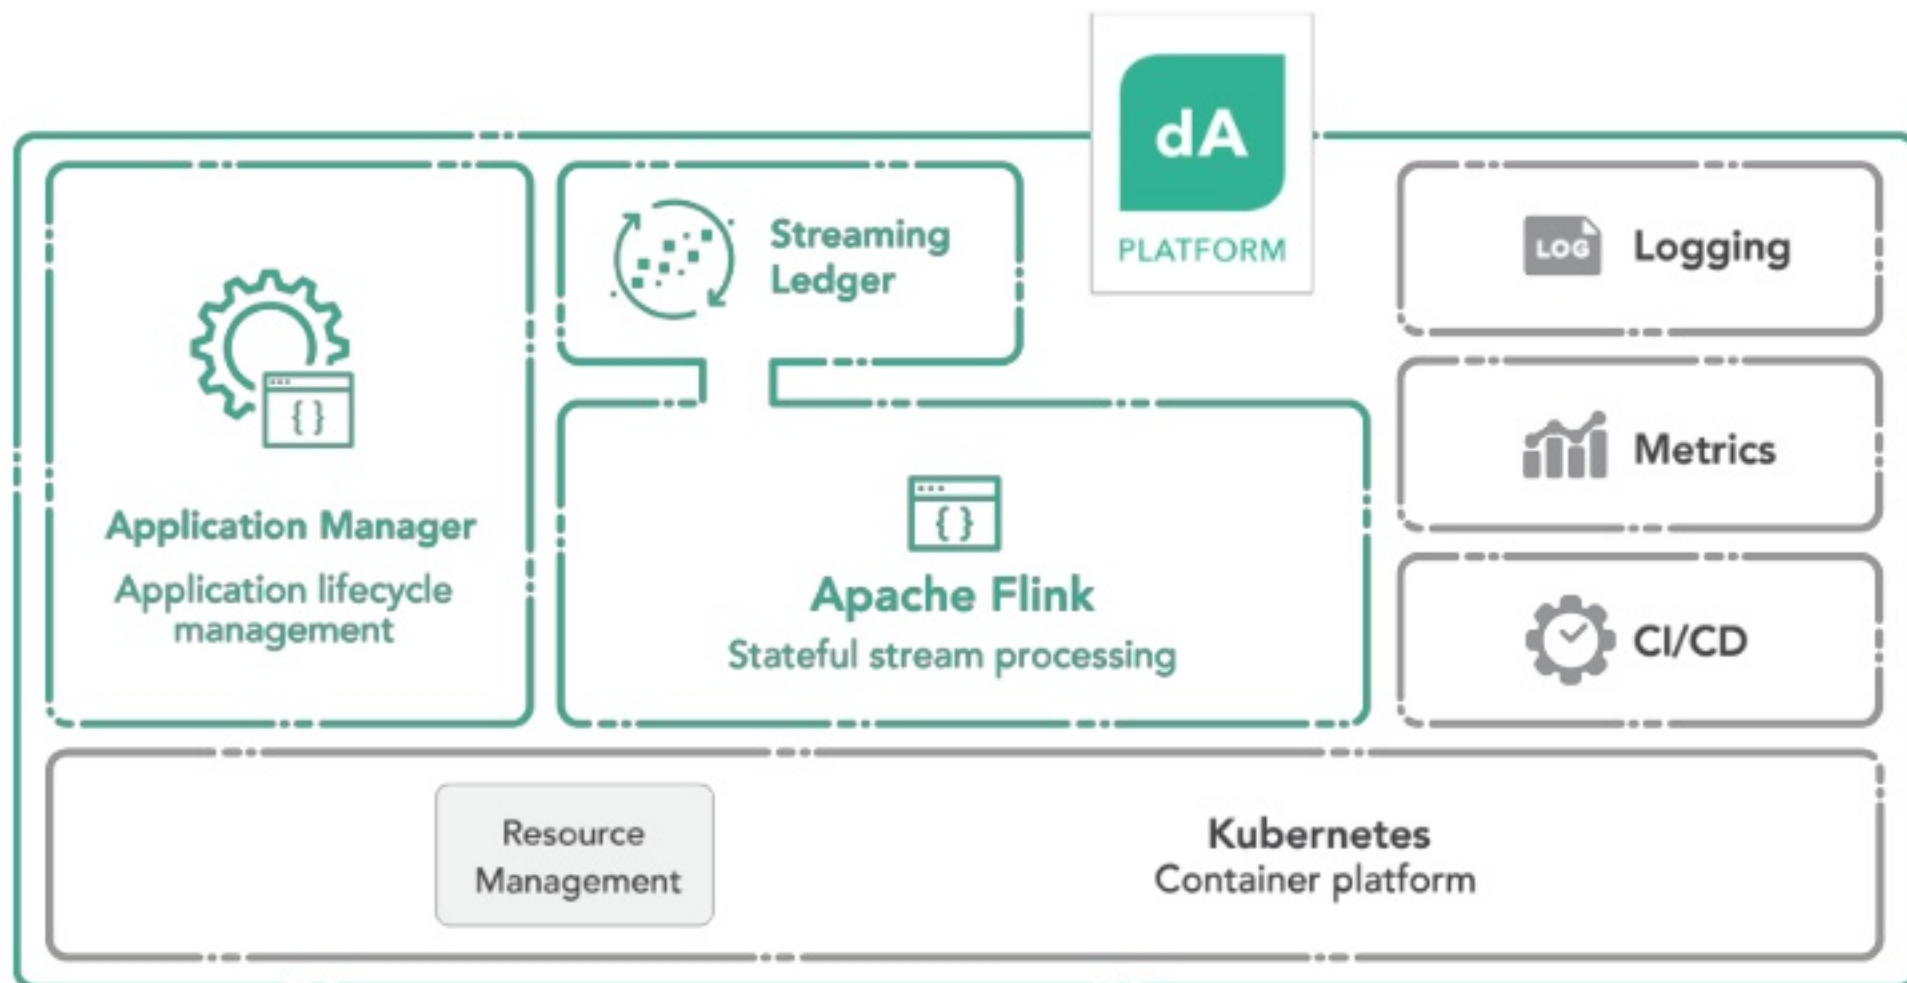


PLATFORM

Real Time Stream Processing Enterprise
Ready



FREE TRIAL DOWNLOAD



data-artisans.com/download



MOTIVATION



FLINK DEPLOYMENT OPTIONS & LIMITATIONS

Before Apache Flink 1.5


| | Session cluster | Job cluster |
|------------|-----------------|-------------|
| Standalone | ✓ | |
| YARN | ✓ | ✓ |
| Mesos | ✓ | |
| Kubernetes | ✓ | |

- Shortcomings:
 - Fixed number of TaskManagers
 - Difficult to support autoscaling
 - Supporting new cluster managers difficult
 - Akka for client—cluster communication
 - Undisciplined about Job clusters



FLINK AS A LIBRARY

- Make deployments easier
- Application-oriented instead of cluster-oriented
- Just a start one or more Job jars, and automatically get distributed execution
- Remove distinction between different roles in the cluster (JM vs. TM)
- Participants of the cluster should determine roles on their own



RESOURCE MANAGEMENT & AUTOSCALING



- Short-lived & high latency
 - Resources only exist for the duration of the job
 - Different stages in a job may require different resources
 - Flink should be in control of allocating/releasing resources
- Long-lived & low latency
 - Dependent on external systems & exposed to varying load
 - External system should drive resource allocation to meet SLO/SLAs
 - Flink should react to available resources



CONTAINERIZED DEPLOYMENT

Why Flink on containers?

- Self-contained runtime environment that includes application code, libraries, dependencies, and configuration files
- Ease of operations by clean separation of concerns
- Enables dynamic resource allocation
- Fits the "*Flink as a Library*" model



CONTAINERIZED DEPLOYMENT (CONT.)

Brief introduction to Kubernetes



- De-facto industry standard for container orchestration
- Resource-oriented with declarative configuration
 - You tell K8s the desired state, and a background process asynchronously makes it happen
 - *"3 replicas of this container should be kept running"*
 - *"a load balancer should exist, listening on port 443, backed by containers with this label"*
- Core resource types:
 - Pod: a group of one or more containers running on a node
 - Deployment: keep n pods running



REWORKING FLINK'S RUNTIME

FLIP-6 – Flink Deployment and Process Model (Flink 1.5 and up)

- Revamped distributed architecture
 - Resource elasticity
 - Support for different environments
 - RESTful API for client-cluster communication
- Building blocks:

ResourceManager

- Cluster manager specific
- Acquires/releases resources

Dispatcher

- Touch point for job submissions
- Spawns JobManagers

JobManager

- Single job only
- Deploys/monitors job execution

TaskManager

- Offers resources to ResourceManager
- Gets tasks from JobManagers




UPCOMING CHANGES

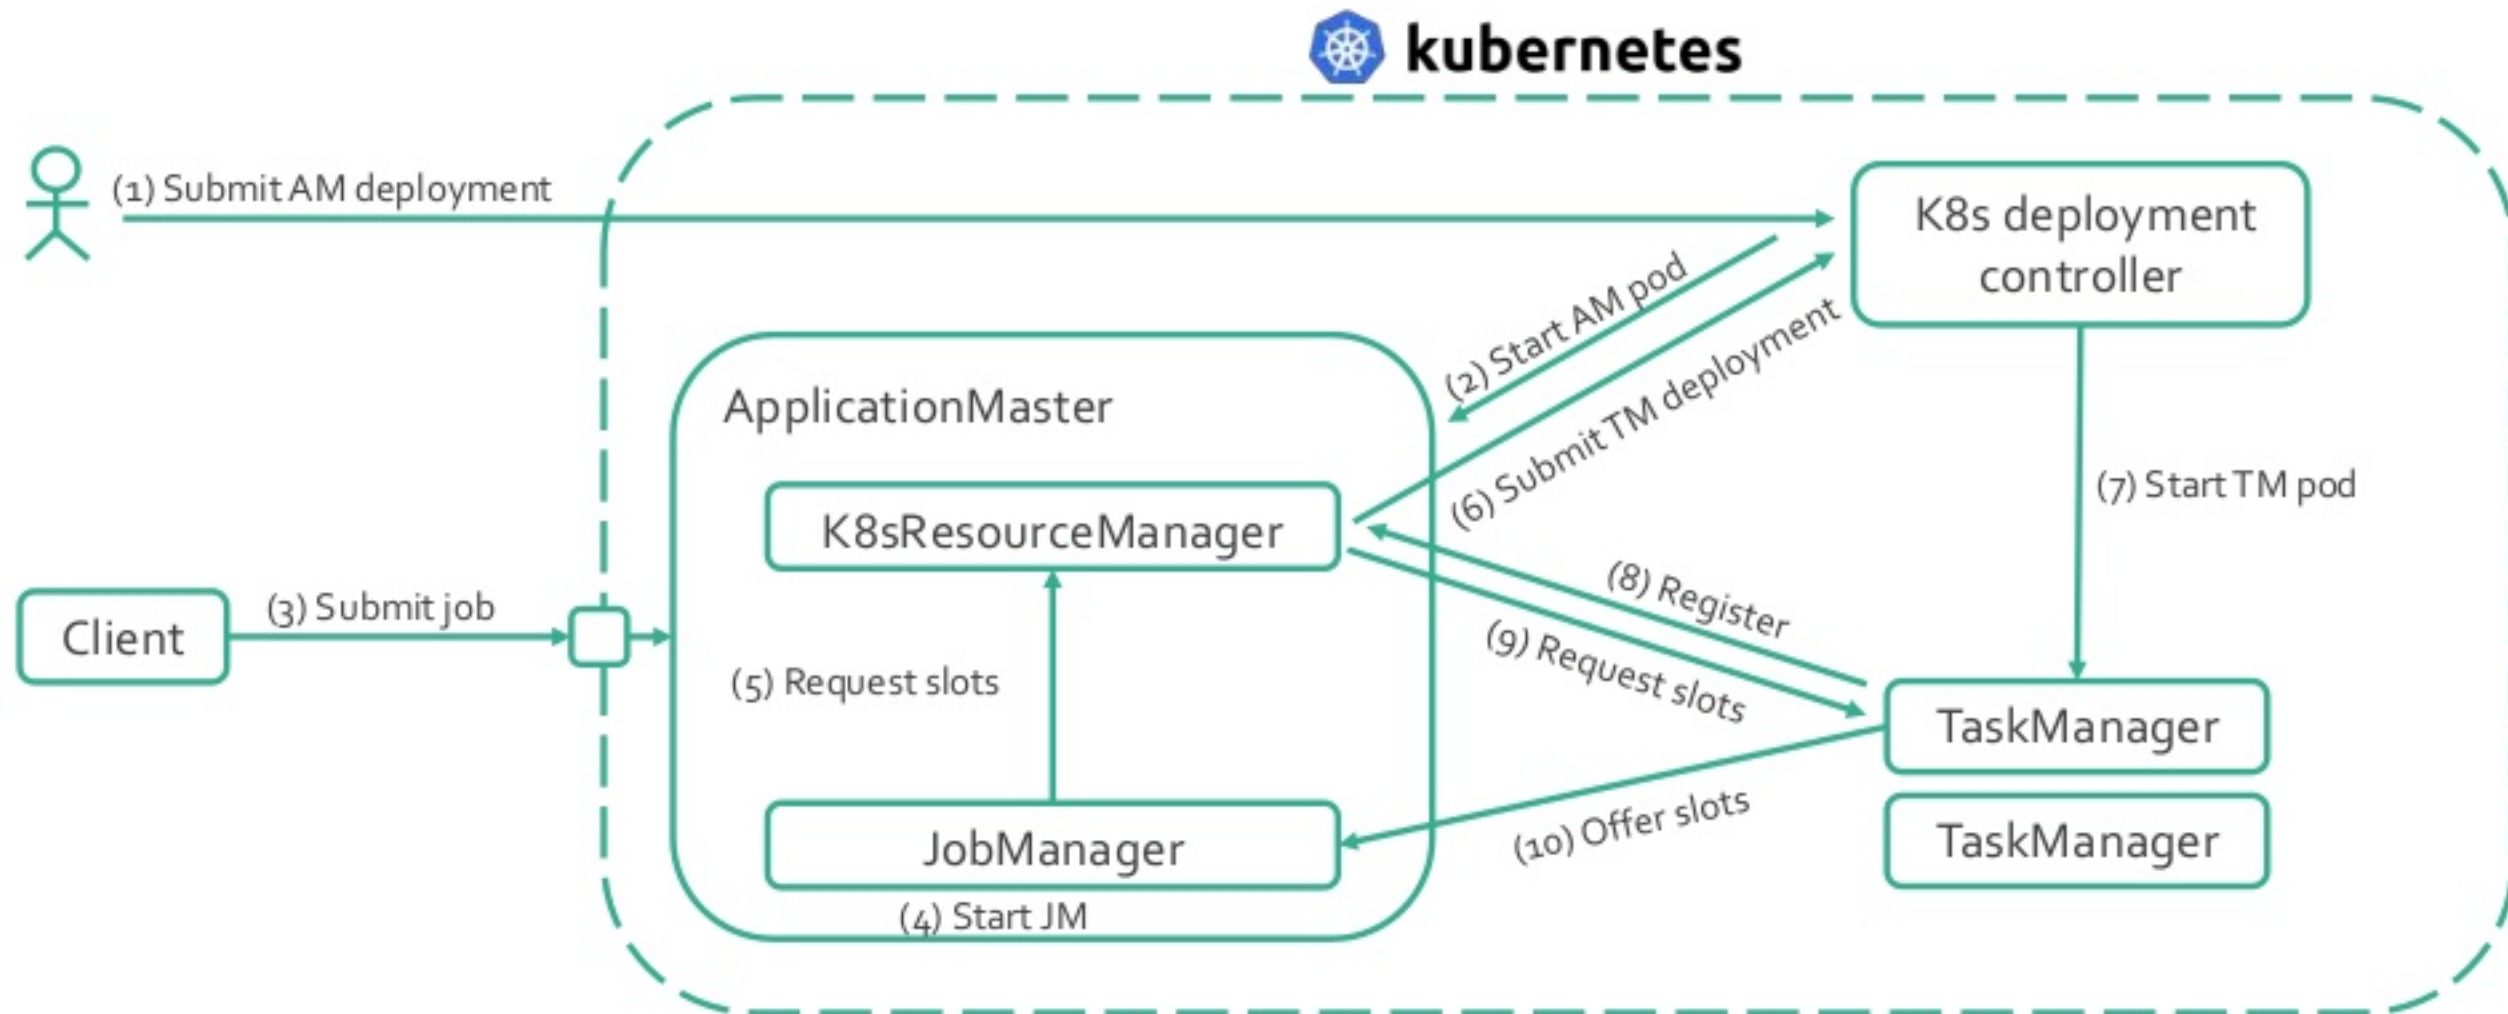


CONTAINER MODES

- Different levels of interaction with cluster manager
- Active Container Mode
 - Flink is aware about the cluster manager that it is running on, and interacts with it
 - Already exists, e.g., FLIP-6 YARN
- Reactive Container Mode
 - Flink is oblivious to its environment
 - Similar to standalone deployment
 - Flink may react on new resources by scaling up job

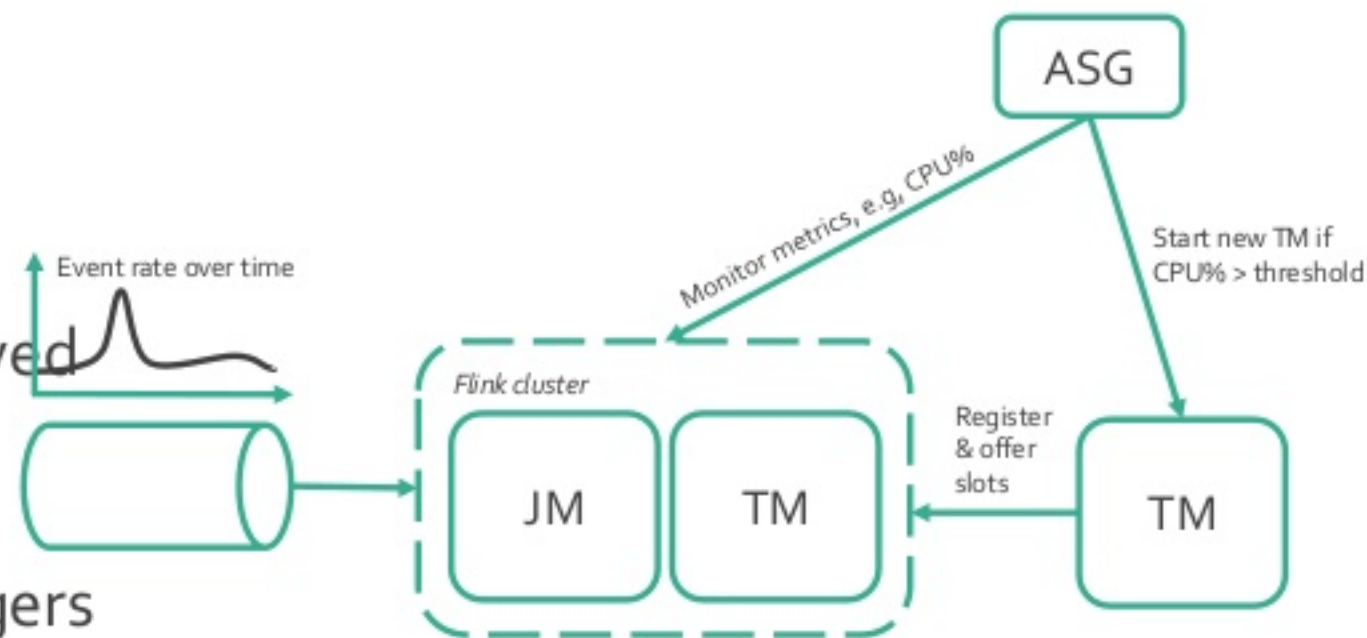


ACTIVE K8S INTEGRATION



REACTIVE CONTAINER MODE

- Relies on external system to start/release TaskManagers, e.g.,
 - Kubernetes Horizontal Pod Autoscaler
 - GCP Autoscaling
 - AWS Auto Scaling Group
- Re-scale job if resources are added/removed (take savepoint and resume job with new parallelism automatically)
- By definition works with all cluster managers



DOES ANY OF THIS WORK?

Demo: Autoscaling in Active & Reactive Container Mode



CURRENT STATE

- Recommendations on how to deploy job and session cluster on K8s
 - Docker image build script
 - K8s resource configs
- [StandaloneJobClusterEntryPoint](#) to start a (containerized) job cluster
- Work in progress/planned:
 - Reactive Container Mode with automatically rescaling
 - Design documents will be published soon
 - Active K8s integration: prototype is available
 - `git clone -b nativeKubernetes git@github.com:tillrohrmann/flink.git`



SUMMARY

- Thinking in terms of jobs & application code as an alternative to deploying Flink clusters
- Active K8s integration & Reactive Container Mode
- Existing code already enables job and session clusters on K8s
- Call to action:
 - Umbrella tickets: FLINK-9953, FLINK-7087
 - Follow dev@flink.apache.org to stay informed about latest state of development



THANK YOU!

@gjyao

@dataArtisans

@ApacheFlink

WE ARE HIRING

data-artisans.com/careers

dataArtisans