




# Assisting Millions of User in Real-Time

Flink Forward Berlin 2018



# The Speakers

Who are these guys?



Alexey Brodovshuk

@alexeybrod



Krzysztof Zarzycki

@k\_zarzyk



# About Kcell

Kcell JSC is a part of the largest Scandinavian telecommunications holding – TeliaCompany

## Largest GSM operator in Kazakhstan

> 10 000 000 subscribers



## Great network coverage

4G (40%), 3G (73%), 2G (96%) population

## We like innovations

Kcell has a strong software development team and lots of experience in building services and products

## Not only telco

There is the ongoing process of company digital transformation

# Business needs

Assisting Millions of Users in Real-Time


## Input



## Process



## Actions






# Use Cases

Use case scenarios. Just few of many.



## Balance Top Up Case

If subscriber top-ups her balance too often in short period of time. We can offer her a less expensive tariff or auto-payment services.



Create new trigger

Inform Marketing Platform whenever balance is topped up during some period

Aggregation period

7 days

Channels

☒ Payment GW  
☒ Payment GW (Terminal)  
☒ Balance Credit  
☒ Money Transfer  
☒ Voucher (via SMS)  
☒ Voucher (via USSD)

Topup condition

$\geq 1000$

Tracked parameter

☐ Number of topups  
☒ Sum of topup amounts

Tracked parameters value

30000

Save

Close



### Roaming case

Trigger to Marketing Platform if subscriber visited X country OR/AND registered in Y visited mobile network and his device's type is Z

Roaming

Fraud



### Fraud case in roaming

Send an email to the anti-fraud unit if subscriber registered in roaming but his balance at the moment is equal to 0.  
This situation is impossible in standard case.

# Old System

Why did we start to look for the new solution?



External Vendor  
Solution

1

## Blackbox Solution

Kcell Developers can't fix, tweak or optimize it

2

## Scalability issues

Limited to ~2000 events / sec

Can't support all needed data sources

3

## Not reliable

Multiple accidents which took too much time to resolve

# Scale

Required system throughput

---

**10M**  
Subscribers

**165K**  
Events / second

**22.5**  
TB / month



# About GetInData

Big Data. Passion. Experience.



Roots at  
Spotify



Focus on  
Big Data  
from Day 1



Production  
Experience

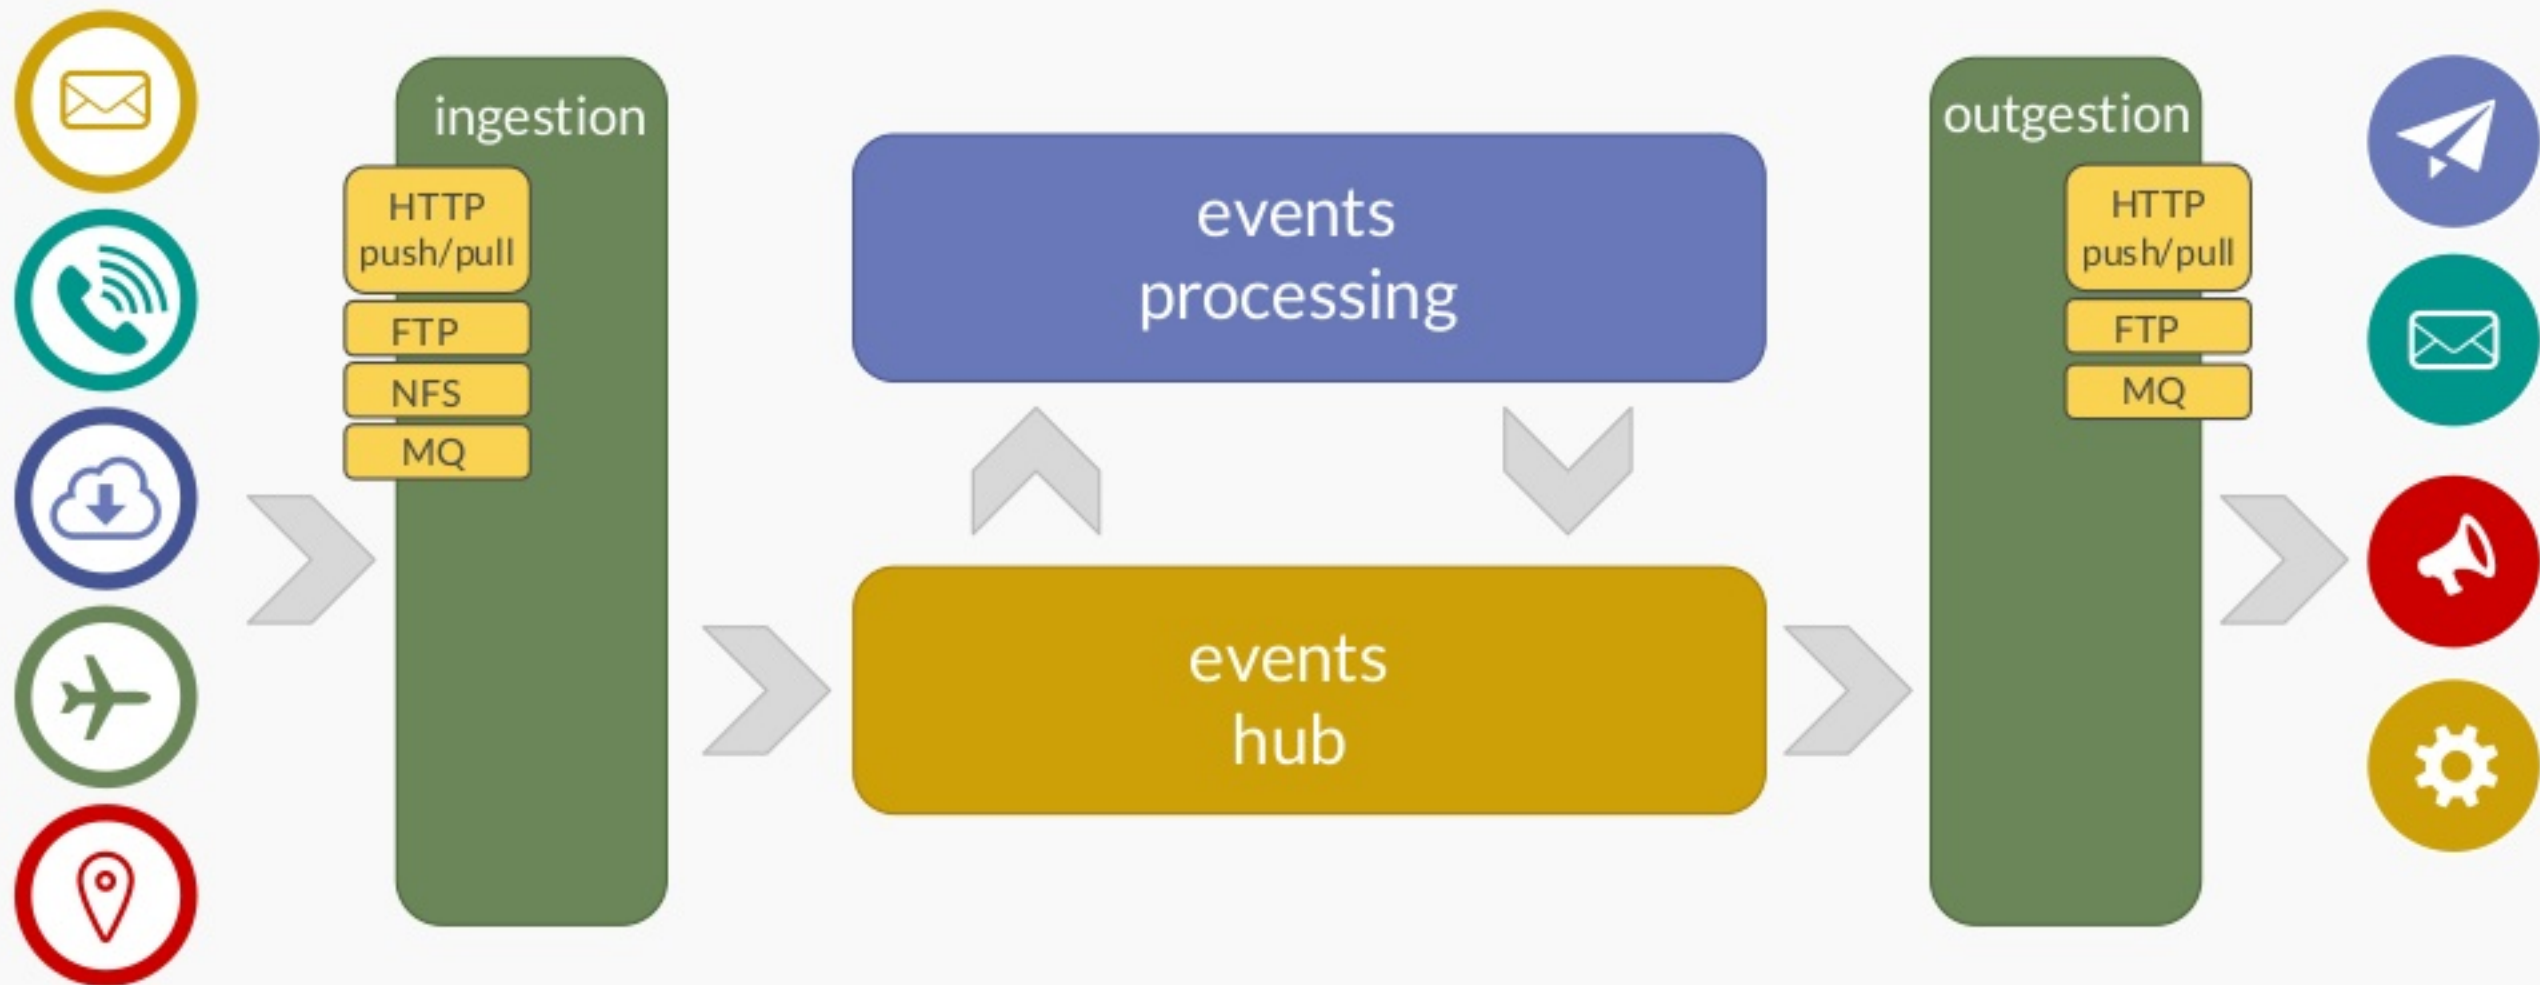


Contributions  
to  
Apache Flink



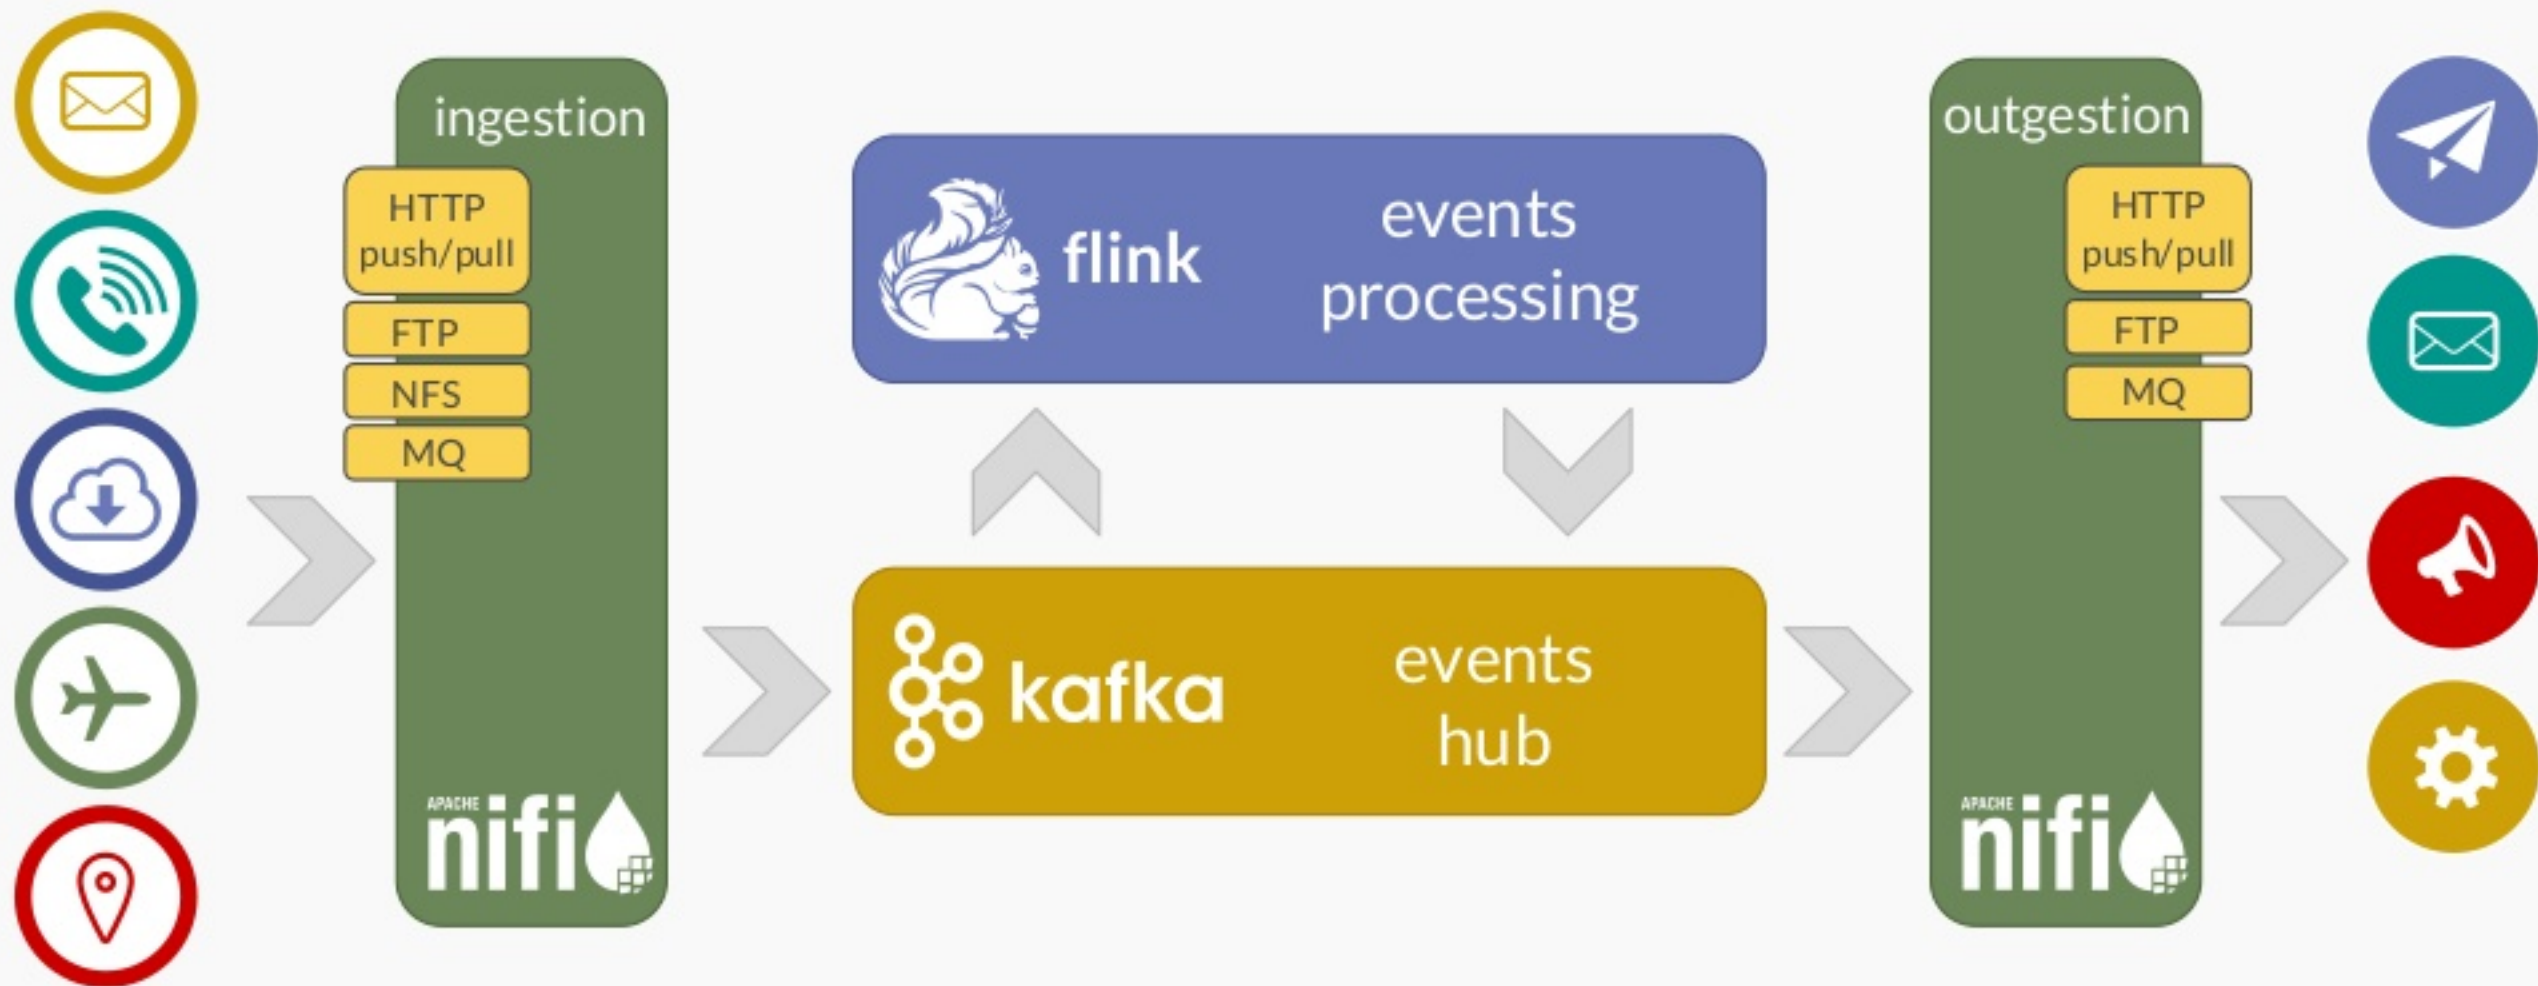
# New Solution

Real-time Stream Processing



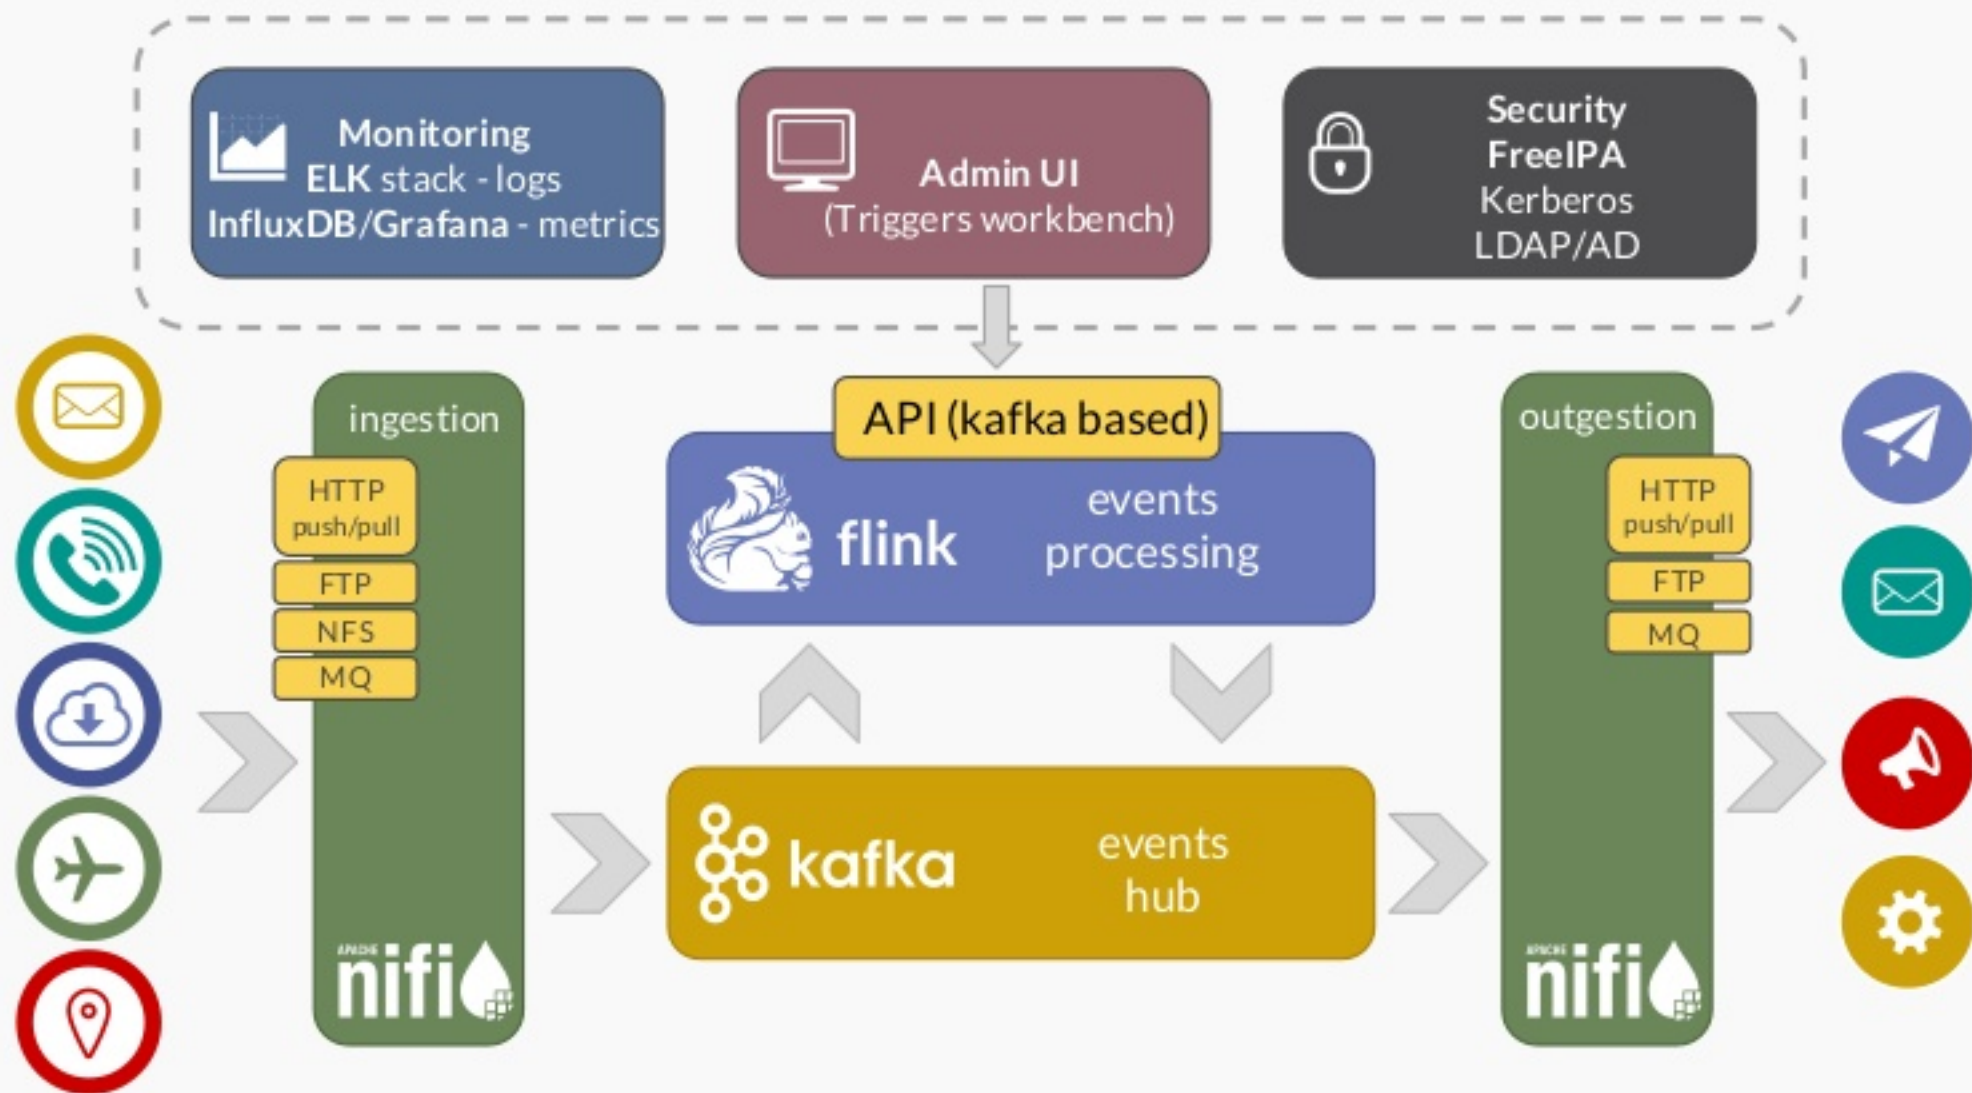
# New Solution

Real-time Stream Processing



# New Solution (Operations)

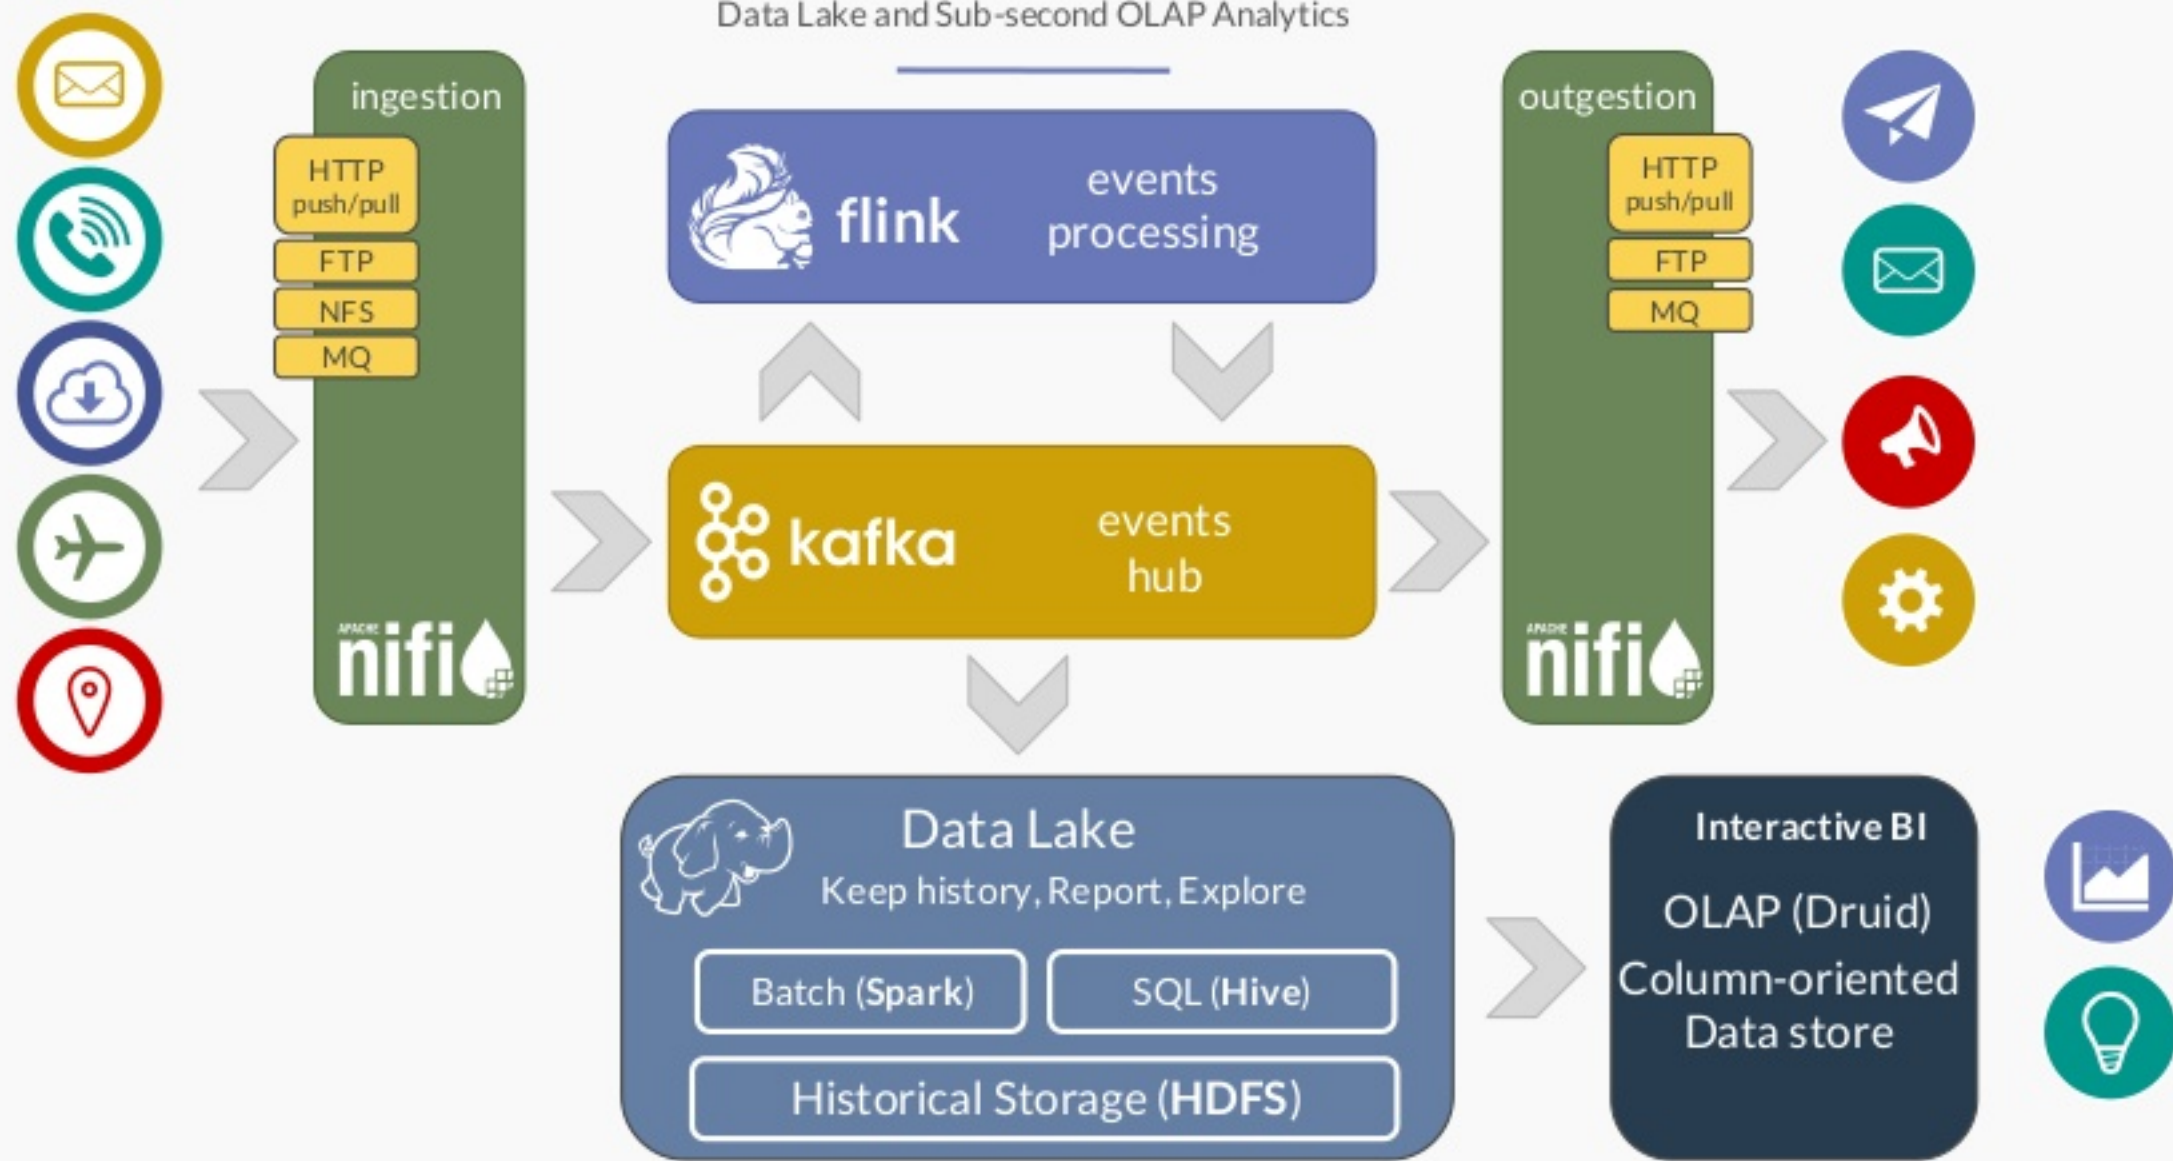
Web UI, Monitoring, Security





# New Solution (Data Lake)

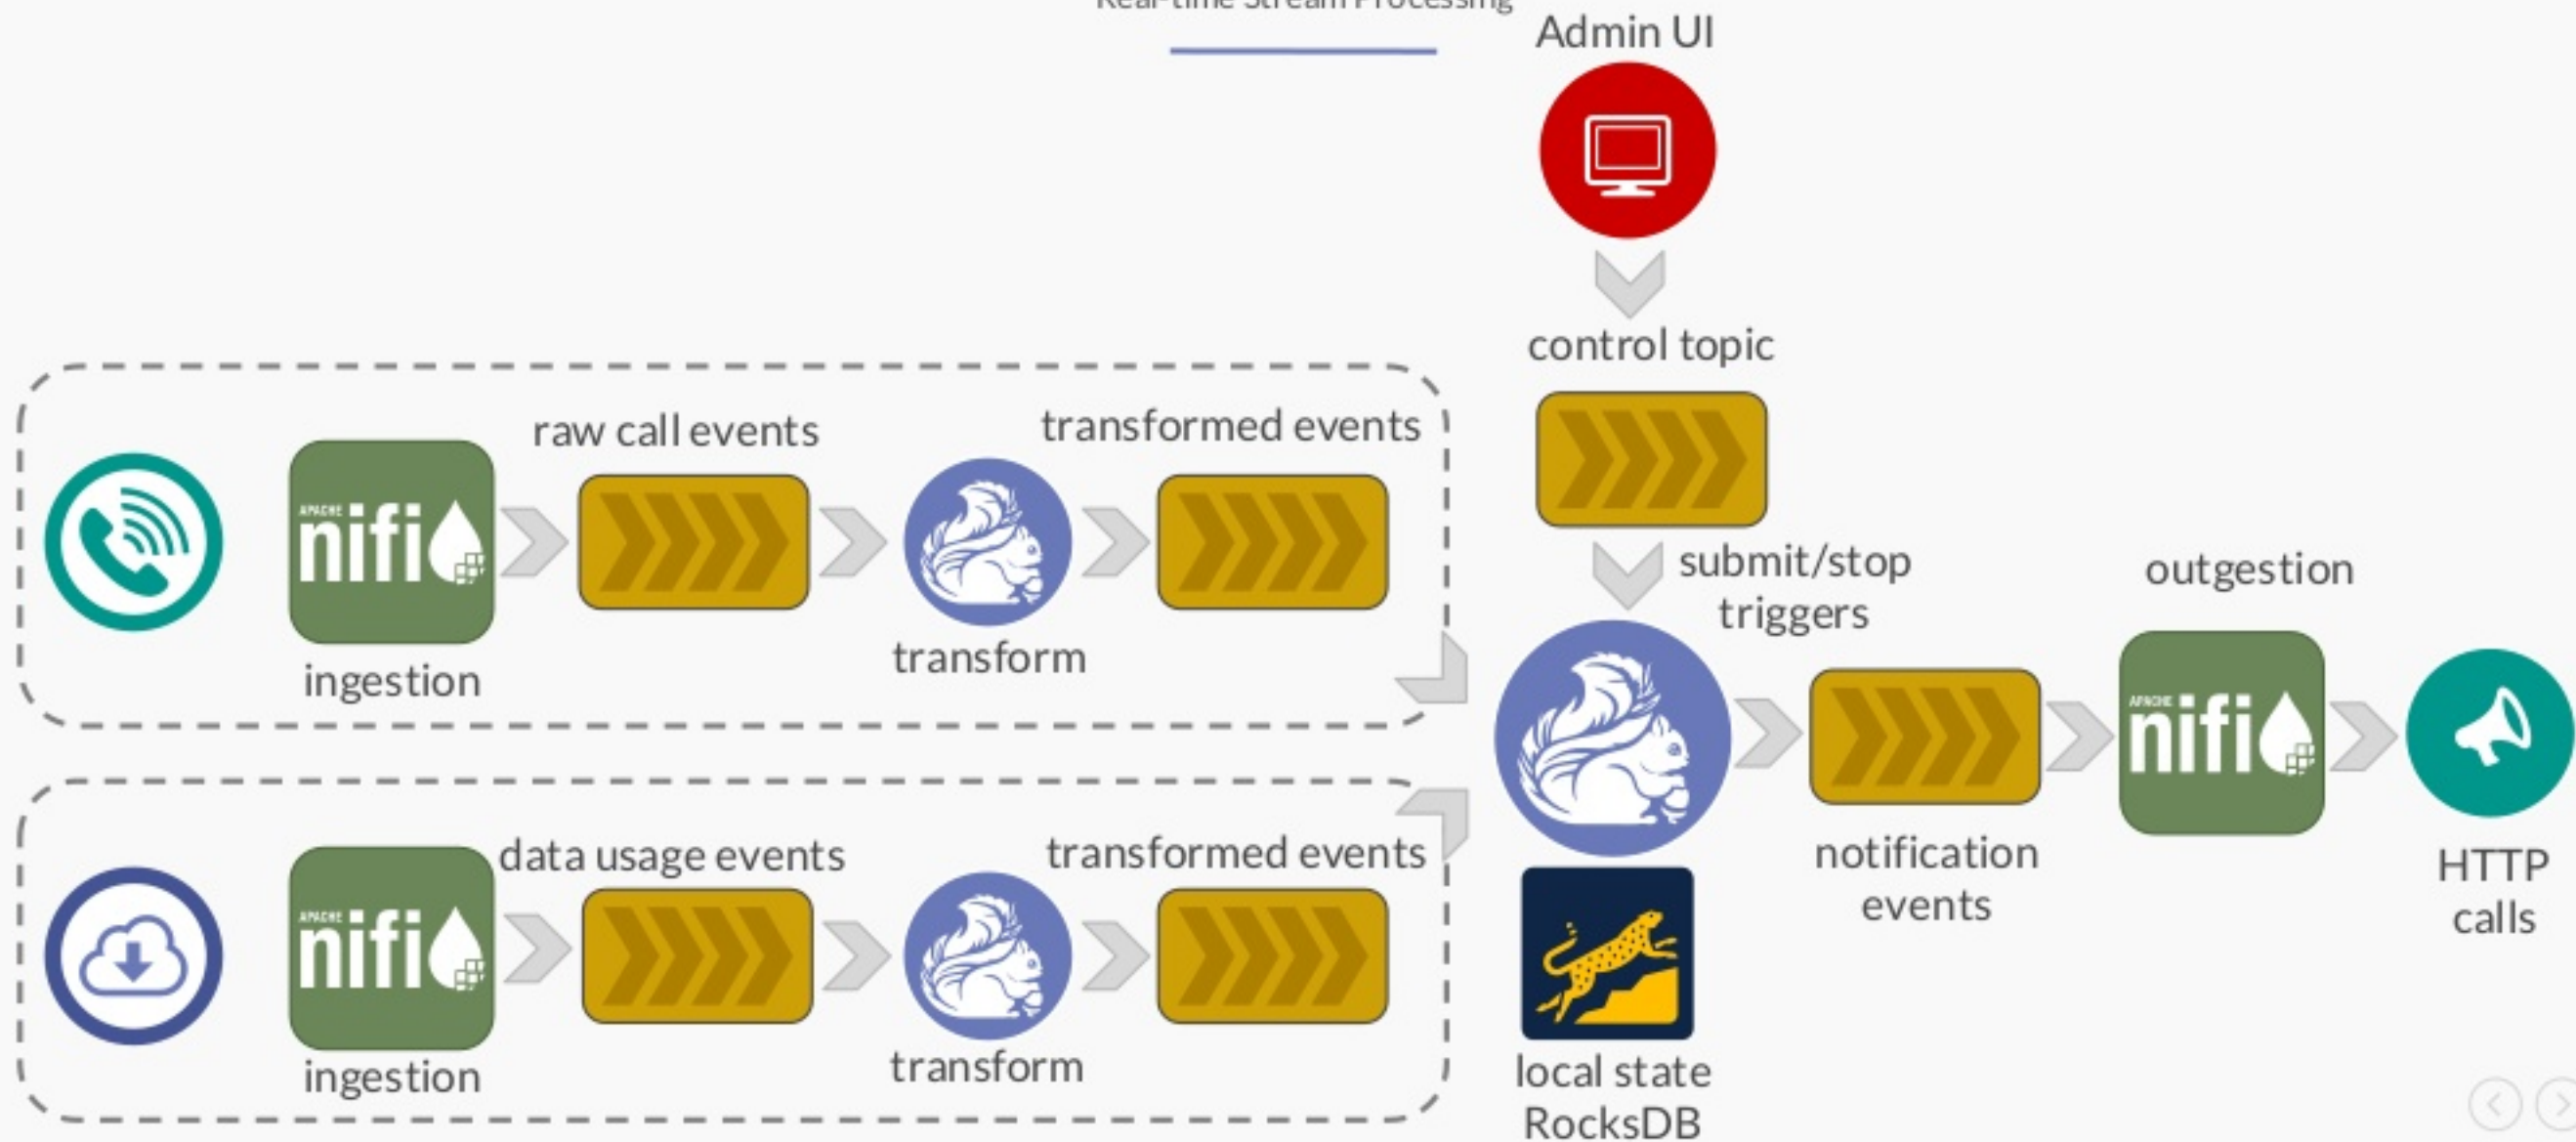
Data Lake and Sub-second OLAP Analytics





# Processing Flow

Real-time Stream Processing





# Dynamic Rules Design

---

Some treats for Squirrels

# Dynamic Rules Design

## Key Points

---

- We want to run 100s of triggers/business rules
- A typical approach: job per rule
- Won't work in our case:
  - Run 100s of topologies/jobs = multiplied resources cost
  - Pull data from Kafka 100s of times
  - State (user features) replicated 100s times
  - Starting rule requires deployment of the job

# Dynamic Rules Design

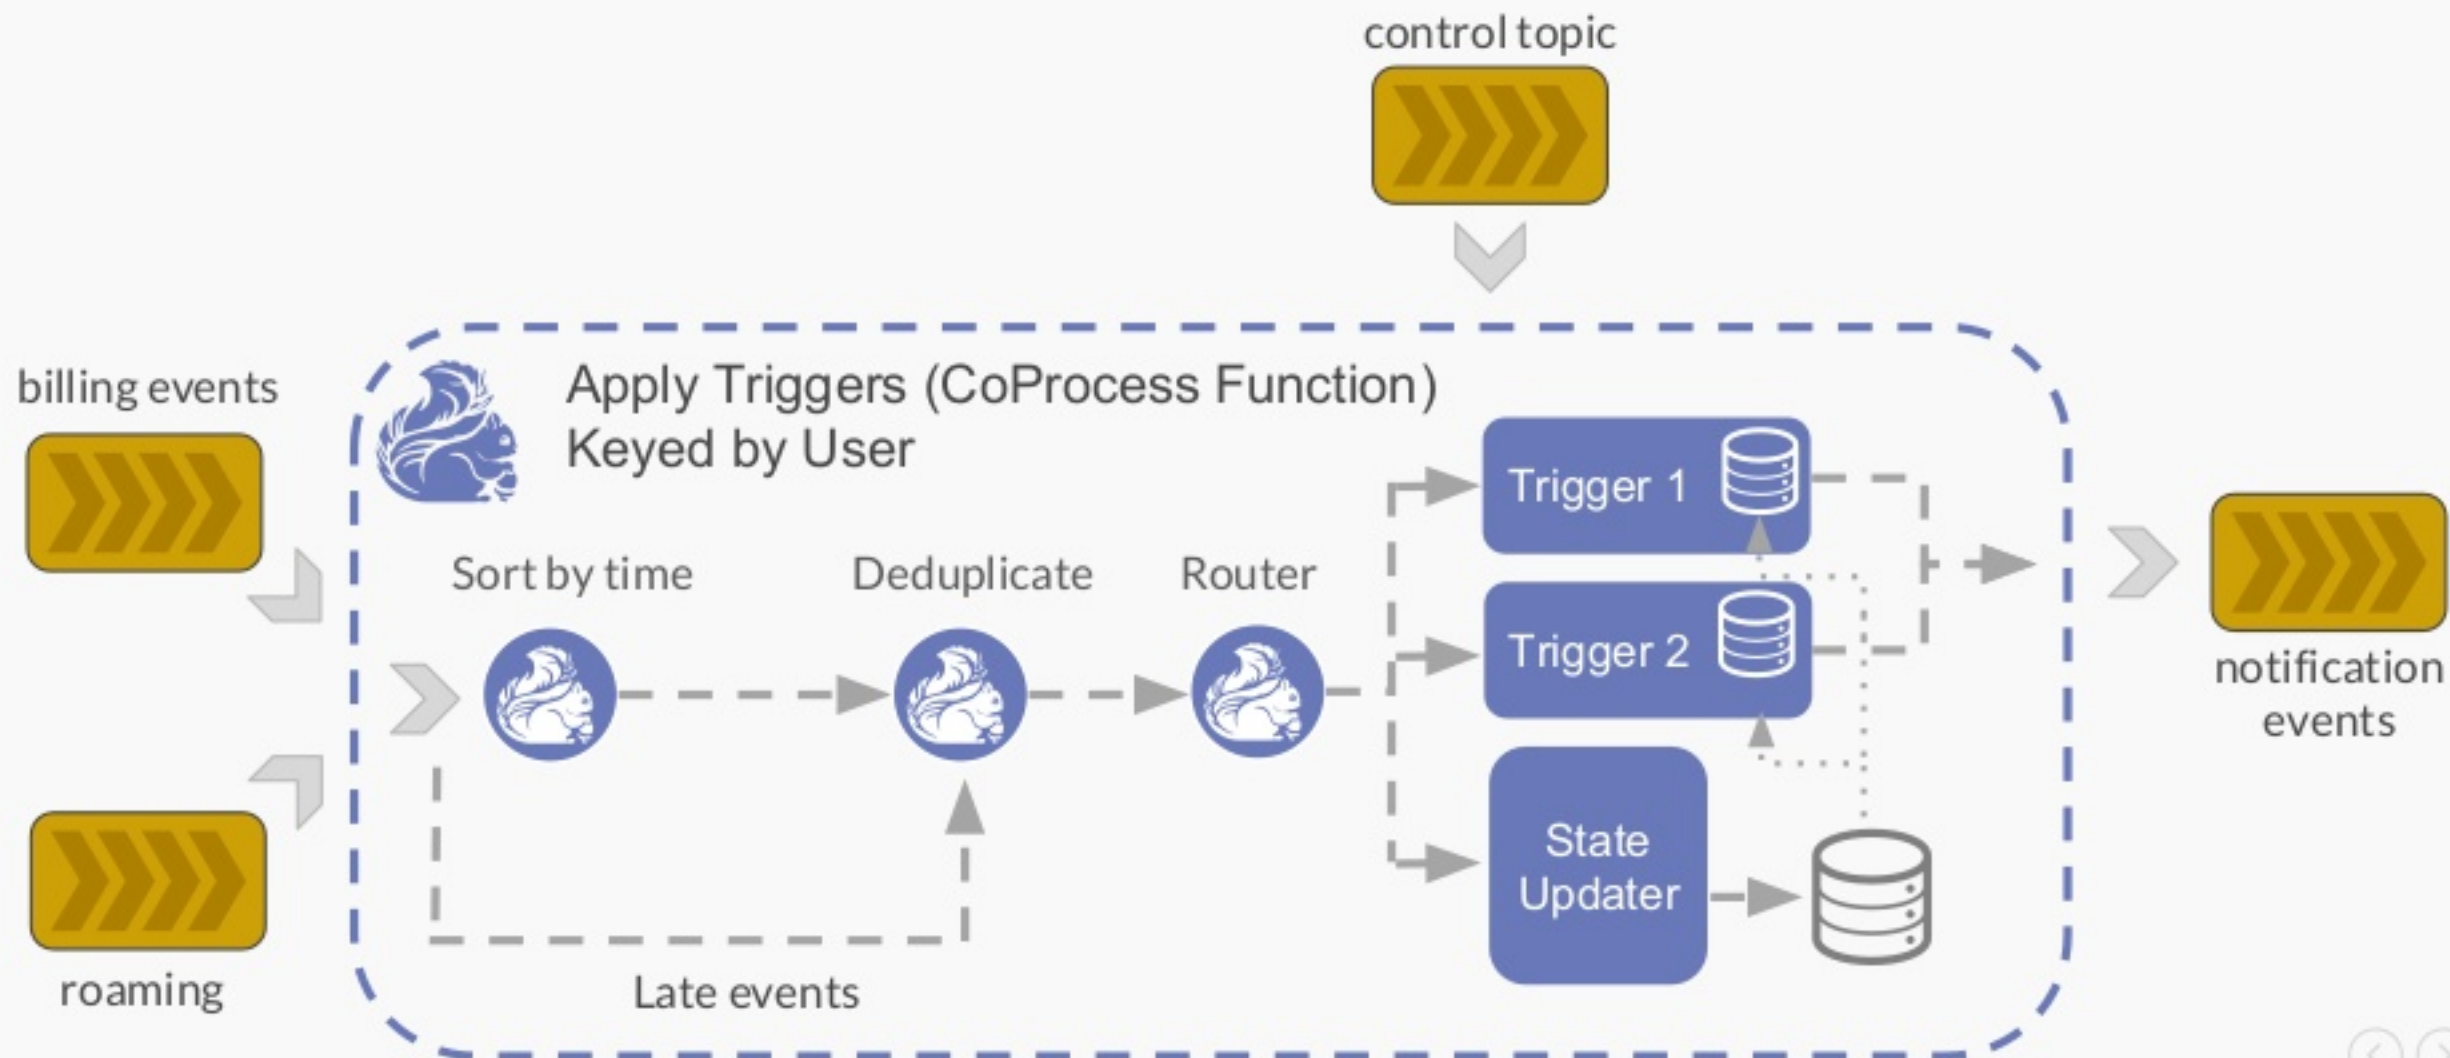
## Key Points

---

- Our approach: One job to run all triggers/rules
  - And to consume all the sources
- Trigger “templates” still coded with java
- adding/removing rules without restarting application
- 100s of rules running efficiently

# Dynamic Rules Design

The Overview





# Dynamic Rules Design

## Pros and Cons



### Shared resources and costs

- CPU, RAM, state, shuffle
- Pulling data from Kafka



### Sharing of state

- Build customer features, that can be seen by all rules



### No job restart on start of new rules

- Rules started by business, no IT involved



### One bad rule affects whole system

- Watermarks are shared
- Failures are shared



### Still need to code rule template in the job

- No way to use SQL, Table API, CEP



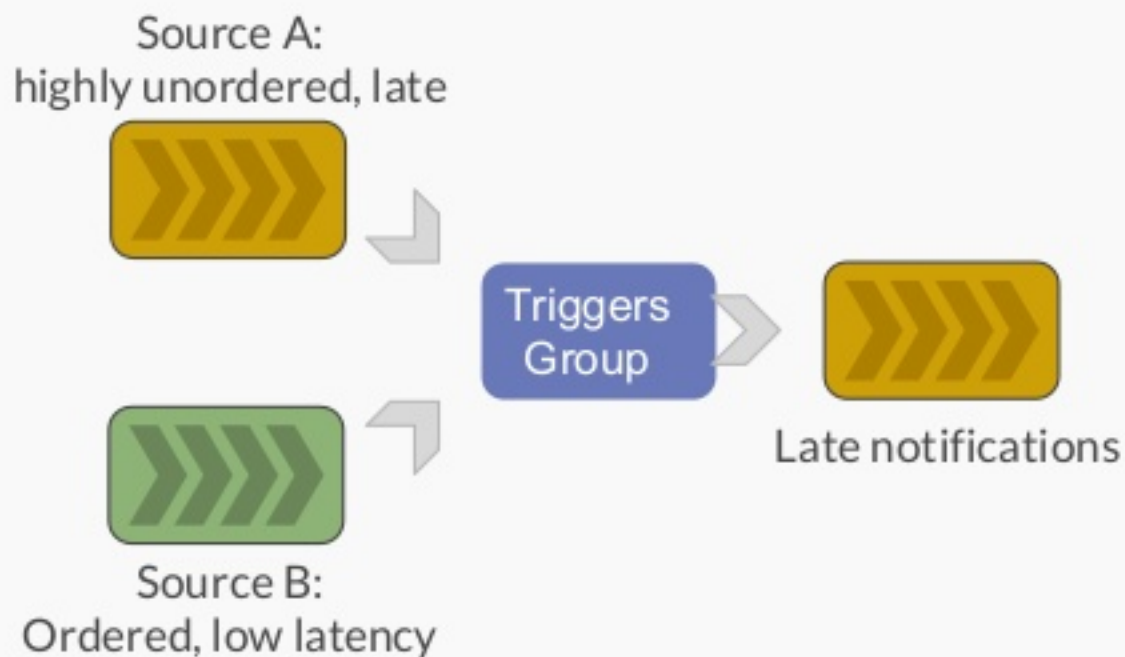
### Can be tricky to debug

- Code is shared
- Code paths enabled externally

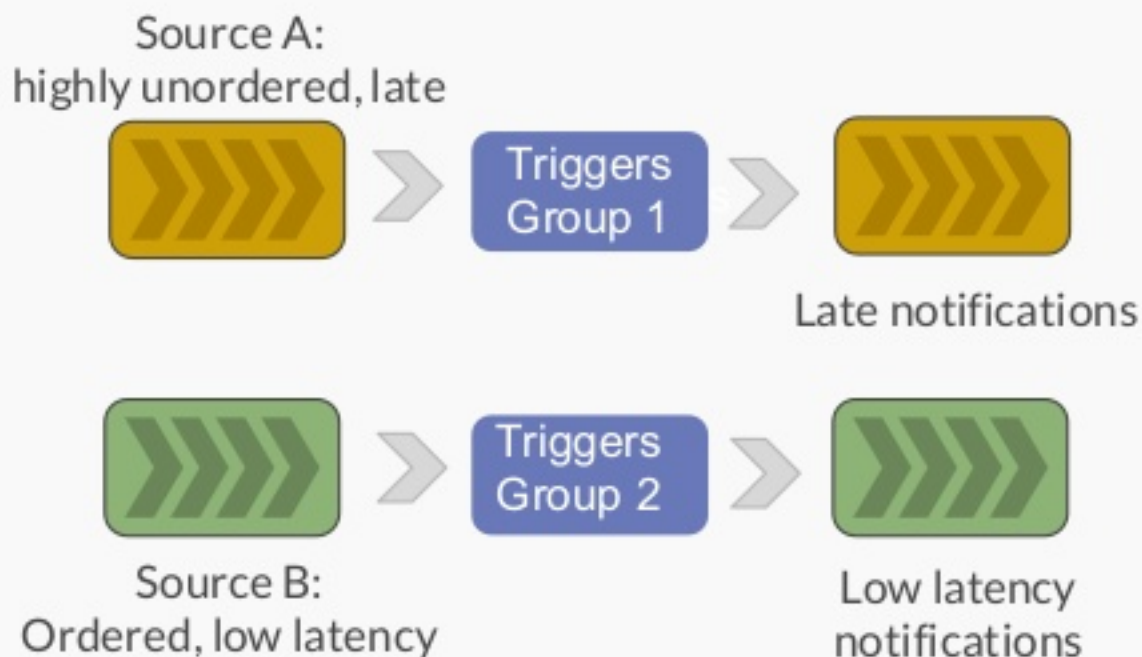
# Dynamic Rules Design

Issue: lagging sources slow down all rules

## Problem



## Solution



## Join Streams

Join slowly changing stream with event stream



Issue: No Data = No Watermark

Stream X



Stream Y



## Solution

Ignore Watermark for slow stream  
(Watermark = Long.MAX\_VALUE)

# Flink Changes Wishlist

What could be even better?



# Decisions made

Some decisions our team made before or during project implementation

---



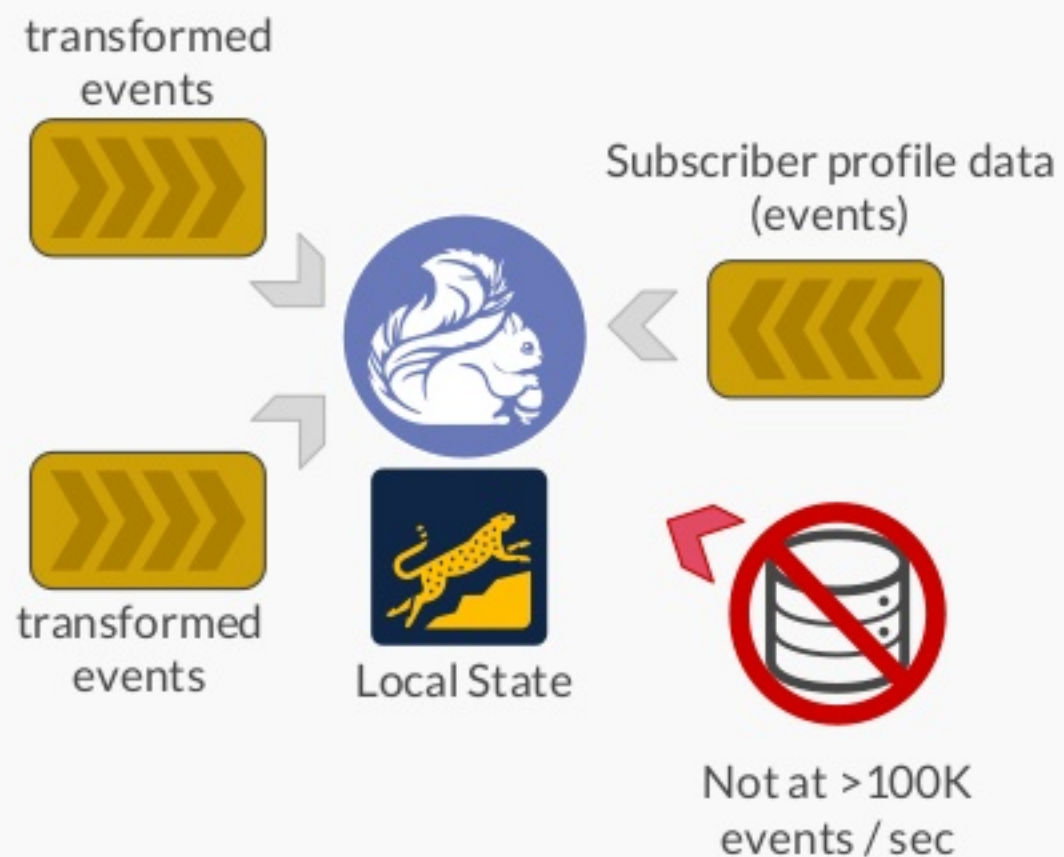
## Powerful Real-Time Analytics

Streaming-first approach

Apache Kafka for event hub

Apache Flink





## Performance

Apache Avro

Keep state local to the process

Ingest reference data for local joins and enrichment

- No need to query external systems while processing
- Data time correlation correctness





## Reliability and battle-tested techniques

Flink on YARN, with HDFS

HA for redundancy and running ~24/7

InfluxDb & Grafana for monitoring & alerting

ELK for logs collection and aggregation



## Security

Kerberos and AD thanks to FreeIPA

Apache Ranger for authorization



**Cost-Efficiency**  
Open-source technologies  
HDP as a licence-free distribution  
Just start with a bunch of servers



## Extensiveness

One platform for the whole Enterprise

Batch (adhoc) queries too

- Spark, Hive/Presto

Online analytics

- OLAP



# Our Collaboration

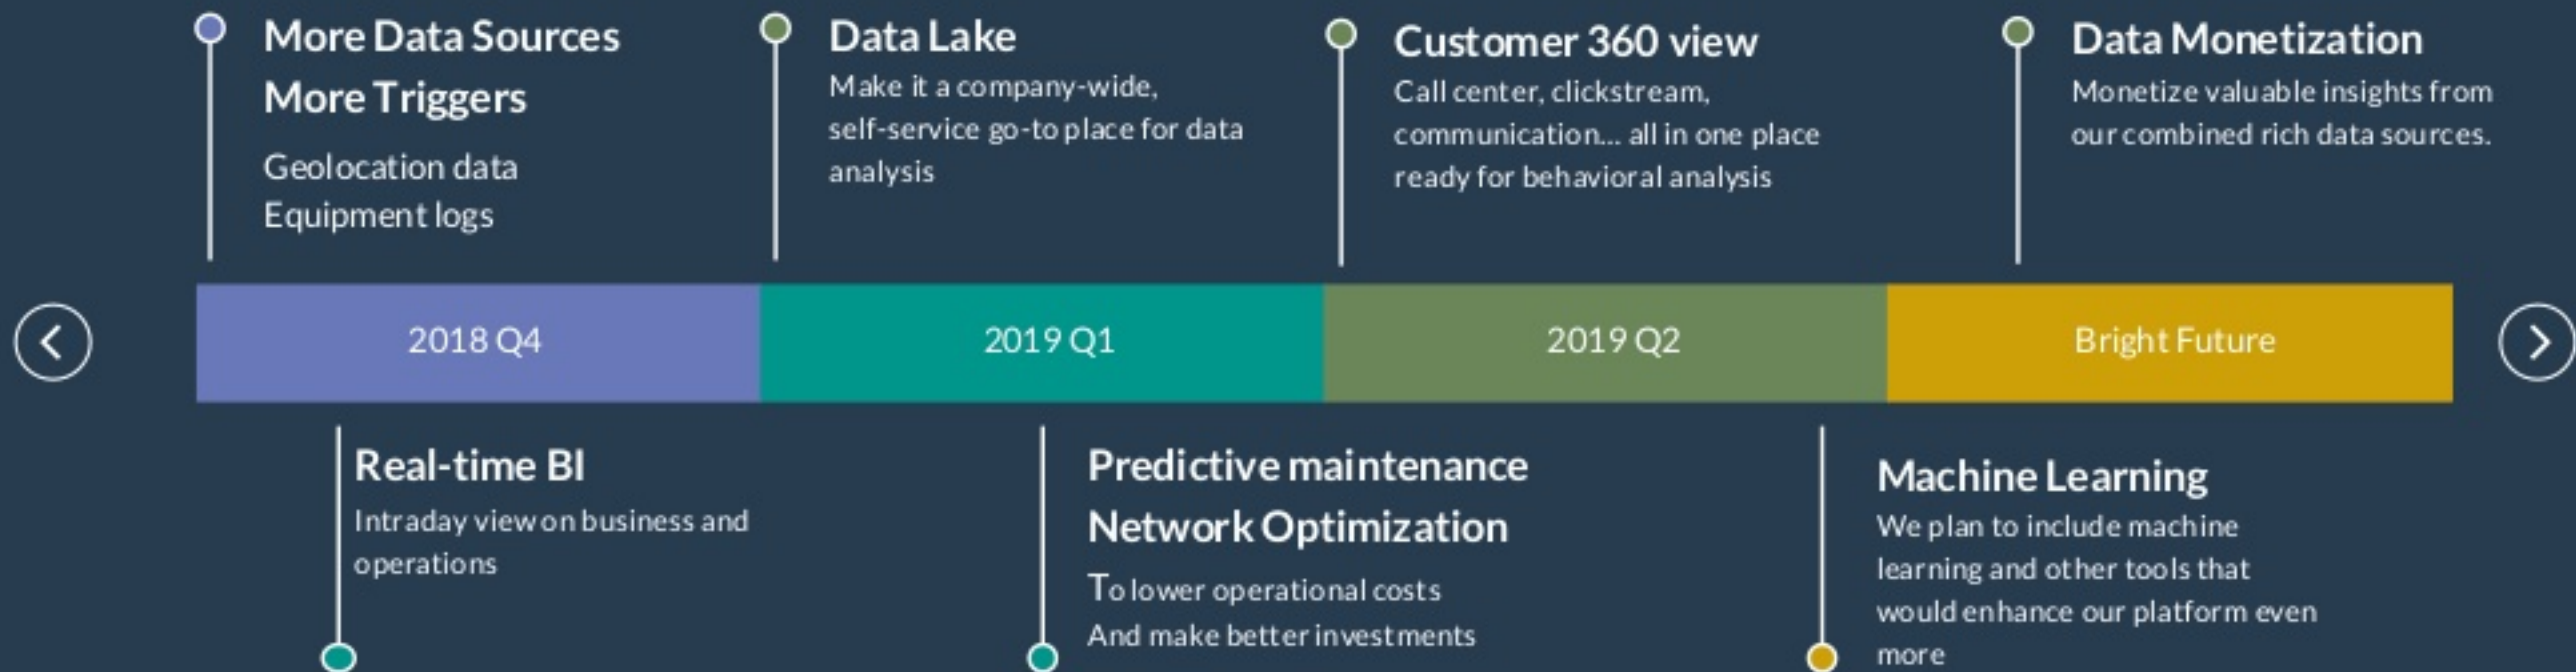
Two heads are better than one





# Future Work

We have already done a lot. But more great things are coming.



And many more...



# Questions?

---

Flink Forward Berlin 2018

Contact us:



[zarz@getindata.com](mailto:zarz@getindata.com)

[alexey.brodovshuk@gmail.com](mailto:alexey.brodovshuk@gmail.com)