

IMPROVING THROUGHPUT AND LATENCY WITH FLINK'S NETWORK STACK

NICO KRUBER

SOLUTION ARCHITECT / SOFTWARE ENGINEER @ DATA ARTISANS,
APACHE FLINK COMMITTEE

FLINK DATA TRANSPORT (LOGICAL)




Abstraction over:

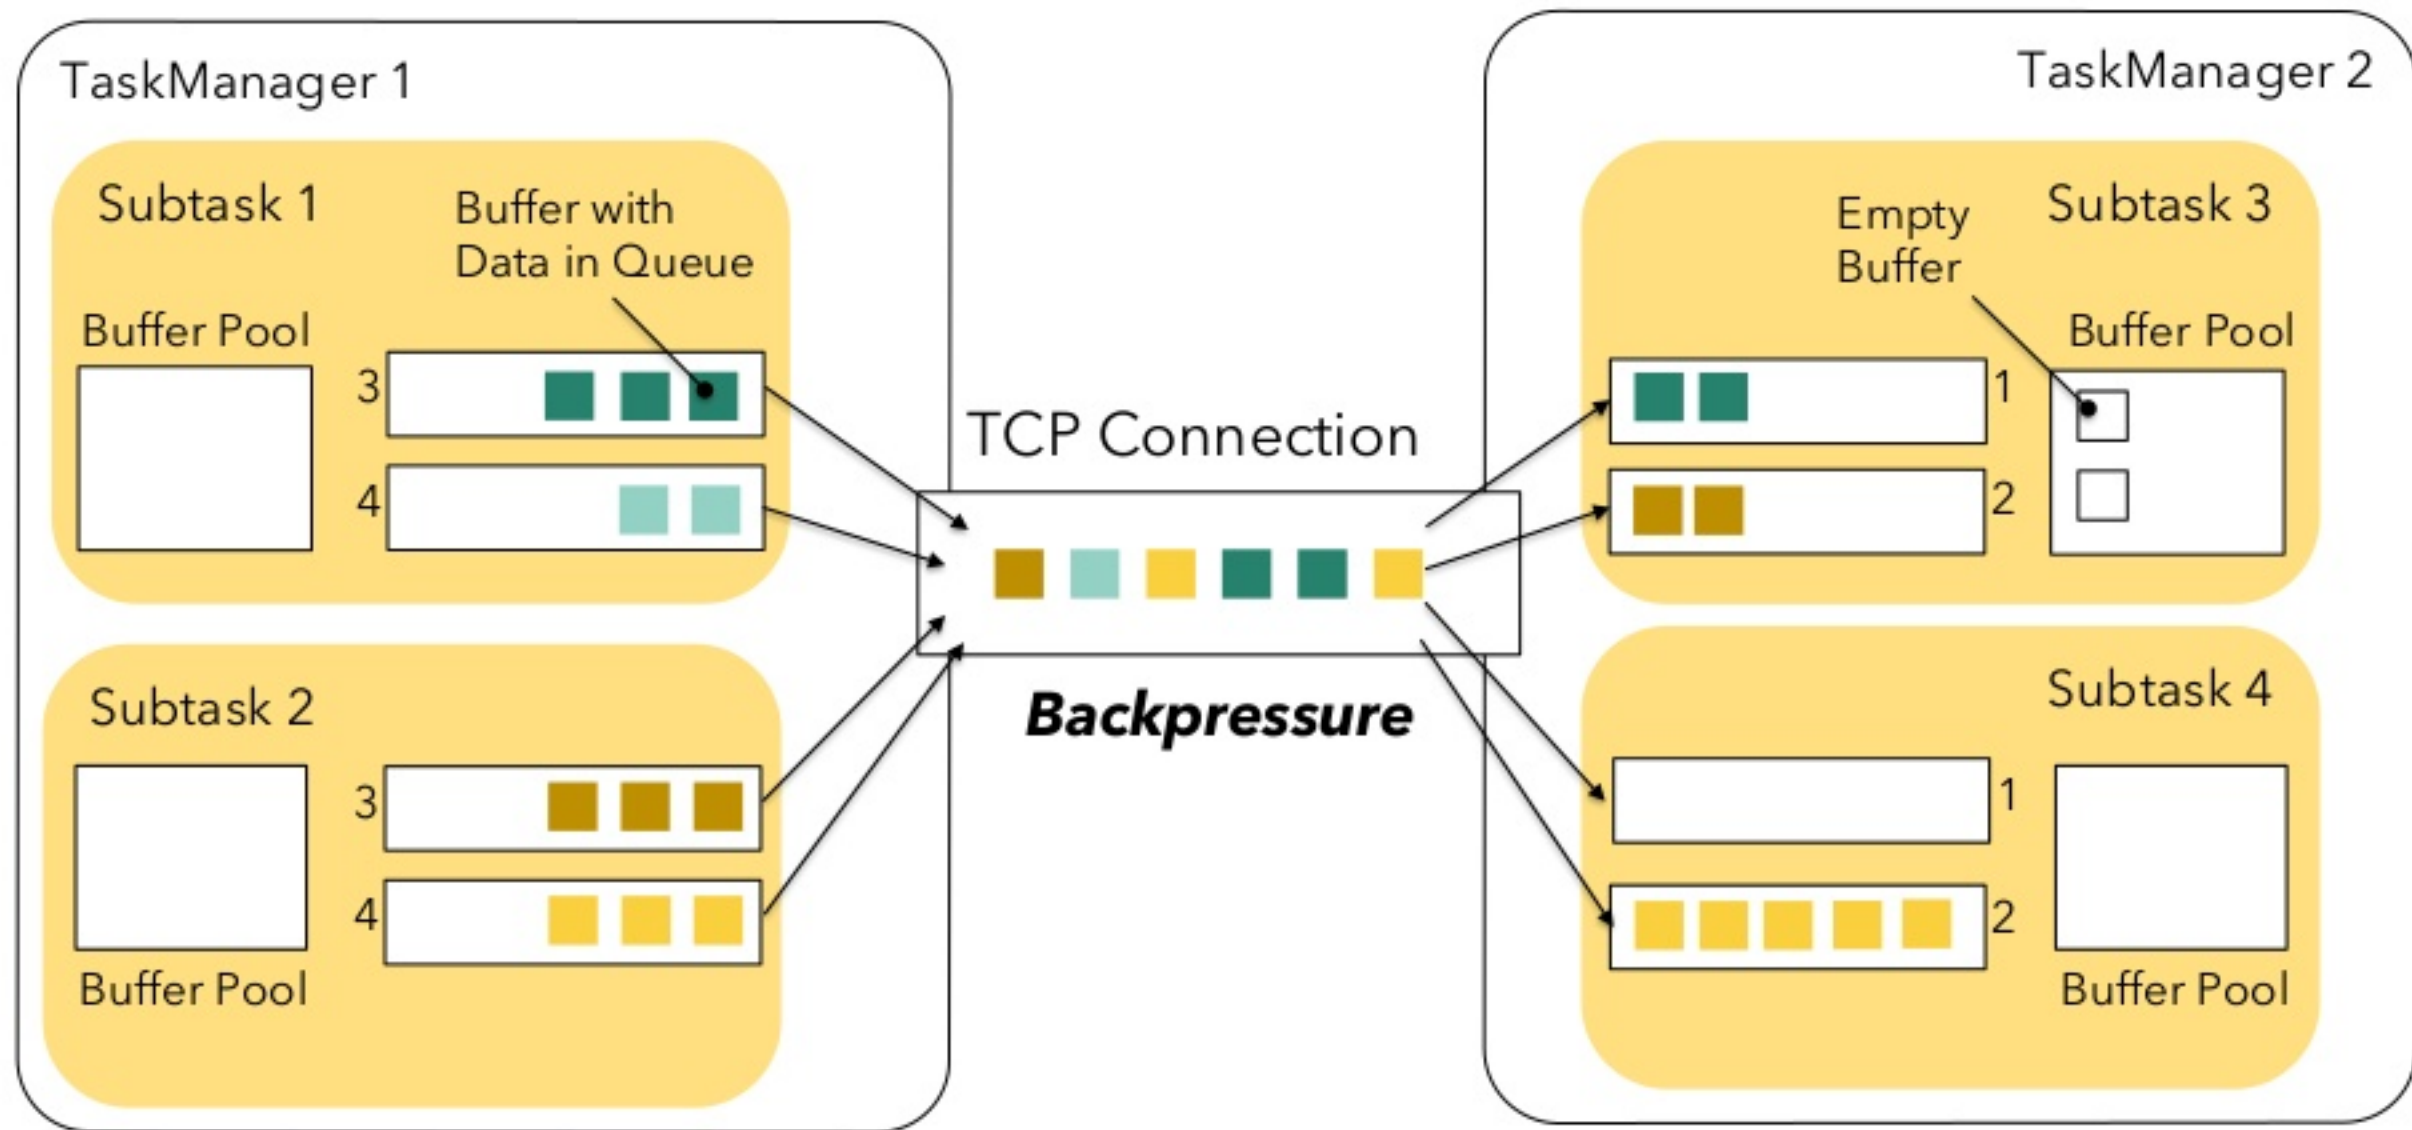
- Subtask output
 - pipelined-bounded
 - pipelined-unbounded
 - Blocking
- Scheduling type
 - all at once
 - next stage on complete output
 - next stage on first output
- Transport
 - high throughput via buffers
 - low latency via buffer timeout



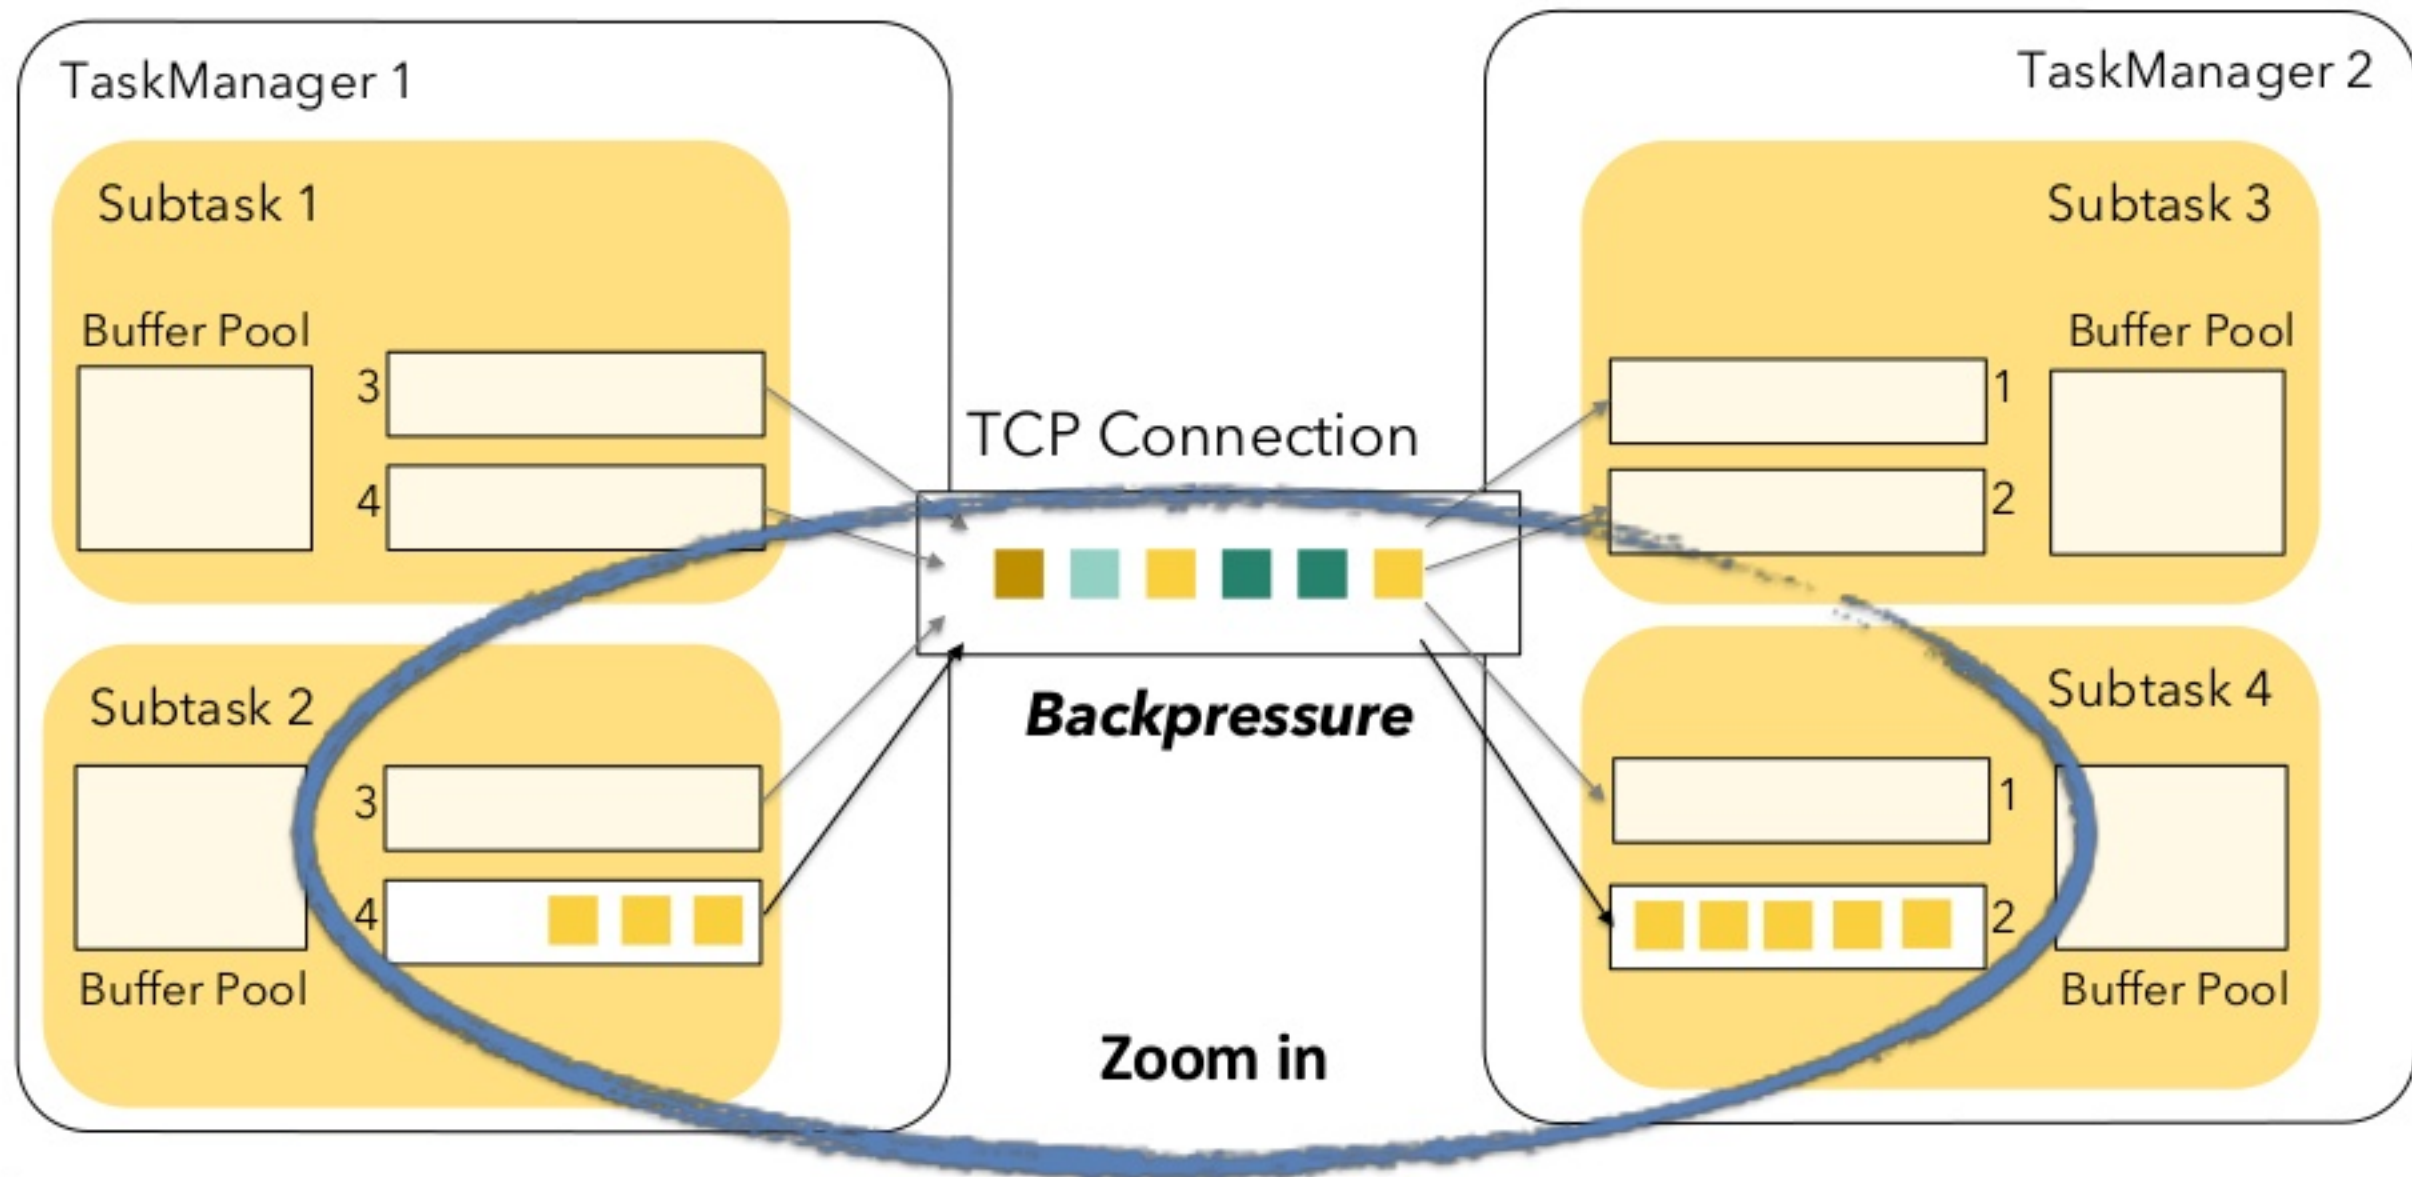
FLINK DATA TRANSPORT (PHYSICAL)



FLINK DATA TRANSPORT (PHYSICAL)




FLINK DATA TRANSPORT (PHYSICAL)




CREDIT-BASED FLOW CONTROL



CREDIT-BASED FLOW CONTROL (FLINK 1.5+)




CREDIT-BASED FLOW CONTROL (FLINK 1.5+)



CREDIT-BASED FLOW CONTROL (FLINK 1.5+)

- Never blocks the TCP connection
 - Better resource utilization with data skew in multiplexed connections
- Avoids overloading of slow receivers (direct control over amount of buffered data)
 - Improves checkpoint alignment
- cost: additional announce messages (piggy-bagged), potential round-trip latency

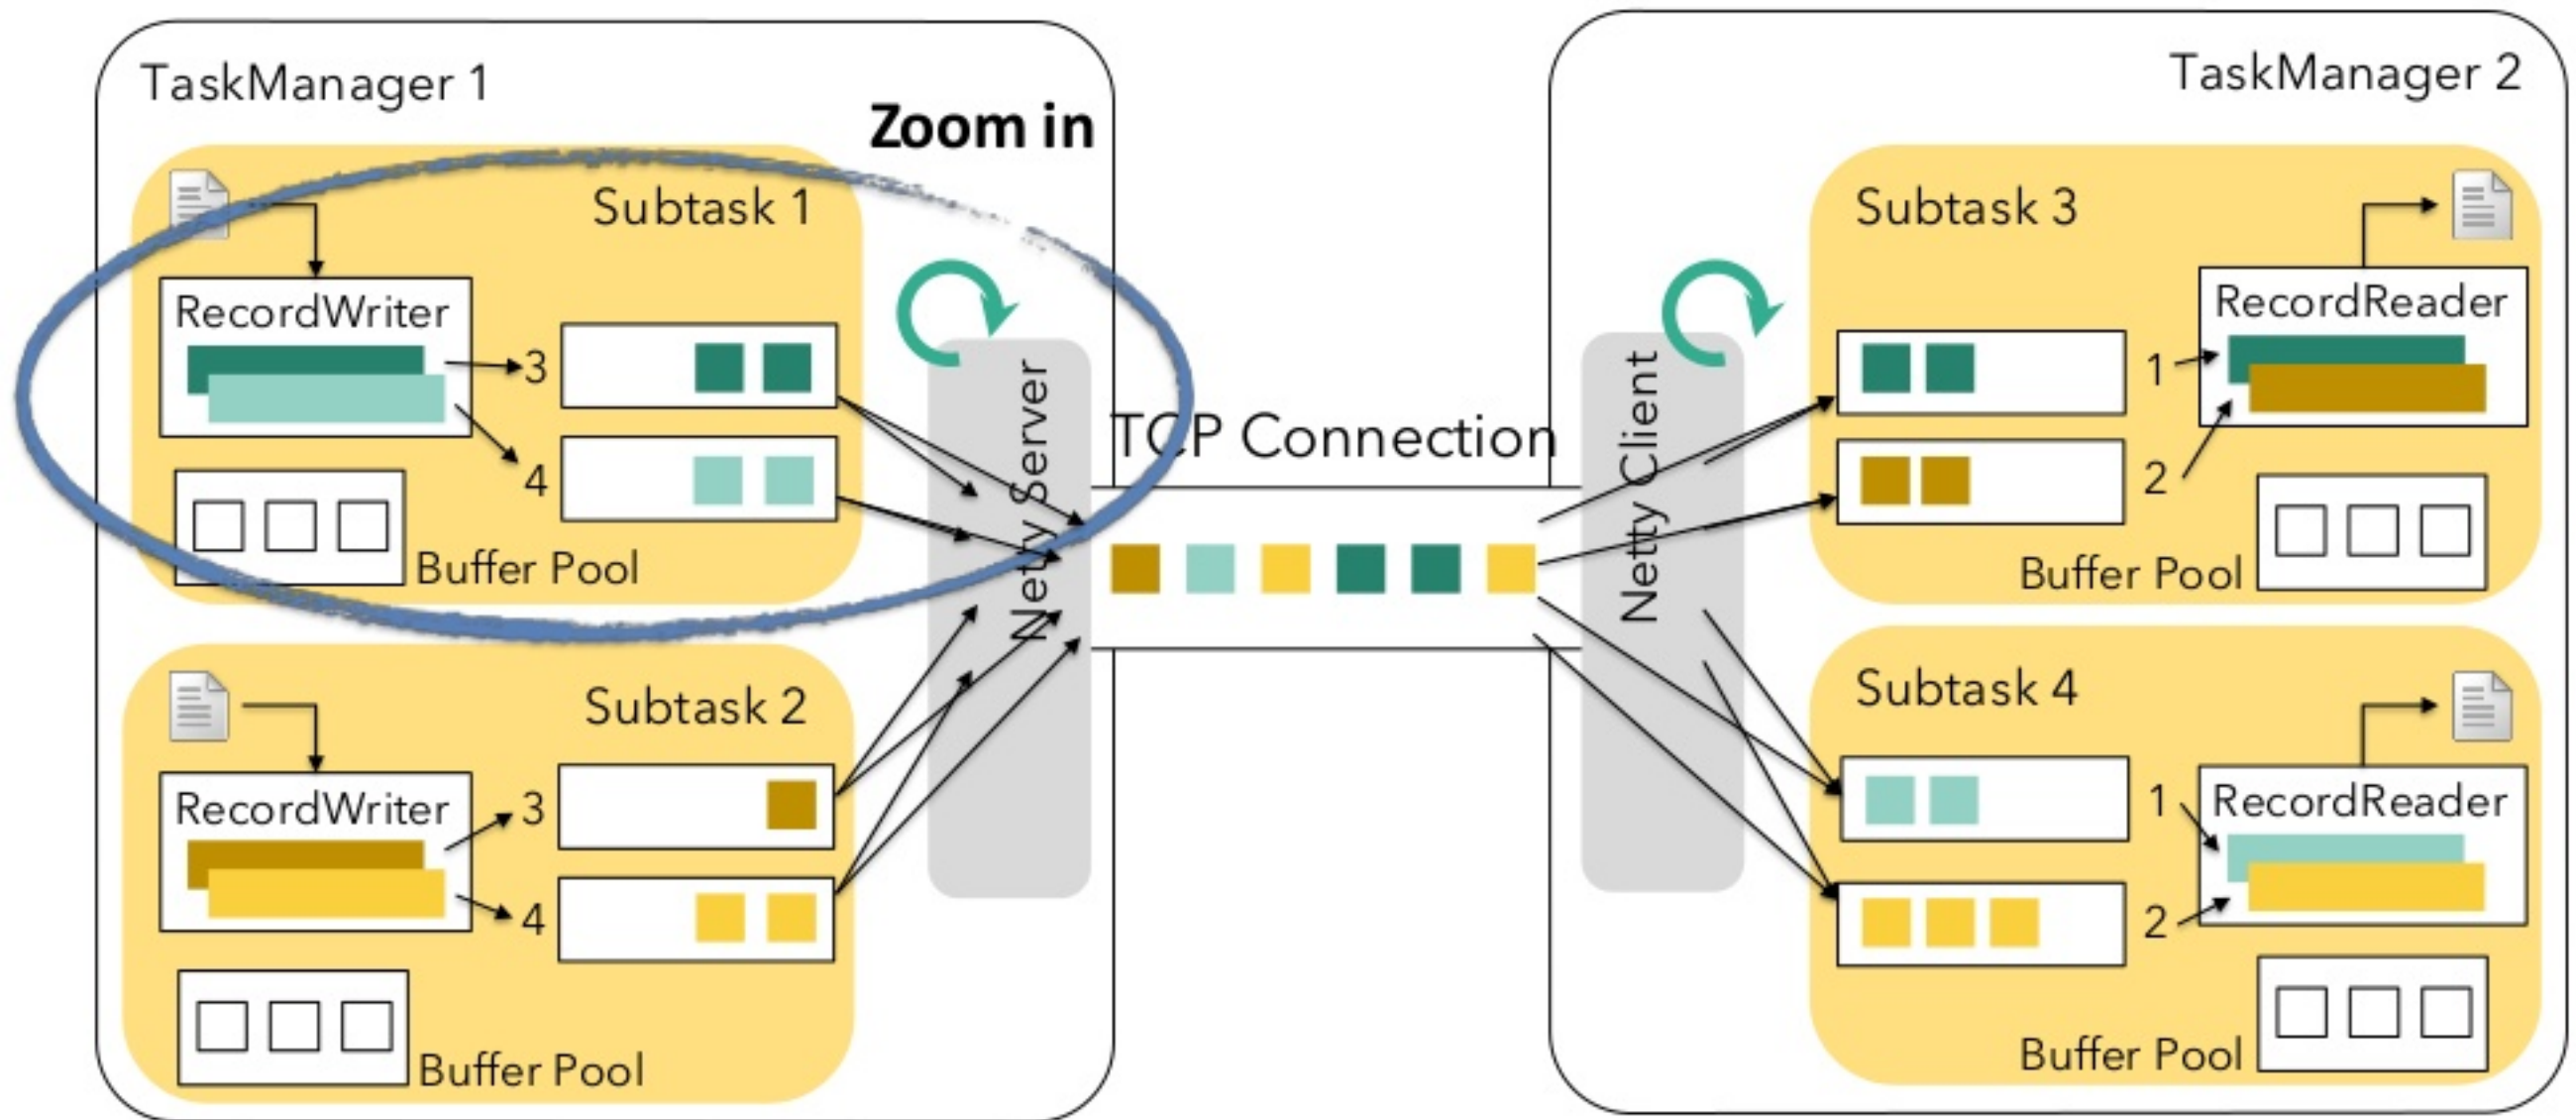
Checkpoint Duration



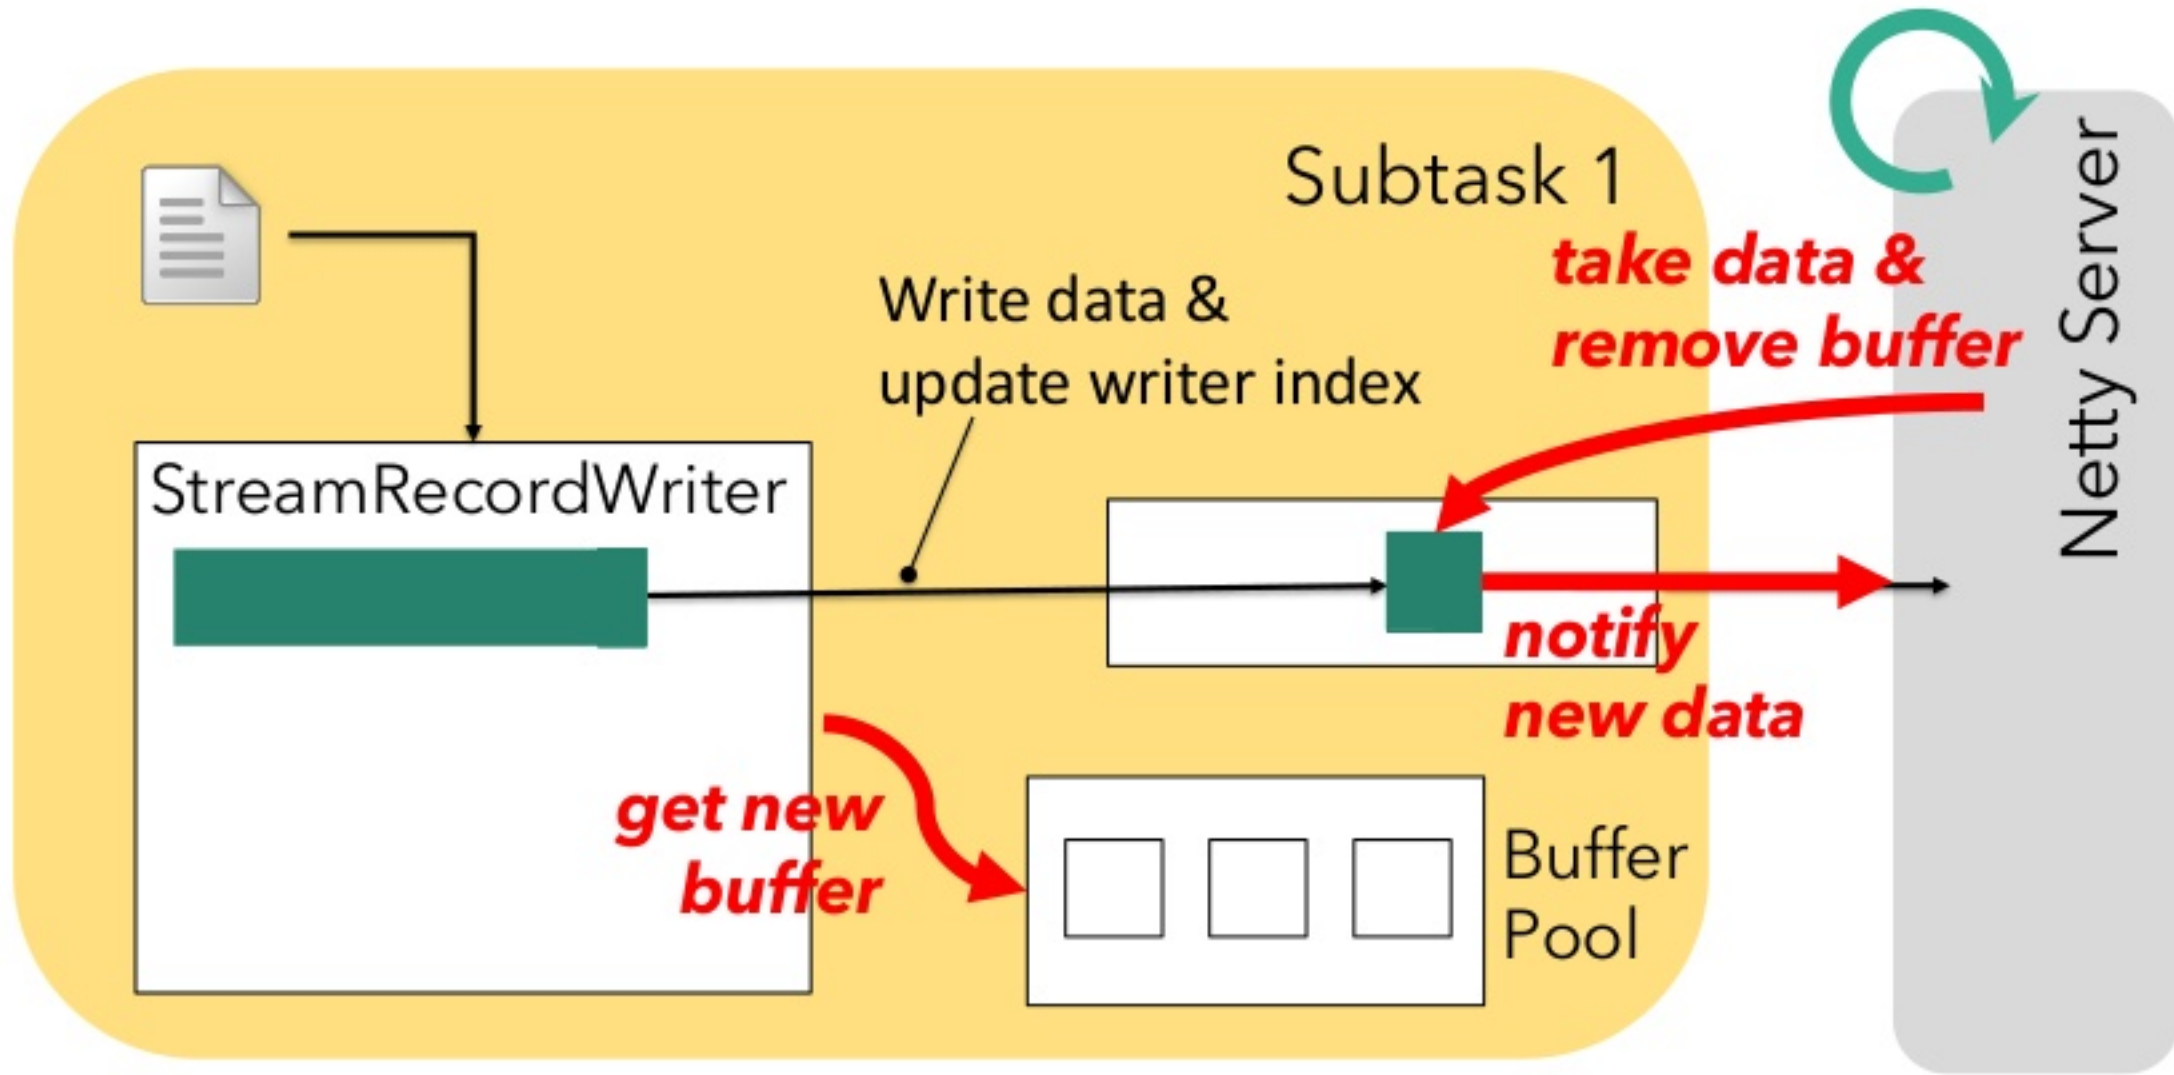
LOW LATENCY IMPROVEMENTS



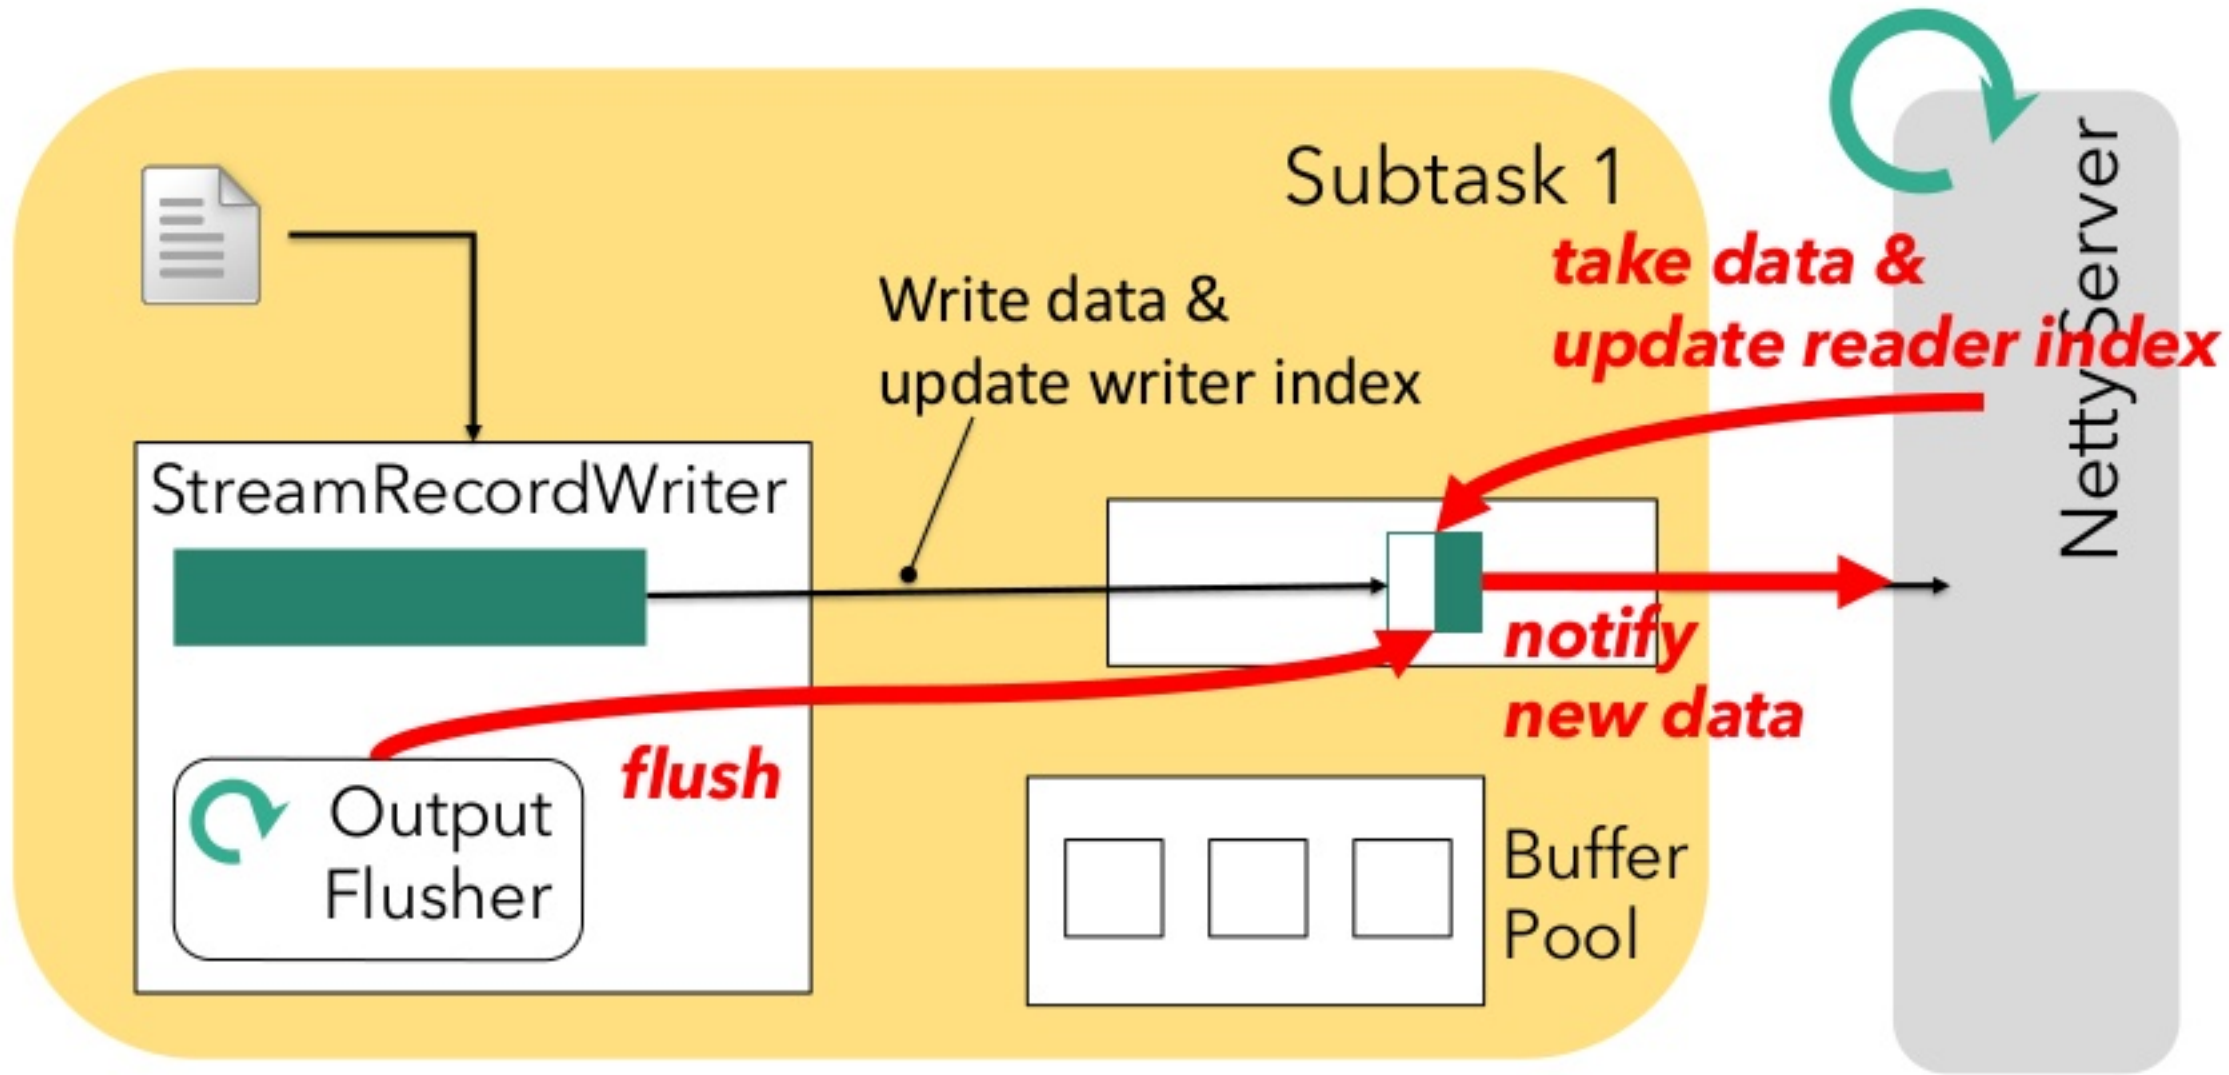
NETWORK STACK (EXTENDED)



FROM RECORD TO NETWORK

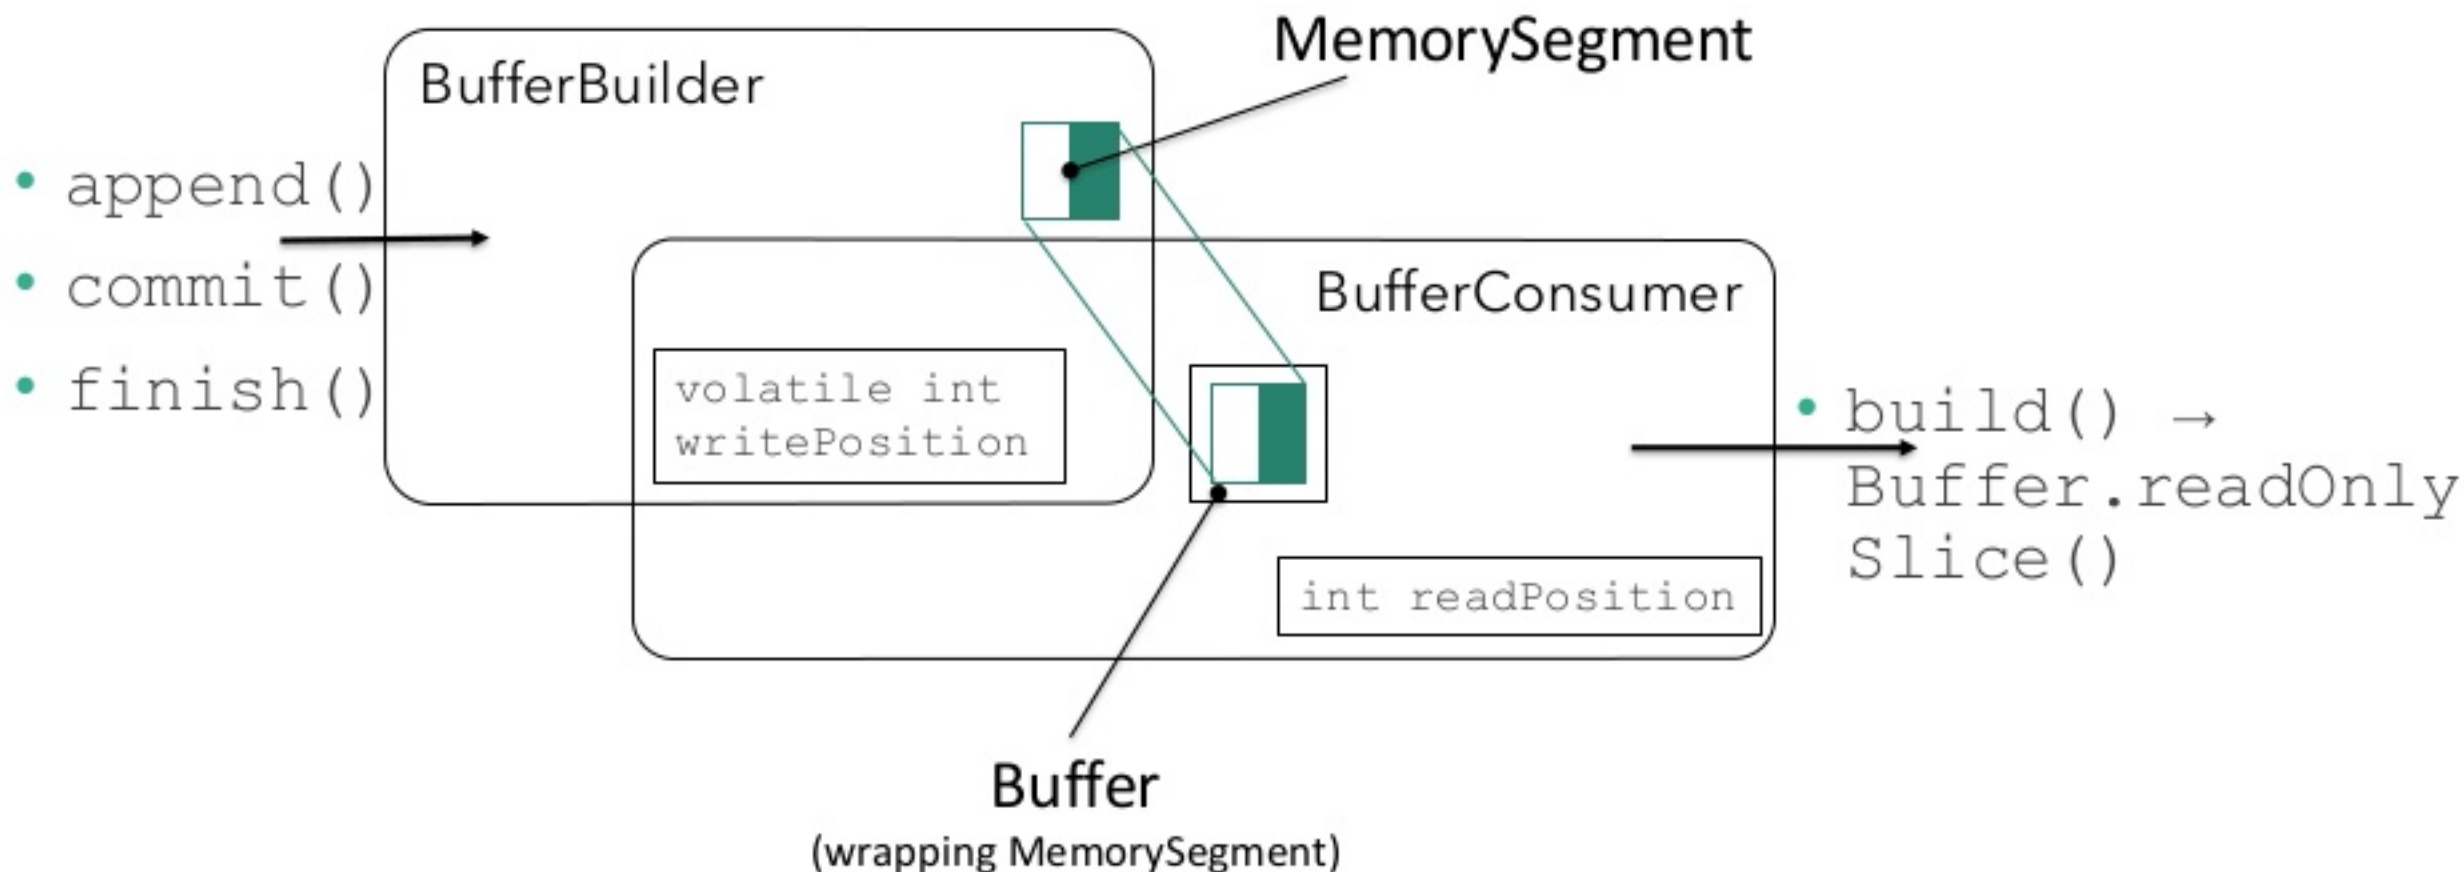


FROM RECORD TO NETWORK




BUFFER BUILDER & CONSUMER

- Producer-Consumer structure with lightweight synchronization



LATENCY VS. THROUGHPUT

- low latency via buffer timeout
- high throughput through buffers



CONNECTION TYPES



LOCAL VS. REMOTE CONNECTIONS

- Every (unchained) connection:
 - Requires serialization
 - Assembles serialized records into buffers
 - Forwards a buffer when it is full or the buffer timeout hit
- Remote connection:
 - Sent via multiplexed Netty TCP connections (one per pair of tasks and task managers)
 - As soon as a buffer is on the wire, it can be re-used
 - Allows credit-based flow control to control amount of buffered data
- Local connection:
 - Direct connection between sender and receiver: buffers are shared
 - No need for further flow control (buffered data = sender buffers)




TUNING OPTIONS




CREDIT-BASED FLOW CONTROL

- `taskmanager.network.credit-model: true/false`
- `taskmanager.network.memory.buffers-per-channel: 2`
- `taskmanager.network.memory.floating-buffers-per-gate: 8`
- Number of exclusive buffers should be enough to saturate the network for a full round-trip-time (2 x network latency)
 - $\#exclBuffers * segmentSize = round-trip-time * throughput$



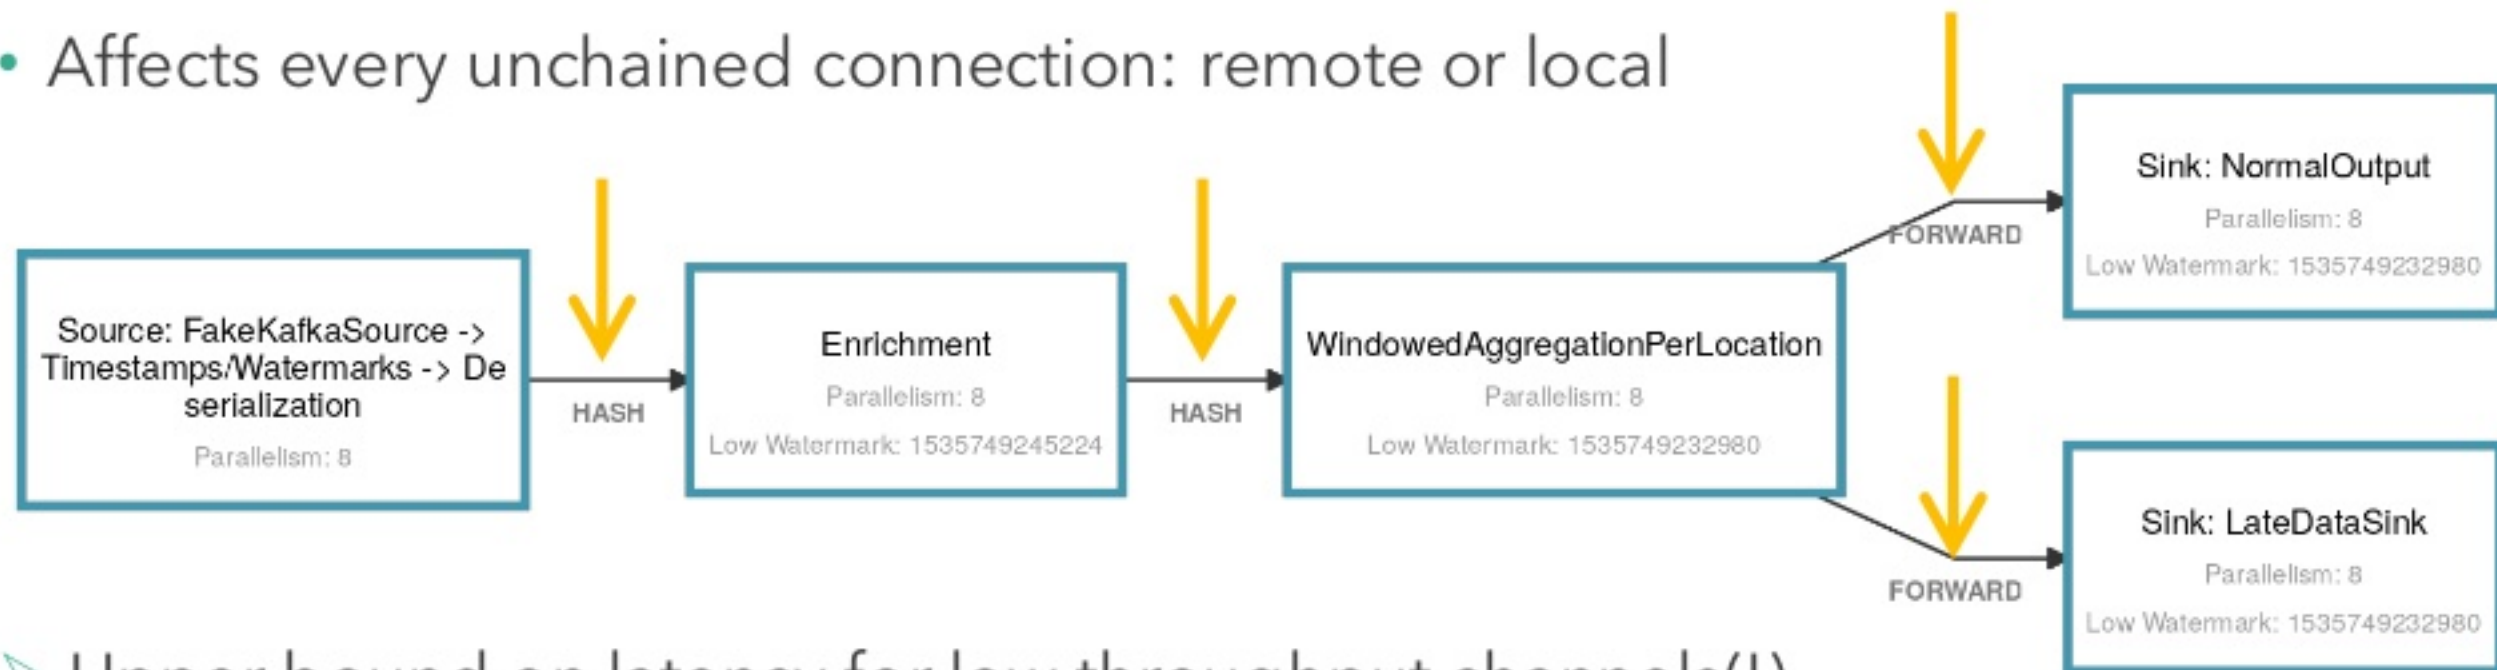
CREDIT-BASED FLOW CONTROL

- Number of exclusive buffers too high
 - higher number of required network buffers
 - buffering more during checkpoint alignment
 - BUT: faster ramp-up (before floating buffers kick in)
- Number of exclusive buffers too low
 - times of in-activity during ramp-up



BUFFER TIMEOUT

- `StreamExecutionEnvironment#setBufferTimeout()`
- Affects every unchained connection: remote or local



- Upper bound on latency for low throughput channels(!)
- Trade-off throughput vs. latency (see [earlier](#))



NETWORK THREADS

- `netty.client.numThreads` (default: number of slots)
 - `netty.server.numThreads` (default: number of slots)
 - May become a bottleneck if thread(s) are overloaded
 - BUT: may also become an overhead if too many
- Do your own benchmarks and verify for your job!



USE LINUX-NATIVE EPOLL (FLINK 1.6+)

- `taskmanager.network.netty.transport:` AUTO | **NIO** | EPOLL
 - EPOLL may reduce the channel polling overhead between user space and kernel/system space
 - There should be no downside in activating this or at least AUTO.
- Do your own benchmarks for your job!
- Please give feedback in [FLINK-10177](#) so that we can decide whether to use AUTO by default.



METRICS



NETWORK STACK METRICS

- Backpressure monitor
 - Web/REST UI, /jobs/:jobid/vertices/:vertexid/backpressure)
- [input, output]QueueLength
- numRecords[In, Out]
- numBytesOut, numBytesIn[Local, Remote]
- numBuffersOut, numBuffersIn[Local, Remote] (Flink 1.5.3+, 1.6.1+)



LATENCY MARKERS

- `ExecutionConfig#setLatencyTrackingInterval()` (default: every 2s)
- Sources periodically emit a `LatencyMarker` with a timestamp
- These flow with the stream and properly queue behind records
- Latency markers bypass operators, e.g. windows
- Once received, they will be re-emitted onto a random output channel
- We create one histogram per source \leftrightarrow operator pair (window size: 128)
- `source_id.<sourceId>.source_subtask_index.<subtaskIdx>.
operator_id.<operatorId>.operator_subtask_index.<subtaskIdx>`
- 10 operators, parallelism 100 = $9 * 100 * 100 = 90,000$ histograms!

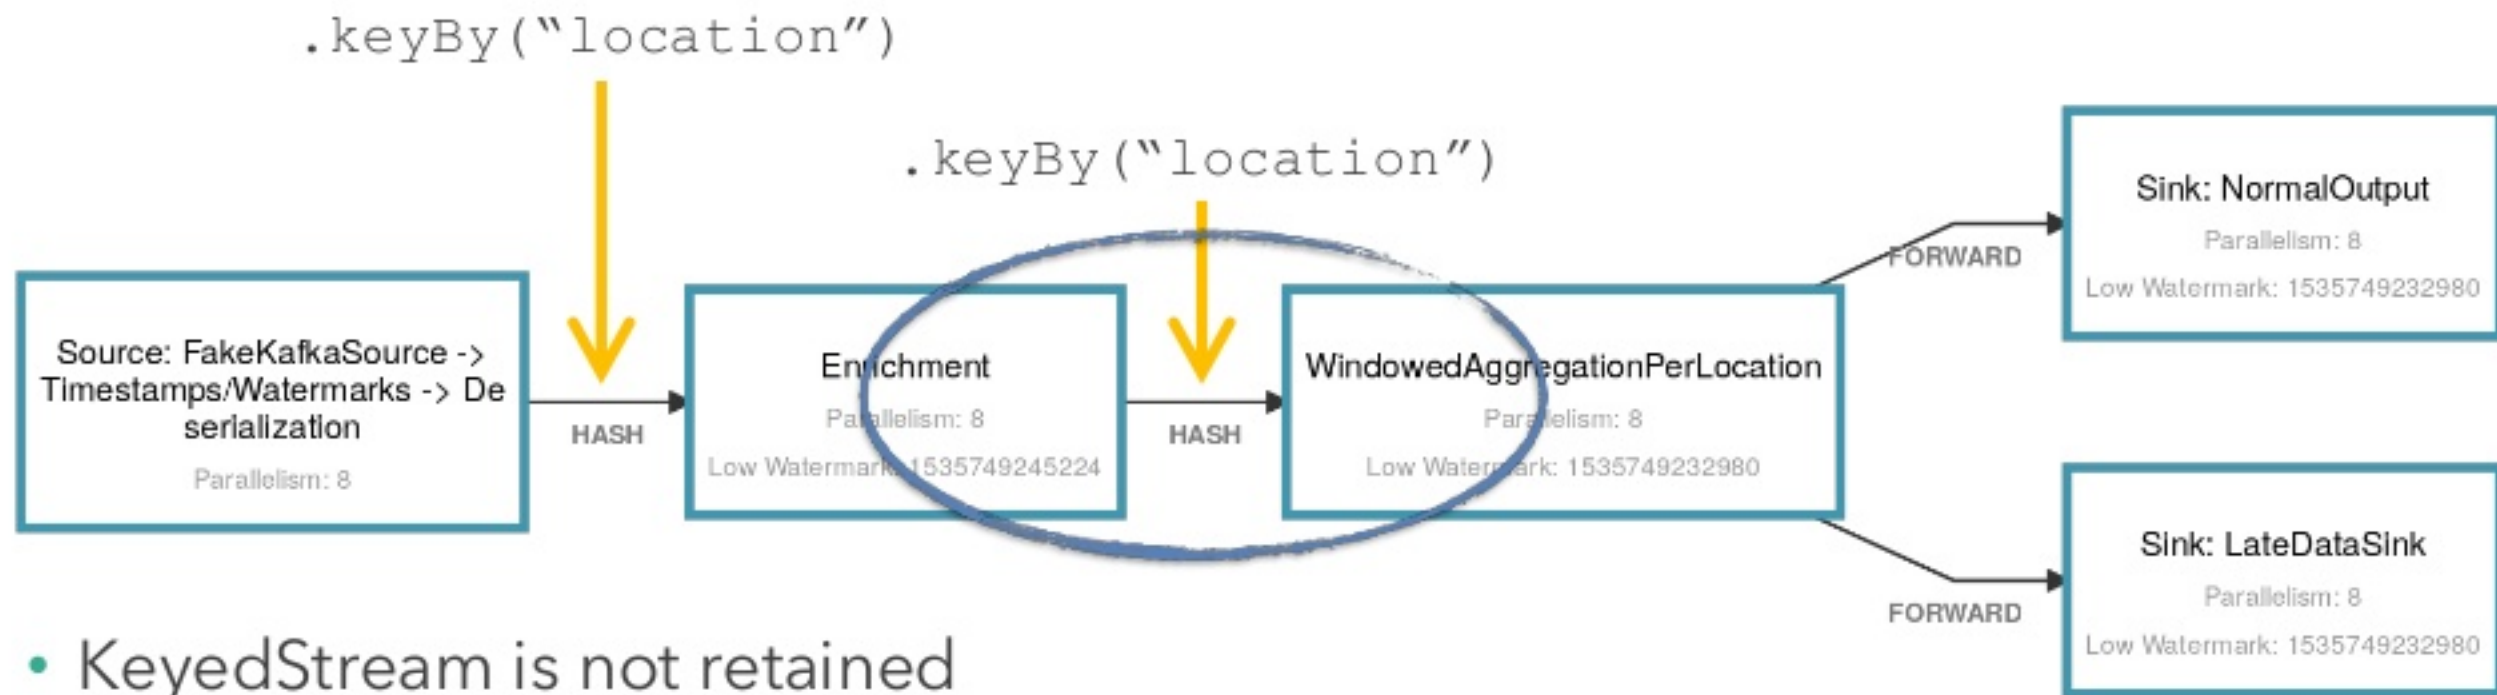
<https://ci.apache.org/projects/flink/flink-docs-stable/monitoring/metrics.html#latency-tracking>



COMMON ANTIPATTERNS



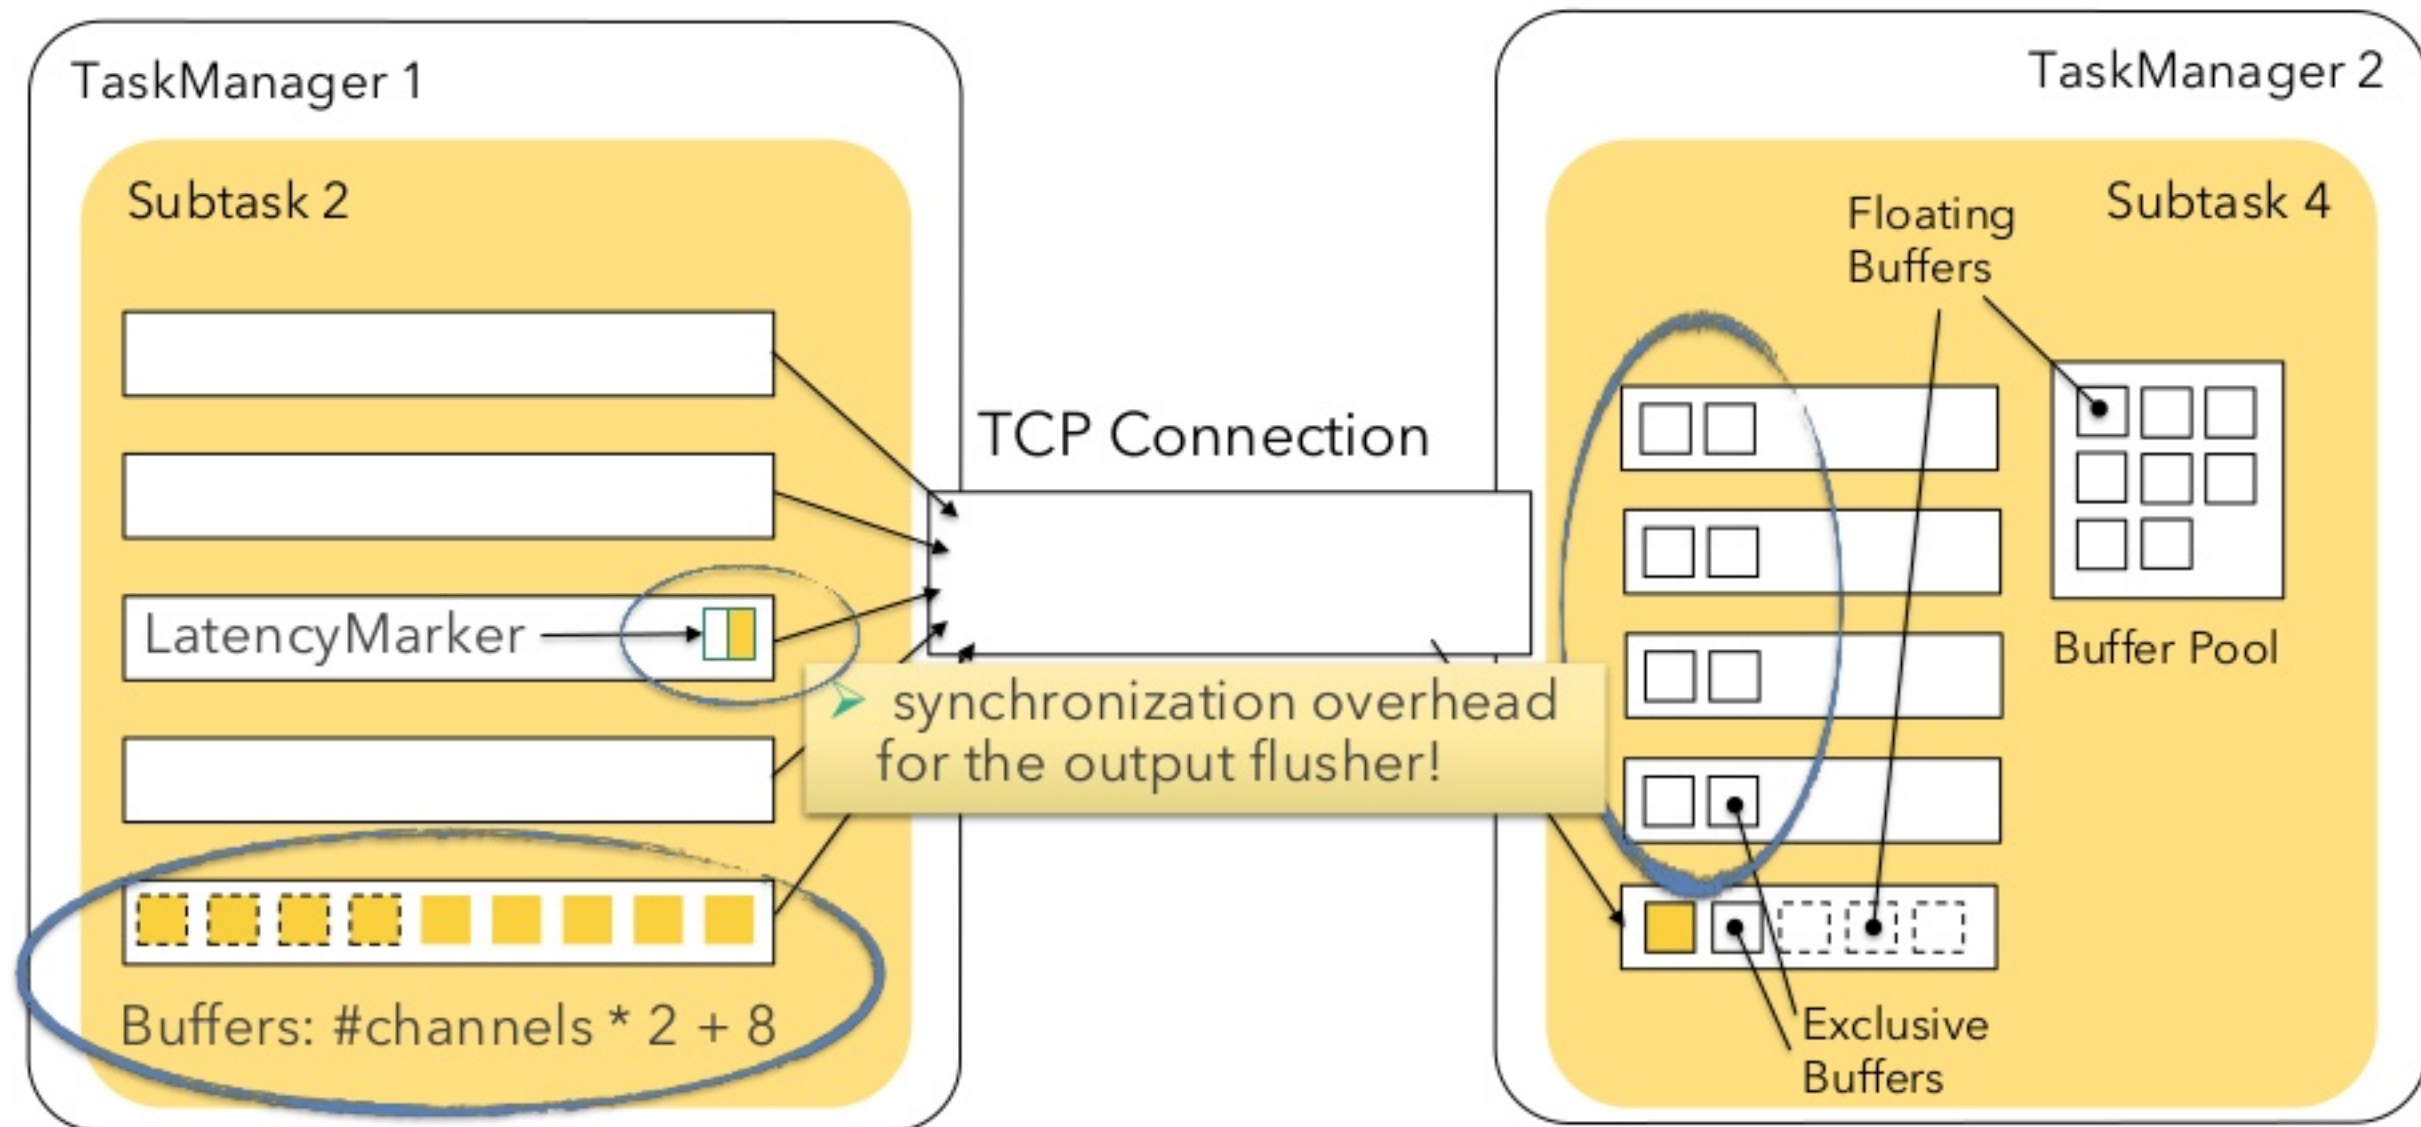
REPEATED KEYBY'S ON THE SAME KEY



- KeyedStream is not retained
 - UDF could have changed the key
 - Additional `keyBy()` is necessary to gain access to keyed state, but:
 - Prevents chaining
 - Adds an additional shuffle
- `DataStreamUtils#reinterpretAsKeyedStream`



CREDIT-BASED FLOW CONTROL (FLINK 1.5+)

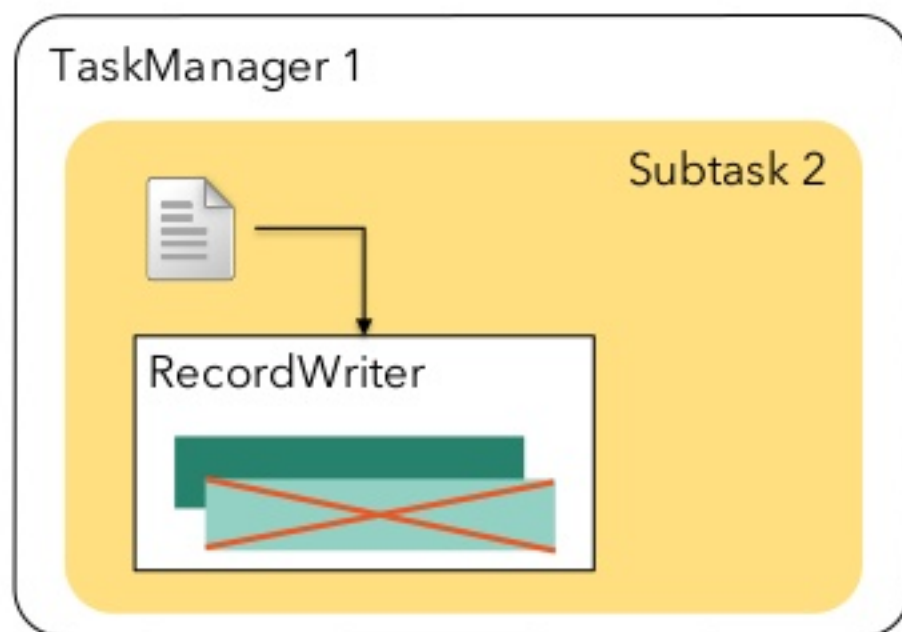


WHAT'S UP NEXT?



NETWORK SERIALIZATION STACK (FLINK 1.7?)

- Serialization for broadcasts once per record, not channel
- Only one intermediate serialization buffer (on heap)
 - significantly reduces the memory footprint
- see [FLINK-9913](#)




OPENSSL-BASED SSL ENGINE (FLINK 1.7?)

- Runs native code
- Uses advanced CPU instruction sets
- May reduce encryption/decryption overhead (needs verification)
- see [FLINK-9816](#)



MOVE OUTPUT FLUSHER TO NETTY

- Current implementation may have (GC) problems with many channels
 - schedule the output flusher inside the Netty event loop



- see [FLINK-8625](#)



THANK YOU!

@dataArtisans
@ApacheFlink

WE ARE HIRING
data-artisans.com/careers

dataArtisans