

Exploiting Apache Flink's Stateful Operators

Gijsbert van Vliet & Olga Slenders
Berlin • 4 September 2018



Agenda

- About ING
- About the Use Case
- Stateful example + code
- Advanced stateful example + code
- Conclusion

About ING

About ING

Almost 9 million customers in the Netherlands*



9 million transactions per month on ING machines*



1,45 million customer logins per day on Mijn ING website*




4,2 million customer logins per day on ING Banking App*



*per january 2018


Use Case

Use Case - Prevent Fraudulent Transactions



Model Input - Features

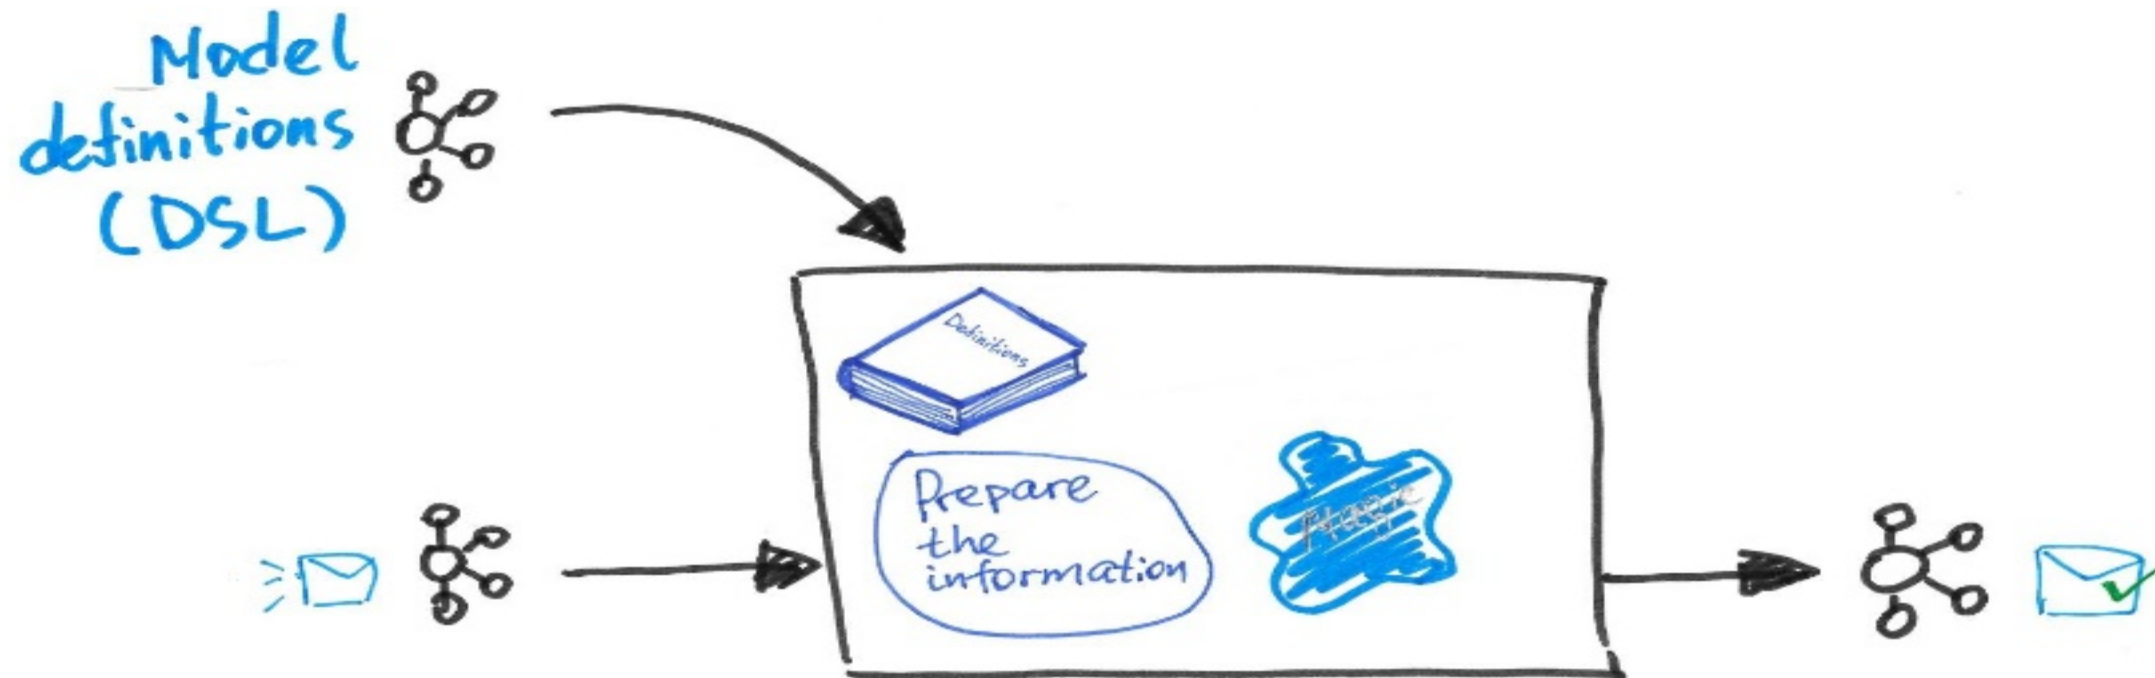
- Speed of the client (between transactions)
- The debit card has been reported broken or stolen
- The customer has been blocked



Use Case – Requirements

- **Dynamic behaviour -> DSL**
- Historical data (state)
- Real-time (< 100 ms)
- High volumes (10K+ tps)
- Scalability

Use Case – Architecture

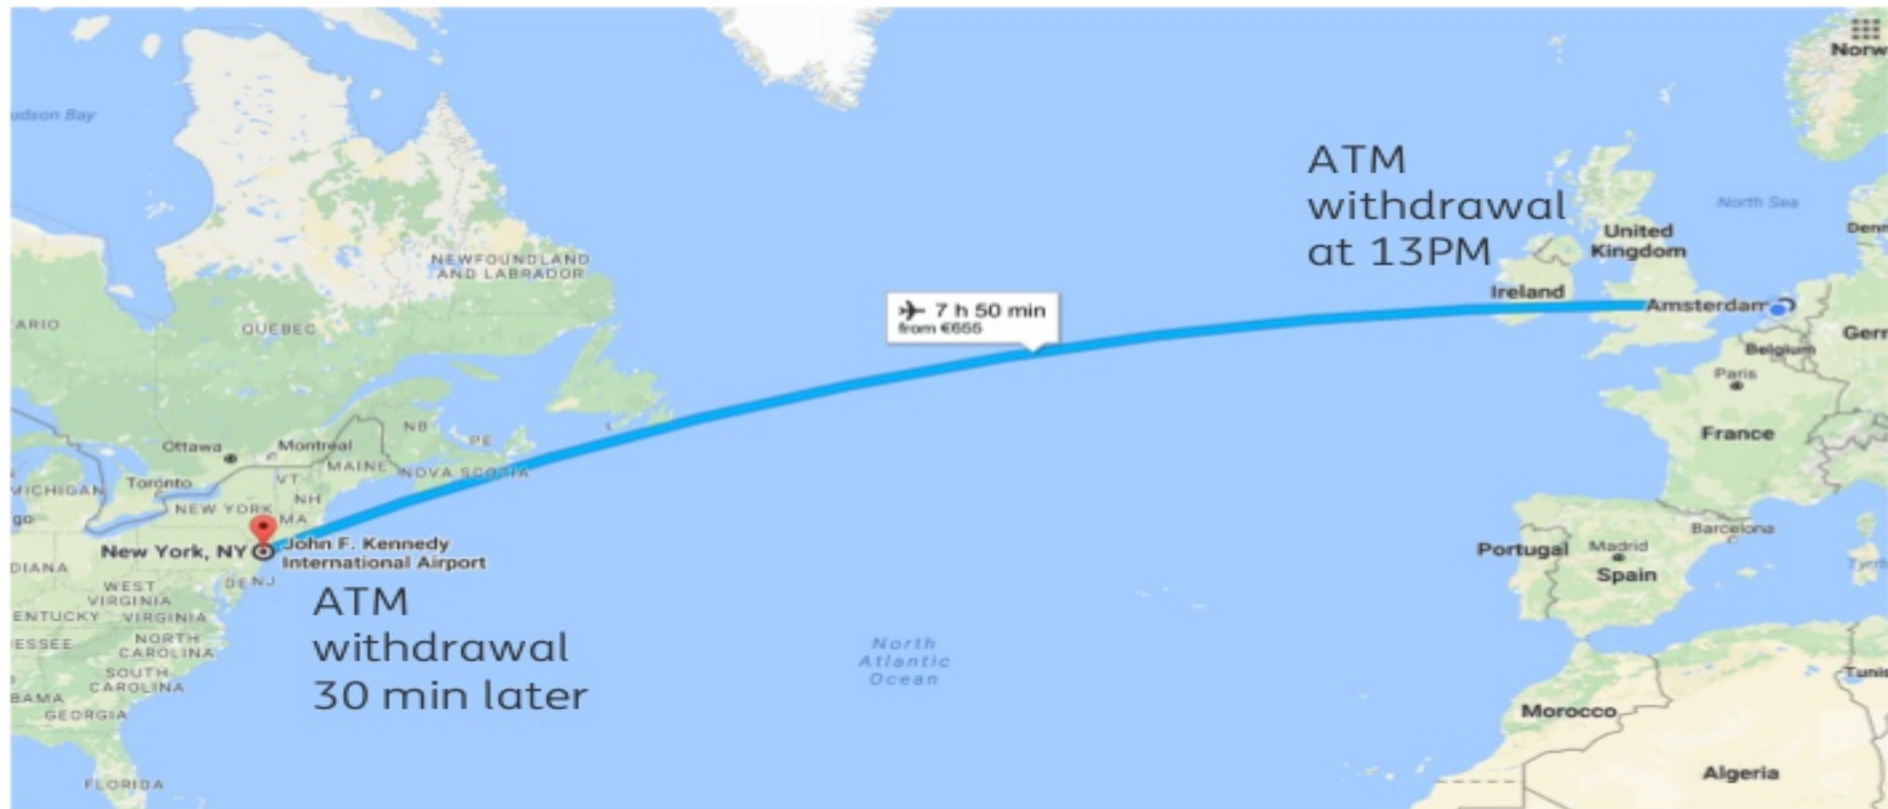


Disclaimer



Scoring Example

Flying Carpet Rule



Need for State



To score we need:
Customer Speed ?

Event: Transaction

customerId
cardId
amount
location
eventTime

Need for State



To score we need:
Customer Speed ?

Event: Transaction

customerId
cardId
amount
location
eventTime

+

prev.Location
prev.eventTime

Need for State



To score we need:
Customer Speed ?

Event: Transaction

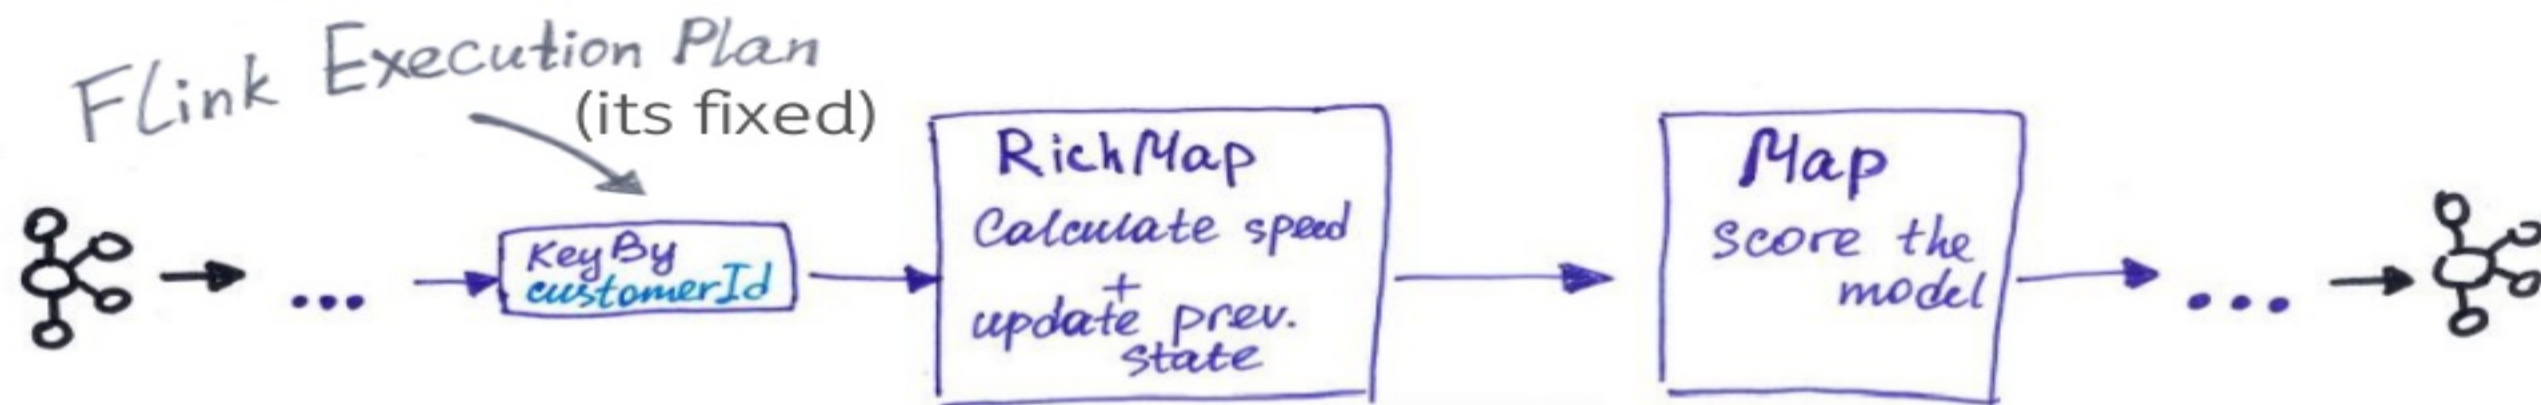
customerId
cardId
amount
location
eventTime

+

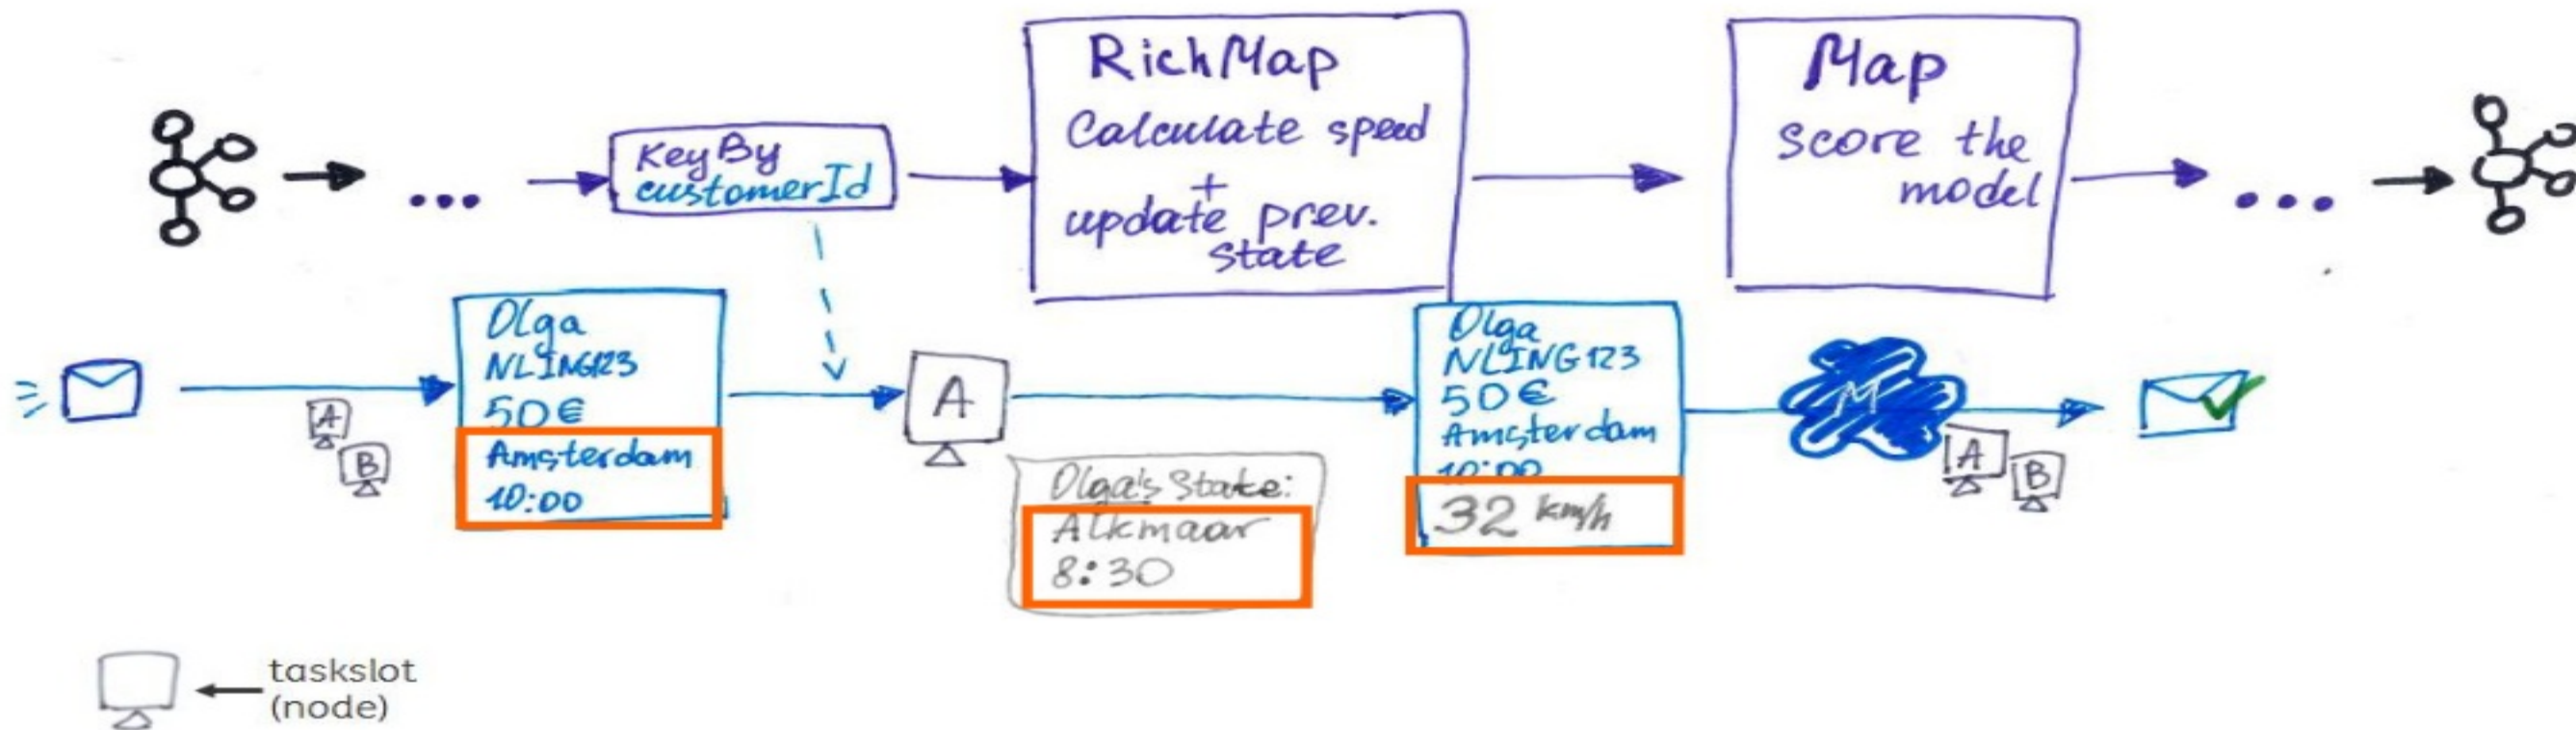
prev.Location
prev.eventTime

 State

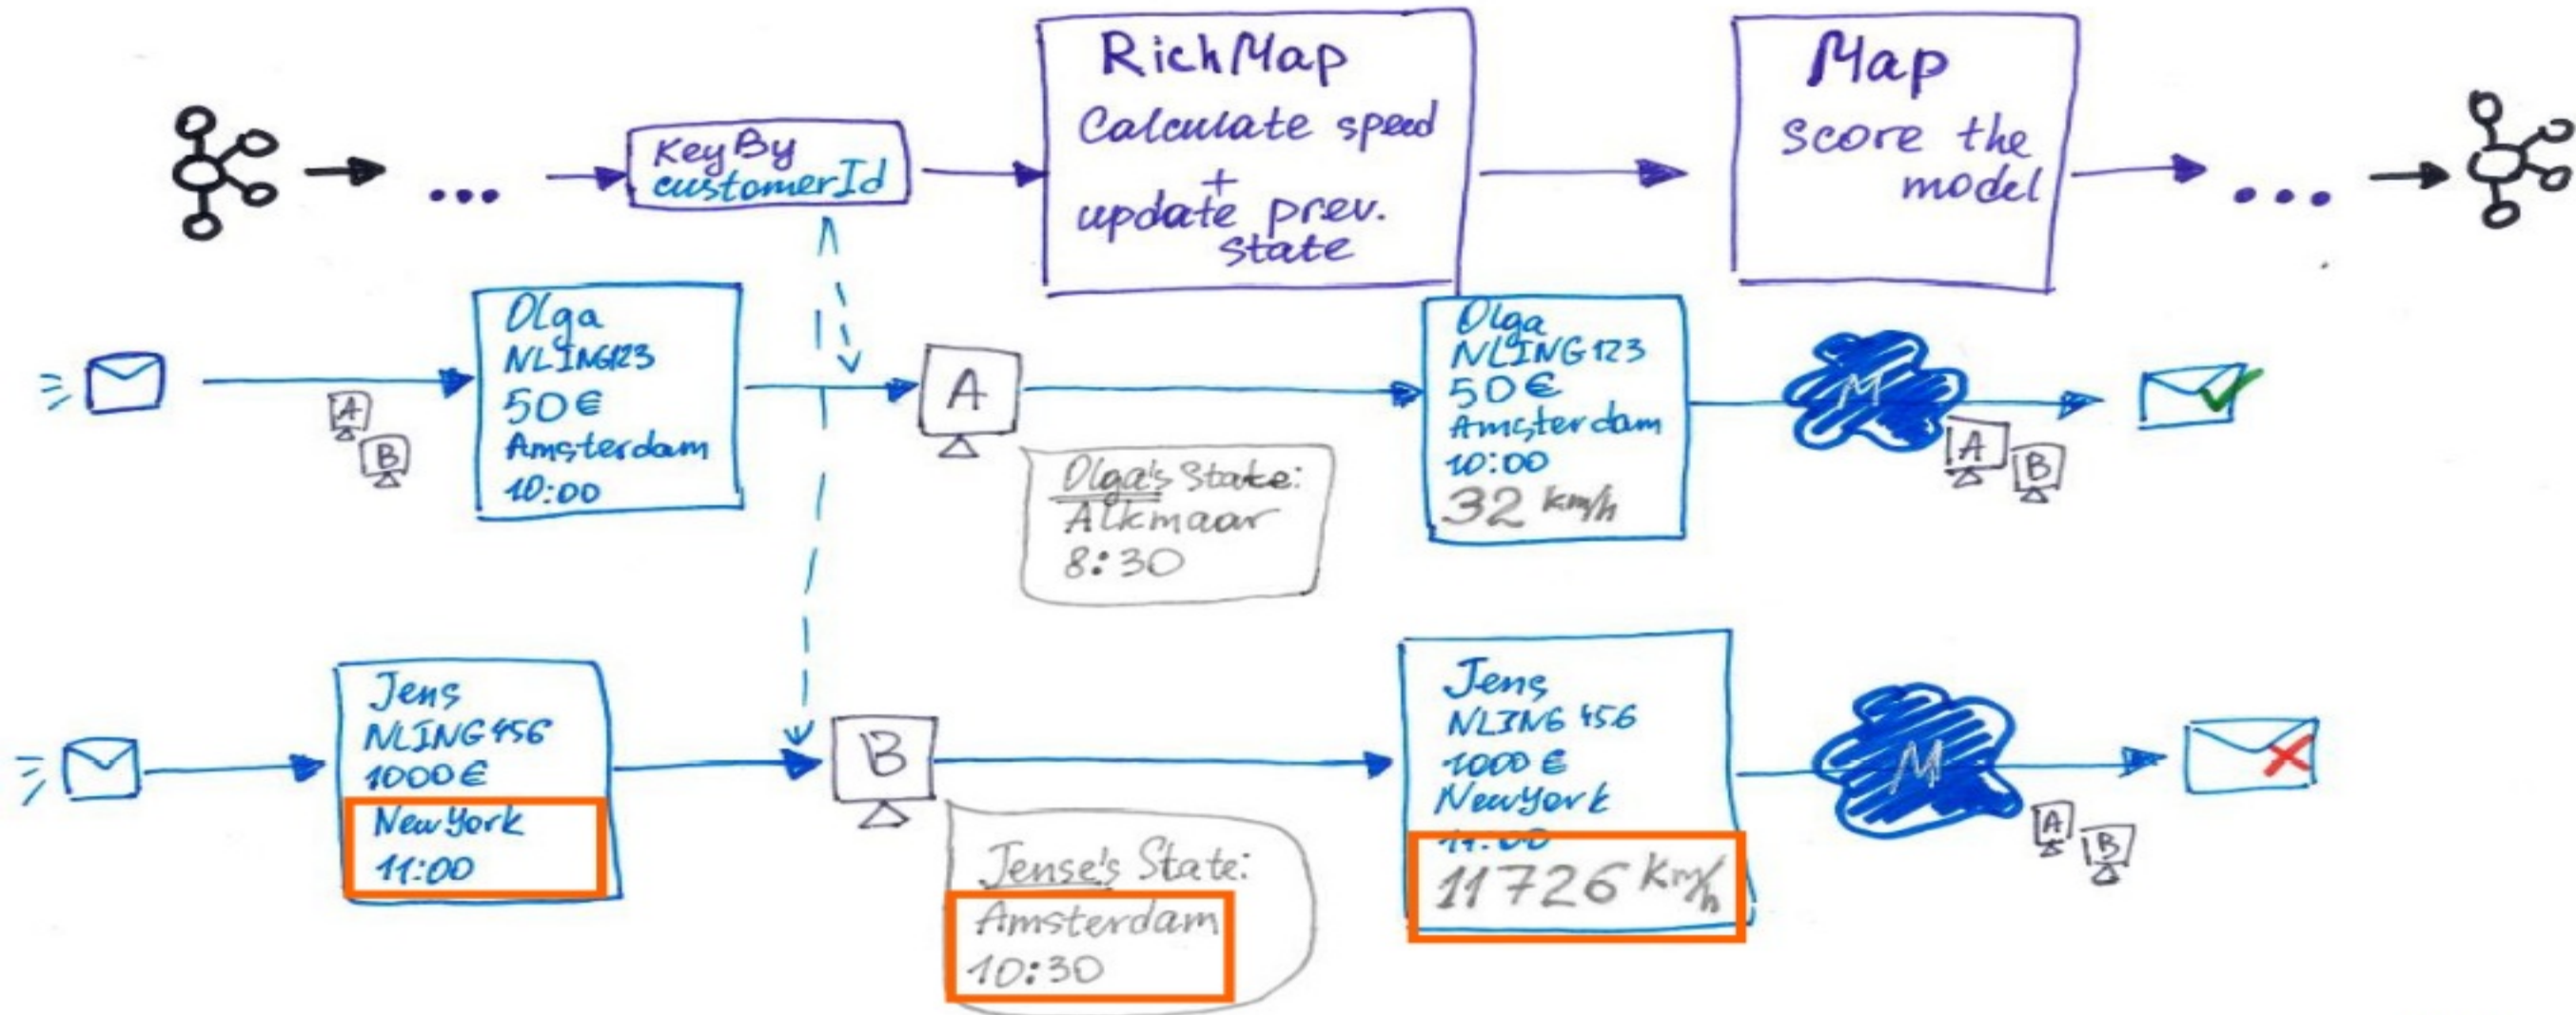
Stateful Flying Carpet



Stateful Flying Carpet



Stateful Flying Carpet



Stateful Flying Carpet

CODE

Improved Stateful Example

Flying Carpet Model Improved



To score *Better* we need:
Customer Speed

Event: Transaction

...

+

prev. Location
prev. event Time

Flying Carpet Model Improved



To score *Better* we need:

Customer Speed
+ *Av. amount per card*

Event: Transaction

...

+

prev. Location
prev. event Time

+

total Amount
total nr. transactions

Flying Carpet Model Improved



To score *Better* we need:

Customer Speed
+ *Av. amount per card*

Event: Transaction

...

+

prev. Location
prev. event Time

 ← *Customer state*

+

total Amount
total nr. transactions

Flying Carpet Model Improved



To score *Better* we need:

Customer Speed
+ *Av. amount per card*

Event: Transaction

...

+

prev. Location
prev. event Time

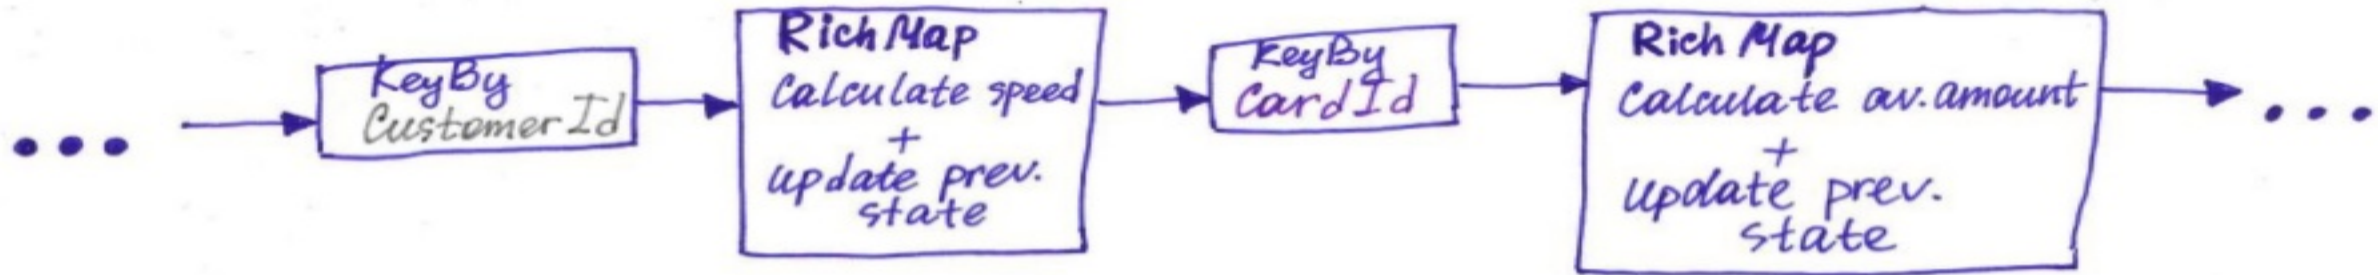
 ← *Customer state*

+

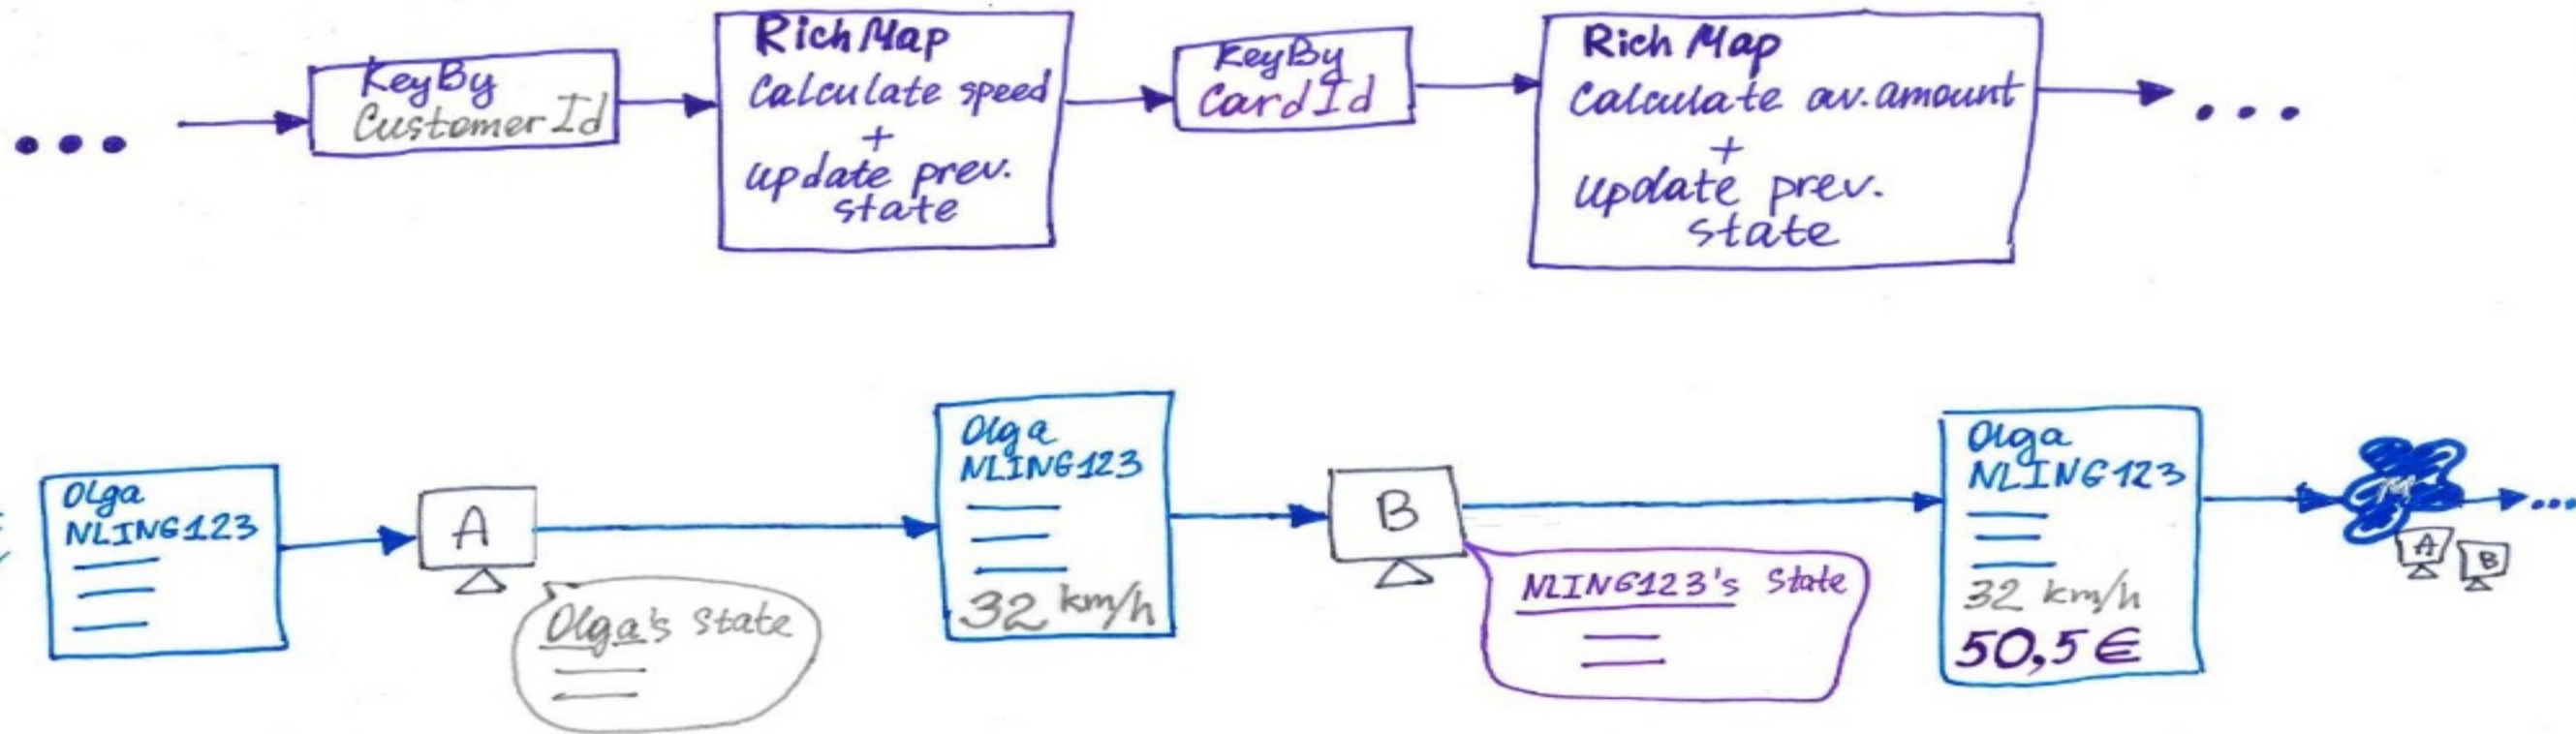
total Amount
total nr. transactions

 ← *Card state*

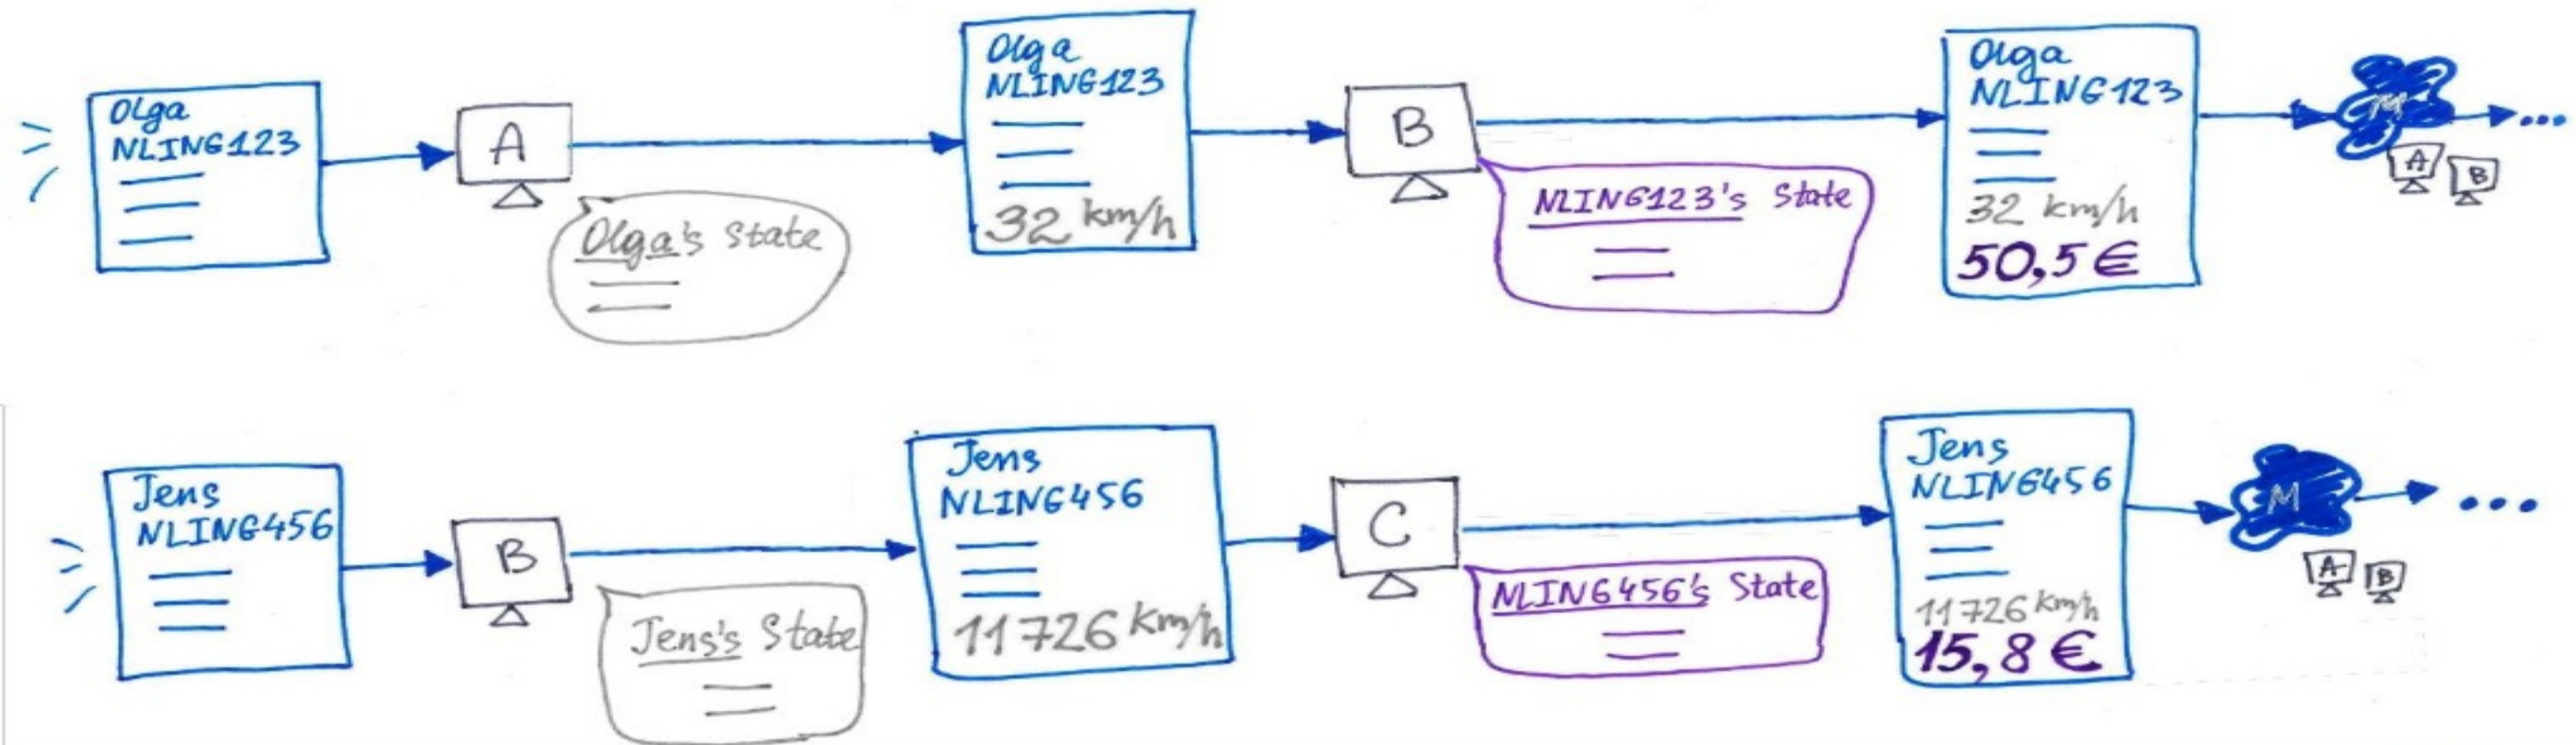
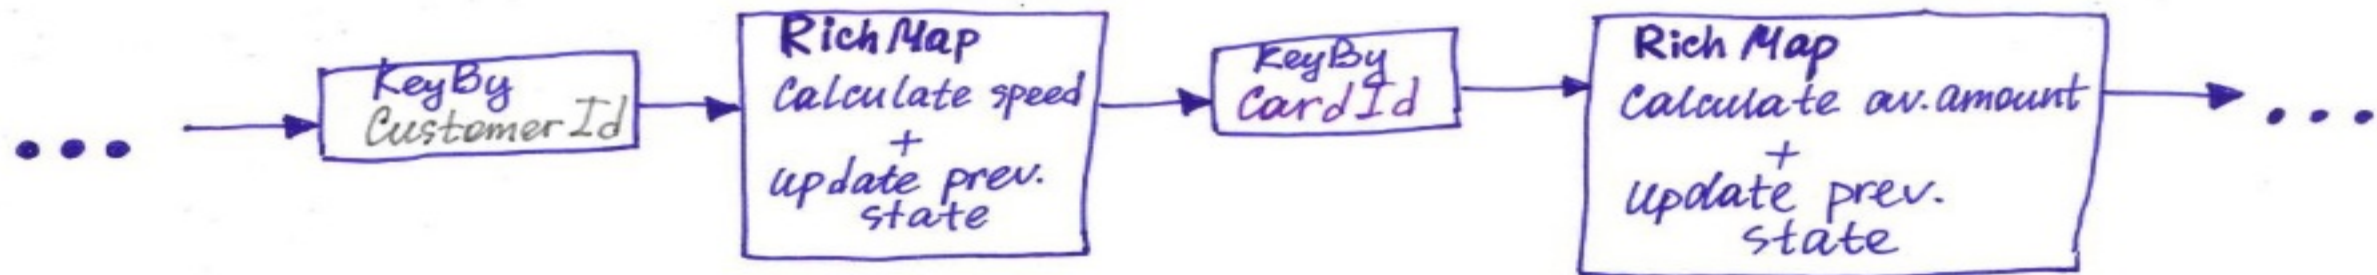
Flink - Sequential Approach



Flink - Sequential Approach



Flink - Sequential Approach



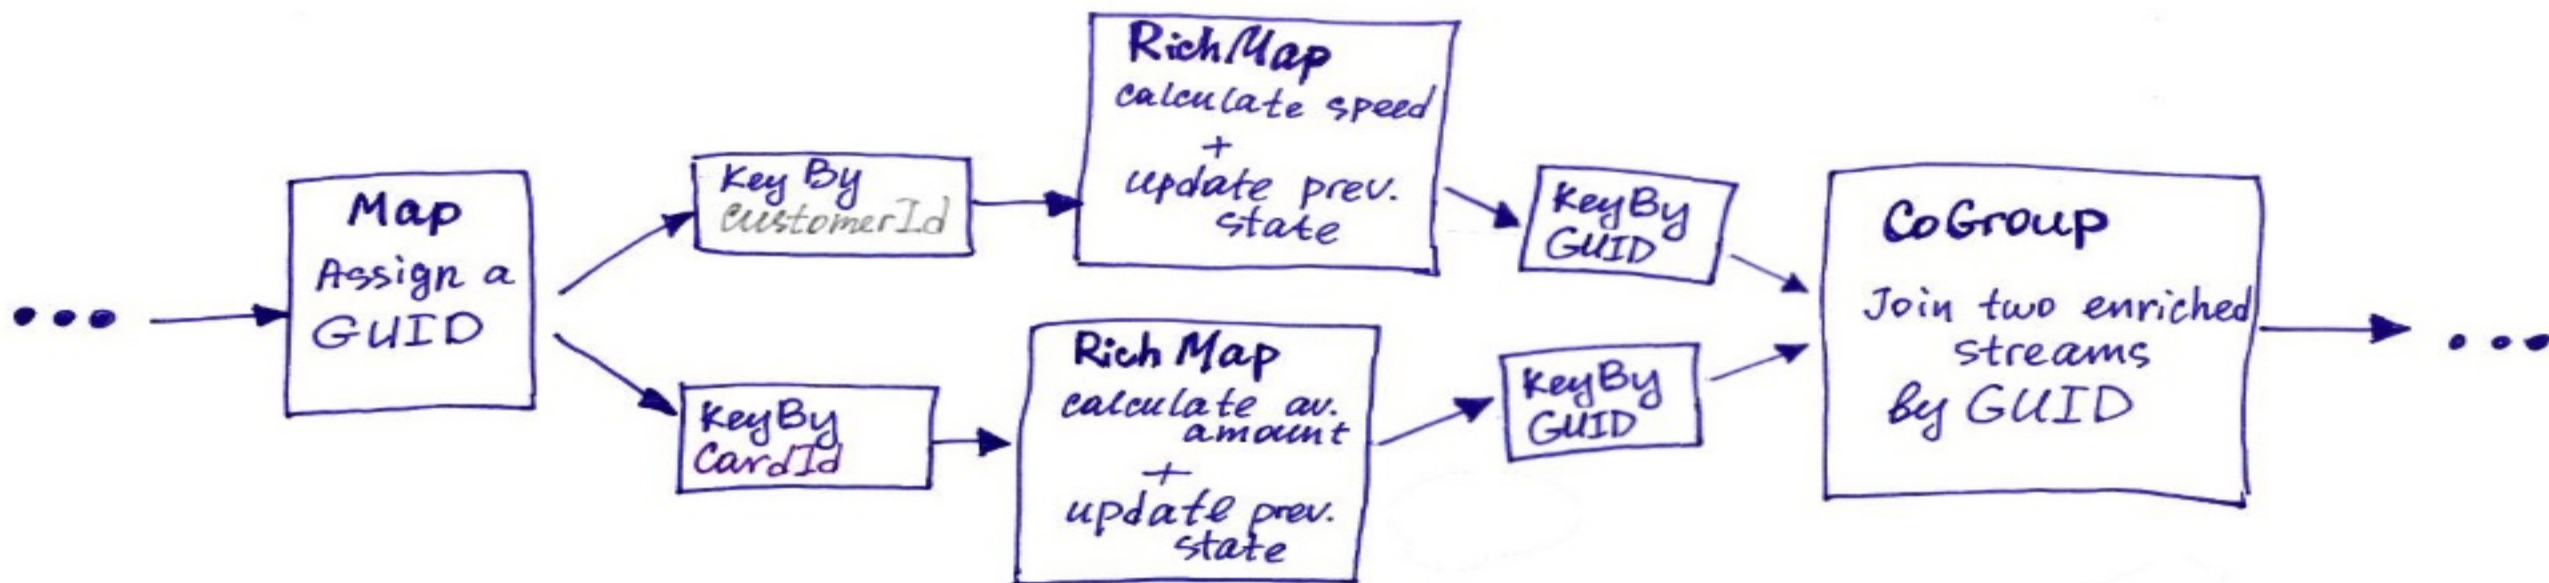
Sequential Approach

CODE

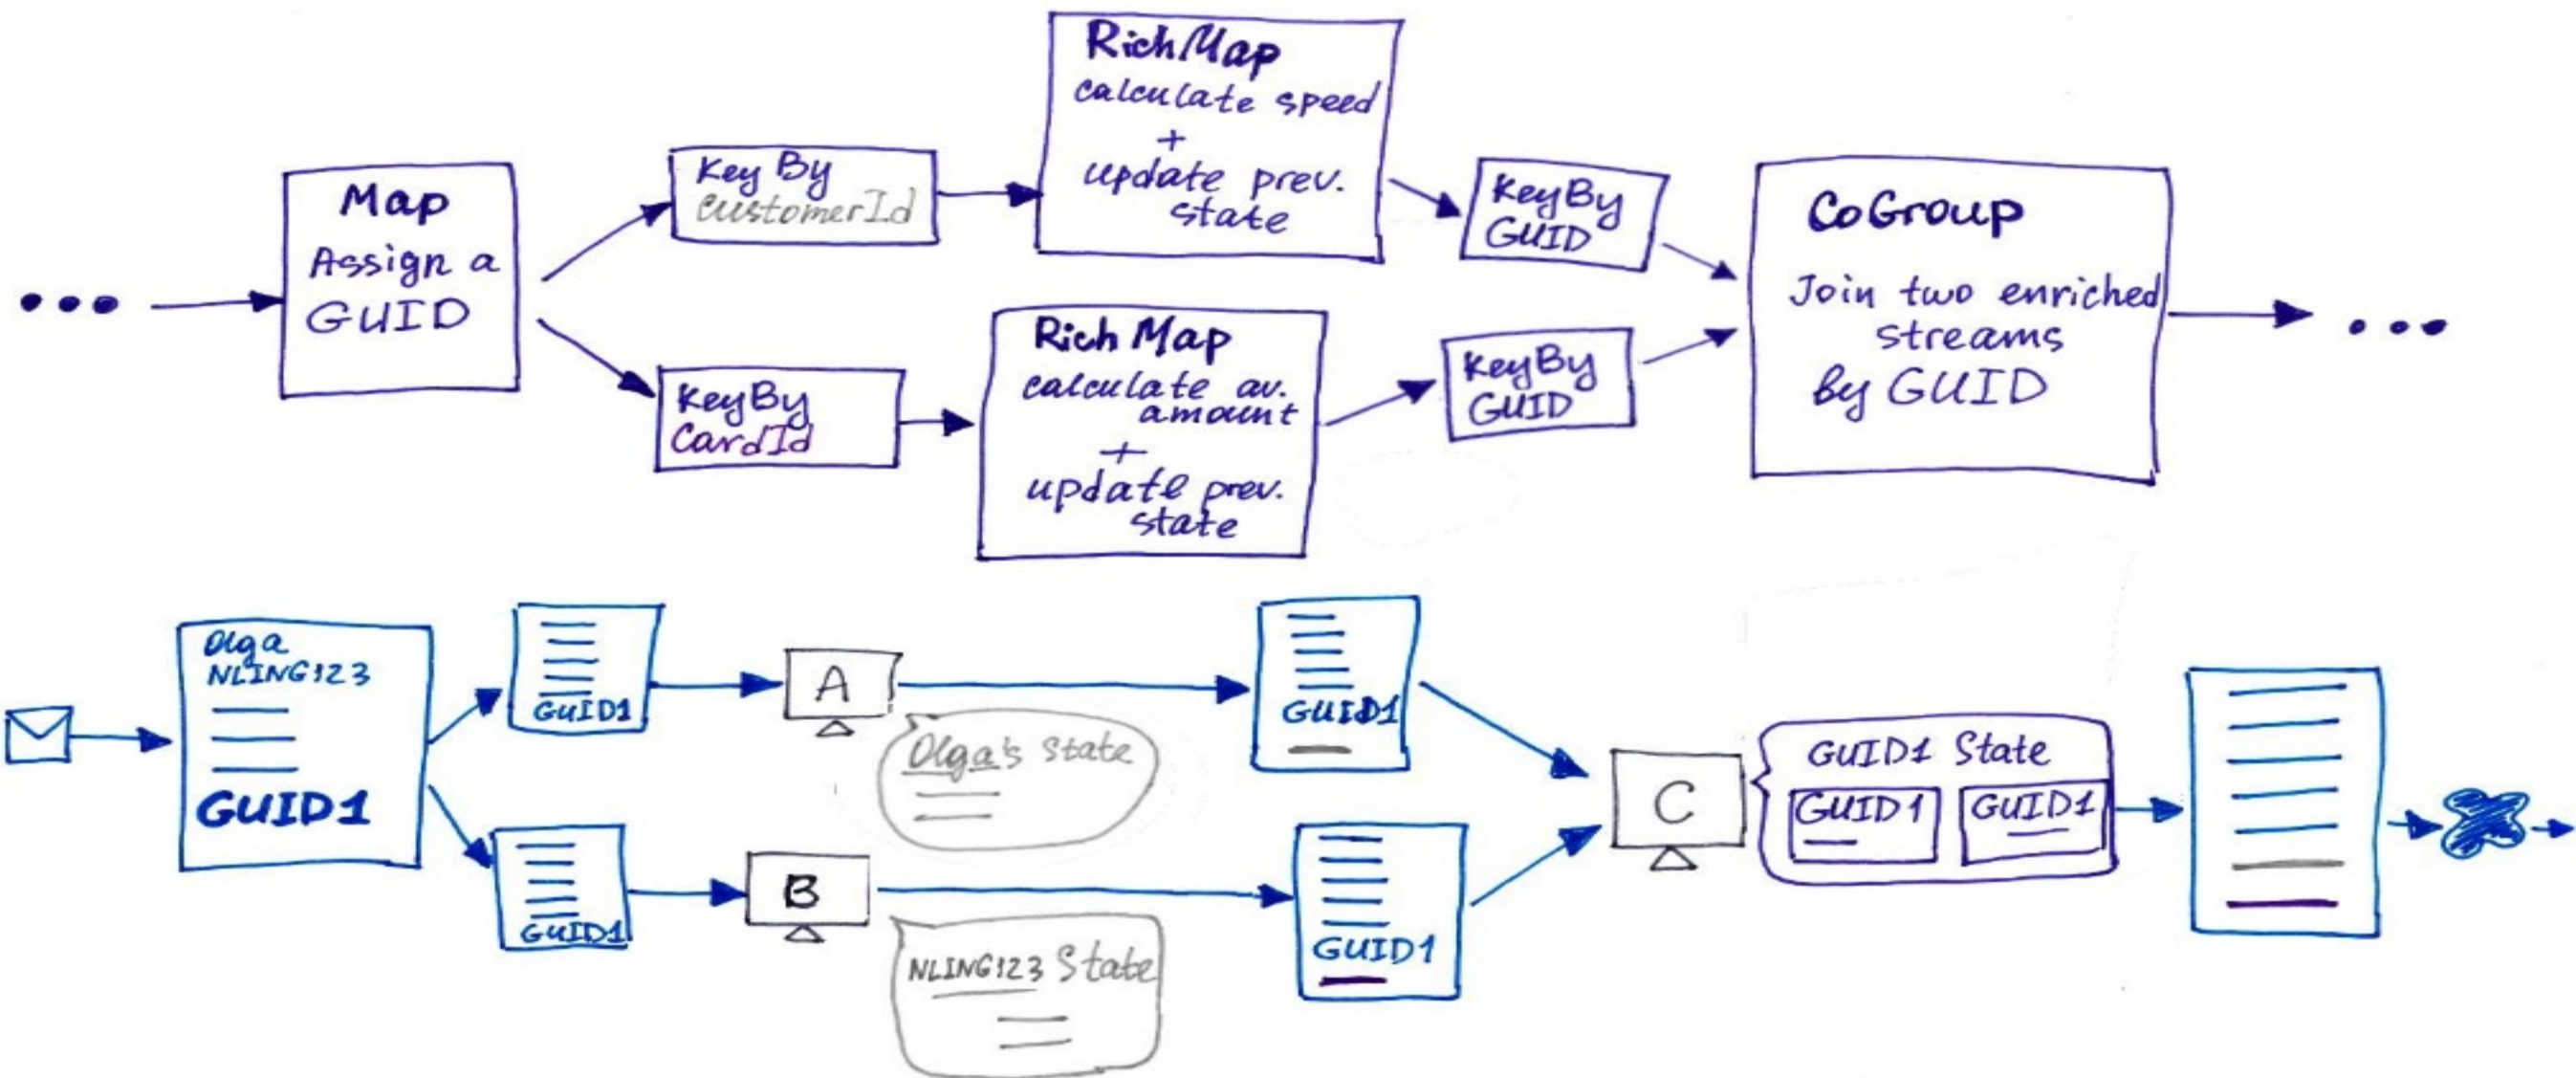
Sequential Approach - Evaluation

- + It works!
- Latency
- Not so scalable
- Event has to pass nodes even if their state is not required
(could be solved with a filter)
- Cannot add new keys without restarting a job
(remember: The execution plan is fixed)

Parallel Approach



Parallel Approach



Parallel Approach

CODE

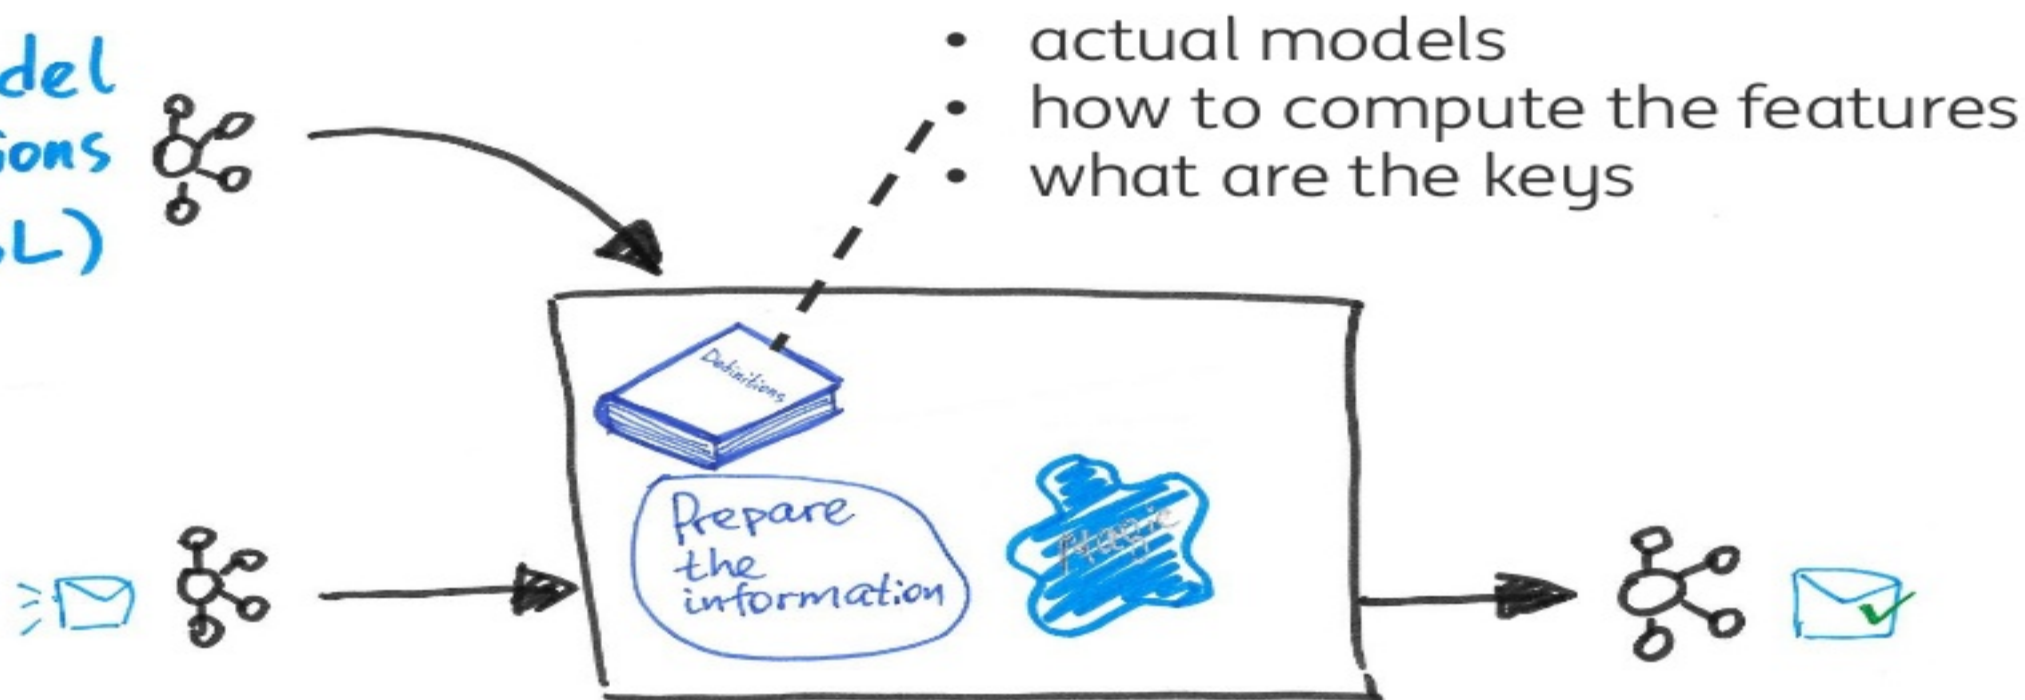
Parallel Approach - Evaluation

- + Less latency than when using sequential approach
- + More scalable
- Event has to pass nodes even if their state is not required
- Complex execution plan
- Cannot add new keys without restarting a job
(remember: The execution plan is fixed)

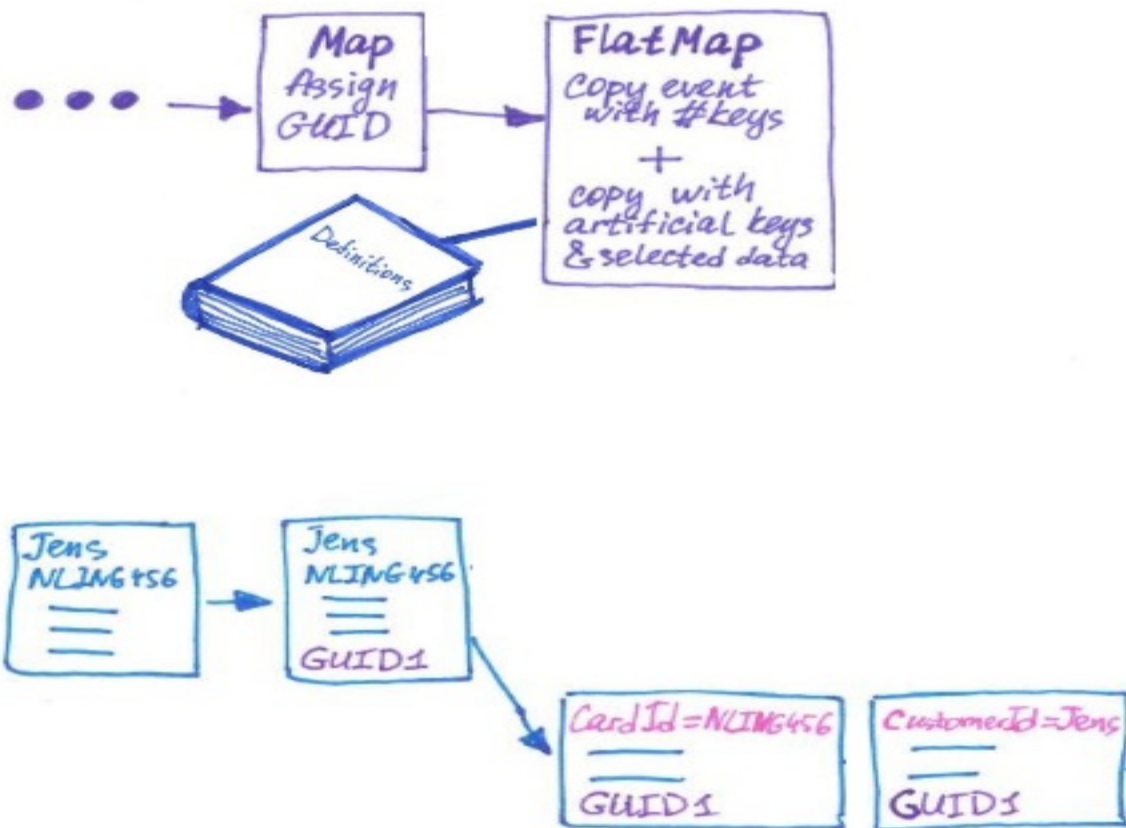
Other Approach to Improved Stateful Example

Architecture recap

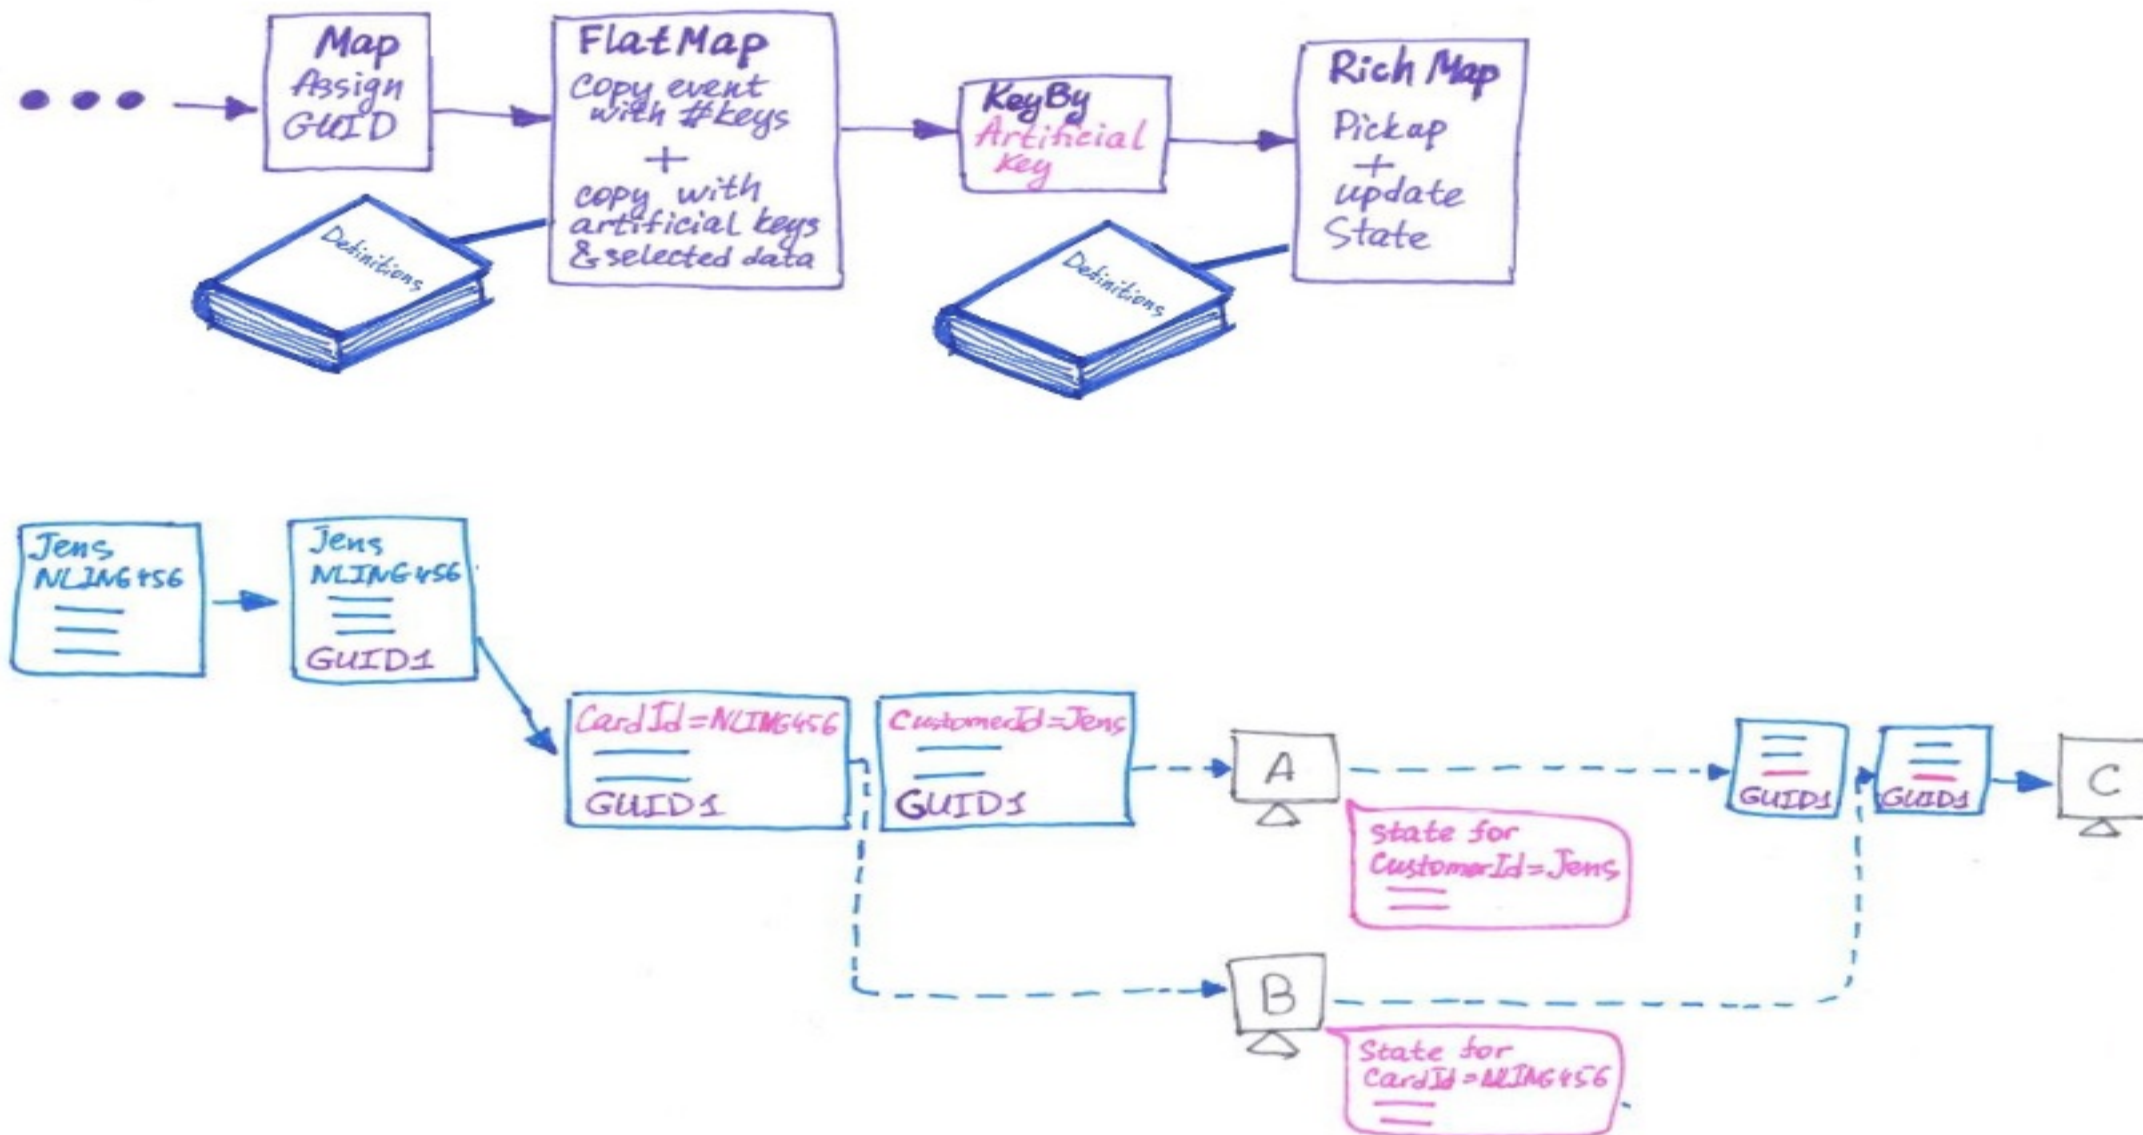
Model
definitions
(DSL)



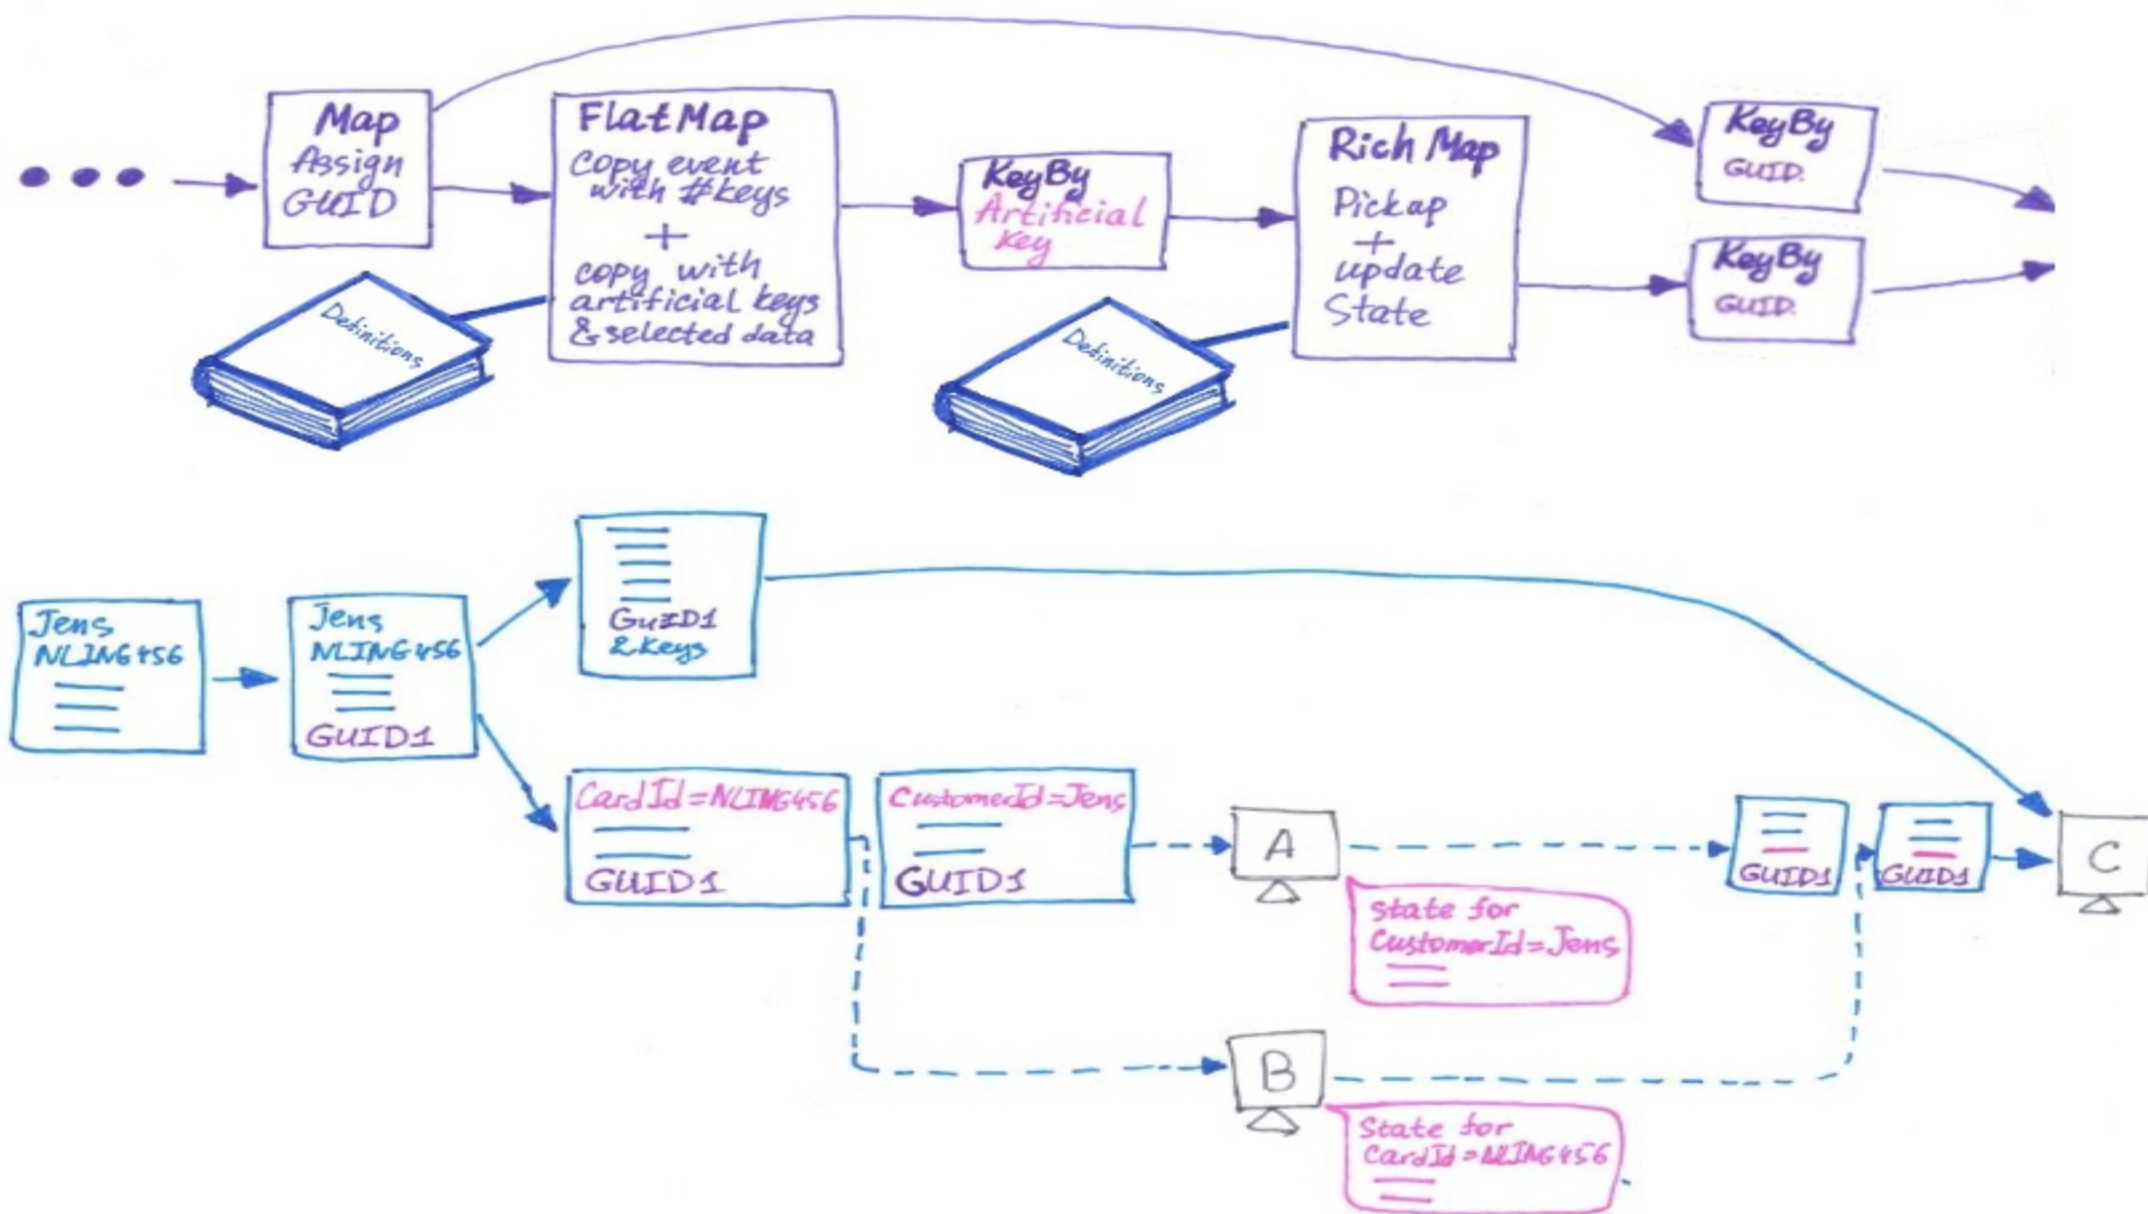
Advanced Approach



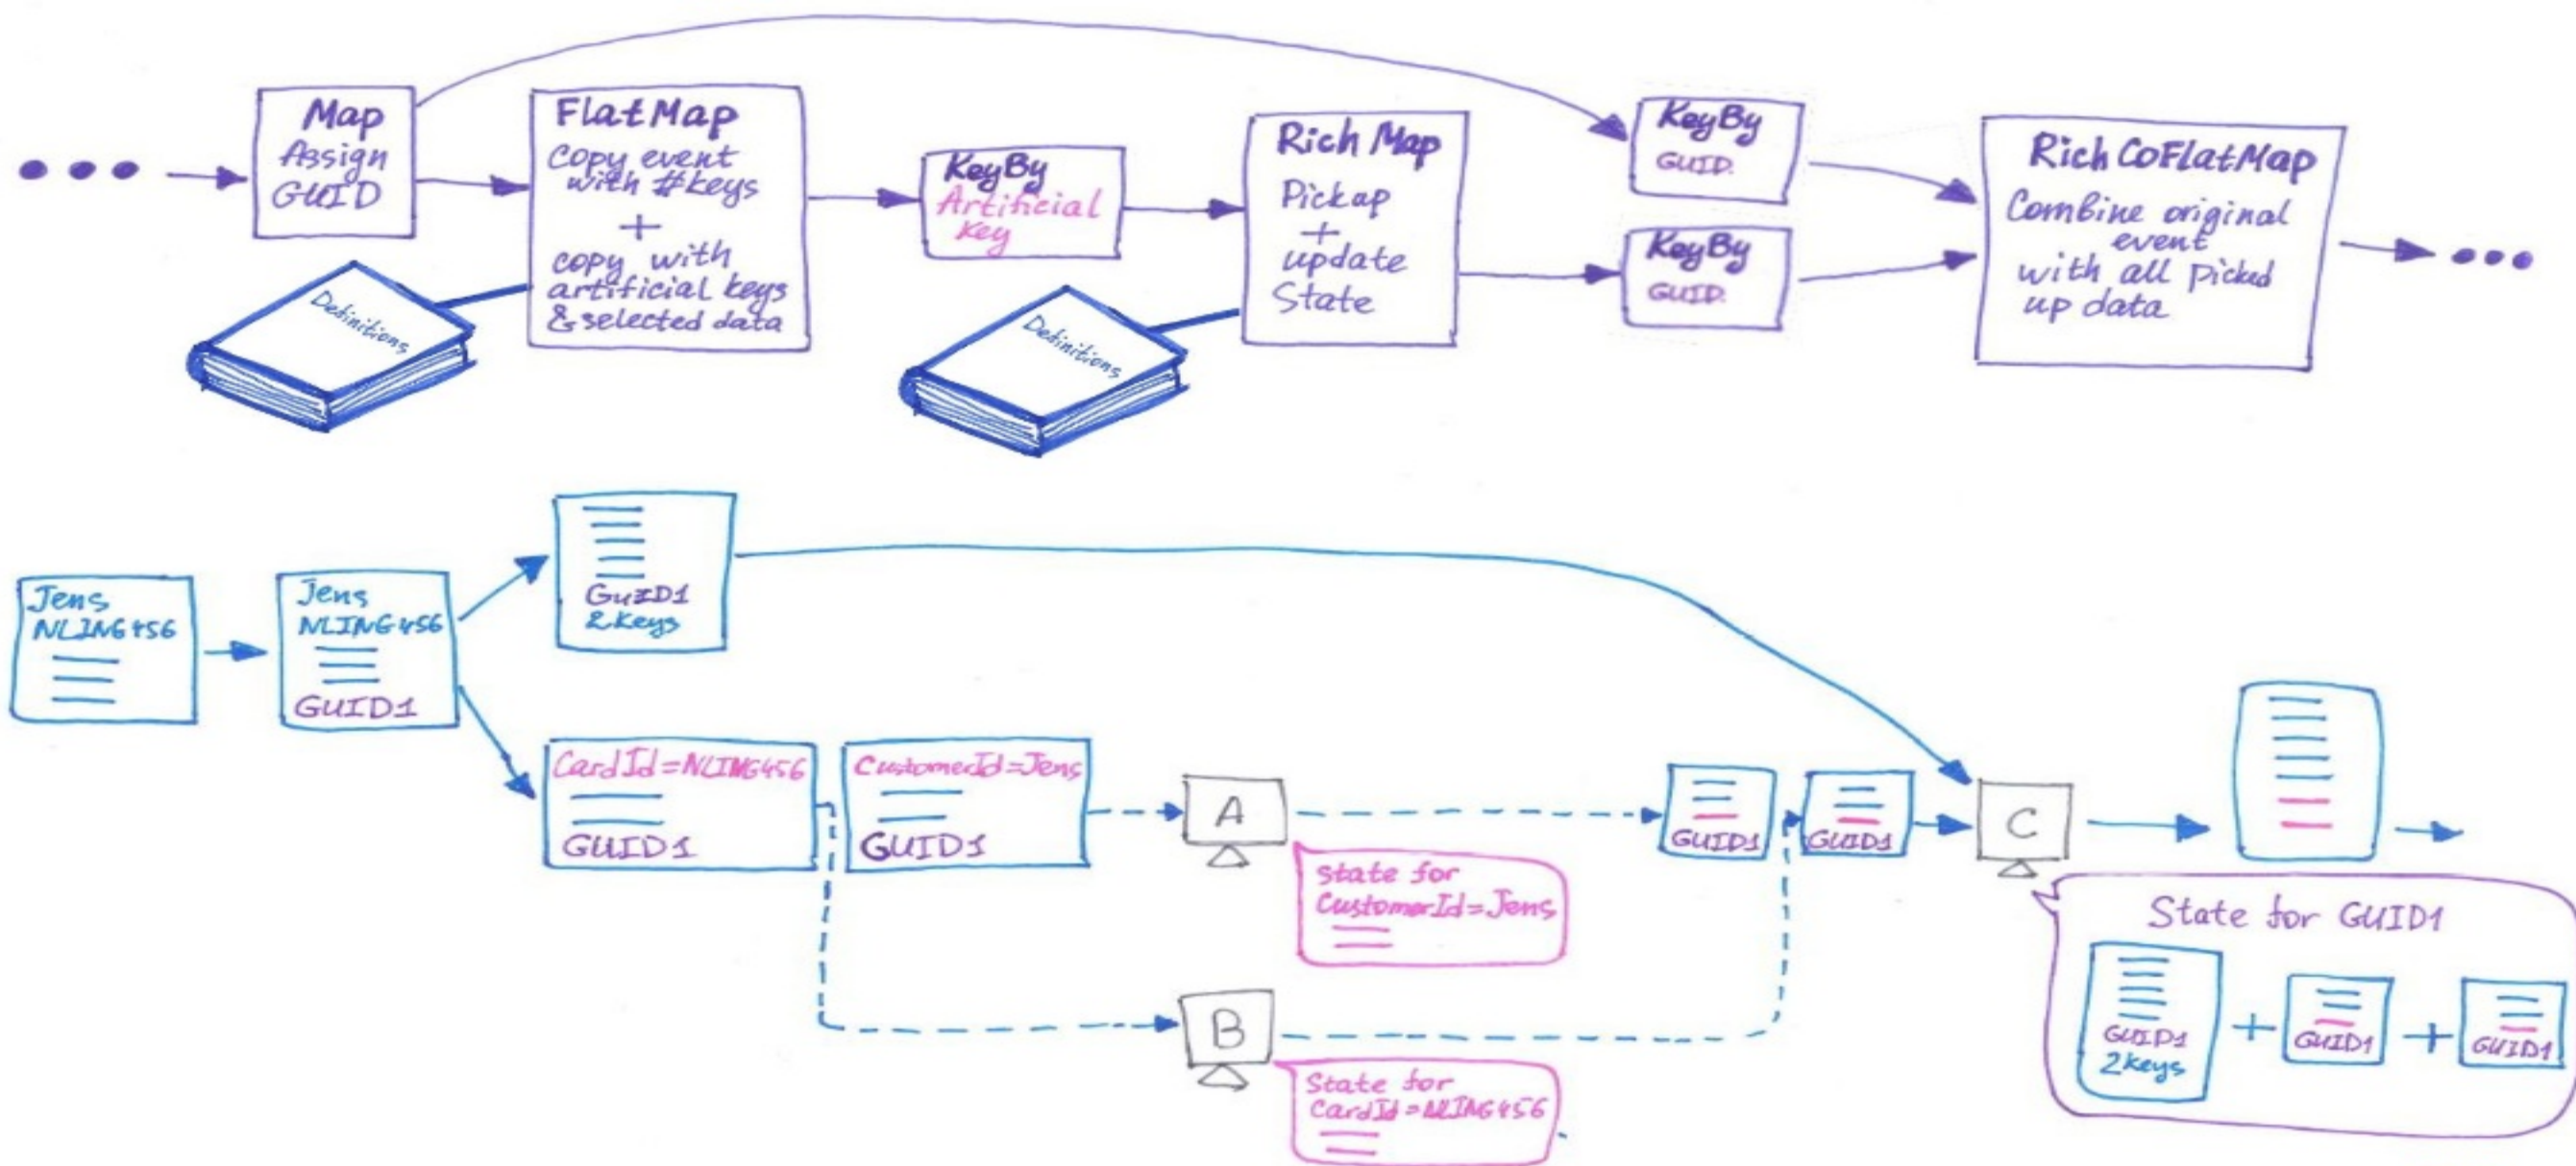
Advanced Approach



Advanced Approach



Advanced Approach



Flink - Advanced Approach

CODE

Advanced approach - evaluation

- + Can add new keys without restarting a job
- + Event is not passing through nodes that are not needed
- + Simple execution plan
- Quite complex implementation
- Requires a way to deliver a description of the behaviour in some way

Conclusion

Conclusion

Although the execution plan is fixed, your business logic does not have to be

The End

Thank you!



Check out the
demo code on GitHub:

github.com/rezolya/exploiting-flinks-stateful-operators