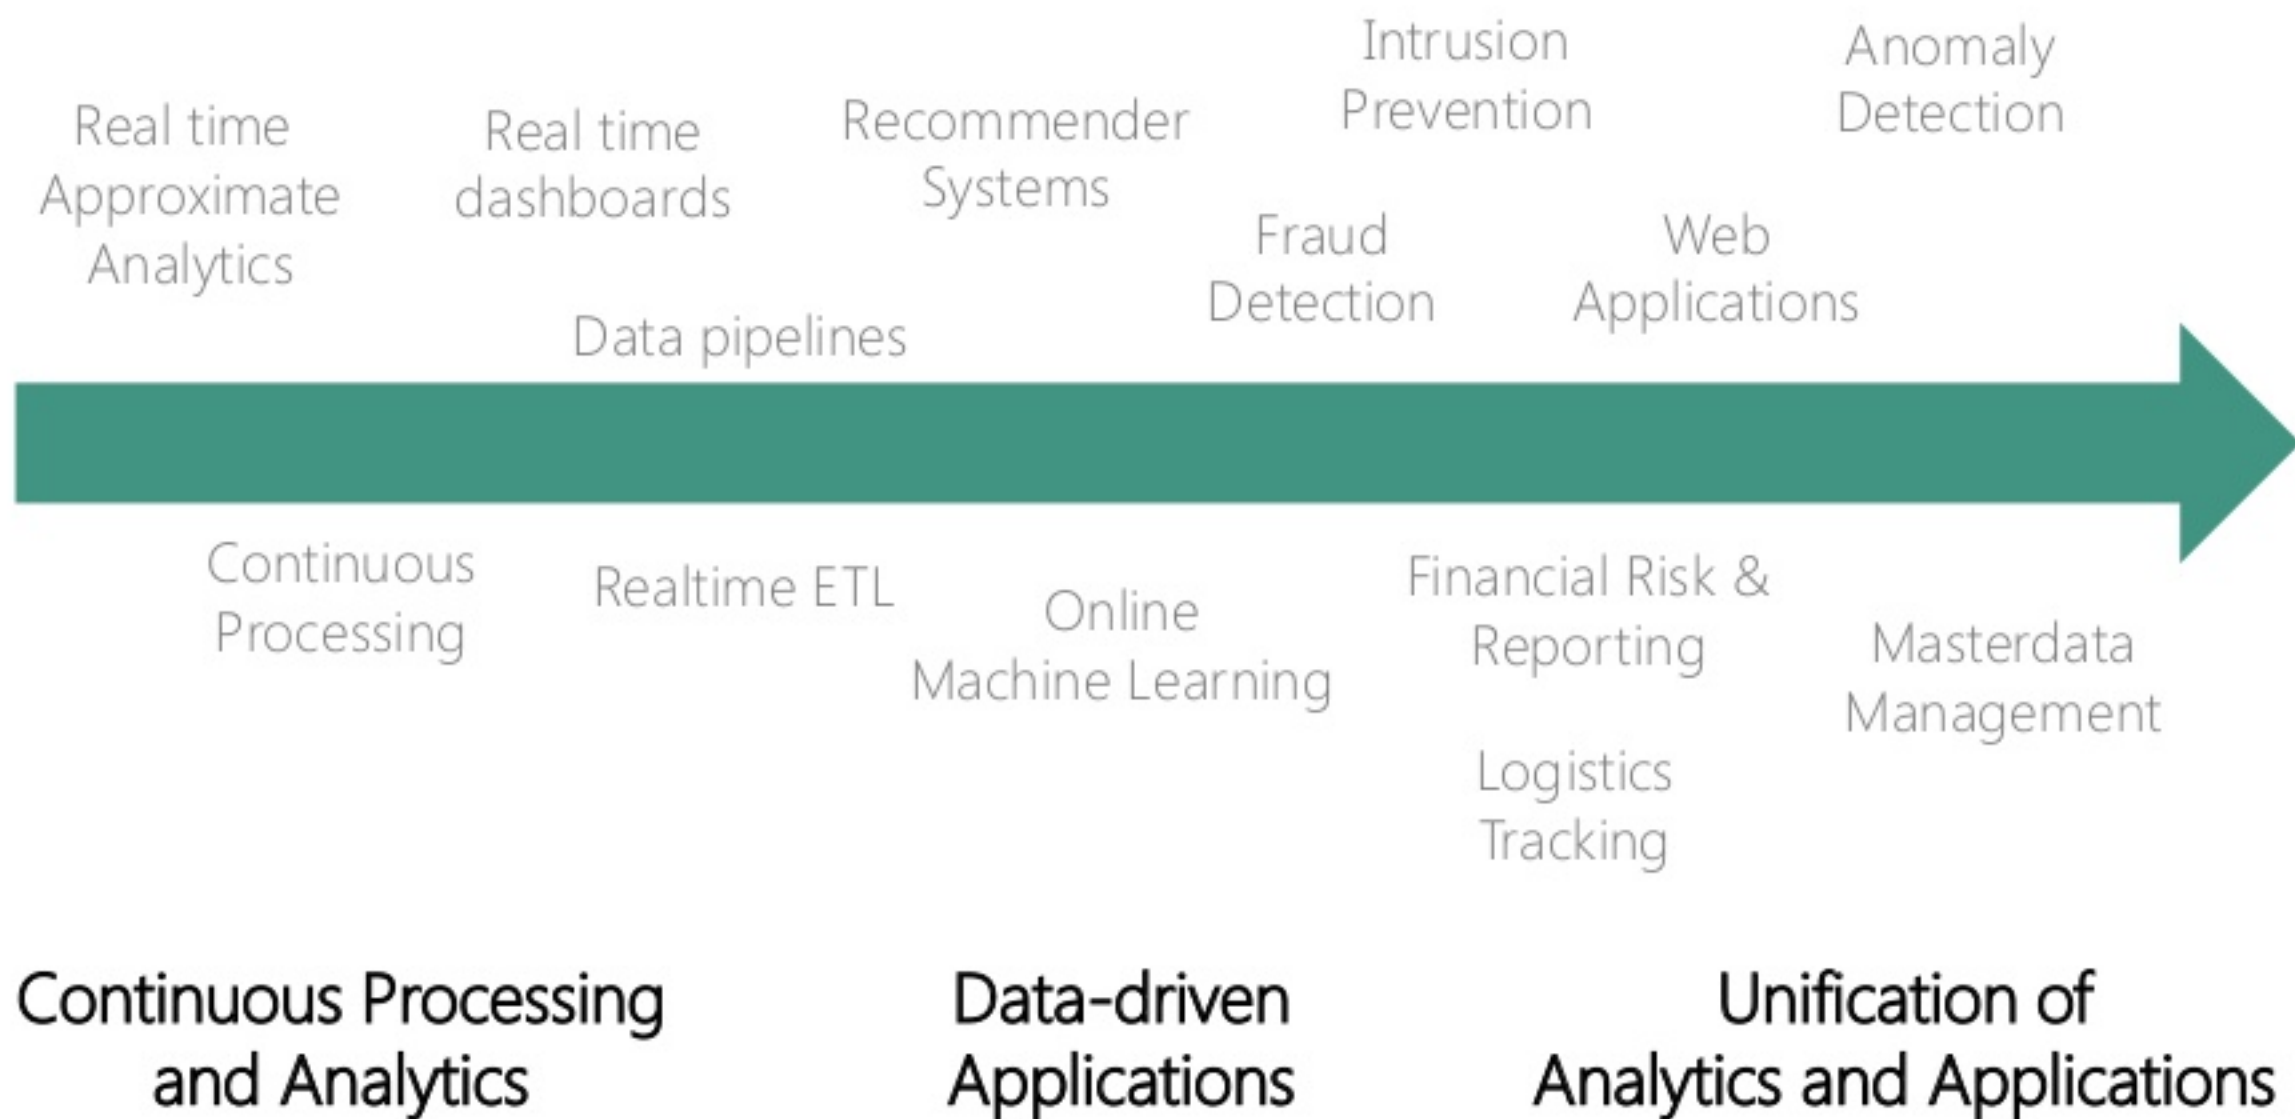

UNLOCKING THE NEXT WAVE OF STREAMING APPLICATIONS

Stephan Ewen – data Artisans CTO, Apache Flink PMC

dataArtisans

Streaming Applications over Time



Enablers of new Applications



Abstractions, APIs

Event-time, streaming SQL,
state & time, CEP



Consistency

Exactly-once, savepoints



Scalability

high parallelism, large state



Interoperability

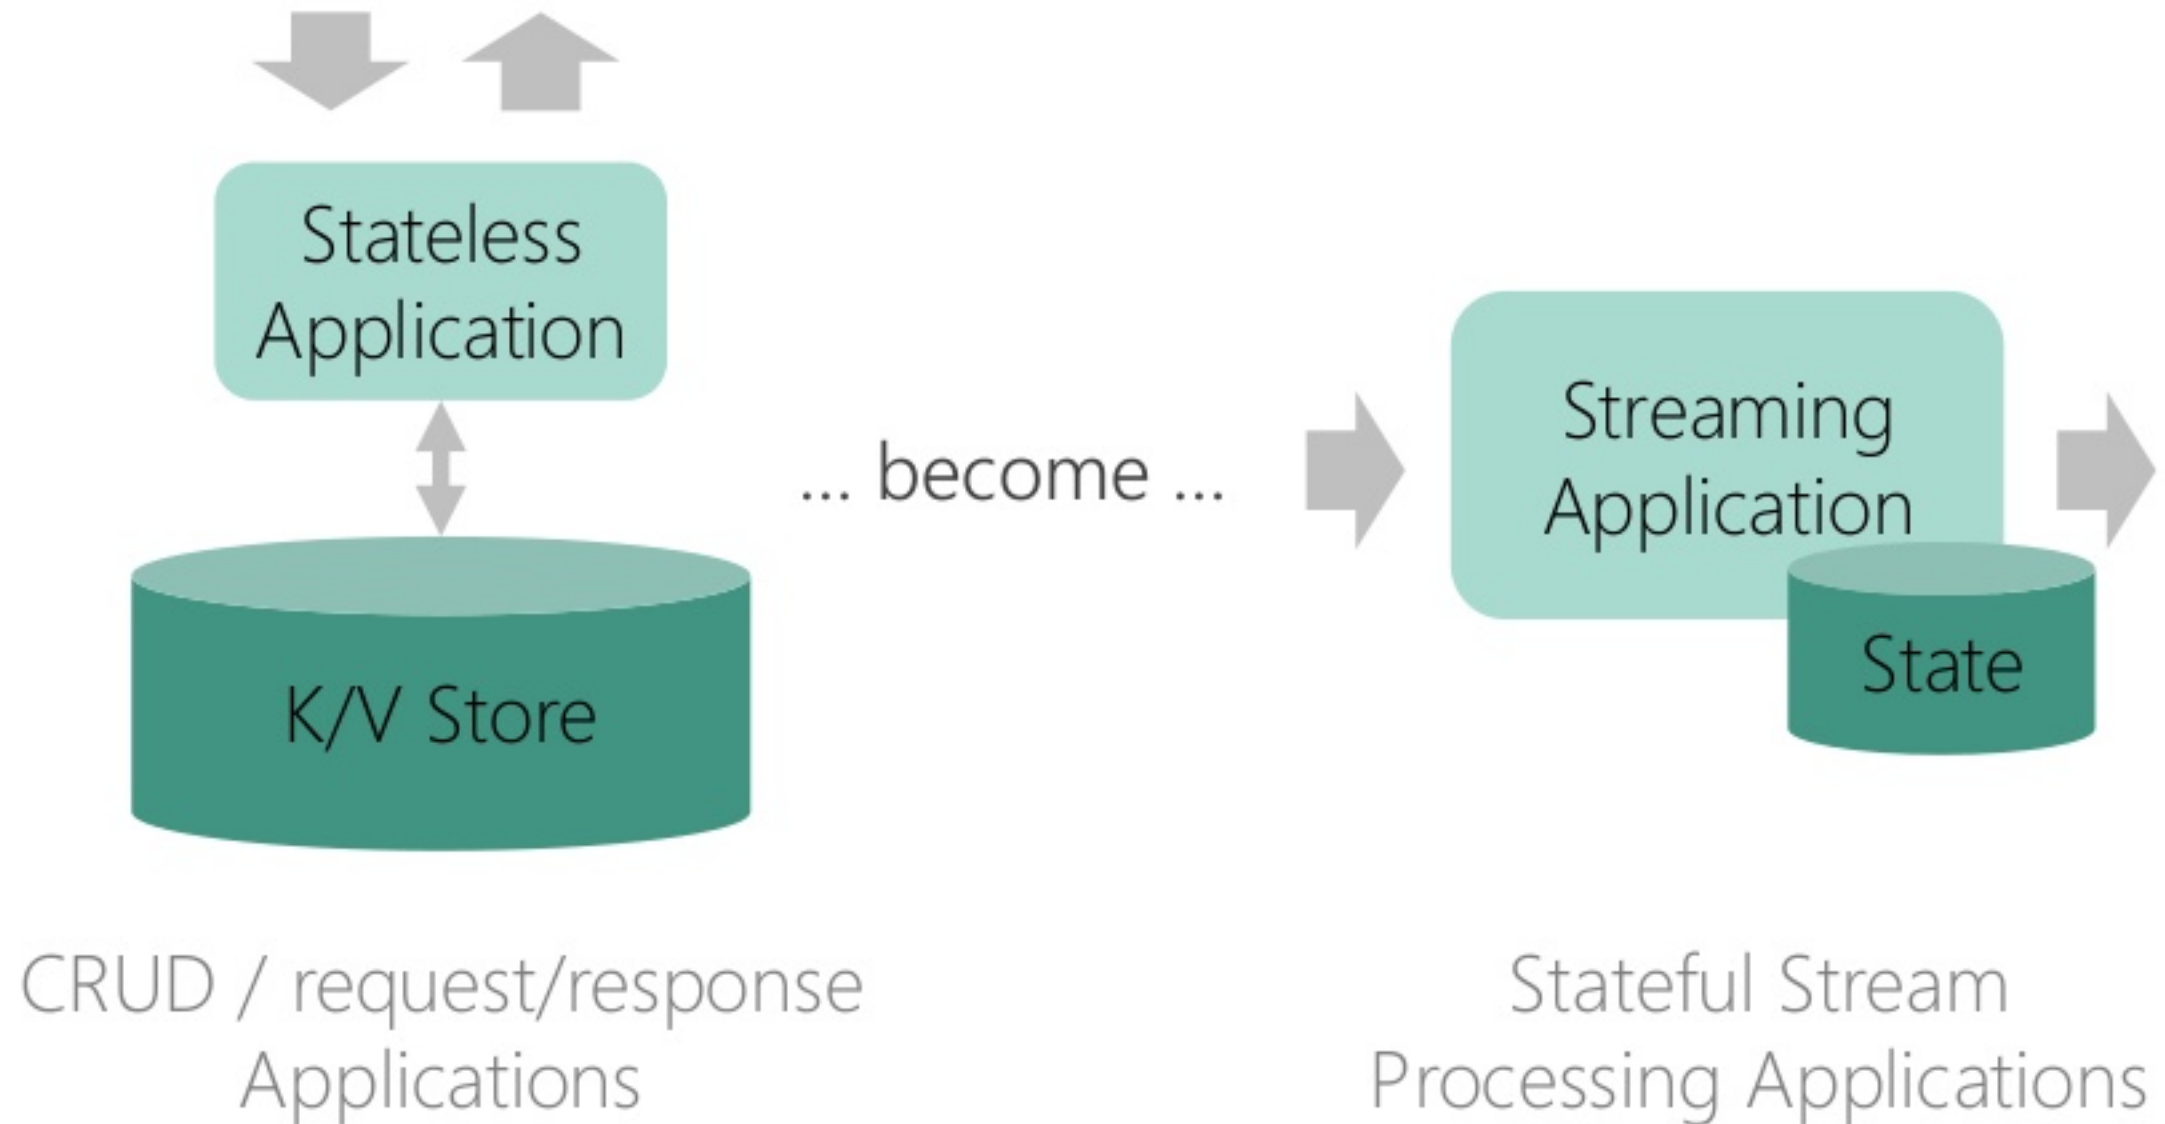
Deployment, Connectors, Operations




Enablers of new Applications



Exactly-once Changed Applications



Some Applications don't move to Stream Processing



Wouldn't it be great if Stream Processors could...

- ... access and update state with multiple keys at the same time
- ... maintain full isolation/correctness for the multi-key operations
- ... operate on multiple states at the same time
- ... share the states between multiple streams



TODAY WE ANNOUNCE

The first system for
serializable multi-row ACID transactions
on streaming data



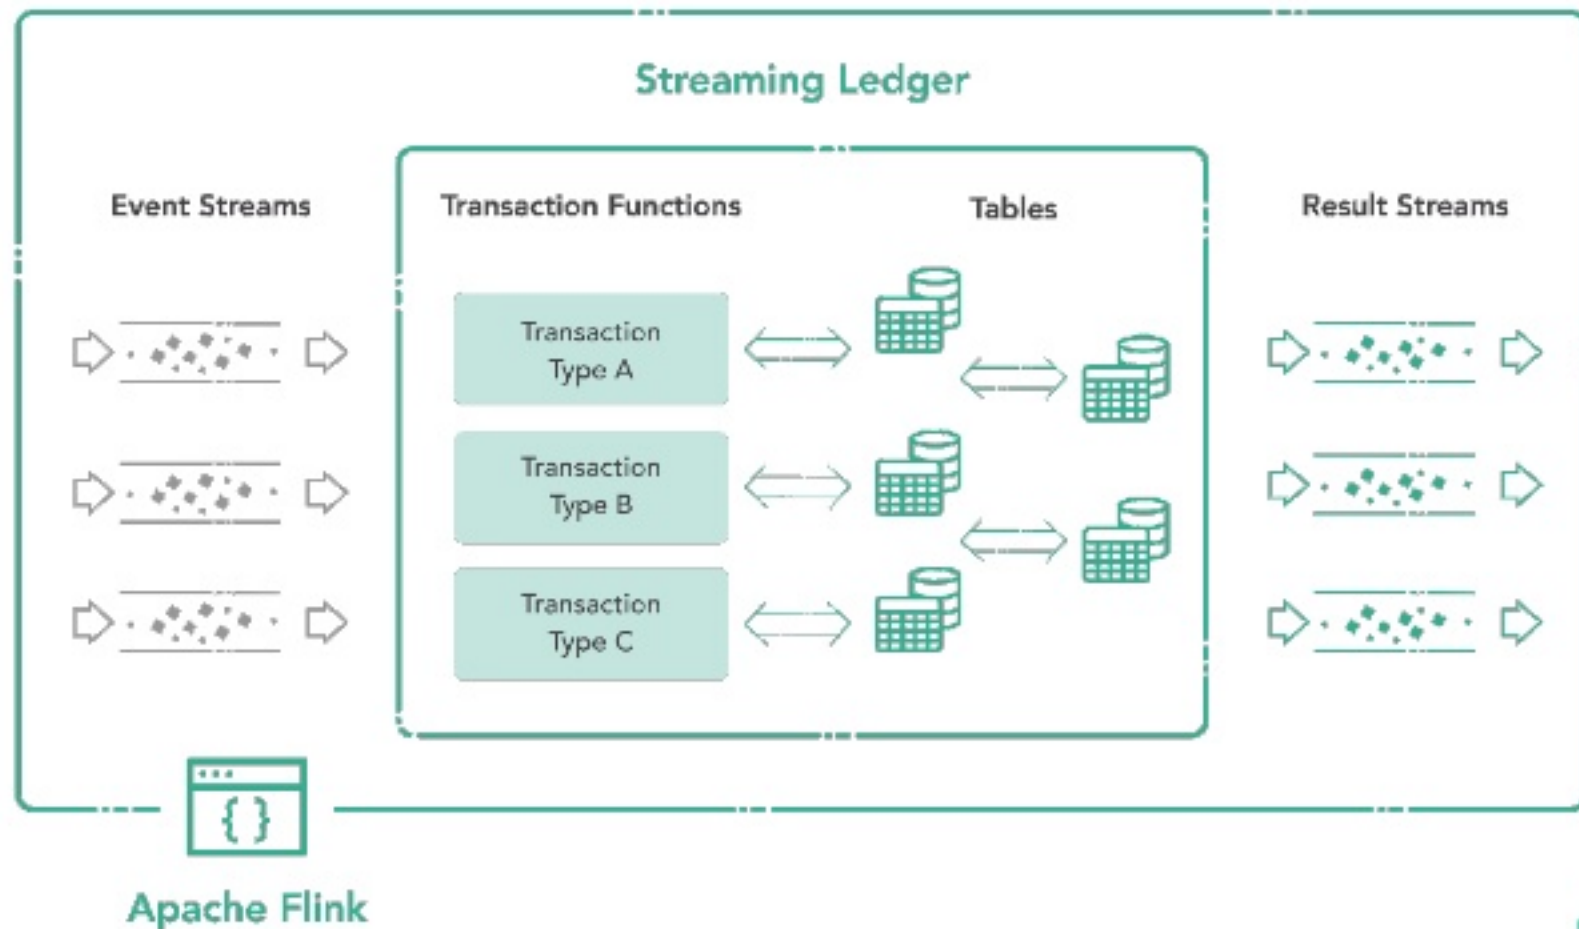
data Artisans Streaming Ledger

Serializable ACID Transactions
on Streaming Data

ACID Transactions for Multi-key Stream Processing

Streaming Ledger provides ACID guarantees across multiple states, rows, and streams

- **Atomicity:** the transfer affects either both accounts or none
- **Consistency:** the transfer must only happen if the account have sufficient funds
- **Isolation:** no other operation can interfere and cause an incorrect result
- **Durability:** the result of the transfer is durable



The Evolution of Stream Processing



ACID
guarantees

consistent general applications

Streaming Ledger provides ACID guarantees supporting applications that read and modify several keys

Exactly-once
guarantees

accurate single-key applications

Flink pioneered exactly once guarantees on true streaming: accurate analytics and applications, consistent single key at a time

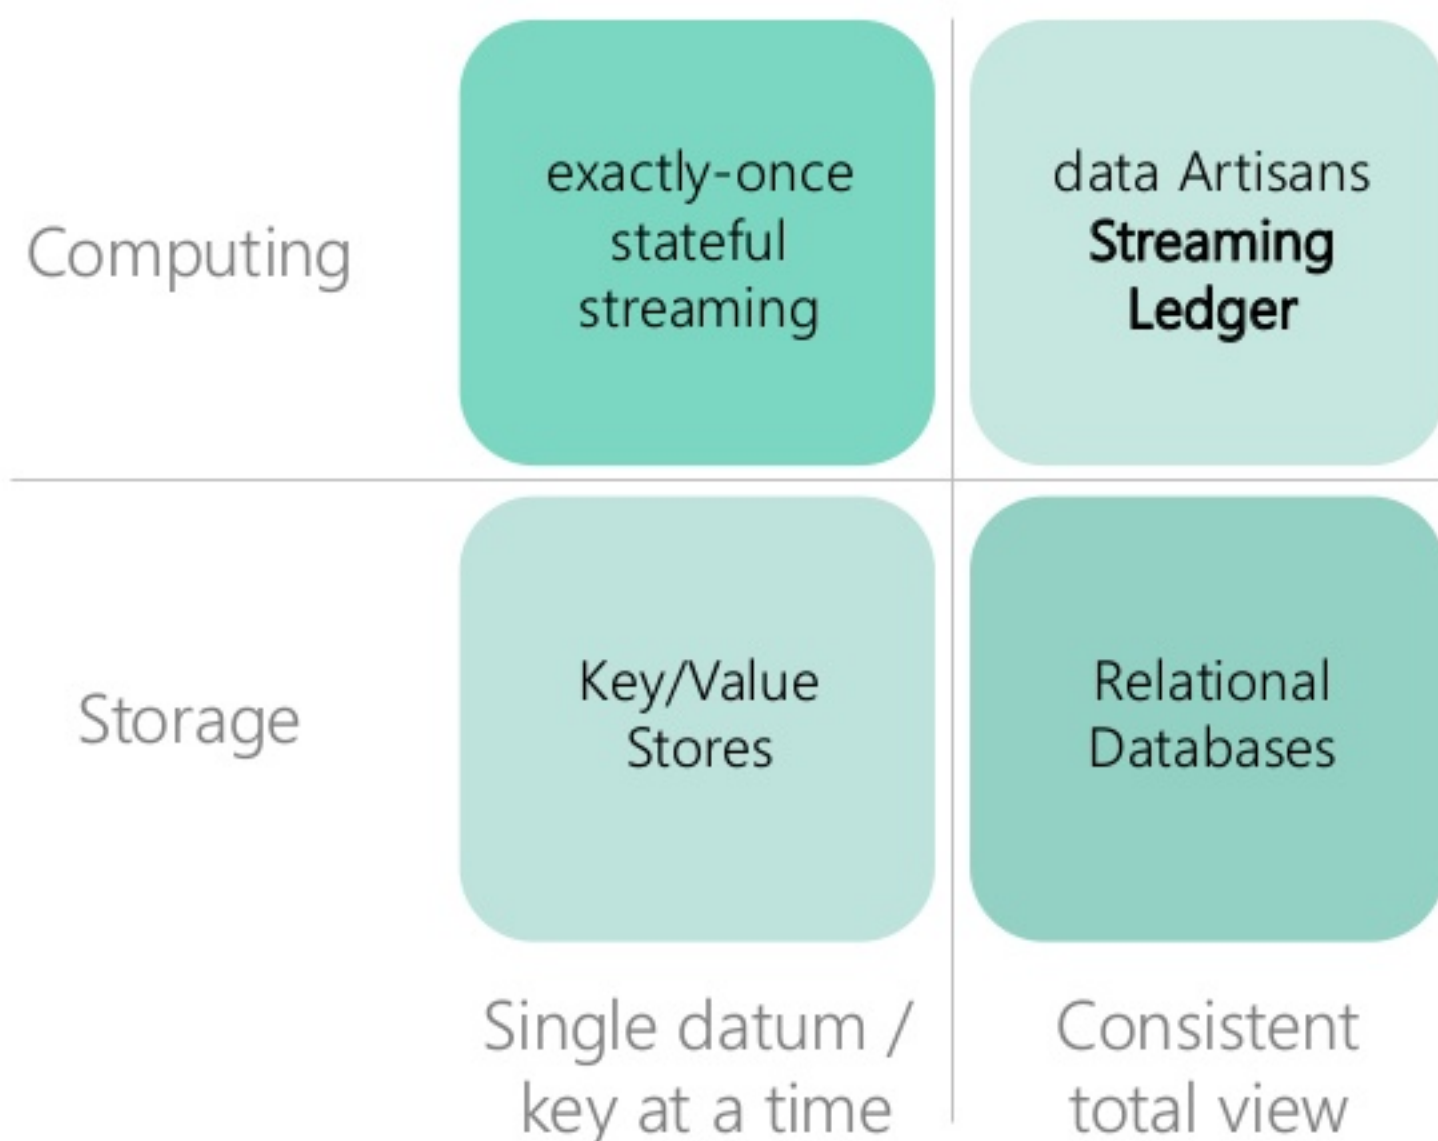
At-least-once
guarantees

approximate real time analytics

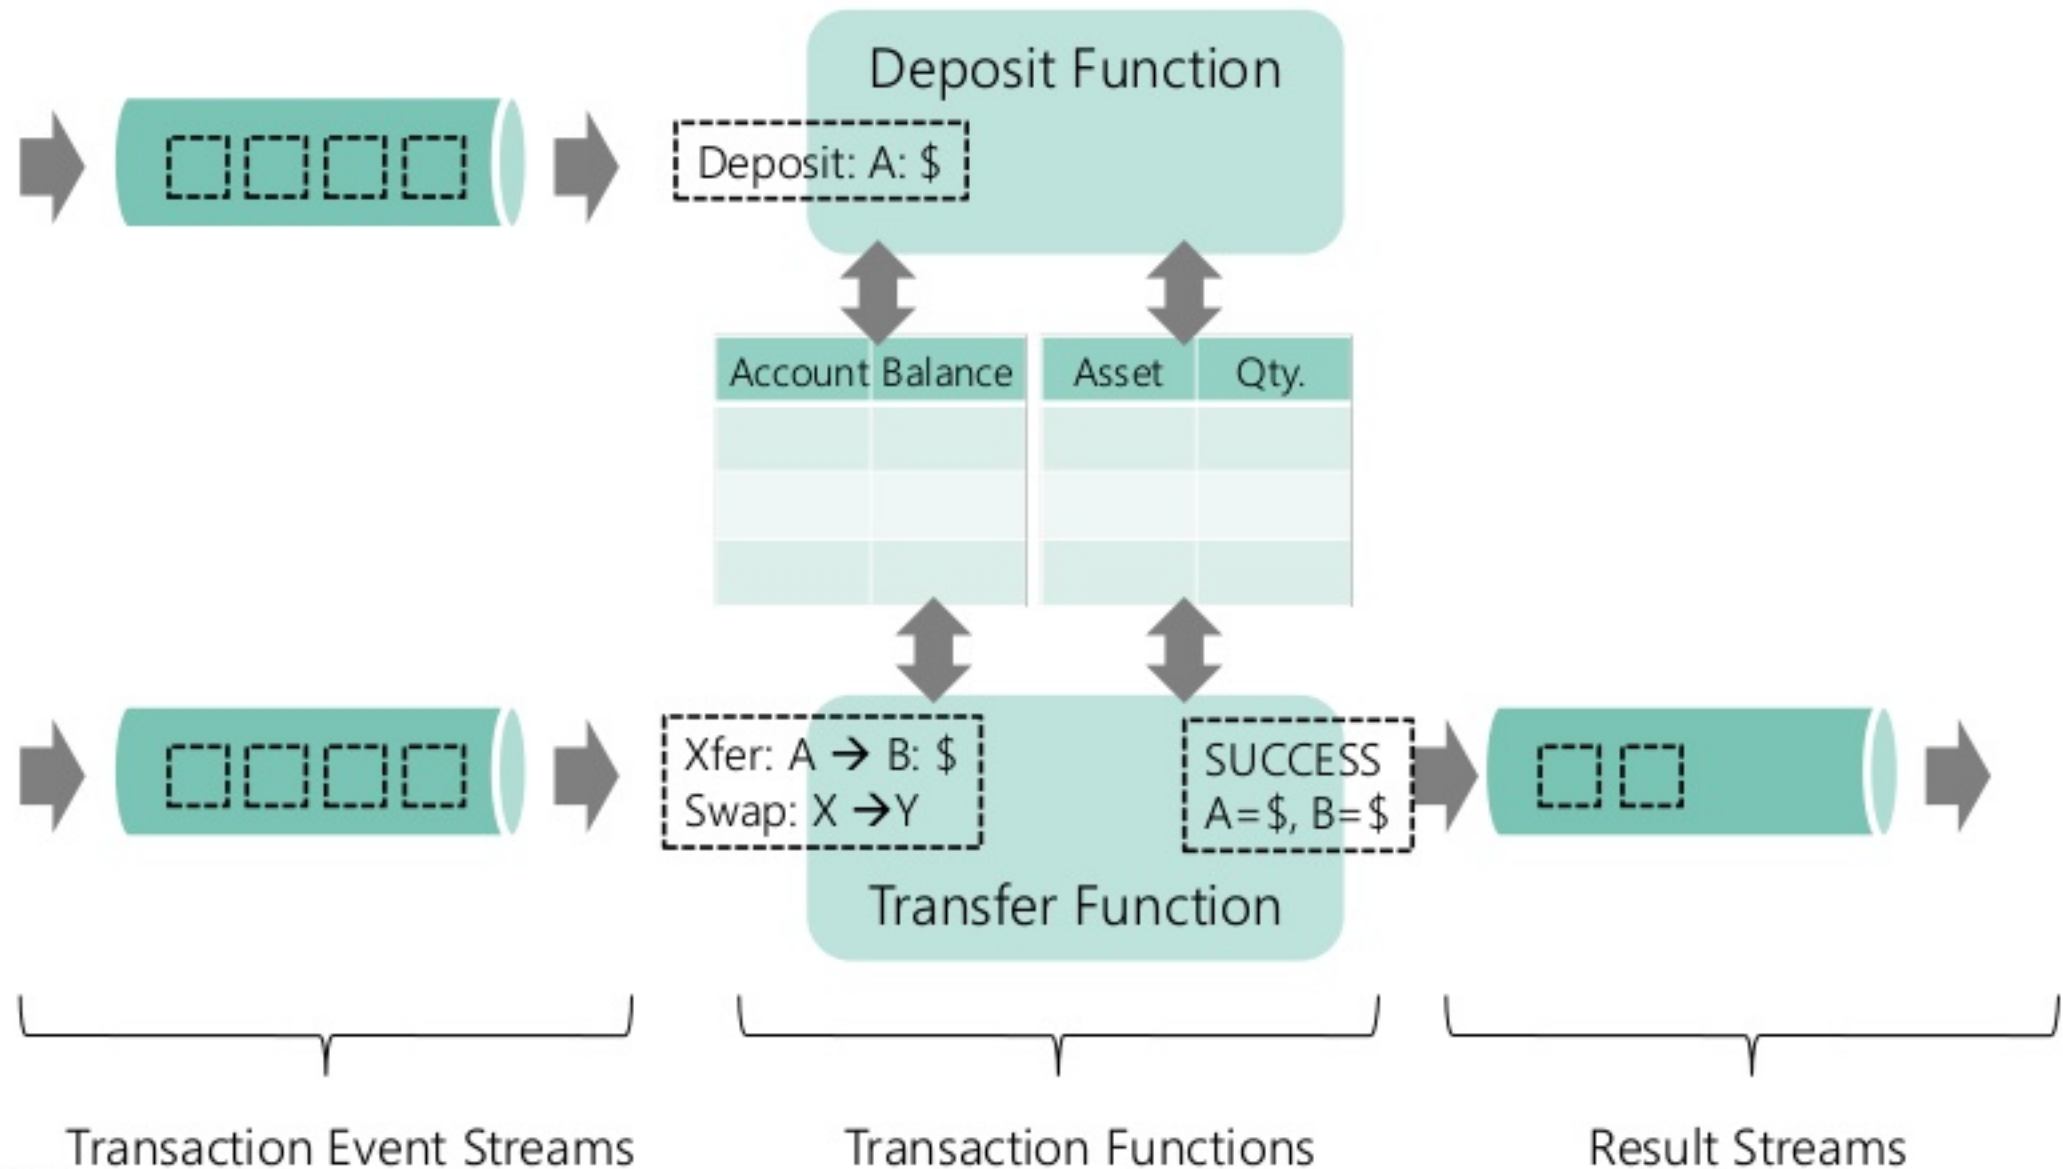
No data loss but possible duplications to support real-time approximate analytics in a "speed layer" (lambda)



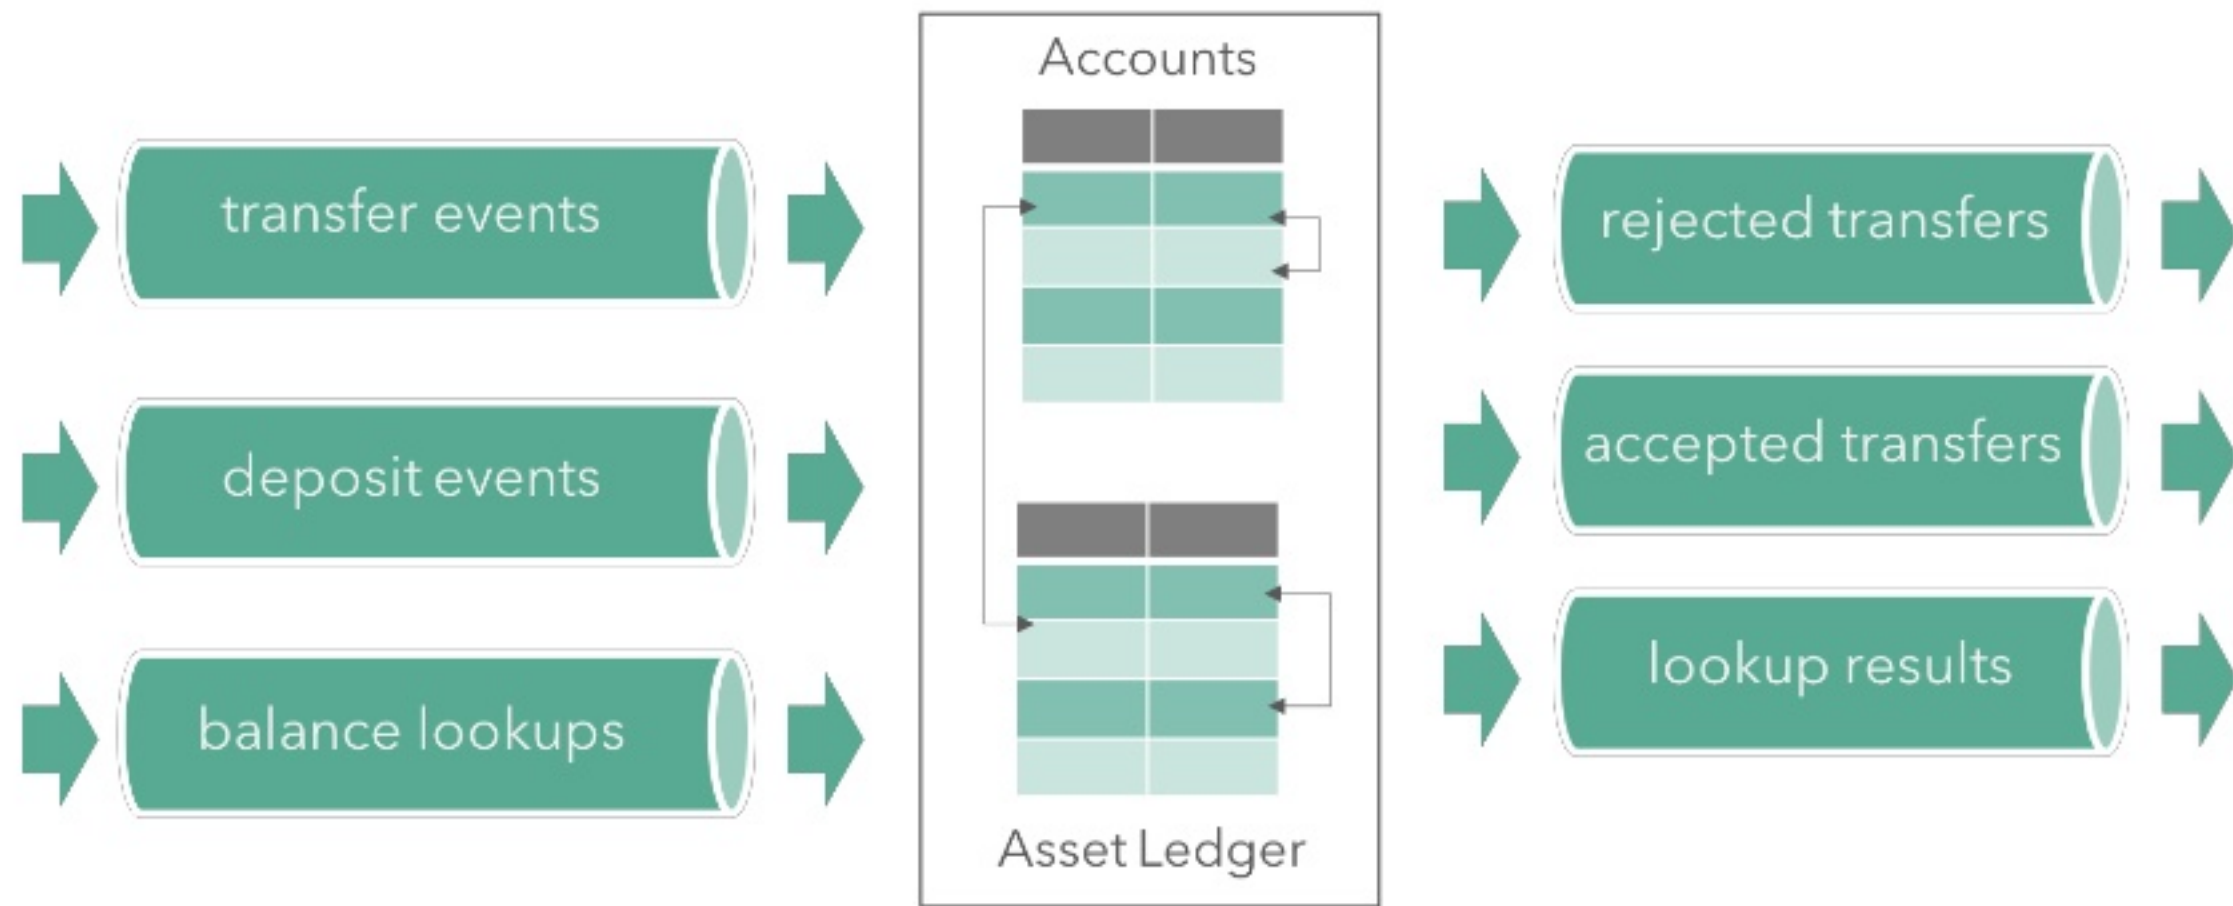
Stateful Streaming vs. Streaming Ledger



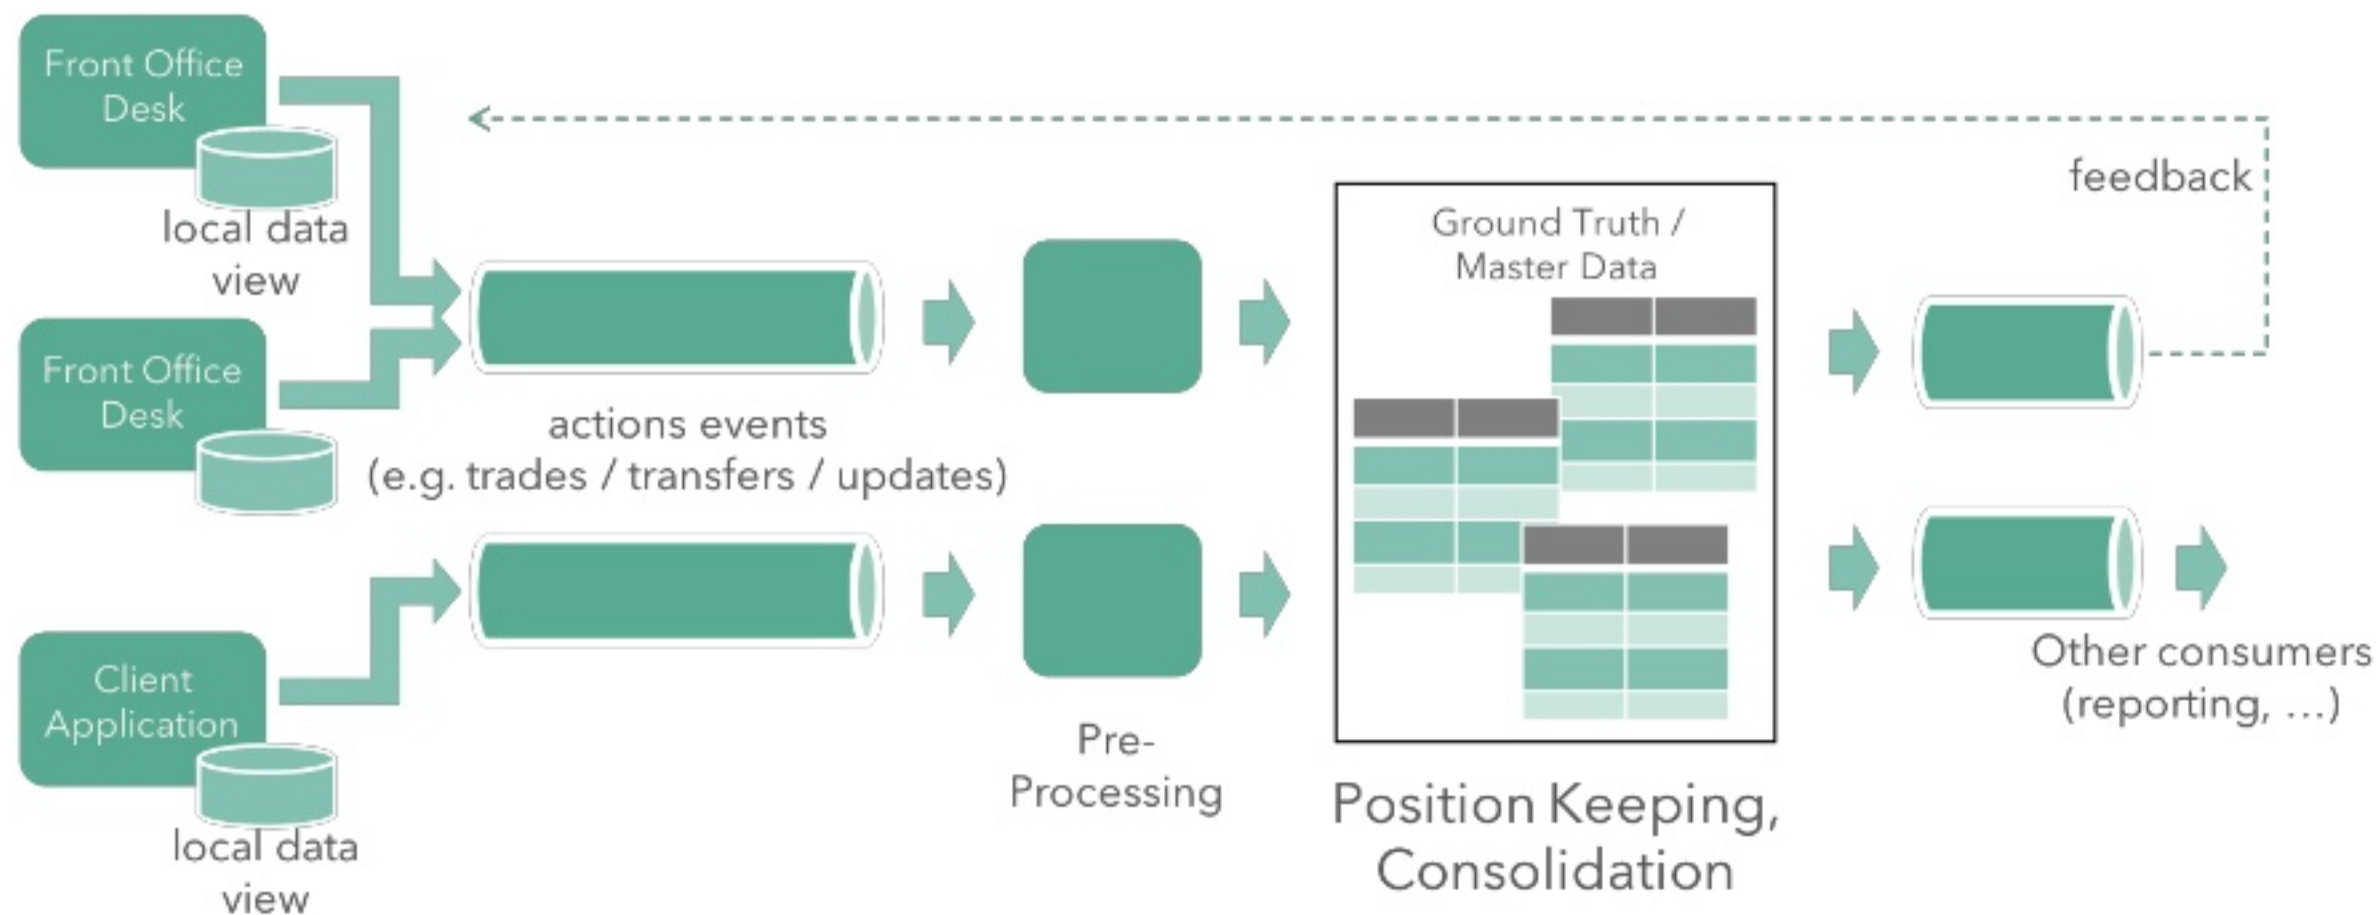
Tables, Streams, Transactions



Example: Transferring Cash/Assets between Accounts




Example: Position-keeping, Reporting, Risk Management in Investment Banking



A Library on top of Apache Flink

- No additional dependencies needed
- Seamlessly integrates and composes with DataStream API and SQL
- Read from- and write to all Flink connectors
- Supports savepoints for upgrades







How does it work?




Relational Database embedded? No!



First-of-a-kind - Unique Approach



Iterative streaming
dataflows



event re-ordering &
out-of-order processing



(0xf876ab78 | 0x0b7cc7a3)

Logical Clocks
to define schedule



Conflict-free Schedule


Txn: (A,B,C) \rightarrow C

Txn: (A,B,C) \rightarrow A

Txn: (D,E) \rightarrow (D,F)

transaction events

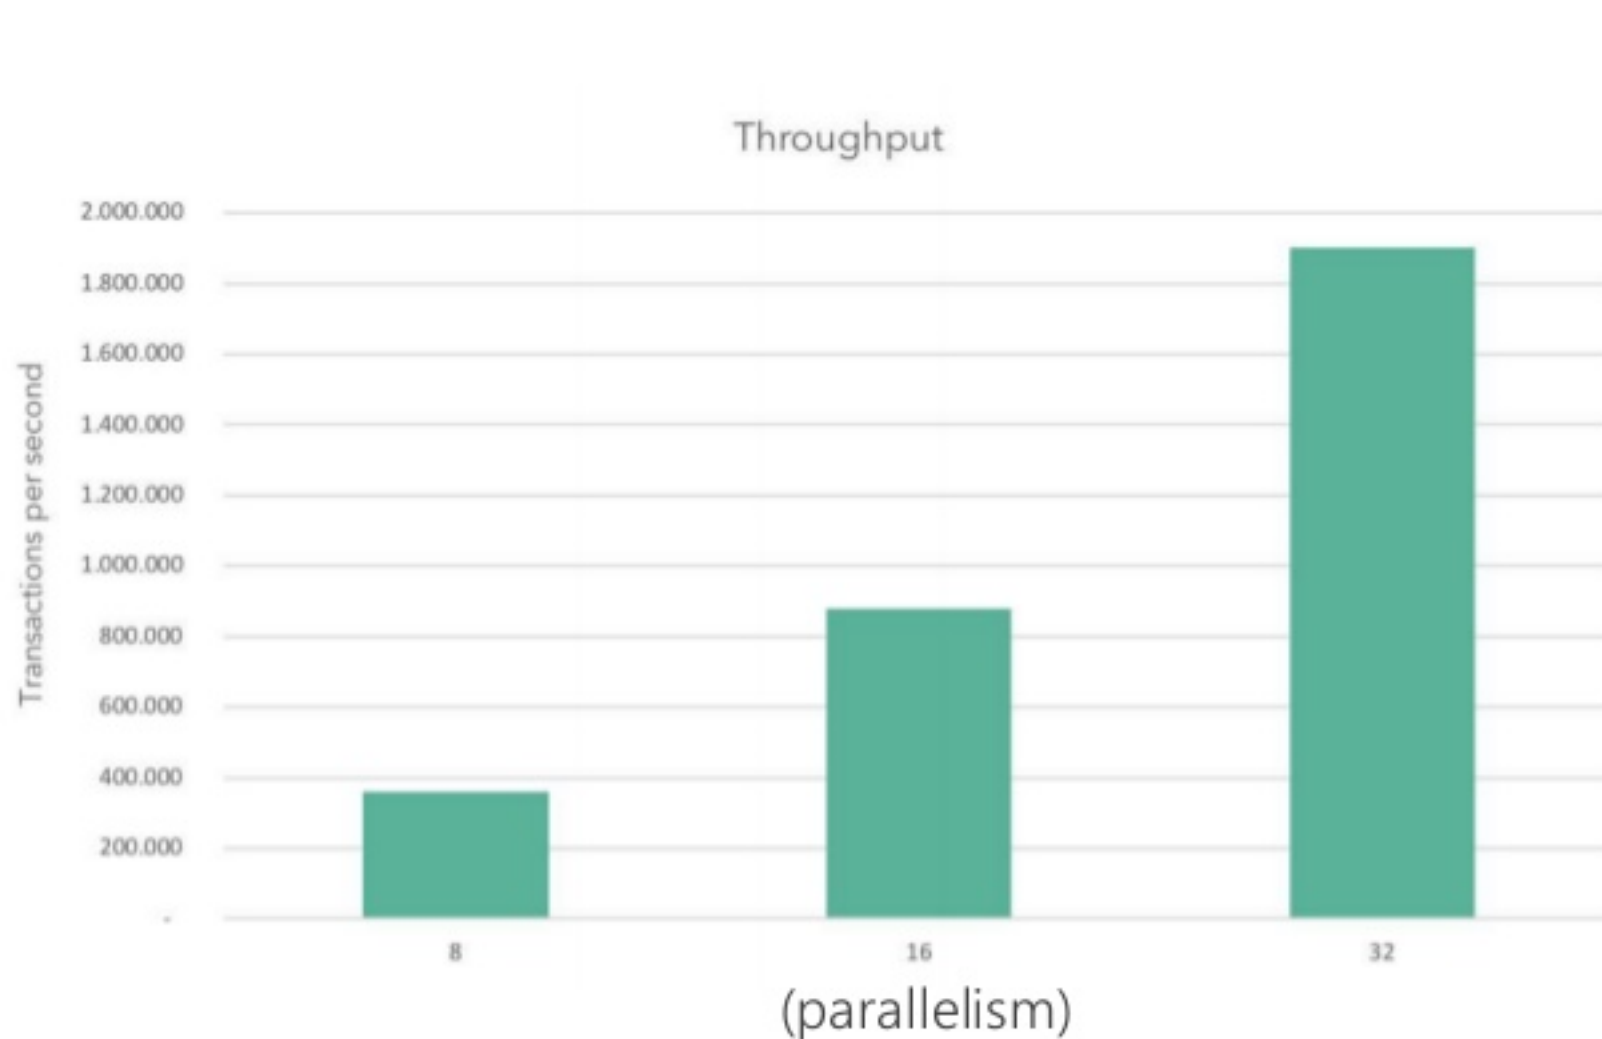
define ordering
for schedule




reorder events to obey schedule



Performance (early results) - Scalability



Performance (early results) – Key Contention

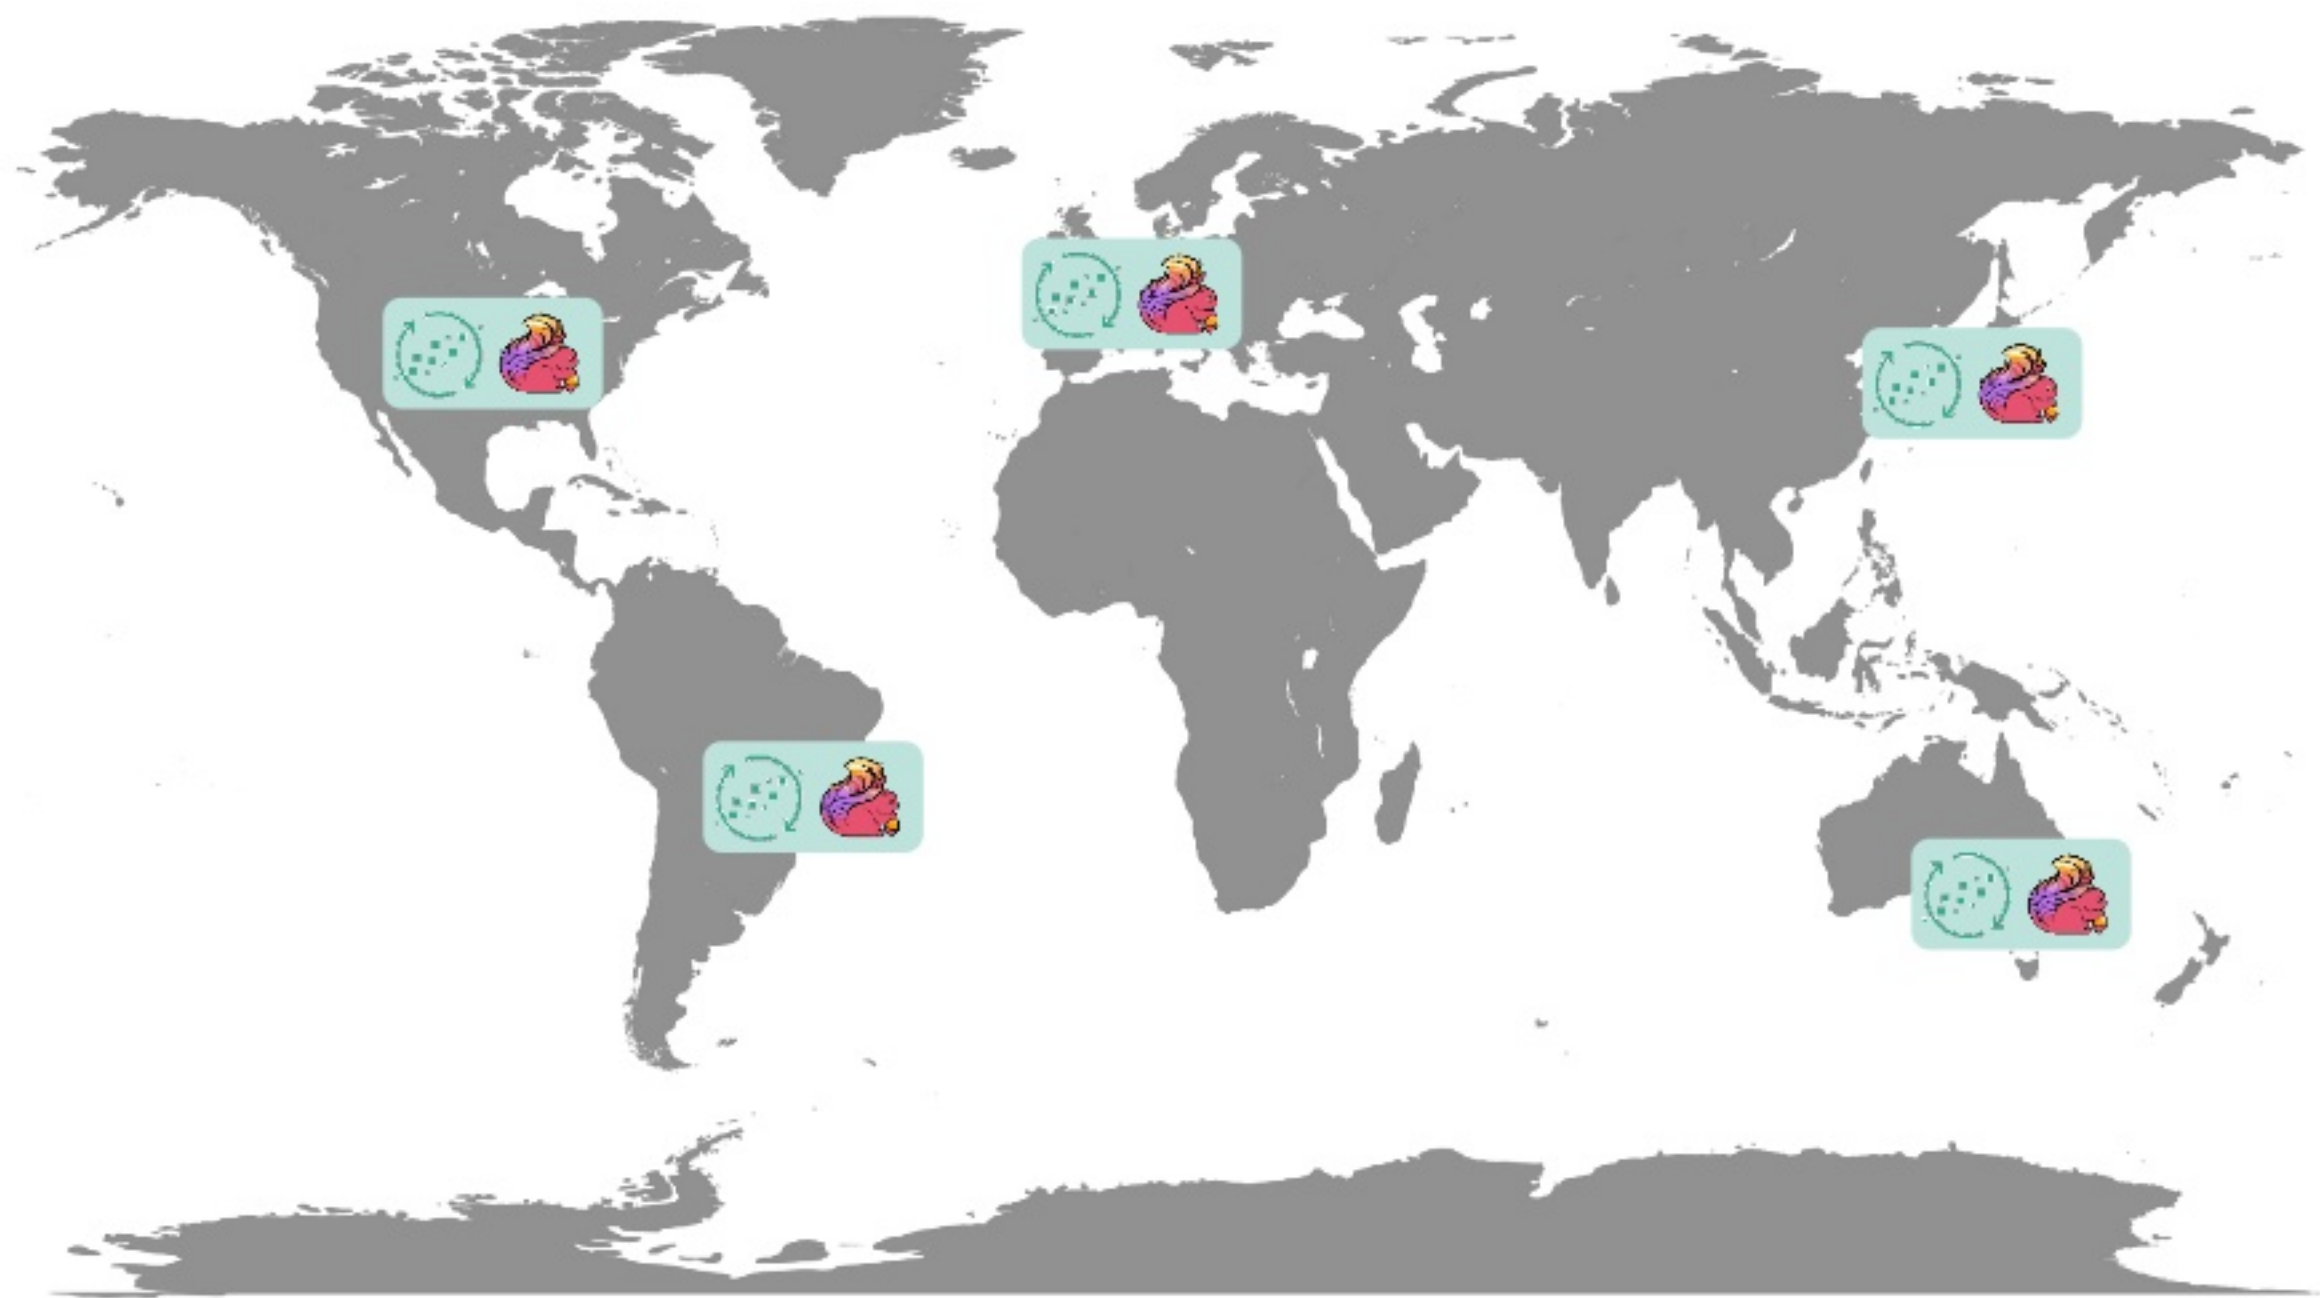


Running a distributed setup

... and because we can...

...we run a globally
distributed setup





Apache Flink: The heavy lifter

This technology is possible, because Apache Flink offers such powerful building blocks

- Continuous processing
- Iterative flows
- Flexible state abstraction
- Asynchronous checkpoints
- Sophisticated event-time/watermarks

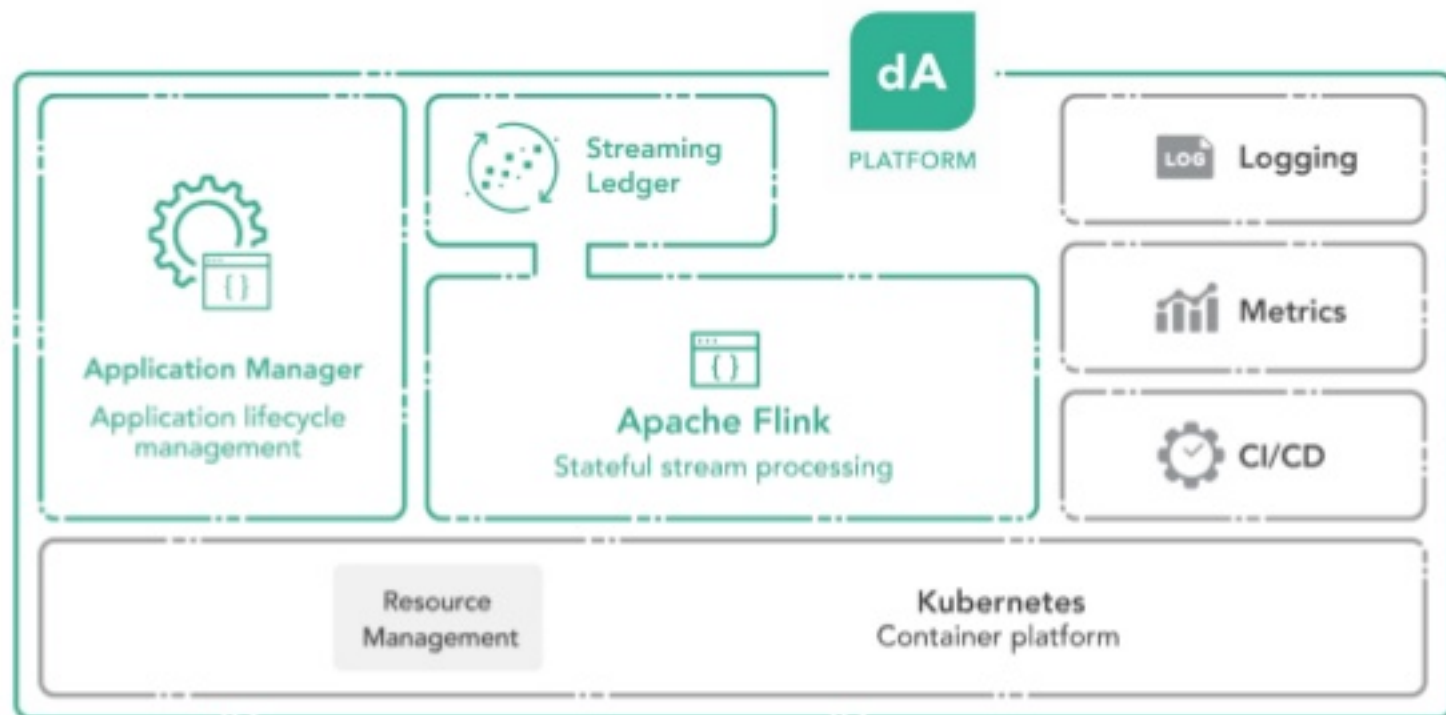


Part of data Artisans Platform River Edition

Streaming Ledger is part of **data Artisans Platform**

- **Stream Edition:** Apache Flink, Application Manager
- **River Edition:** Apache Flink, Application Manager, Streaming Ledger

API and single-node implementation under Apache 2.0 license




Learn more at

**data Artisans Product
Announcement**

Igal Shilman, data
Artisans

4:10pm - 4:30pm

 Maschinenhaus



At the booth

dataArtisans





Thank you!
Enjoy the conference!

dataArtisans