

# Python Streaming Pipelines with Beam on Flink

Flink Forward Berlin, 2018



Apache Beam



Apache Flink

<https://s.apache.org/streaming-python-beam-flink>

**dataArtisans** 

Aljoscha Krettek, Thomas Weise

# Agenda

1. What is Beam?
2. The Beam Portability APIs
3. Executing Pythonic Beam Jobs on Flink
4. The Future

# Problem

- Many of the big data ecosystem projects are Java / JVM based
- Use cases with different language environments
  - Python is the primary option for Machine Learning
- Barrier to entry for teams that want to adopt streaming but have no Java experience
- Cost of too many API styles and runtime environments
- (Currently no good option for native Python + Streaming)


# Multi-Language Support in Beam

- Effort to support multiple languages in Beam started late 2016
- Python SDK on Dataflow available for ~ 1 year
- Go SDK added recently
- At Flink Forward 2017...


**2018: Portable Flink Runner  
MVP near completion  
(~ Beam release 2.8.0)**

## Talk Python to Me

Stream Processing in Your Favourite Language with Beam on Flink



Apache Beam



Apache Flink

Based on work and slides by Frances Perry, Tyler Akidau, Kenneth Knowles & Sourabh Bajaj

Slides by Aljoscha Krettek, September 2017, Flink Forward 2017

What is Beam?

# What is Apache Beam?

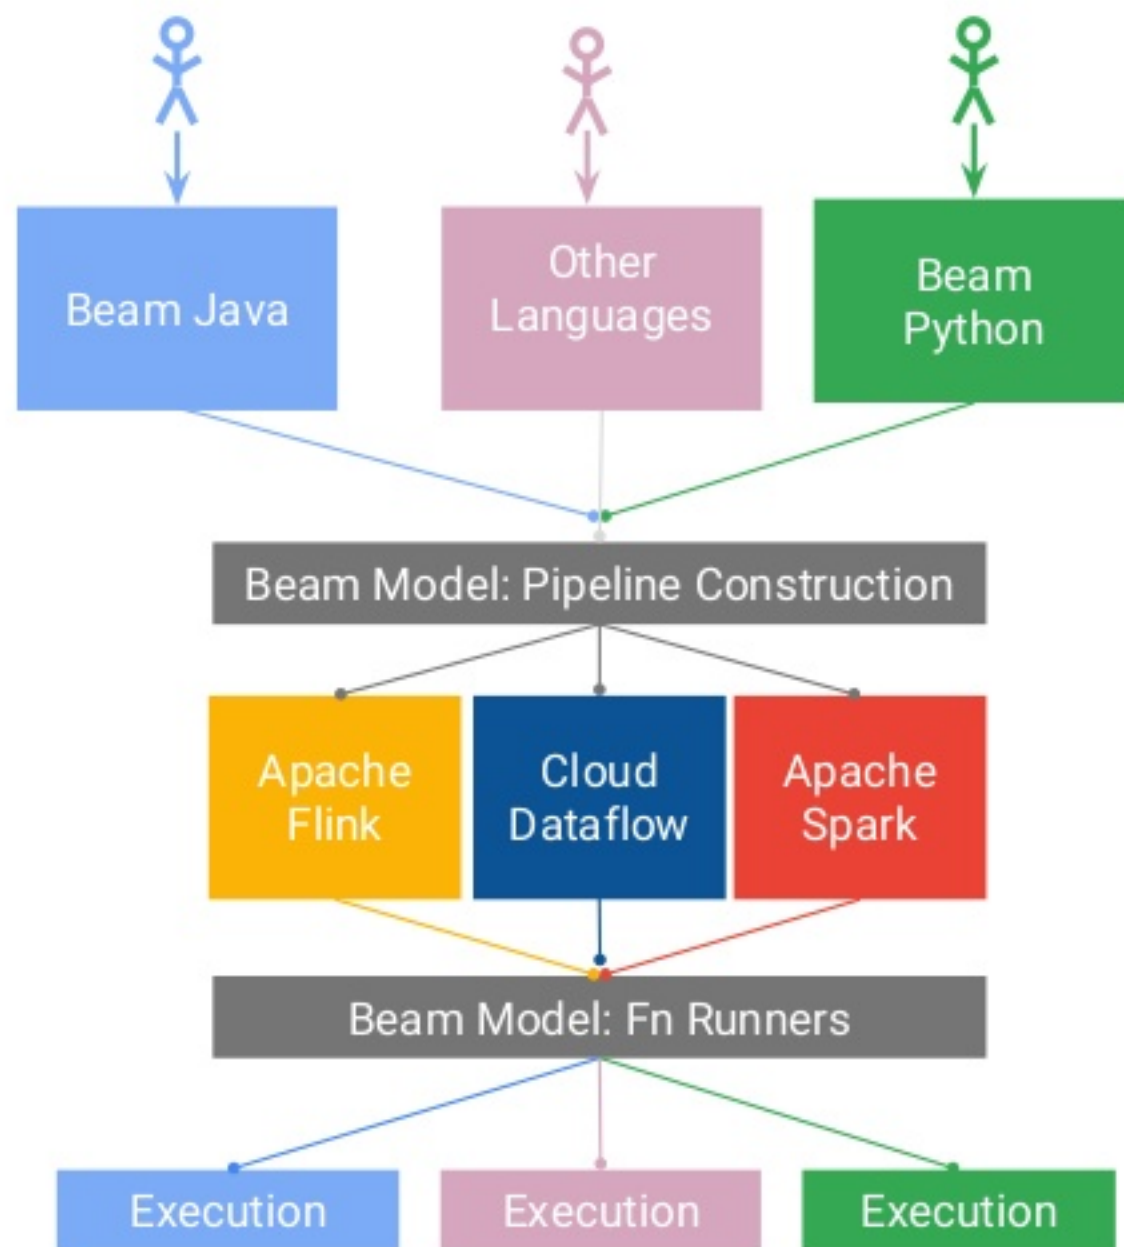
**Apache Beam** is a **unified** programming model designed to provide **efficient** and **portable** data processing pipelines

1. Unified model (**B**atch + str**EAM**)  
**What** / **Where** / **When** / **How**
2. **SDKs** (Java, Python, Go, ...) & **DSLs** (Scala, ...)
3. **Runners** for Existing Distributed Processing Backends (Google Dataflow, Spark, Flink, ...)
4. **IOs**: Data store Sources / Sinks





# The Apache Beam Vision

1. **End users:** who want to write pipelines in a language that's familiar.
2. **SDK writers:** who want to make Beam concepts available in new languages.
3. **Runner writers:** who have a distributed processing environment and want to support Beam pipelines






# The Beam Model



(Flink draws it more like this)



# The Beam Model



# Beam Model: Asking the Right Questions

**What** results are calculated?

**Where** in event time are results calculated?

**When** in processing time are results materialized?


**How** do refinements of results relate?

# The Beam Model: **What** is Being Computed?

```
PCollection<KV<String, Integer>> scores = input  
    .apply(Sum.integersPerKey());
```

```
scores = (input  
    | Sum.integersPerKey())
```

# The Beam Model: **What** is Being Computed?




# The Beam Model: **Where** in Event Time?

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Duration.standardMinutes(2))))
    .apply(Sum.integersPerKey());
```

```
scores= (input
| beam.WindowInto(FixedWindows(2 * 60))
| Sum.integersPerKey())
```

# The Beam Model: **Where** in Event Time?




# The Beam Model: **When** in Processing Time?

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Duration.standardMinutes(2))
        .triggering(AtWatermark())))
    .apply(Sum.integersPerKey());
```

```
scores = (input
    | beam.WindowInto(FixedWindows(2 * 60)
        .triggering(AtWatermark()))
    | Sum.integersPerKey())
```



# The Beam Model: **When** in Processing Time?

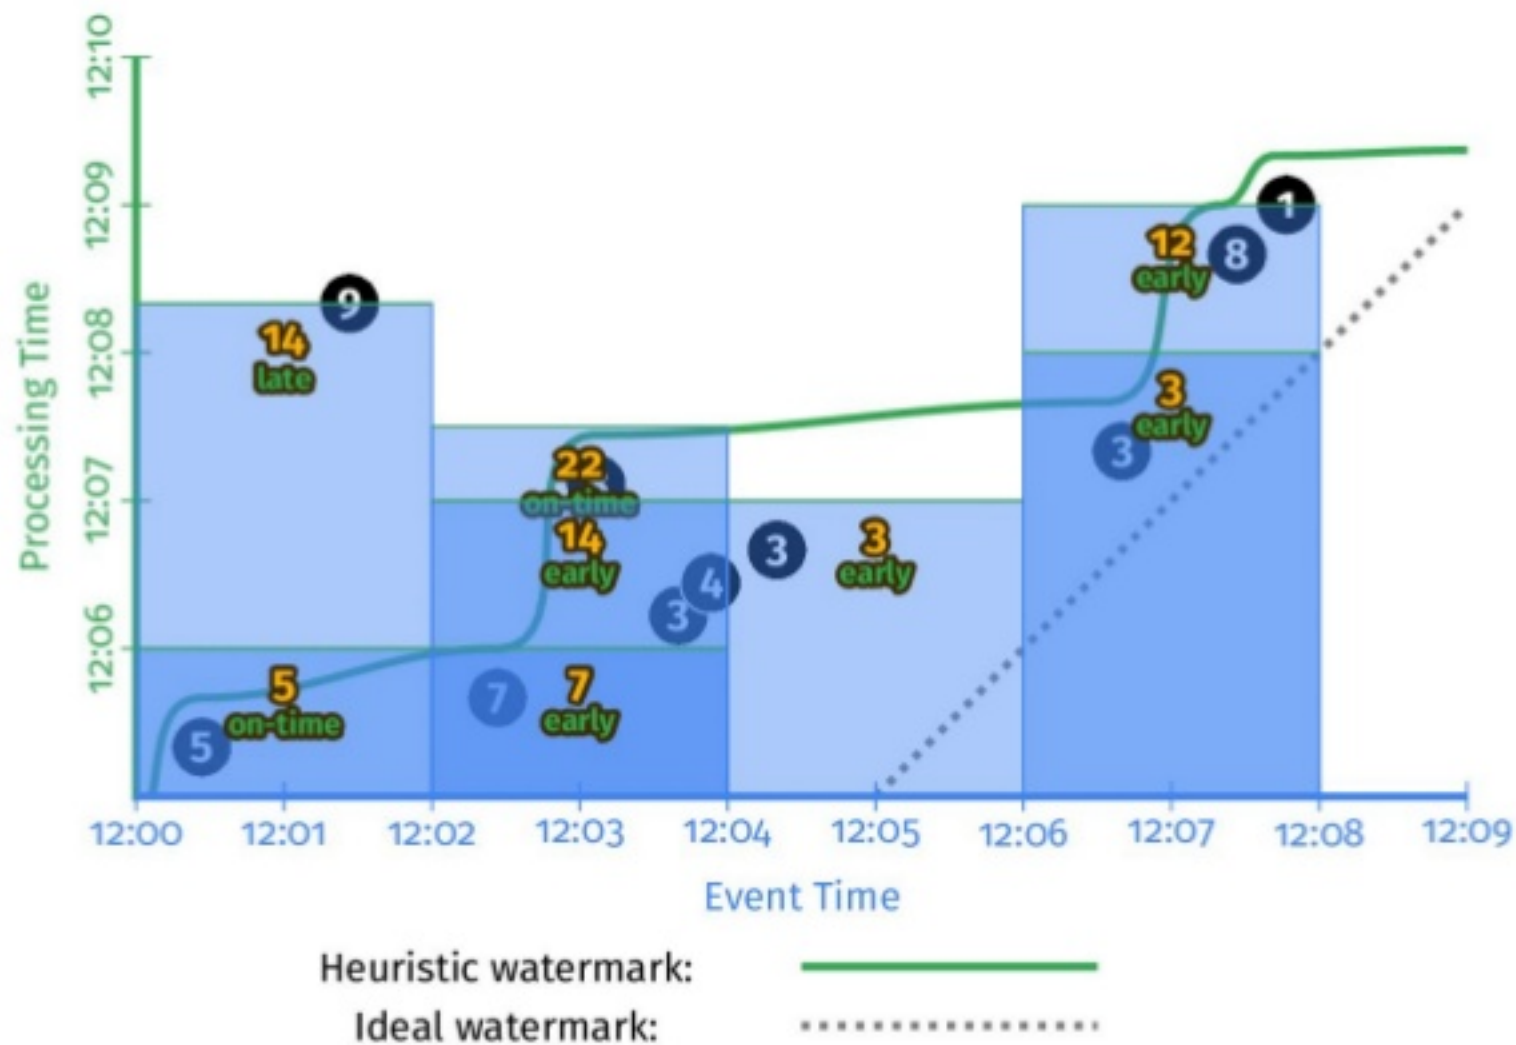


# The Beam Model: **How** Do Refinements Relate?

```
PCollection<KV<String, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(Duration.standardMinutes(2))
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(Duration.standardMinutes(1)))
            .withLateFirings(AtCount(1)))
        .accumulatingFiredPanels()))
    .apply(Sum.integersPerKey());
```

```
scores = (input
    | beam.WindowInto(FixedWindows(2 * 60)
        .triggering(AtWatermark()
            .withEarlyFirings(AtPeriod(1 * 60))
            .withLateFirings(AtCount(1)))
        .accumulatingFiredPanels())
    | Sum.integersPerKey())
```

# The Beam Model: **How** Do Refinements Relate?



# Example of Pythonic Beam Code

```
import apache_beam as beam
with beam.Pipeline() as p:
    (p
     | beam.io.ReadStringsFromPubSub("twitter_topic")
     | beam.WindowInto(SlidingWindows(5*60, 1*60))
     | beam.Map(ParseHashTagDoFn())
     | beam.CombinePerKey(sum)
     | beam.Map(BigQueryOutputFormatDoFn())
     | beam.io.WriteToBigQuery("trends_table"))
```

# Runners

Runners “translate” the code into the target runtime



Apache Beam  
Direct Runner



Apache Apex




Apache Spark



Apache Flink



Apache Gearpump



Apache Samza



Google Cloud  
Dataflow




IBM Streams



Apache Storm  
WIP



Ali Baba  
JStorm



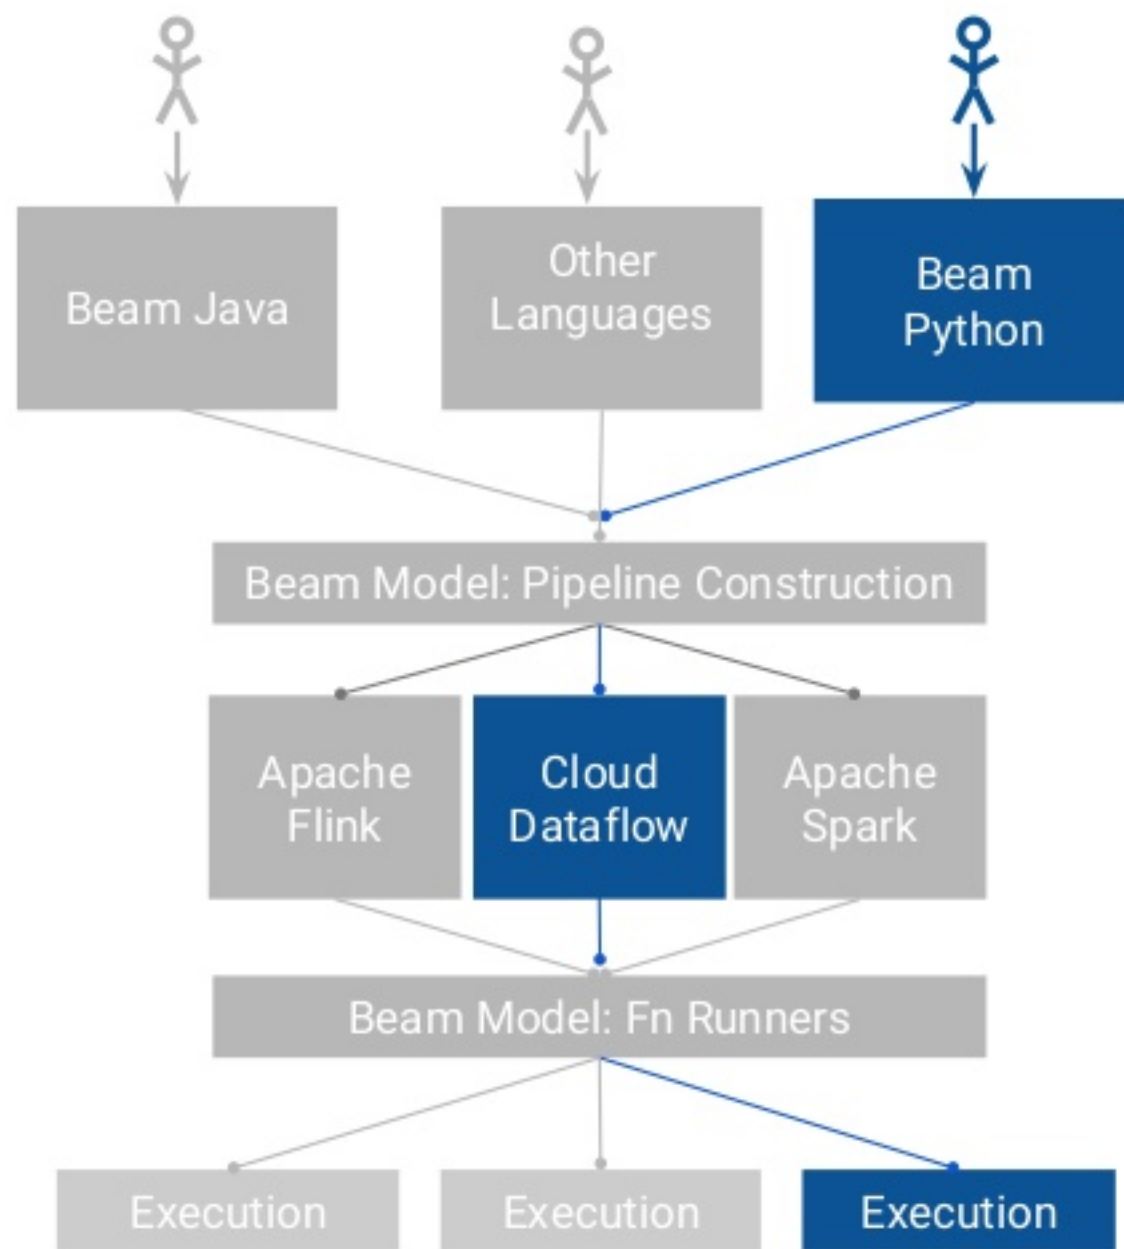
Hadoop  
MapReduce

\* Same code, different runners & runtimes



# Awesome but...

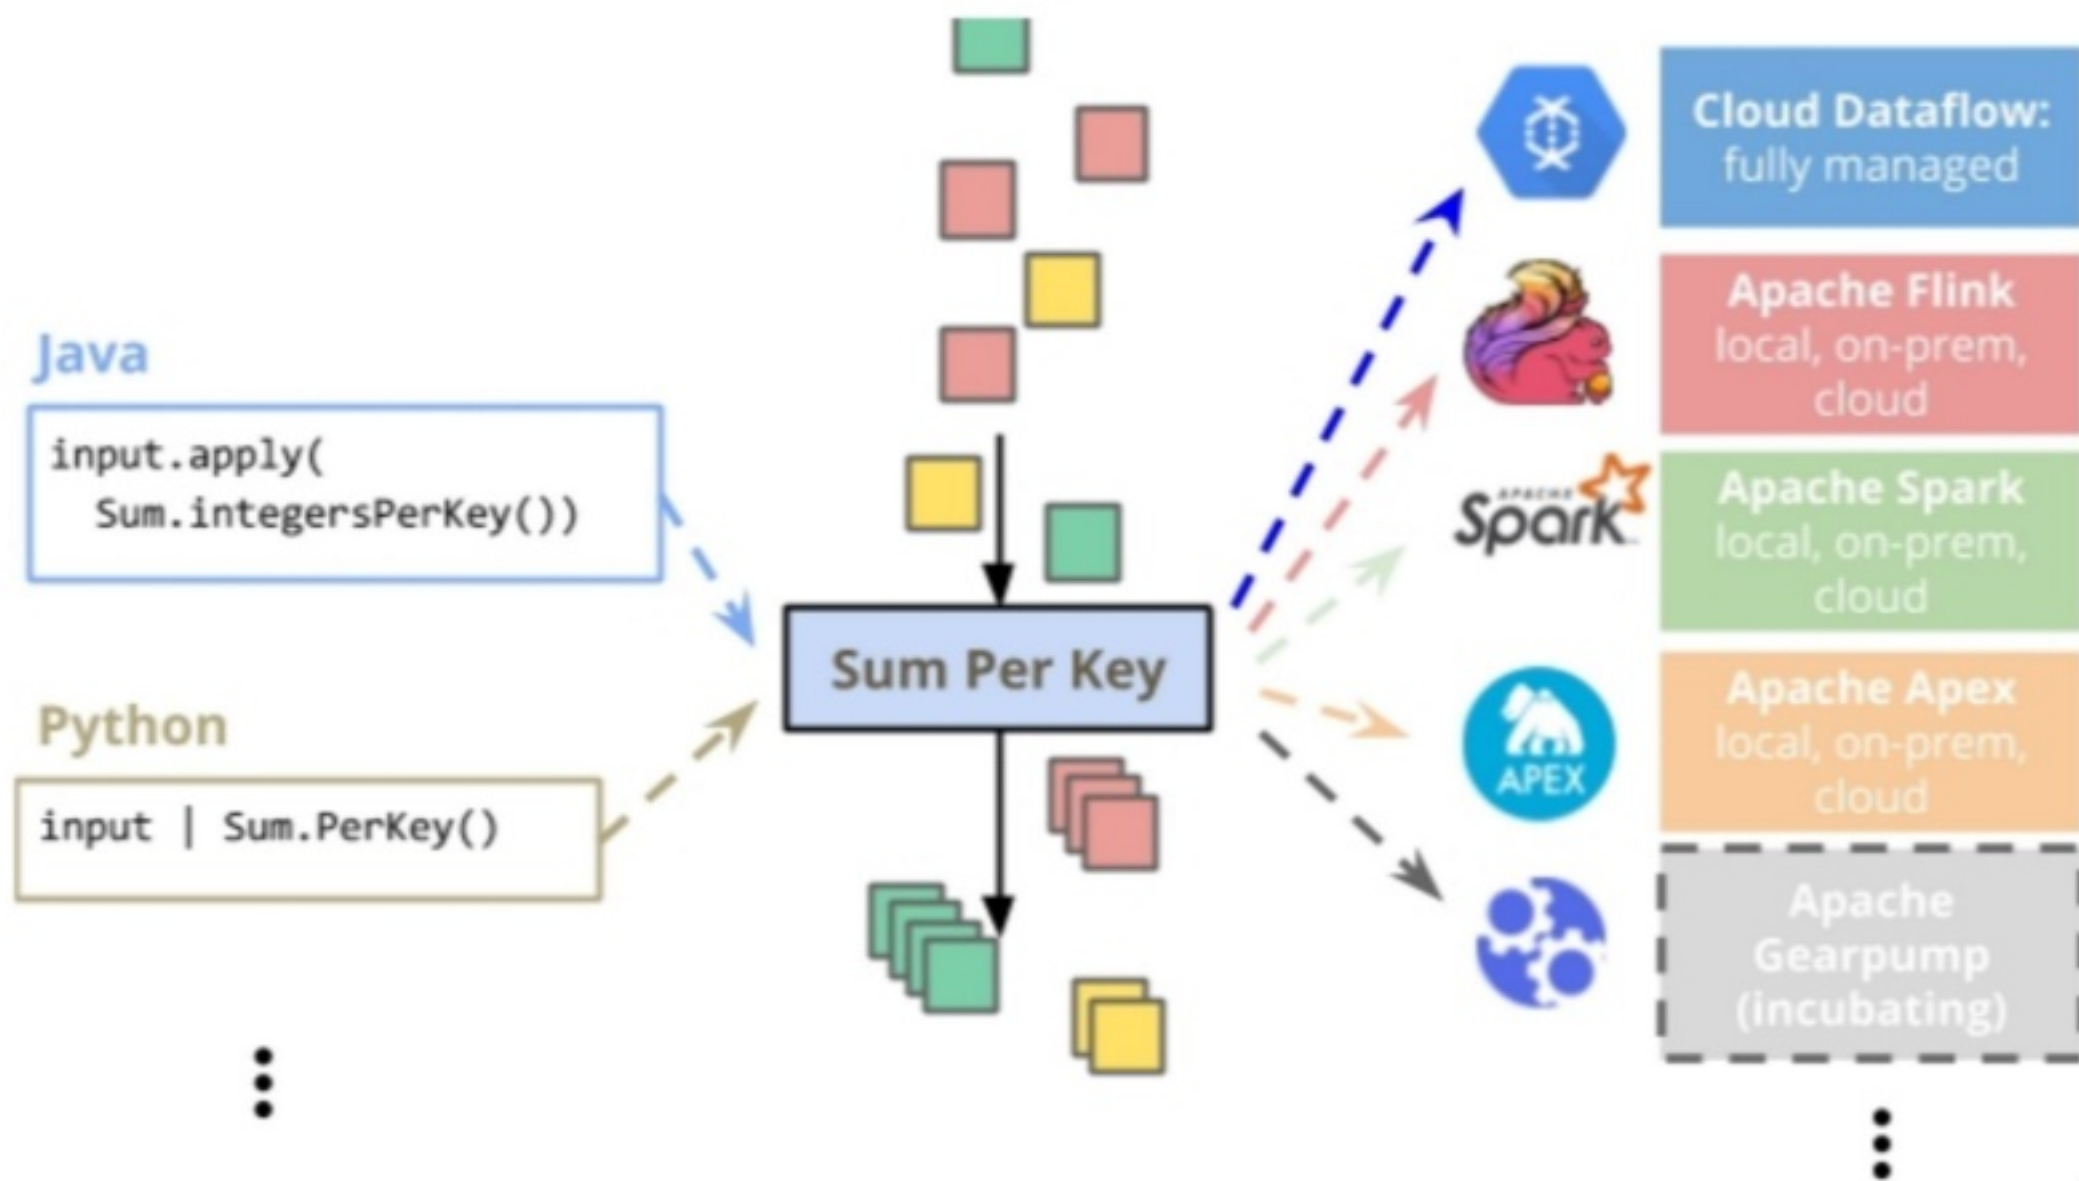
- Can a Python pipeline run on any of the Java/JVM based runners?
- Can I use the Python Tensorflow transform from a Java pipeline?
- I want to read from Kafka in my Python pipeline, but there is no connector - can I use the Java implementation?



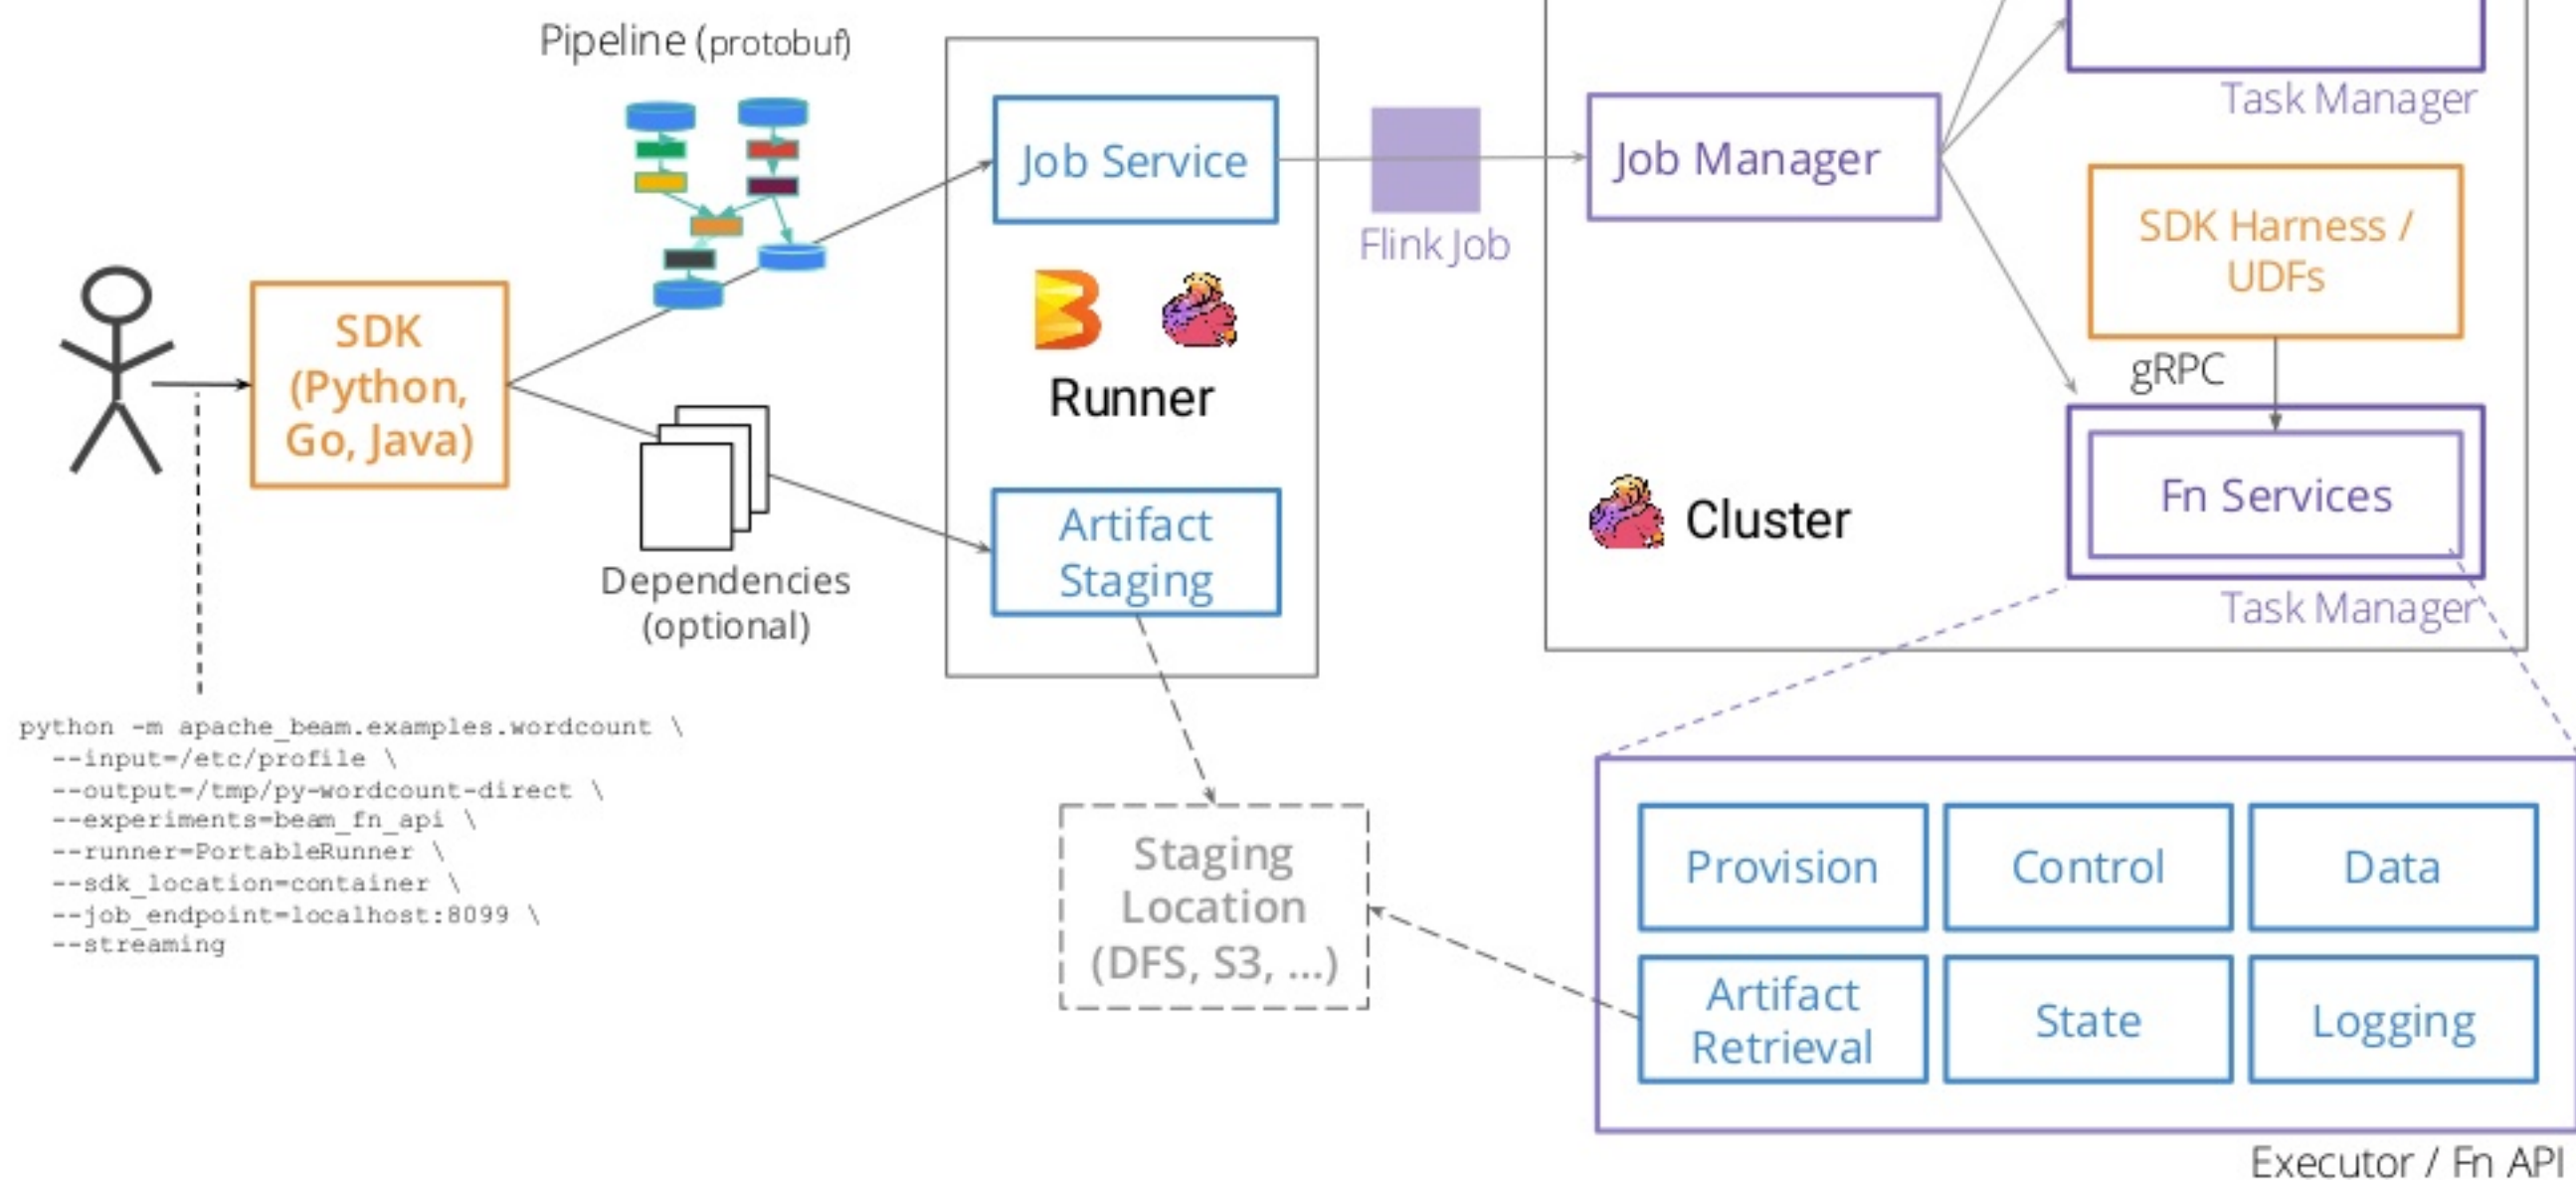
# Beam Portability



# What are we trying to solve?



# Portability Framework



# APIs for Different Pipeline Lifecycle Stages

## Pipeline API

- Used by the SDK to construct SDK-agnostic Pipeline representation
- Used by the Runner to translate a Pipeline to runner-specific operations

## Job API

- Launching and interacting with a running Pipeline

## Fn API

- Used by an SDK harness for communication with a Runner
- User by the Runner to push work into an SDK harness

# Pipeline API (simplified)

- Definition of common primitive transformations (Impulse, ExecutableStage, Flatten, AssignWindow, GroupByKey, Reshuffle)
- Definition of serialized Pipeline (protobuf)

```
Pipeline = {PCollection*, PTransform*,  
            WindowingStrategy*, Coder*}
```

```
PTransform = {Inputs*, Outputs*, FunctionSpec}
```

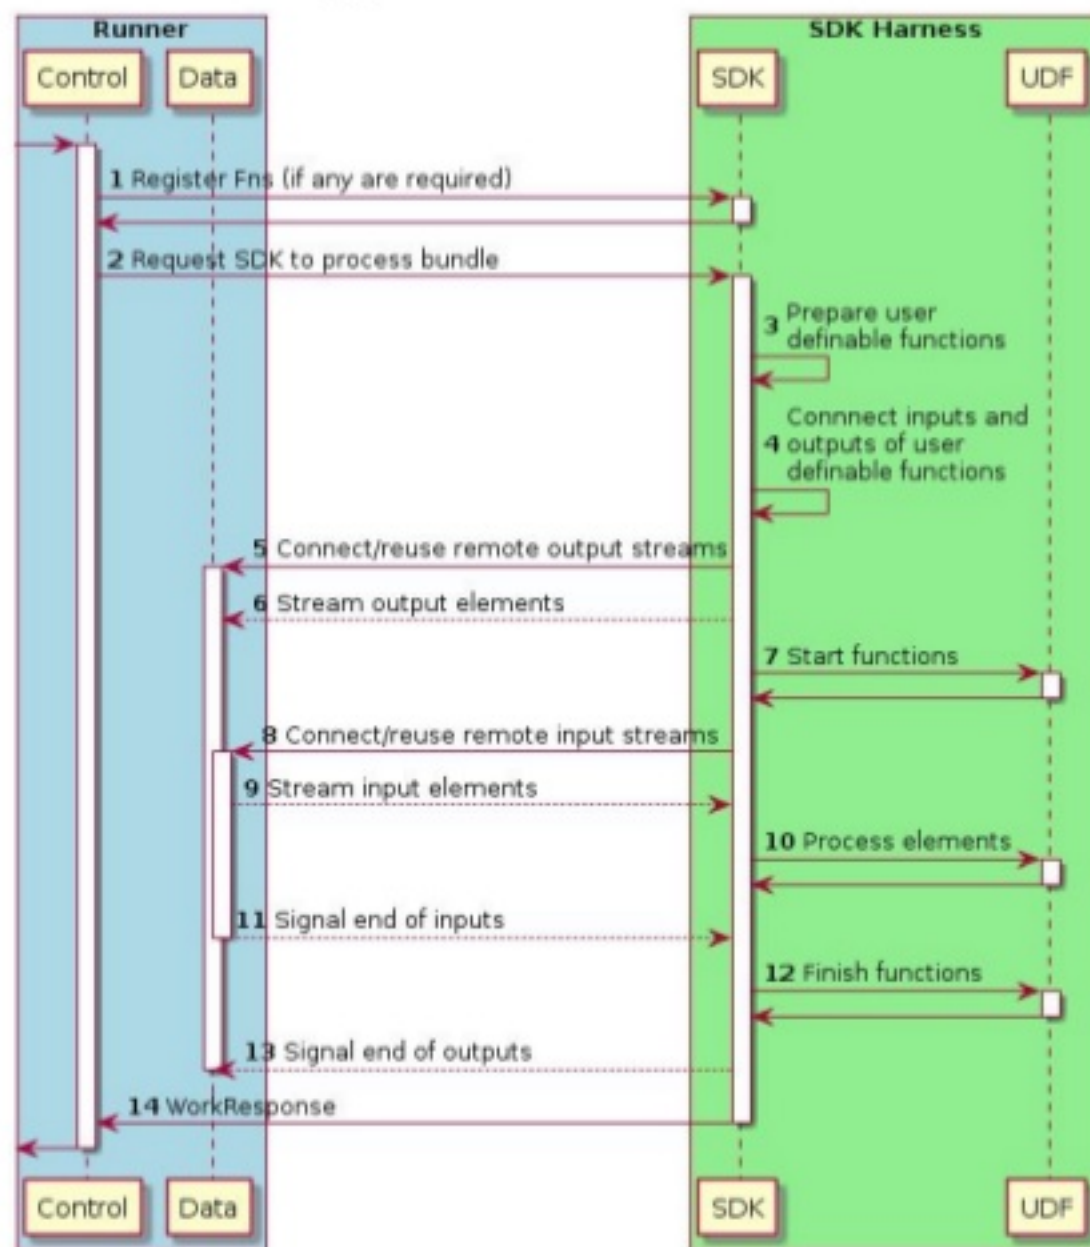
```
FunctionSpec = {URN, payload}
```

# Fn API

gRPC interfaces for communication between SDK harness and Runner

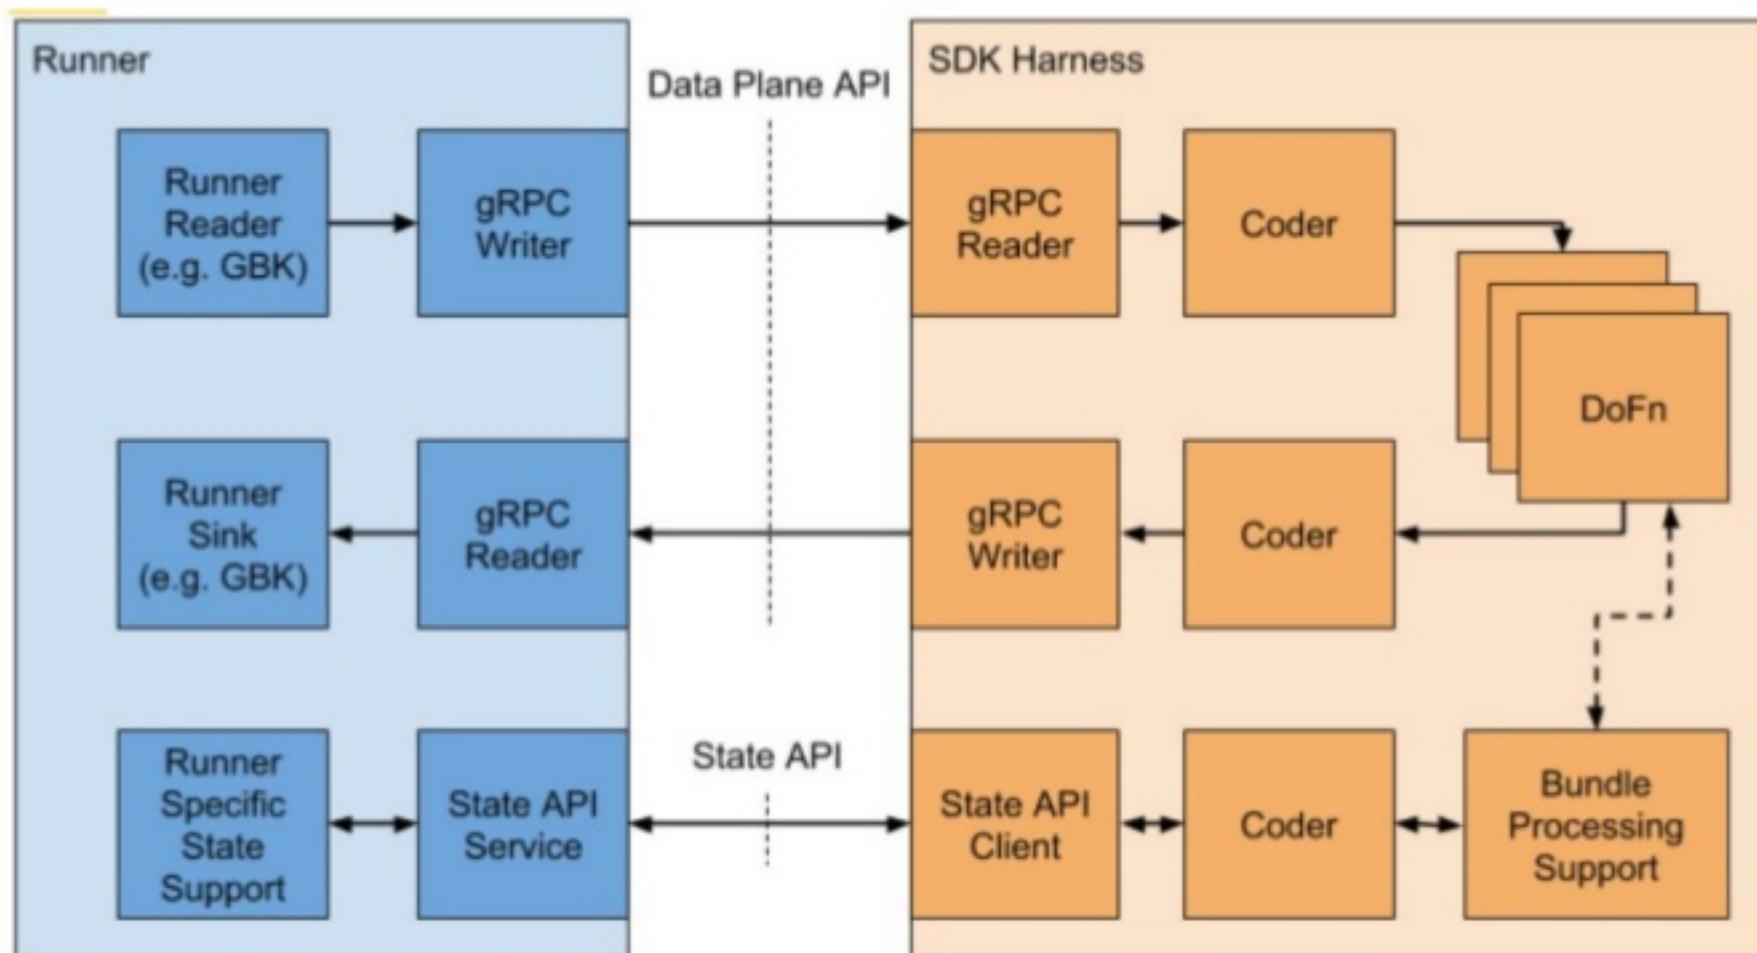
- **Control:** Used to tell the SDK which UDFs to execute and when to execute them.
- **Data:** Used to move data between the language specific SDK harness and the runner.
- **State:** Used to support user state, side inputs, and group by key reiteration.
- **Logging:** Used to aggregate logging information from the language specific SDK harness.

# Fn API - Bundle Processing





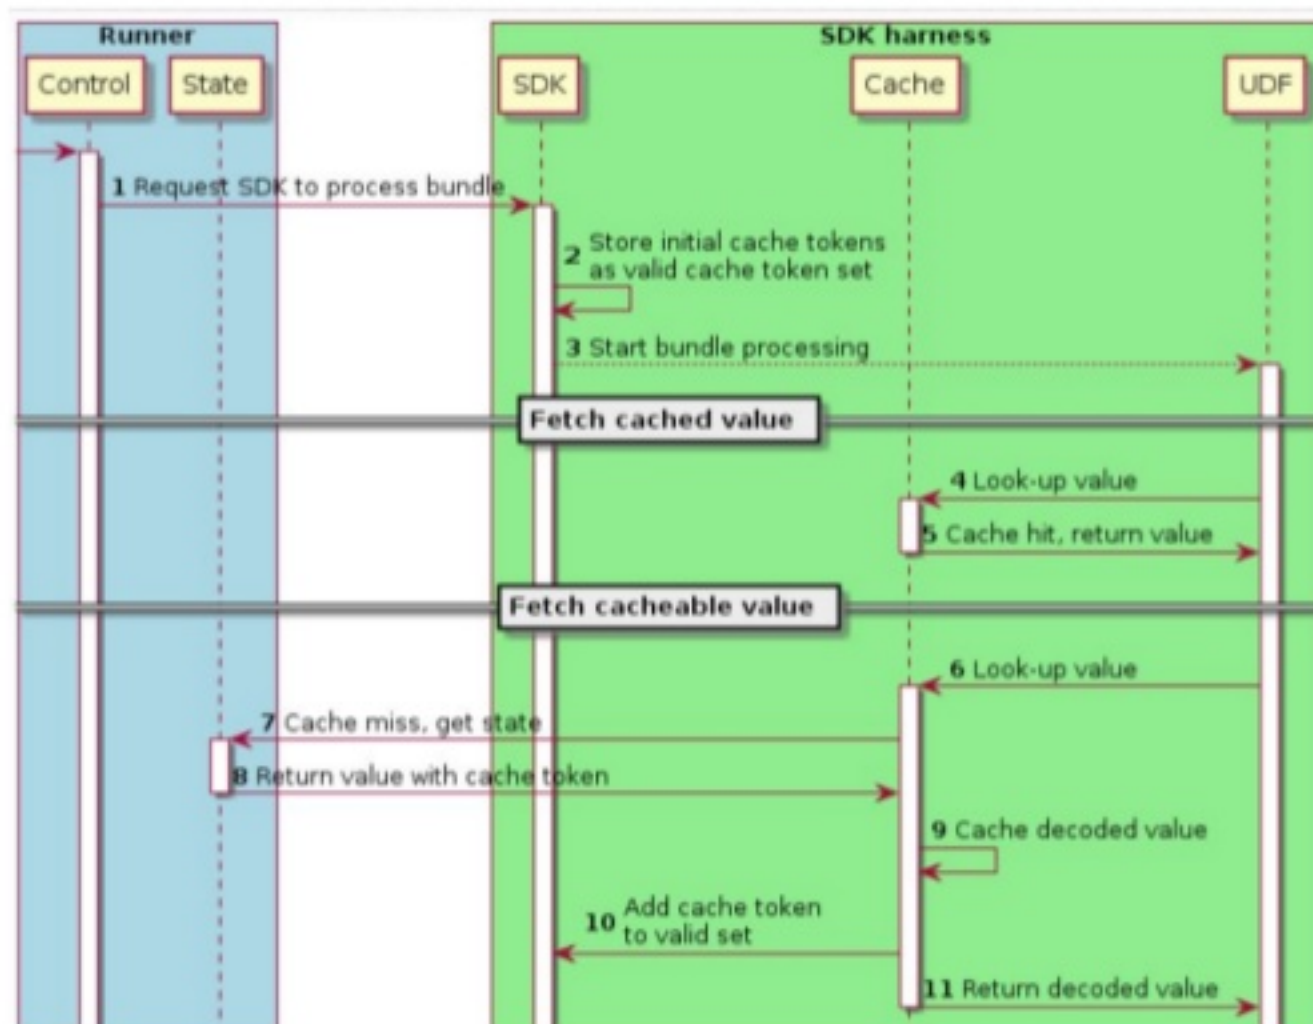
# Fn API - Data






# Fn API - State


User state, side inputs




# Fn API - Timers



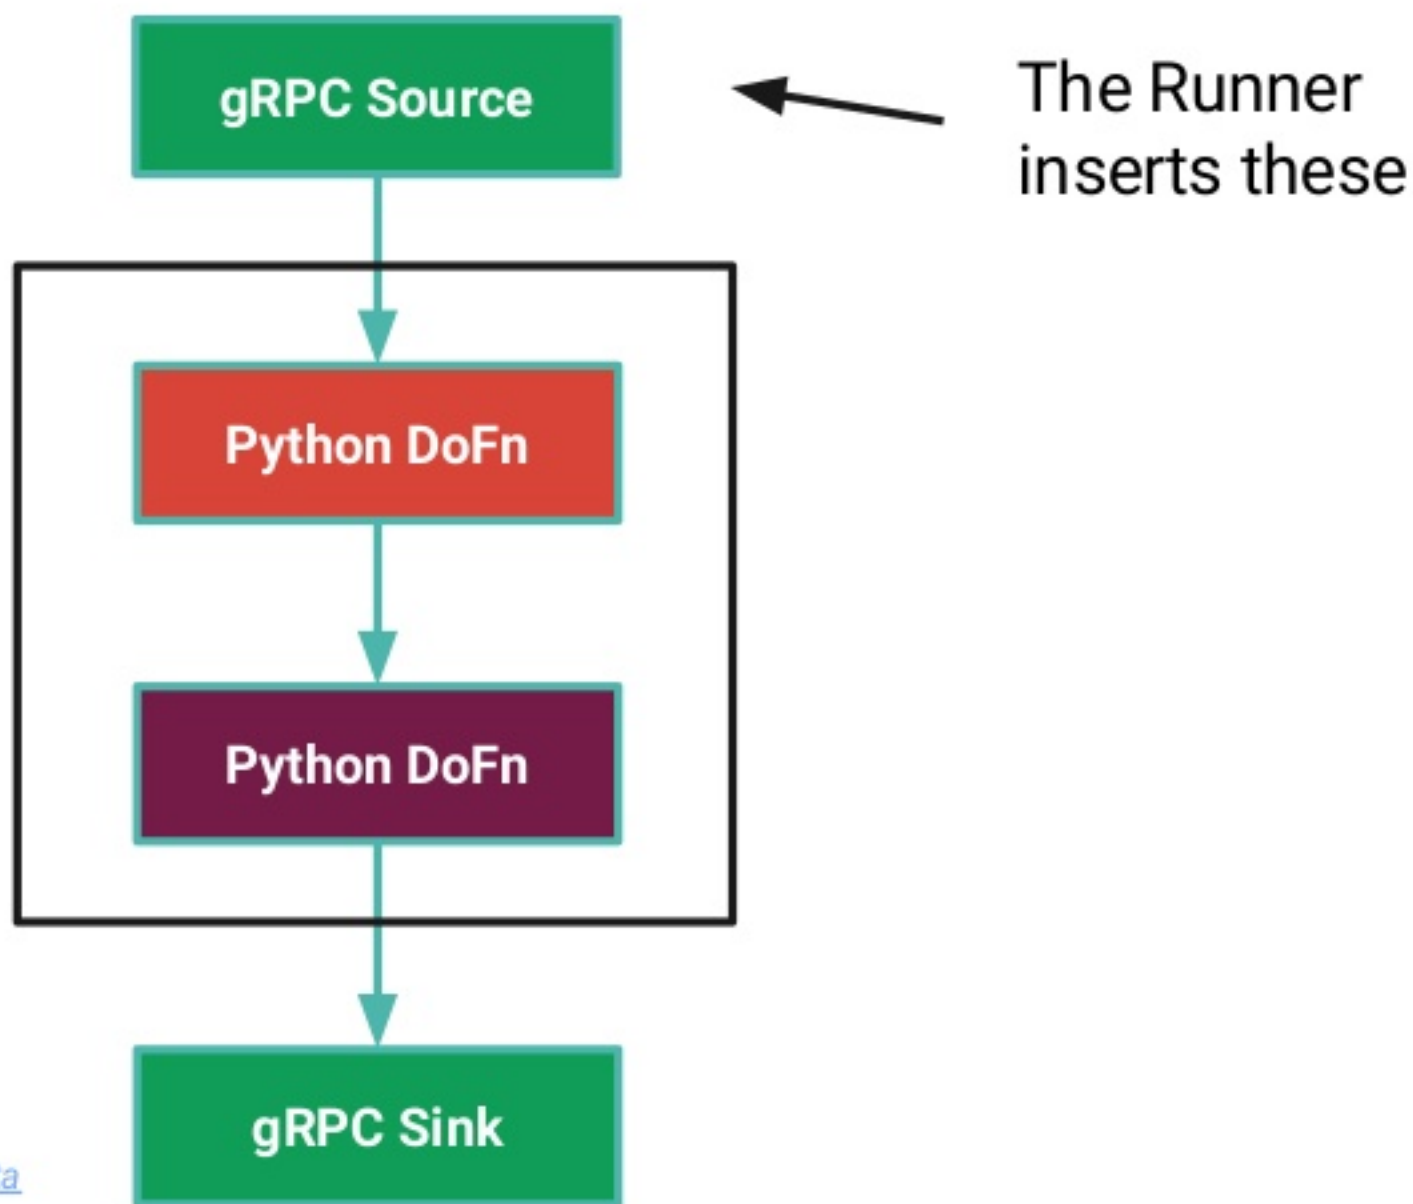
# Fn API - Processing DoFns (Executable Stages)



# Fn API - Processing DoFns

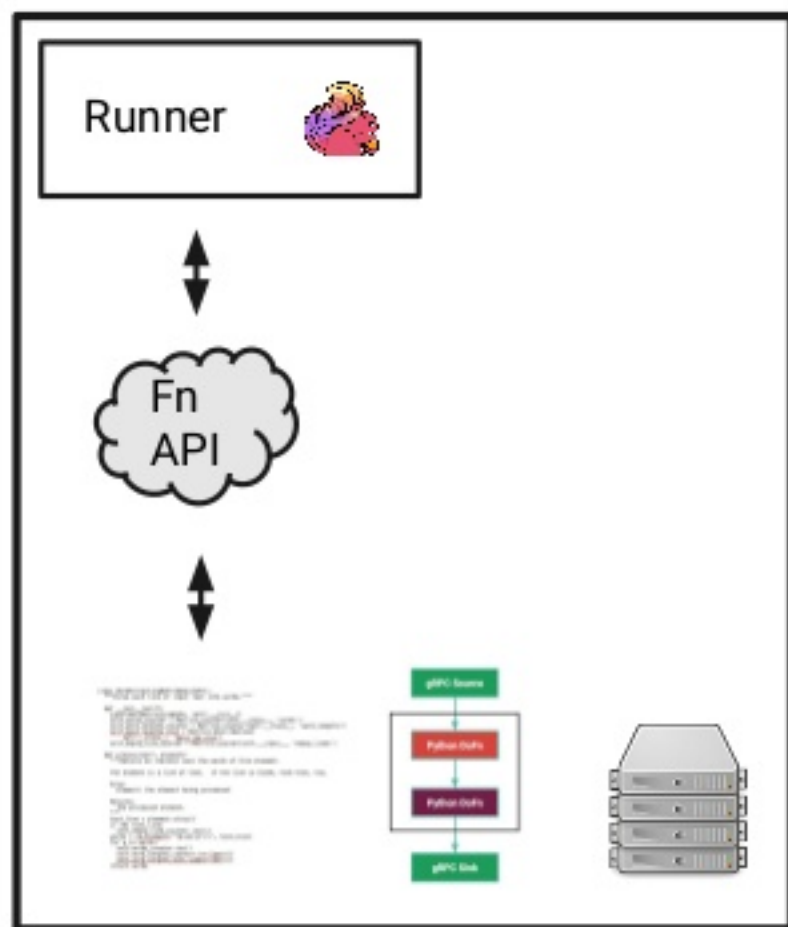


# Fn API - Processing DoFns (Pipeline manipulation)



# Fn API - Executing the user Fn in the SDK Harness

- Environments
  - Docker container
  - Separate process ([BEAM-5187](#))
  - Embedded (SDK and runner same language) - TBD
- Repository of containers for different SDKs
- Container is user-configurable
- User code can be added to container at runtime (artifact retrieval service)



# Executing Pythonic\* Beam Jobs on Fink

\*or other languages



# What is the (Flink) Runner doing in all this?

- Provide Job Service endpoint (Job Management API)
- Translate portable pipeline representation to native API
- Provide gRPC endpoints for control/data/logging/state plane
- Manage SDK Harness processes that execute user code
- Execute bundles (with arbitrary user code) using the Fn API
- Manage state for side inputs, user state/timers

**Reference runner provides common implementation baseline for JVM based runners (/runners/java-fn-execution) and we have a portable Validate Runner integration test suite in Python!**

# What's specific to the Flink Runner?

- Job Server packaging (fat jar and docker container)
- Pipeline translators for batch (DataSet) and streaming (DataStream)
  - Translation/operators for primitive URNs: Impulse, Flatten, GBK, Assign Windows, Executable Stage, Reshuffle
- Side input handlers based on Flink State
- User State and Timer integration (TBD)
- Flink Job Launch (same as old, non-portable runner)

# Advantages/Disadvantages

- Complete isolation of user code
- Configurability of execution environment (Docker, ...)
- Support for code written in non-JVM languages
- Ability to mix code written in different languages in a single pipeline (future)
- Slower (RPC overhead)
- Using Docker requires docker 😊
  - Direct Process Executor WIP
- Early Adoption (complete runner overhaul)

# Roadmap

The proposed project phases are roughly as follows and are not strictly sequential, as various components will likely move at different speeds. Additionally, there have been (and continues to be) supporting refactorings that are not always tracked as part of the portability effort. Work already done is not tracked here either.

- **P1 [MVP]:** Implement the fundamental plumbing for portable SDKs and runners for batch and streaming, including containers and the ULR [BEAM-2899]. Each SDK and runner should use the portability framework at least to the extent that wordcount [BEAM-2896] and windowed wordcount [BEAM-2941] run portably.
- **P2 [Feature complete]:** Design and implement portability support for remaining execution-side features, so that any pipeline from any SDK can run portably on any runner. These features include side inputs [BEAM-2863], User state [BEAM-2862], User timers [BEAM-2925], Splittable DoFn [BEAM-2896] and more. Each SDK and runner should use the portability framework at least to the extent that the mobile gaming examples [BEAM-2940] run portably.
- **P3 [Performance]:** Measure and tune performance of portable pipelines using benchmarks such as Nexmark. Features such as progress reporting [BEAM-2940], combiner lifting [BEAM-2937] and fusion are expected to be needed.
- **P4 [Cross language]:** Design and implement cross-language pipeline support, including how the ecosystem of shared transforms should work.

# Feature Support Matrix as of Beam 2.7.0

A	B	C	K	L	M	N	O	P	Q	R	S	T	U	V
			Flink (master)	<a href="#">instructions</a>					Dataflow					
			Java		Python		Go		Java		Python		Go	
FEATURE			Batch	Streaming	Batch	Streaming	Batch	Streaming	Batch	Streaming	Batch	Streaming	Batch	Streaming
	Impulse													
	ParDo													
		w/ side input					BEAM-3286	BEAM-3286					BEAM-3286	BEAM-3286
		w/ multiple output												
		w/ user state	BEAM-2918/BEA	BEAM-2918/BEA	BEAM-2918/BEA	BEAM-2918/BEA	BEAM-2918/BEA	BEAM-2918/BEA	BEAM-2902/BEA	BEAM-2902/BEA	BEAM-2902/BEA	BEAM-2902/BEA	BEAM-2902/BEA	BEAM-2902/BEA
		w/ user timers												
		w/ user metrics												
	Flatten													
		w/ explicit flatten					BEAM-3300	BEAM-3300					BEAM-3300	BEAM-3300
	Combine													
		w/ first-class rep					BEAM-4276	BEAM-4276	BEAM-3513	BEAM-3513			BEAM-4276	BEAM-4276
		w/ lifting					BEAM-4276	BEAM-4276	BEAM-3711	BEAM-3711			BEAM-4276	BEAM-4276
	SDF						BEAM-3301	BEAM-3301					BEAM-3301	BEAM-3301
		w/ liquid sharding												
	GBK													
	CoGBK													
	WindowInto													
		w/ sessions					BEAM-4152	BEAM-4152					BEAM-4152	BEAM-4152
		w/ custom windowfn												
EXAMPLE			Batch	Streaming	Batch	Streaming	Batch	Streaming	Batch	Streaming	Batch	Streaming	Batch	Streaming
	WordCap													
	WordCount													
		w/ write to Sink												
		w/ write to GCS												

Demo

# The Future



# Future work

- Support for user state and timers
- Mixing and matching connectors written in different languages
- Wait for new SDKs in other languages, they will just work 😊
- Unified batch and streaming API in Flink?
  - currently 4 Flink translators (batch + streaming for each, portable and old, Java-only runner)
- Beam Flink Runner compatibility story
  - Flink upgrades
  - Pipeline upgrades / state migration

# Streaming@Lyft

We are hiring! [lyft.com/careers](https://lyft.com/careers)  
<https://goo.gl/RsyLkS>



---

WE ARE HIRING

[data-artisans.com/careers](https://data-artisans.com/careers)

**dataArtisans**



# Learn More!

## Beam Portability Framework

<https://beam.apache.org/contribute/portability/>

<https://beam.apache.org/contribute/design-documents/#portability>

## Apache Beam

<https://beam.apache.org>

<https://s.apache.org/beam-slack-channel> **#beam #beam-portability**

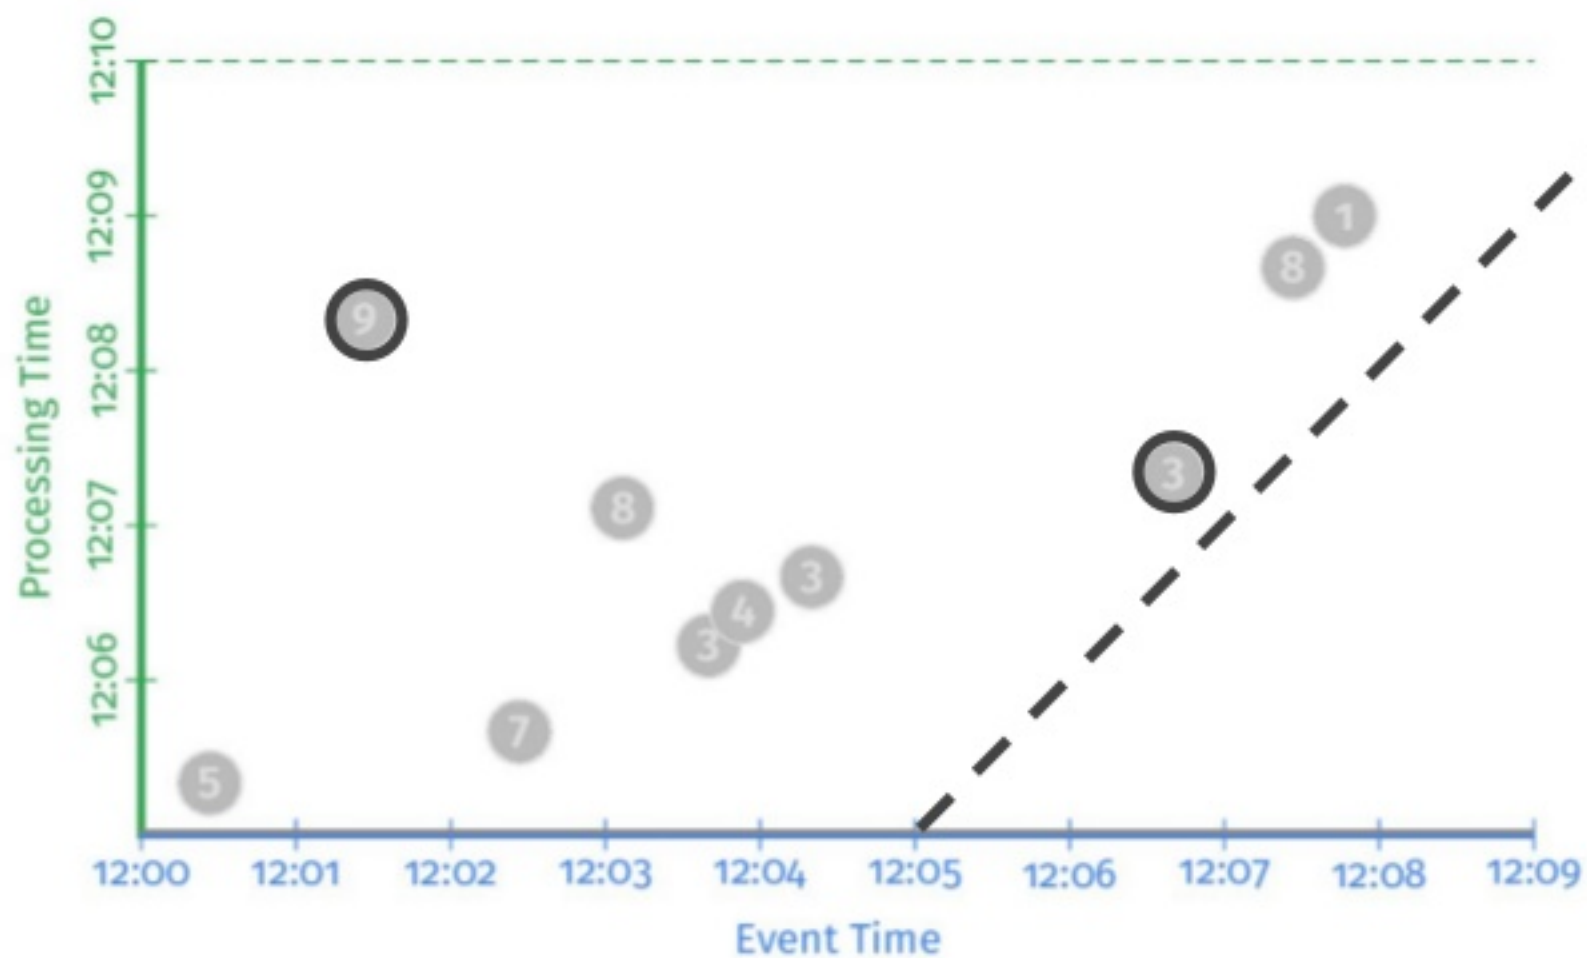
<https://beam.apache.org/community/contact-us/>

**Follow @ApacheBeam on Twitter**

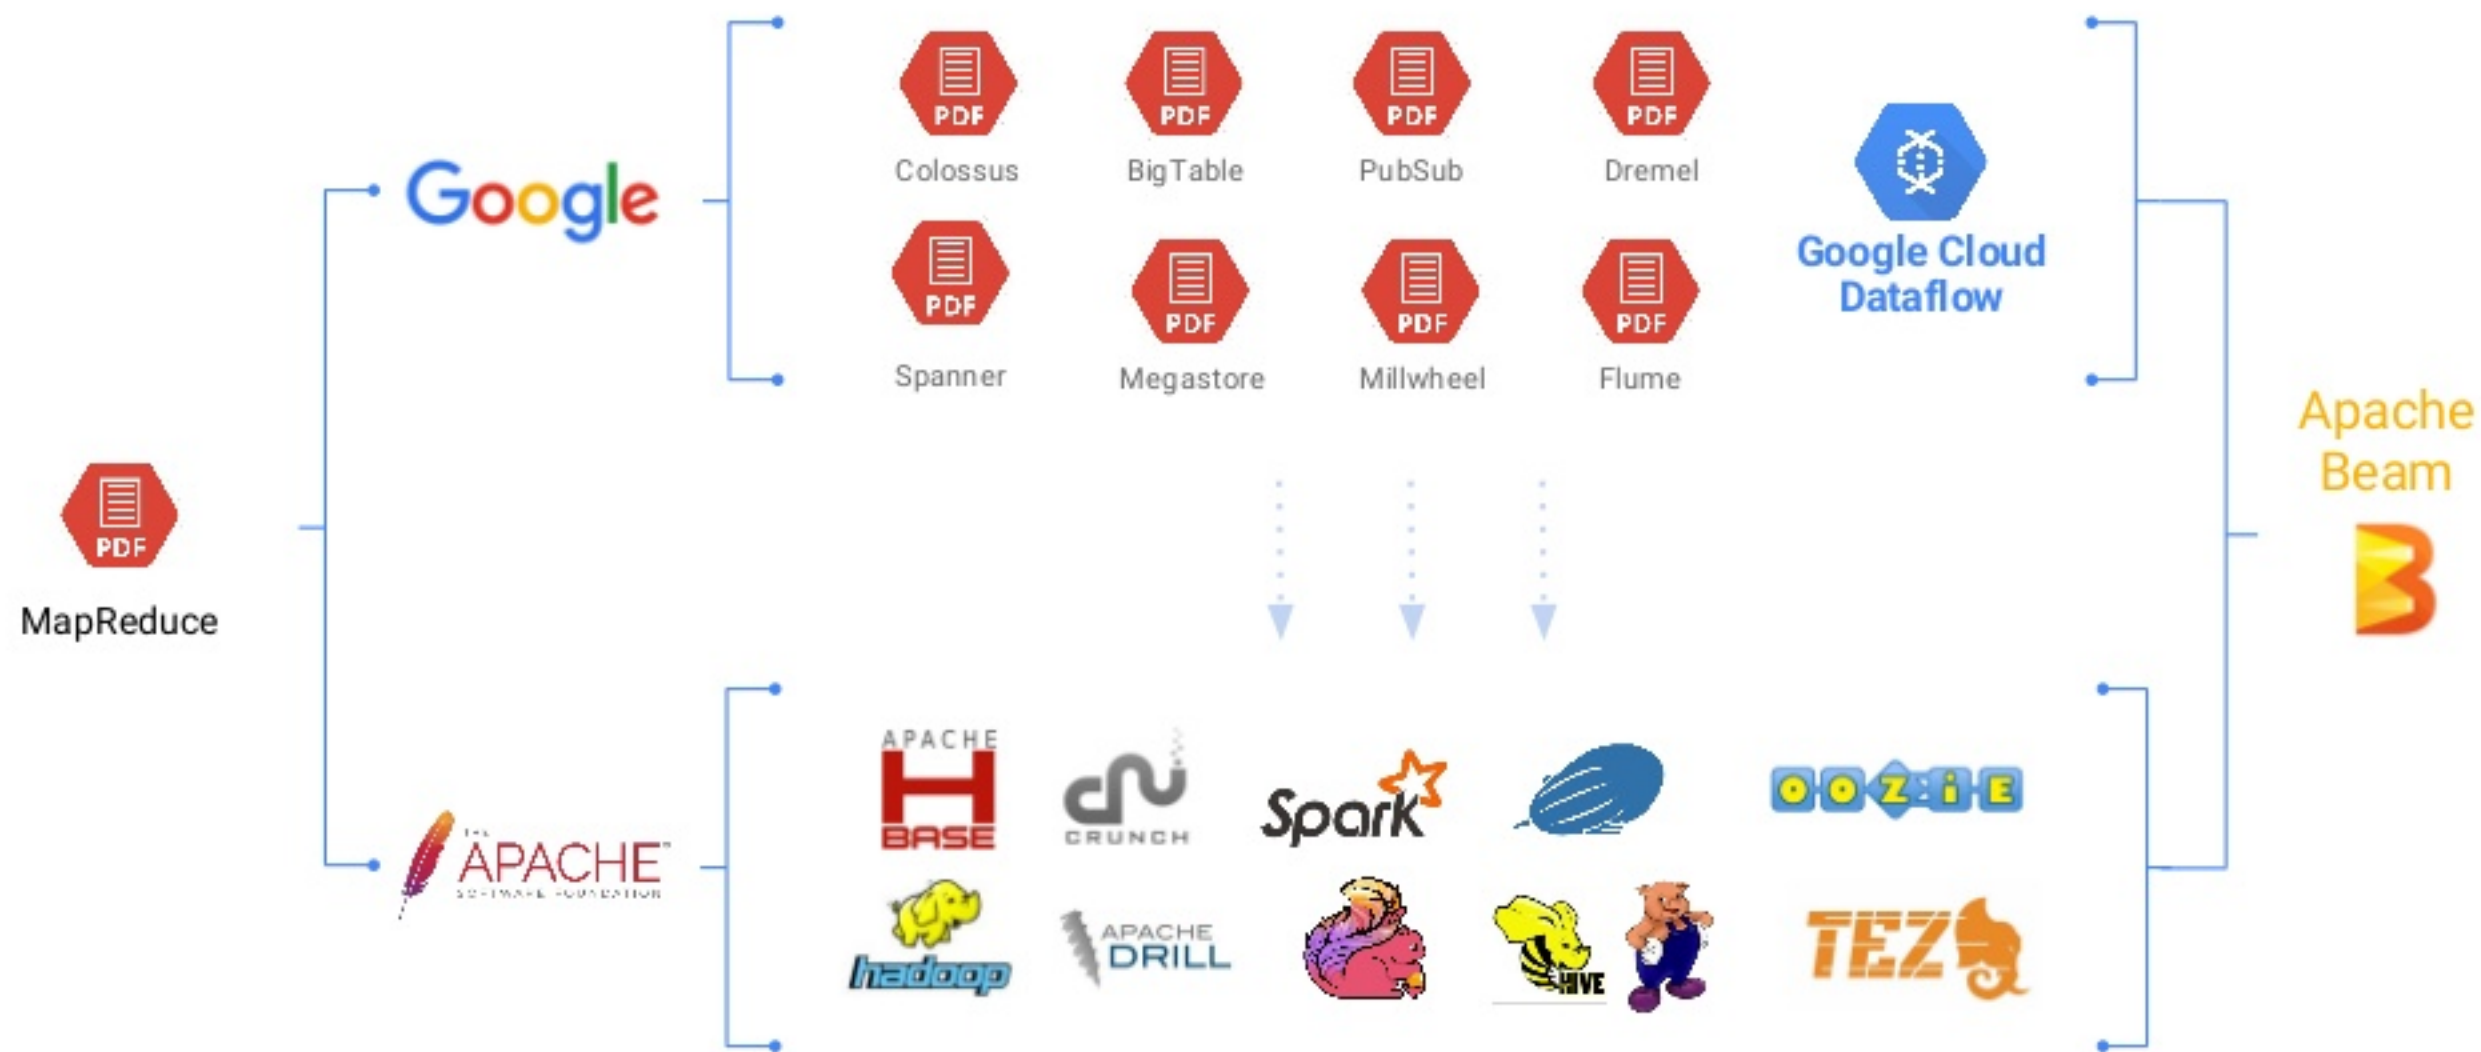
Thank you!

# Backup Slides

# Processing Time vs. Event Time



# The Origin of Apache Beam






# Beam Model: Generations Beyond MapReduce

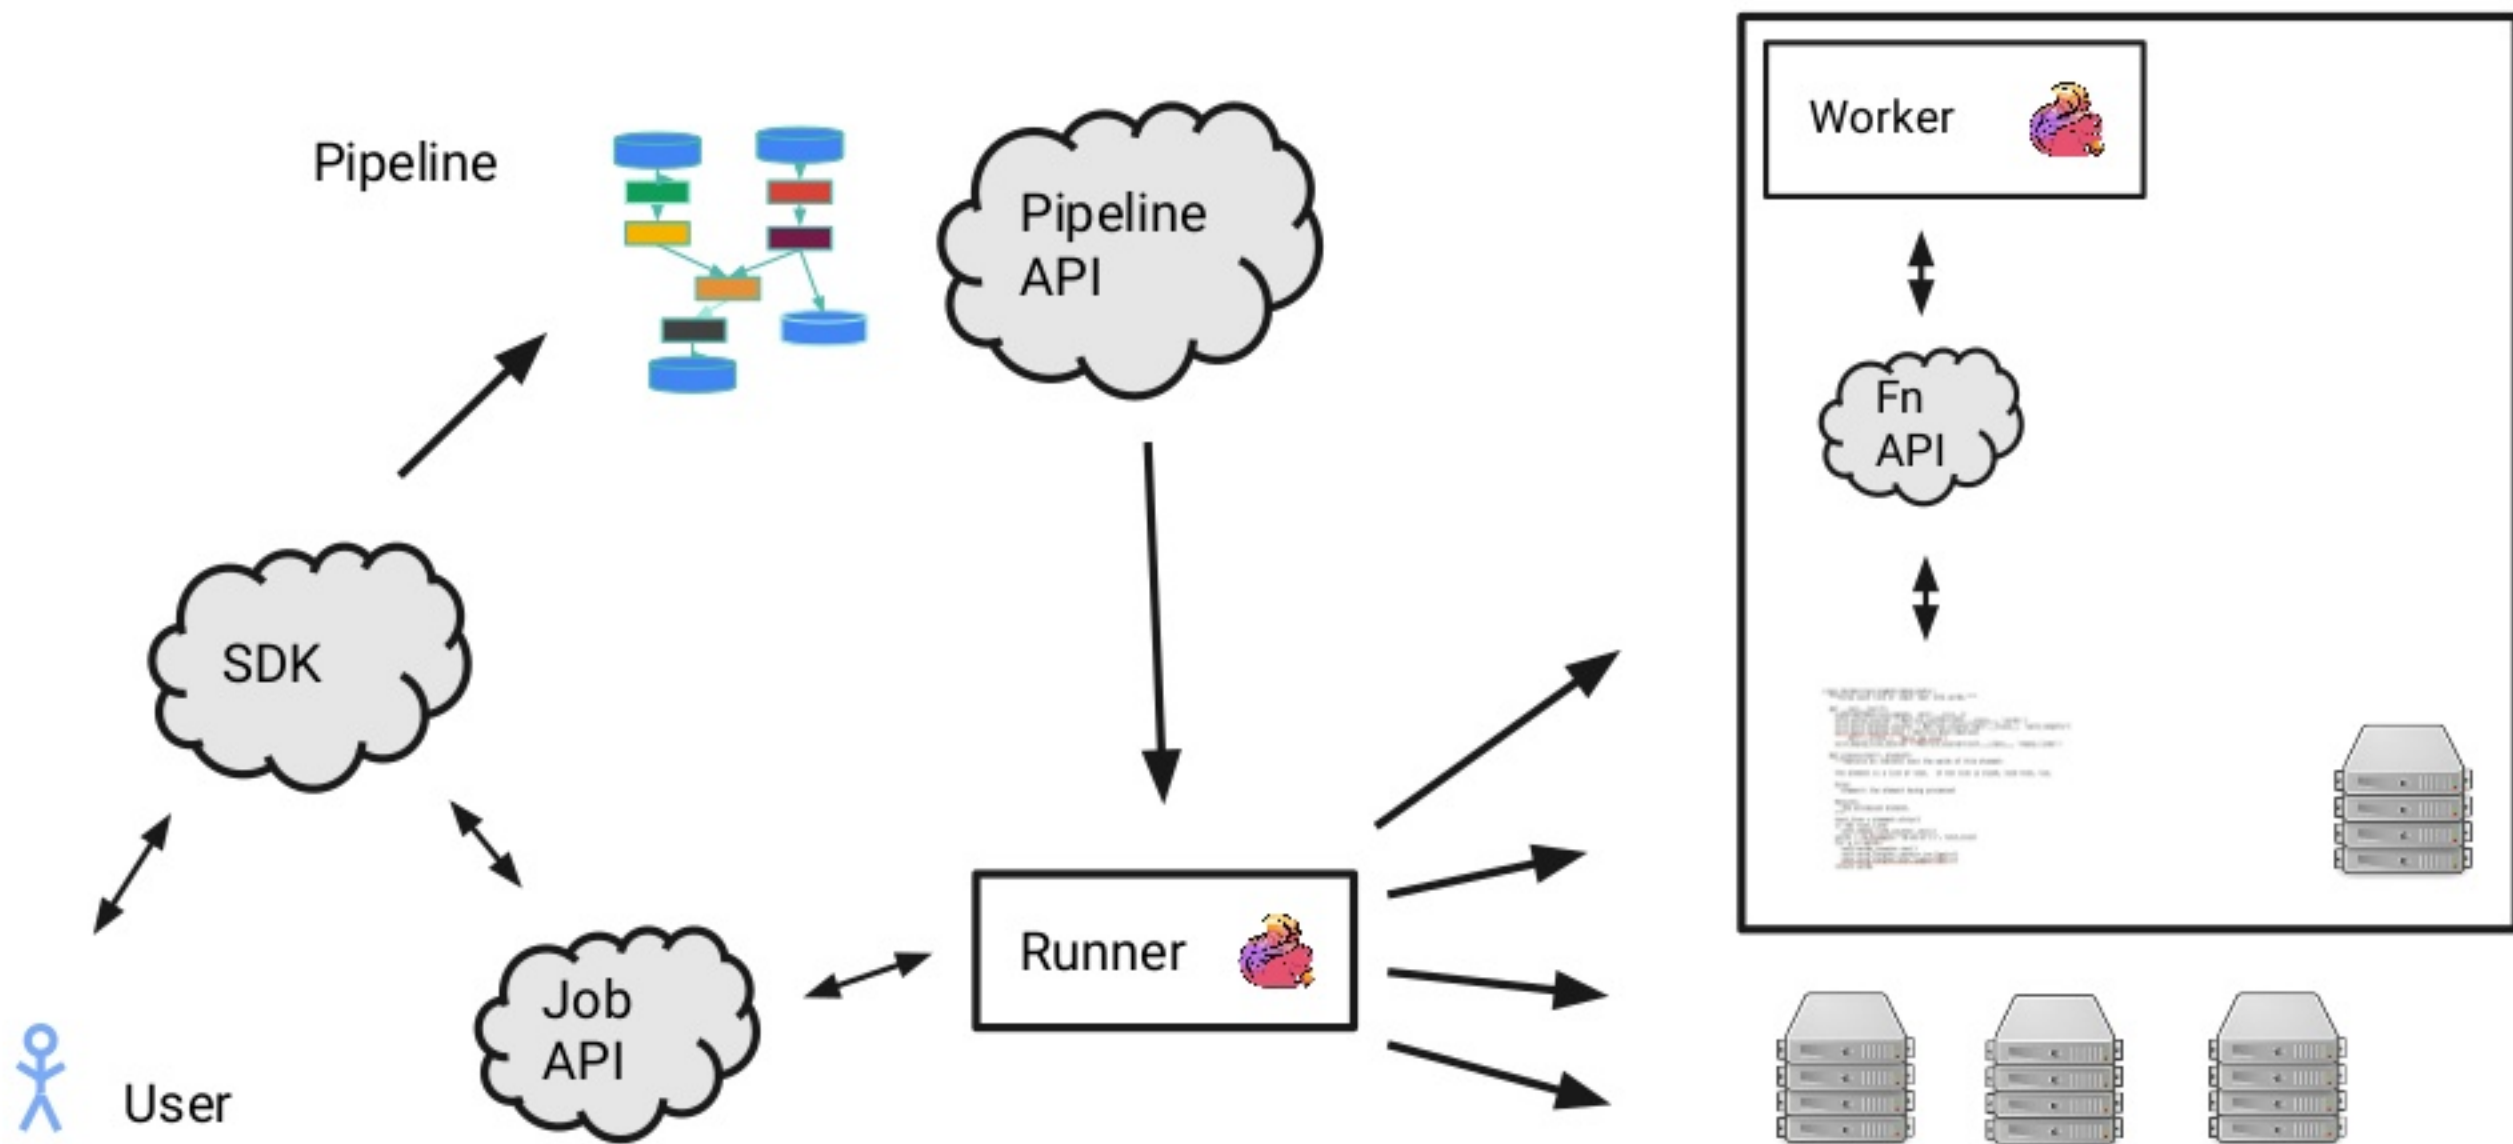
Improved abstractions let you focus on your application logic

Batch and stream processing are *both* first-class citizens -- no need to choose.


Clearly separates event time from processing time.




# Executing a Beam Pipeline - The Big Picture




# Customizing **What** **Where** **When** **How**




**1**  
**Classic  
Batch**



**2**  
**Windowed  
Batch**



**3**  
**Streaming**



**4**  
**Streaming  
+ Accumulation**

For more information see <https://cloud.google.com/dataflow/examples/gaming-example>

# Terminology

## **Beam Model**

Describes the API concepts and the possible operations on PCollections.

## **Pipeline**

User-defined graph of transformations on PCollections. This is constructed using a Beam SDK. The transformations can contain UDFs.

## **Runner**

Executes a Pipeline. For example: *FlinkRunner*.

## **Beam SDK**

Language specific library/framework for creating programs that use the Beam Model. Allows defining Pipelines and UDFs and provides APIs for executing them.

## **User-defined function (UDF)**

Code in Java, Python, Go, ... that specifies how data is transformed. For example *DoFn* or *CombineFn*.

# Job API

```
public interface JobApi {  
    State getState(); // RUNNING, DONE, CANCELED, FAILED ...  
    State cancel() throws IOException;  
    State waitUntilFinish(Duration duration);  
    State waitUntilFinish();  
    MetricResults metrics();  
}
```

# Fn API (continued)

