


@stsffap @joerg_schad

Elastic Streams at Scale



Workload cycles and spikes


Daily cycles



Weekly cycles



Seasonal spikes



Unplanned



Resource adaption




+




Scalable Flink applications

Scaling stateless jobs



- Scale up: Deploy new tasks
- Scale down: Cancel running tasks


Scaling stateful jobs




- Problem: Which state to assign to new task?

Keyed vs. operator state

Keyed



Operator




- State bound to a key
- E.g. Keyed UDF and window state

- State bound to a subtask
- E.g. Source state


Repartitioning keyed state

- Similar to consistent hashing
- Split key space into key groups
- Assign key groups to tasks




Repartitioning keyed state contd.

- Rescaling changes key group assignment
- Maximum parallelism defined by #key groups




Repartitioning operator state

- Breaking operator state up into finer granularity
 - State has to contain multiple entries
 - Automatic re-partitioning wrt granularity




Example: Kafka Source




- Store offset for each partition
- Individual entries are repartitionable

Automatic scaling




- Latency
- Throughput
- Resource utilization
- Connector signals



Resource elasticity

Where to get the resources from?

- Scaling Flink applications changes required number of slots
- Flink needs to start and stop TaskManagers
- ClusterManager is required to start TaskManagers dynamically



Apache Mesos in a nutshell

- Cluster management framework
 - Dynamic resource allocation
 - Running multiple applications
 - 2-level scheduling
- Fault-tolerant, battle-tested
- Scalable to 10,000+ nodes
- Created by Mesosphere founder @ UC Berkeley, used in production by 100+ web-scale companies



Why Flink on Mesos?

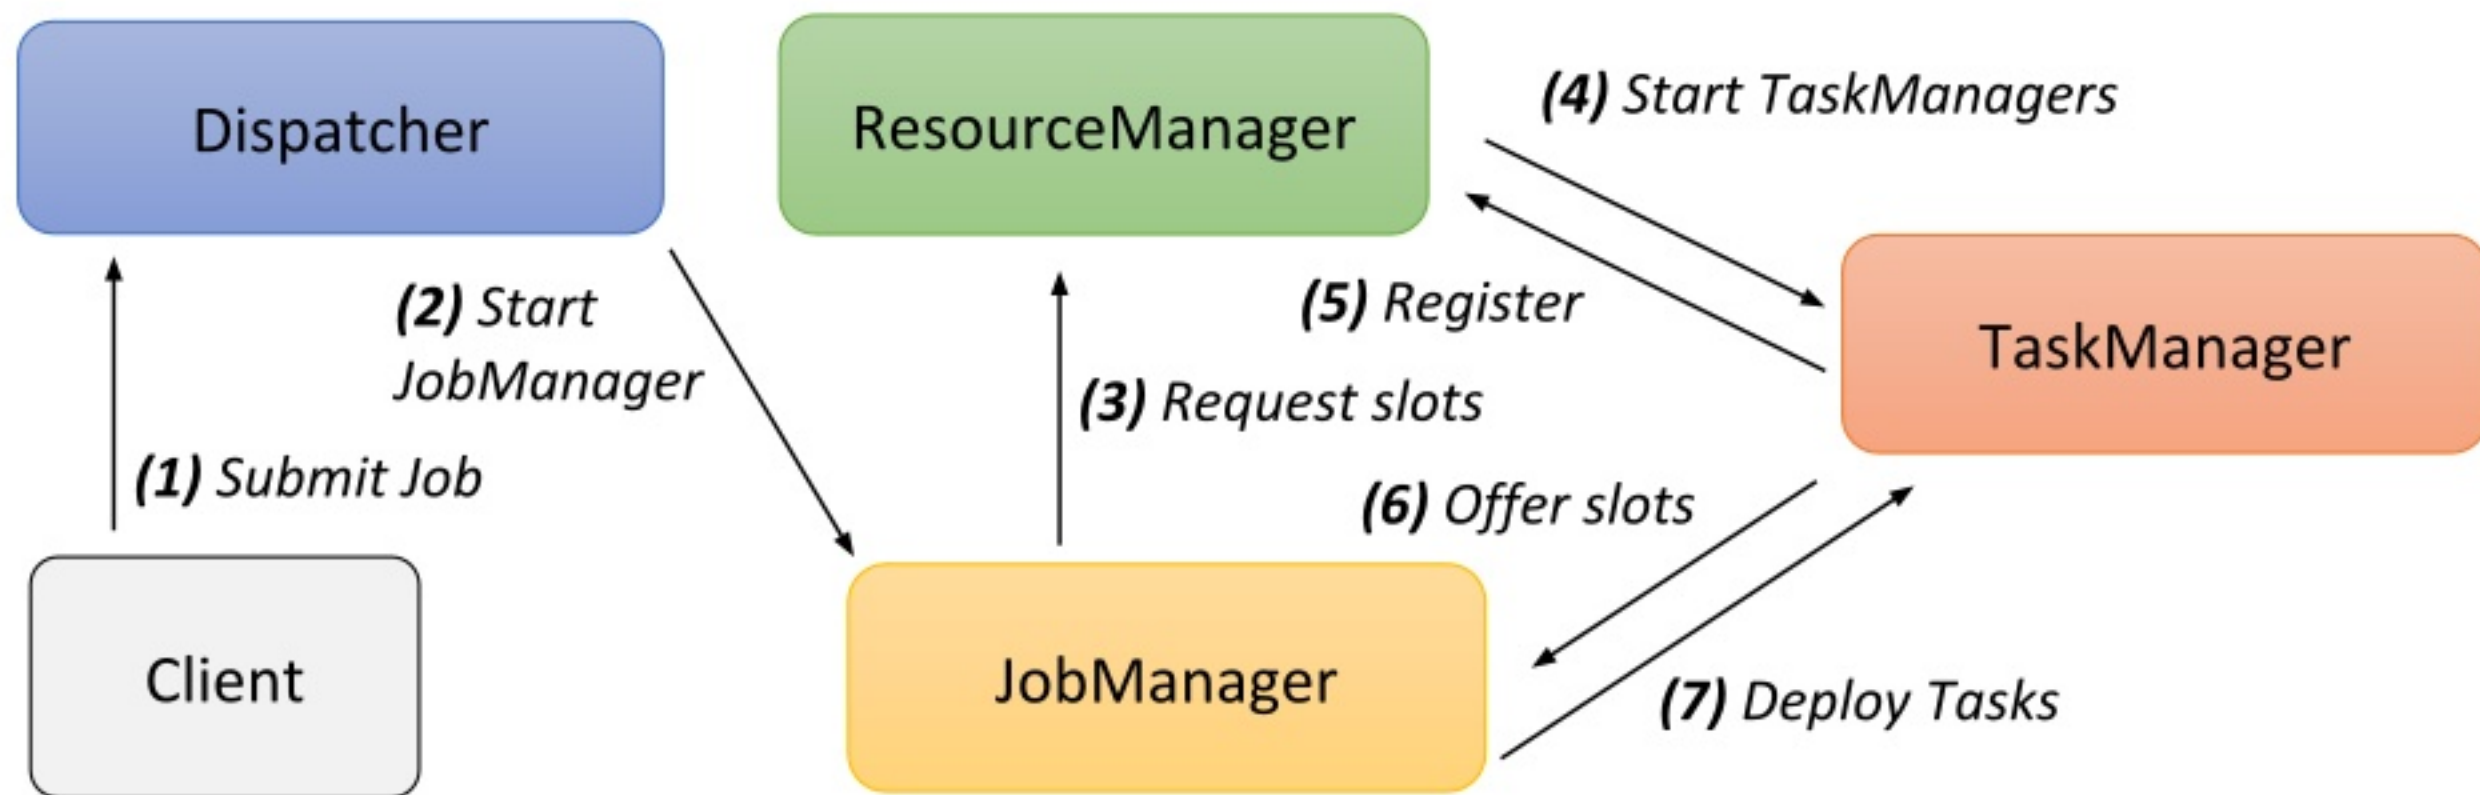
- Mesos offers full functionality to implement fault tolerant and elastic distributed applications
- 30% of survey respondents were running Flink on Mesos (prior to proper Mesos support*, September 2016)
- Other Deployment Models
 - Standalone
 - Yarn
 - Kubernetes

Flink's Revamped Distributed Architecture

- Motivation
 - Resource elasticity
 - Support for different deployments
 - REST interface for client-cluster communication
- Introduce generic building blocks
- Compose blocks for different scenarios



The Building Blocks



The Building Blocks

ResourceManager

- ClusterManager-specific
- May live across jobs
- Manages available Containers/TaskManagers
- Used to acquire / release resources

JobManager

- Single job only, started per job
- Thinks in terms of "task slots"
- Deploys and monitors job/task execution

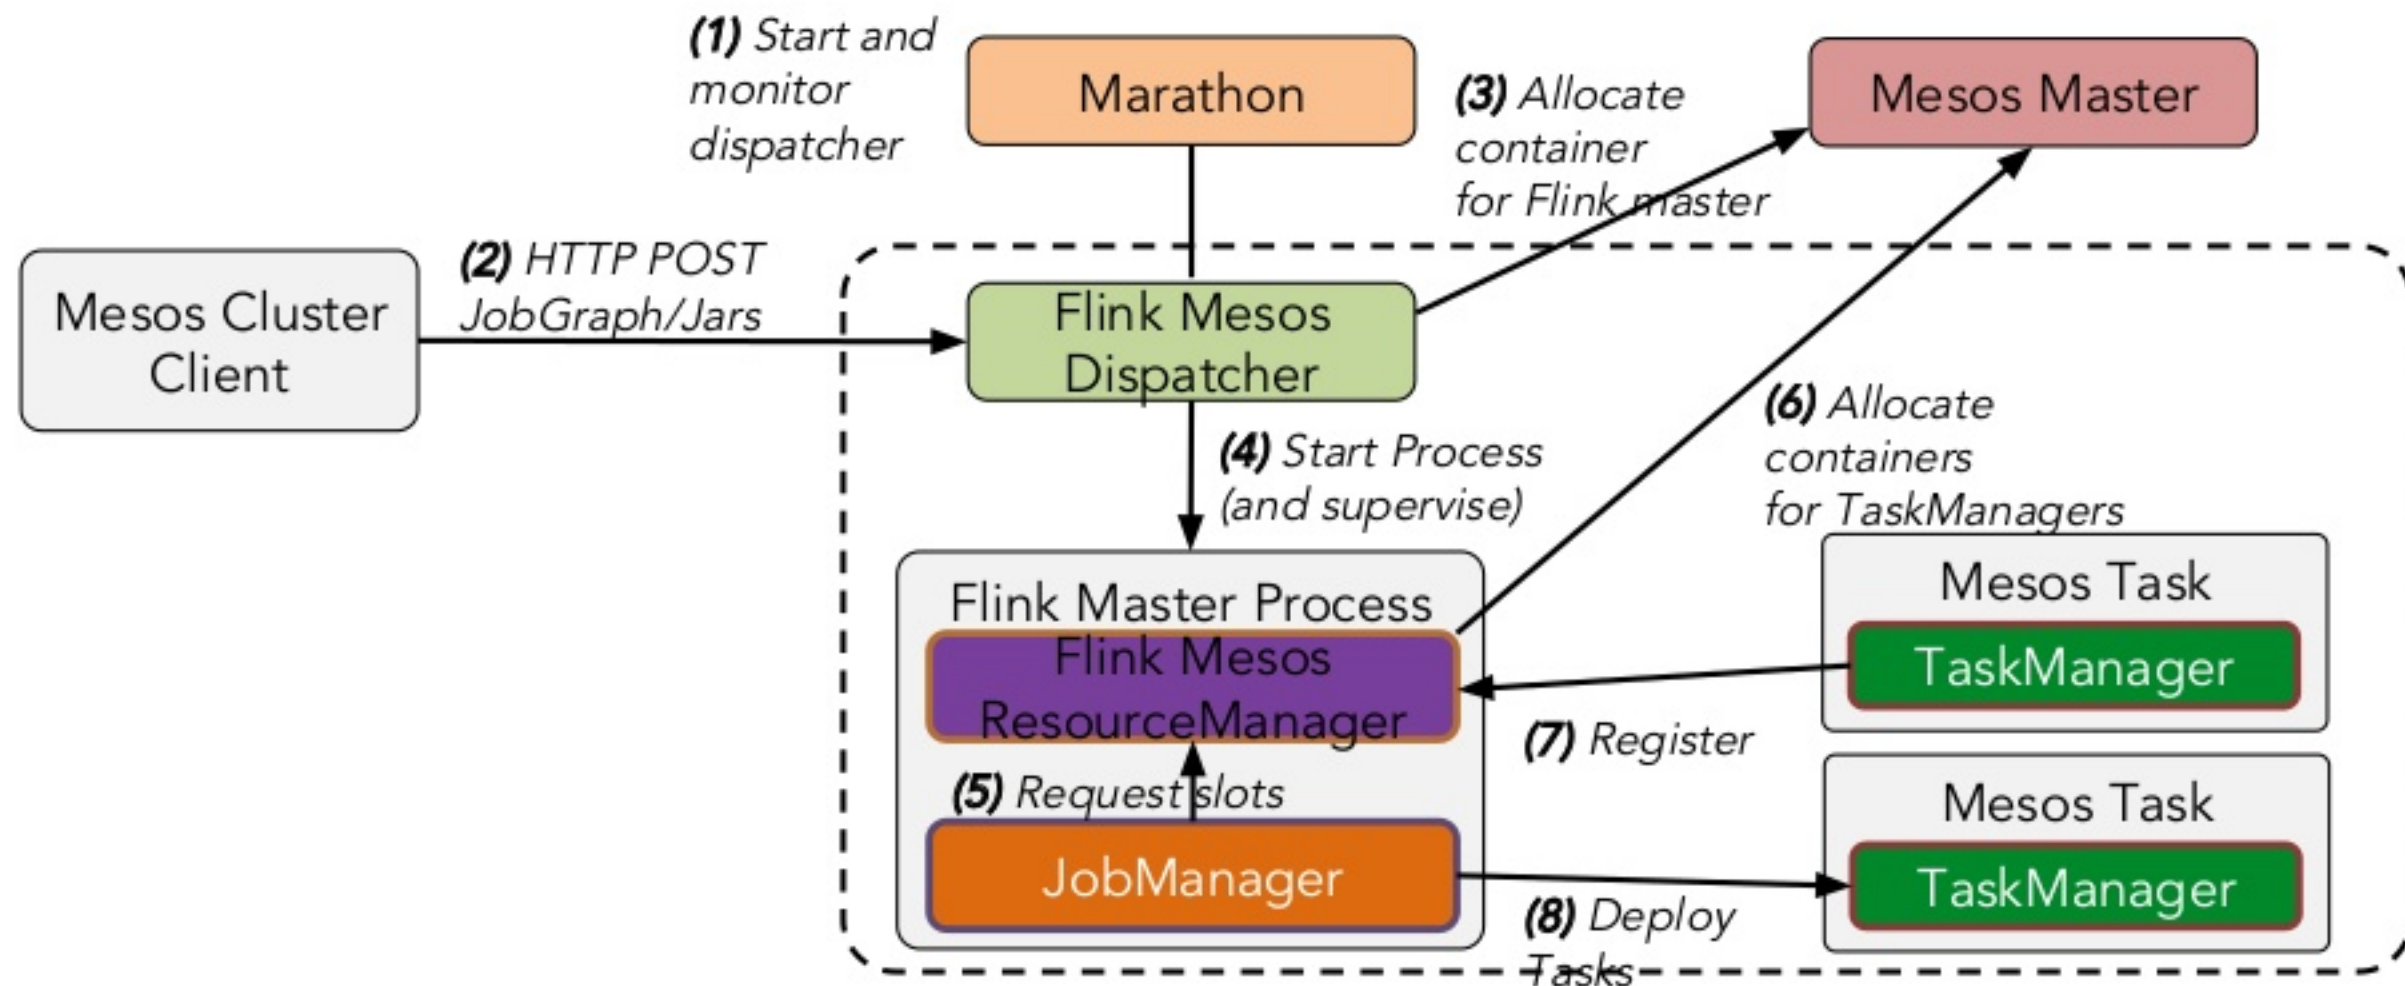
Dispatcher

- Lives across jobs
- Touch-point for job submissions
- Spawns JobManagers

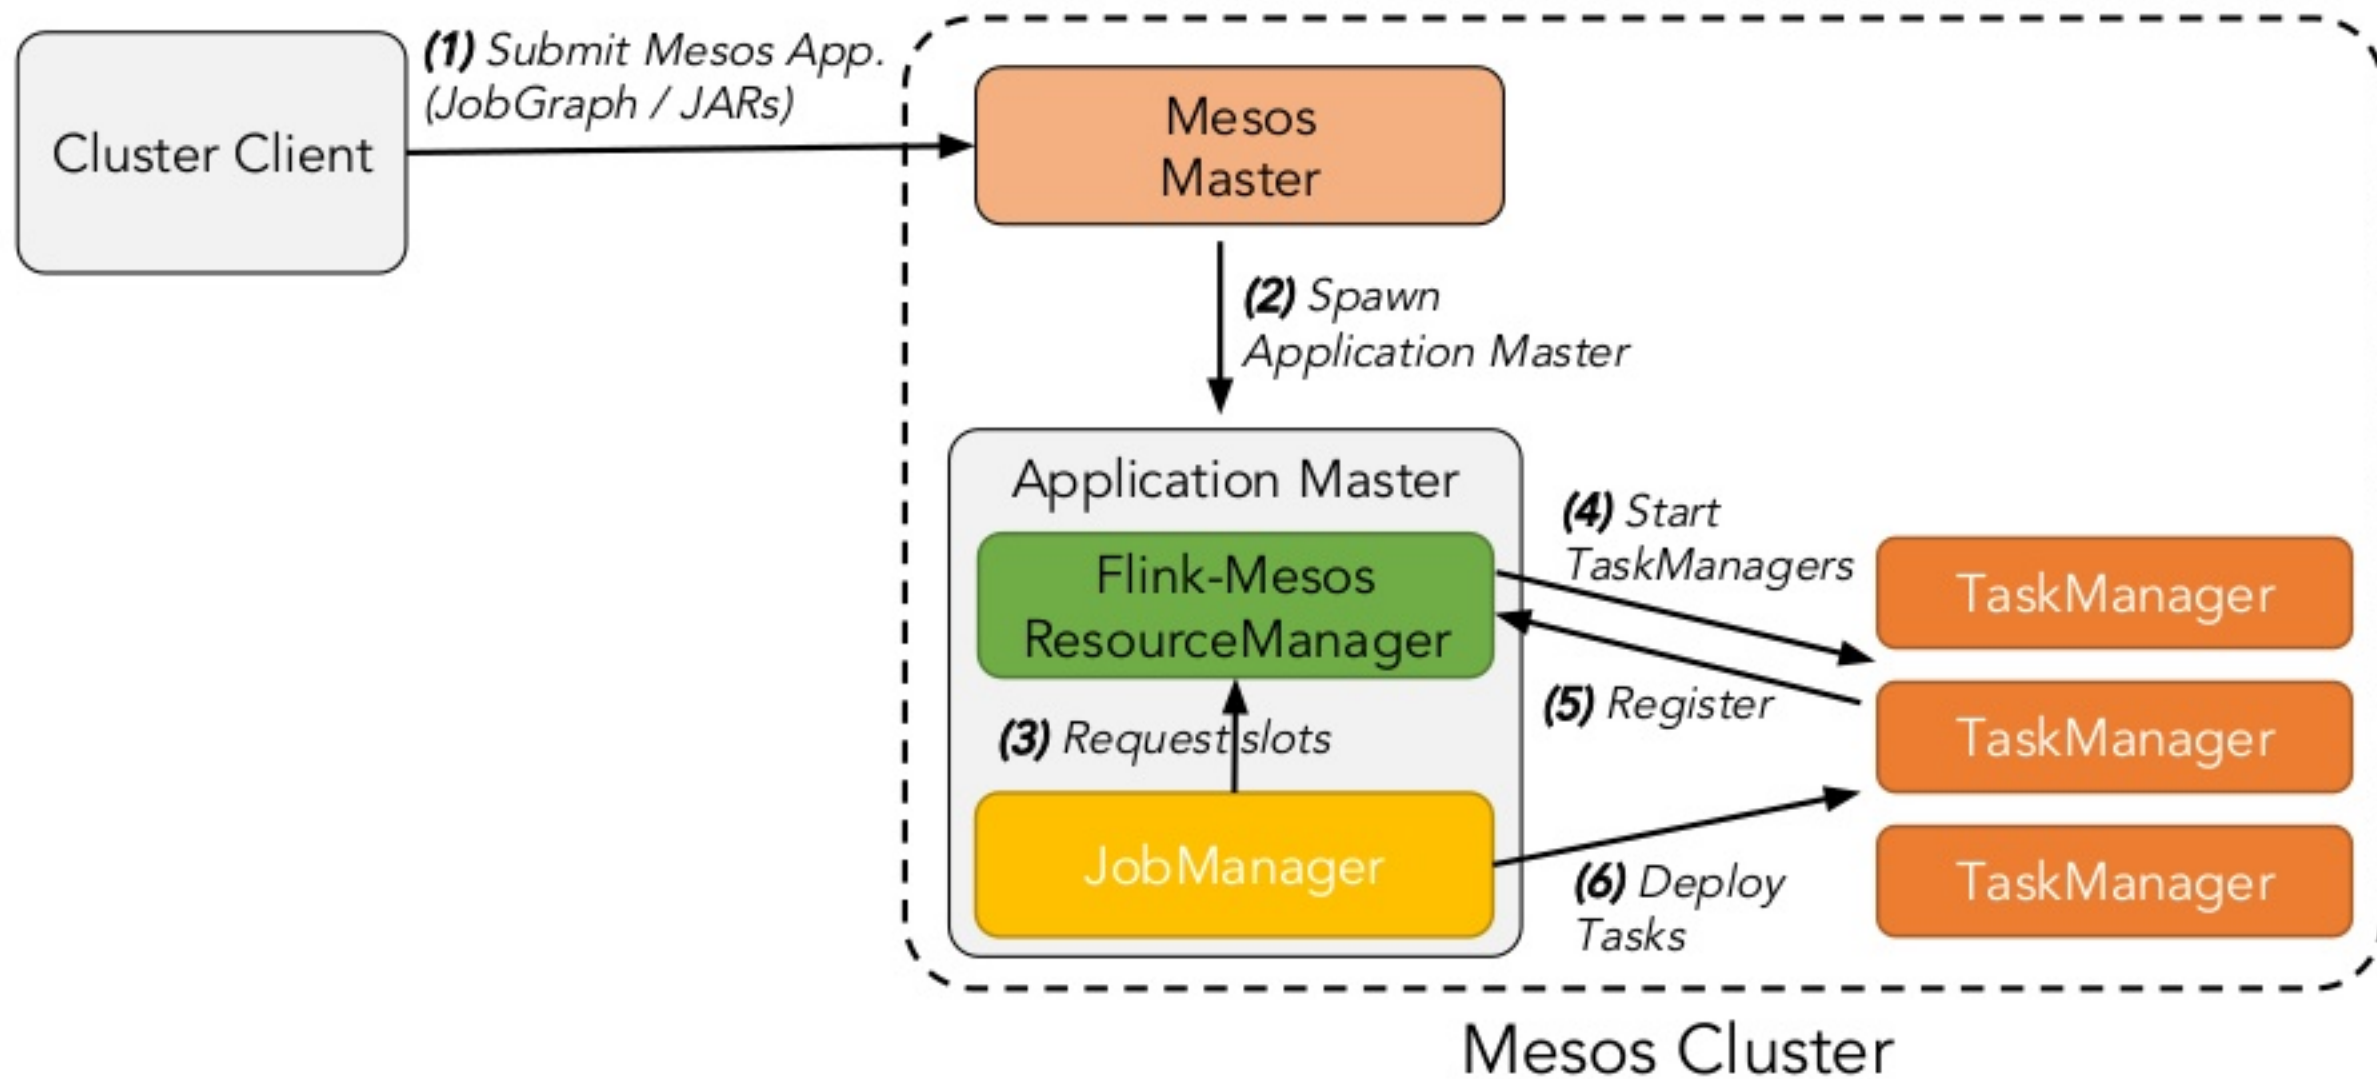
TaskManager

- Registers at ResourceManager
- Gets tasks from one or more JobManagers

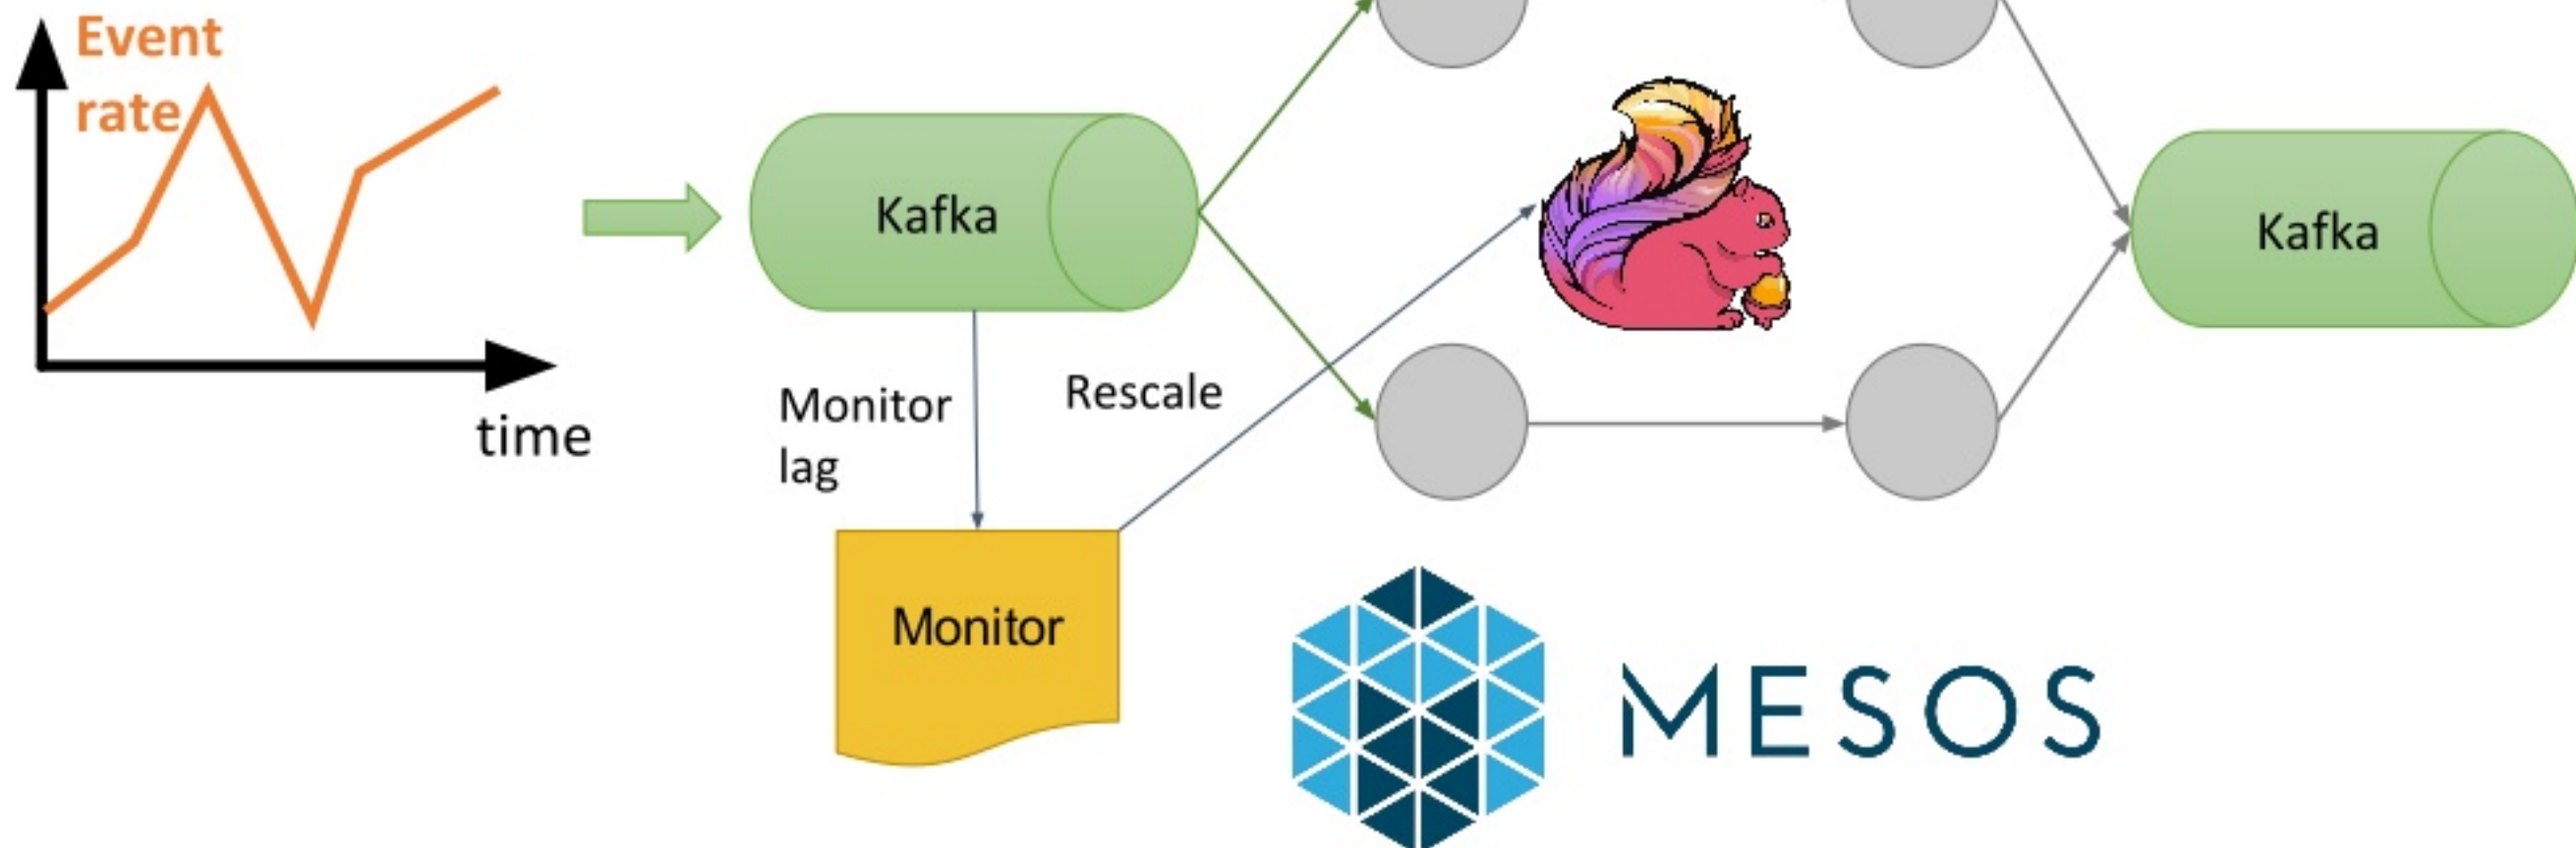
Flink Mesos Integration



Building Flink-on-Mesos (job mode)



Demo Topology



- Executed on Mesos to support dynamic resource allocation

Wrap Up

- Flink supports resource elasticity
- Flink applications can adapt to changing workloads
 - Currently manually
 - Future automatically via re-scaling policies
- Flink integrates natively with Mesos

THANK YOU!

@stsffap

@dataArtisans

@ApacheFlink

WE ARE HIRING

data-artisans.com/careers

dataArtisans