
FLINK SQL IN ACTION

TIMO WALTHER, SOFTWARE ENGINEER

FLINK FORWARD, BERLIN

SEPTEMBER 5, 2018

ABOUT DATA ARTISANS



Original creators of
Apache Flink®



PLATFORM

Open Source Apache Flink
+ dA Application Manager
+ dA Streaming Ledger



BIG APACHE FLINK SQL USERS



UBER



FLINK'S POWERFUL ABSTRACTIONS

Layered abstractions to
navigate simple to complex use cases

High-level
Analytics API

SQL / Table API (dynamic tables)

```
SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp)
FROM sensors
GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```

Stream- & Batch
Data Processing

DataStream API (streams, windows)

```
val stats = stream
  .keyBy("sensor")
  .timeWindow(Time.seconds(5))
  .sum((a, b) -> a.add(b))
```

Stateful Event-
Driven Applications

Process Function (events, state, time)

```
def processElement(event: MyEvent, ctx: Context, out: Collector[Result]) = {
  // work with event and state
  (event, state.value) match { ... }

  out.collect(...) // emit events
  state.update(...) // modify state

  // schedule a timer callback
  ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
}
```



APACHE FLINK'S RELATIONAL APIS

ANSI SQL

```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```

LINQ-style Table API

```
tableEnvironment  
    .scan("clicks")  
    .groupBy('user')  
    .select('user', 'url.count as 'cnt')
```

Unified APIs for batch & streaming data

*A query specifies exactly the same result
regardless whether its input is
static batch data or streaming data.*



QUERY TRANSLATION

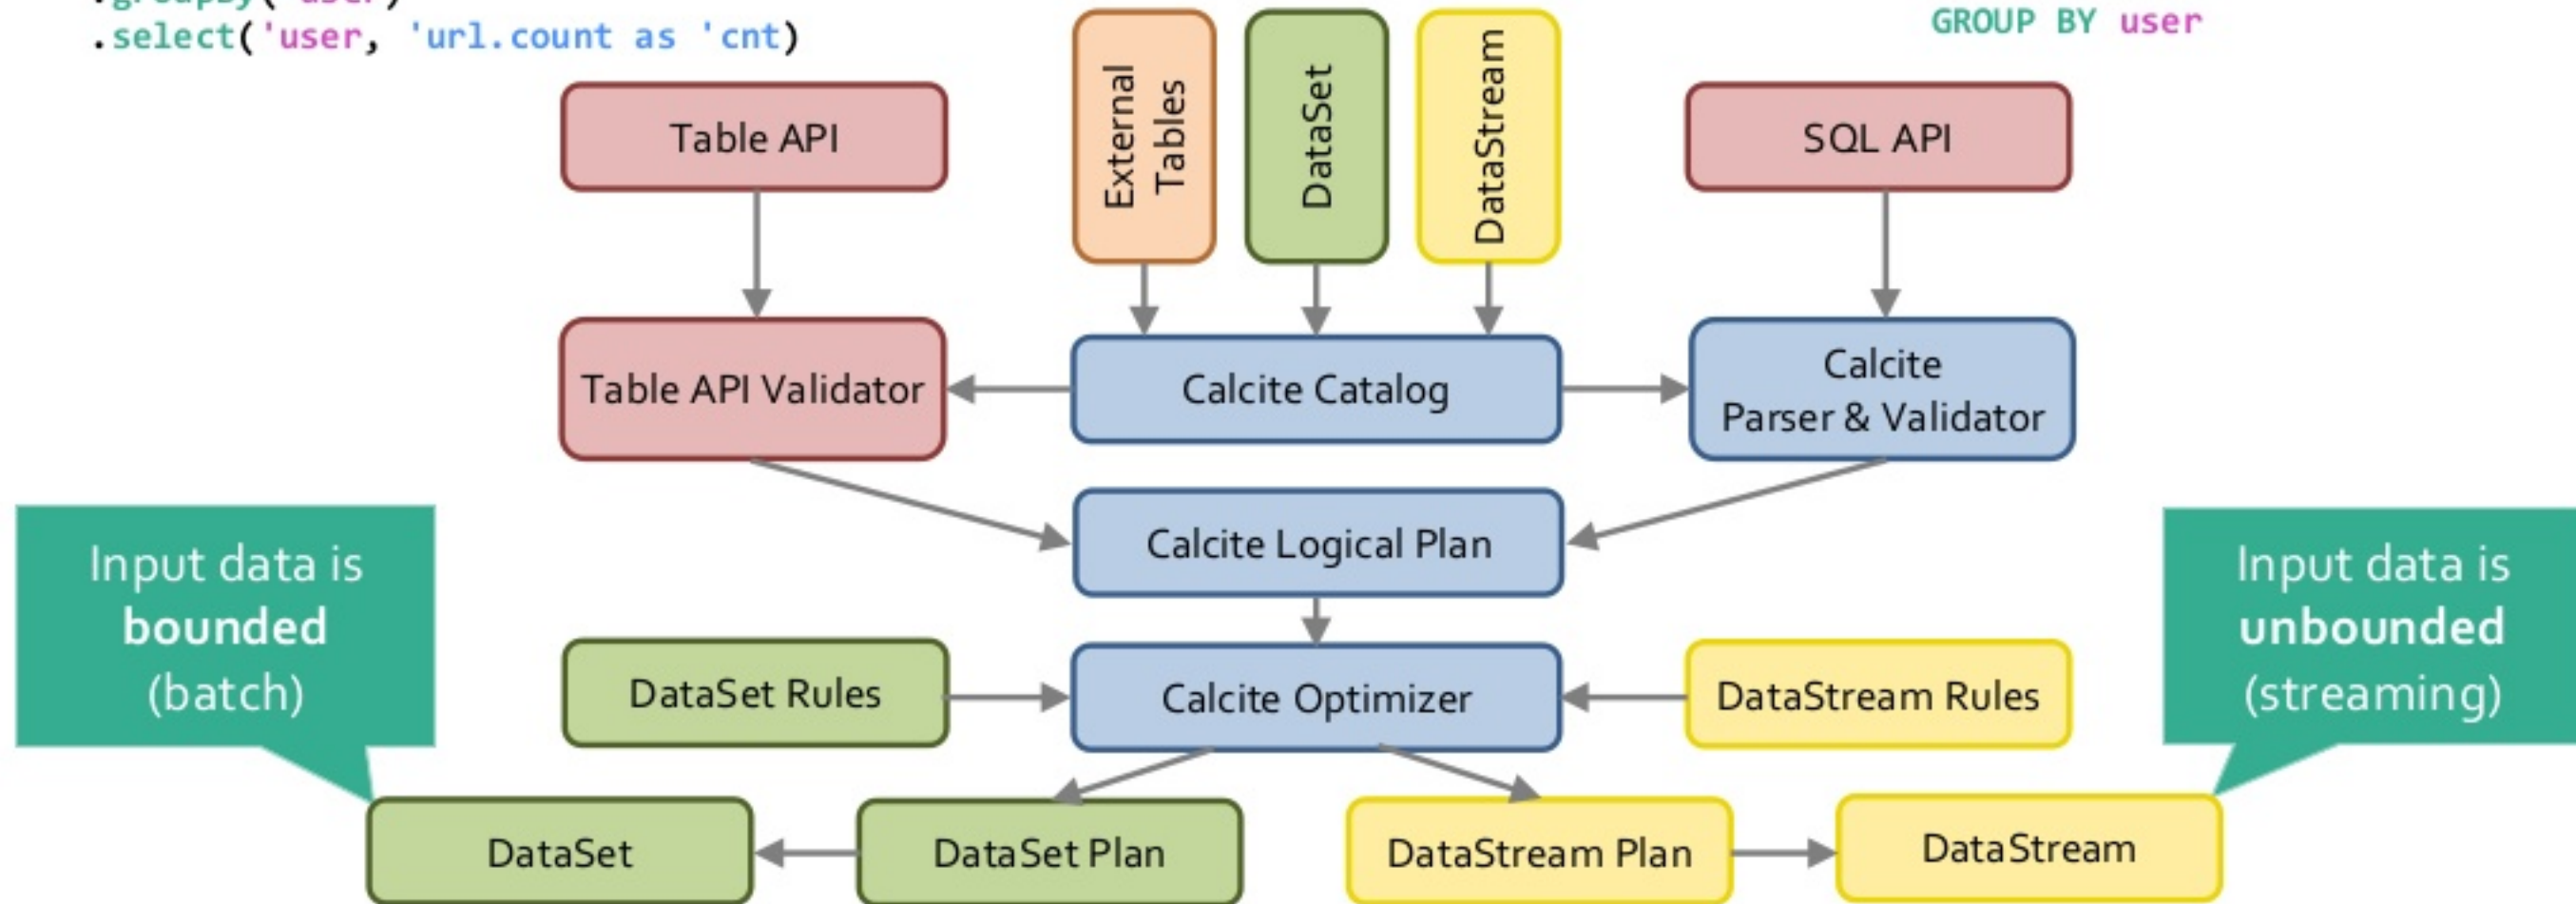
```
tableEnvironment
```

```
.scan("clicks")
```

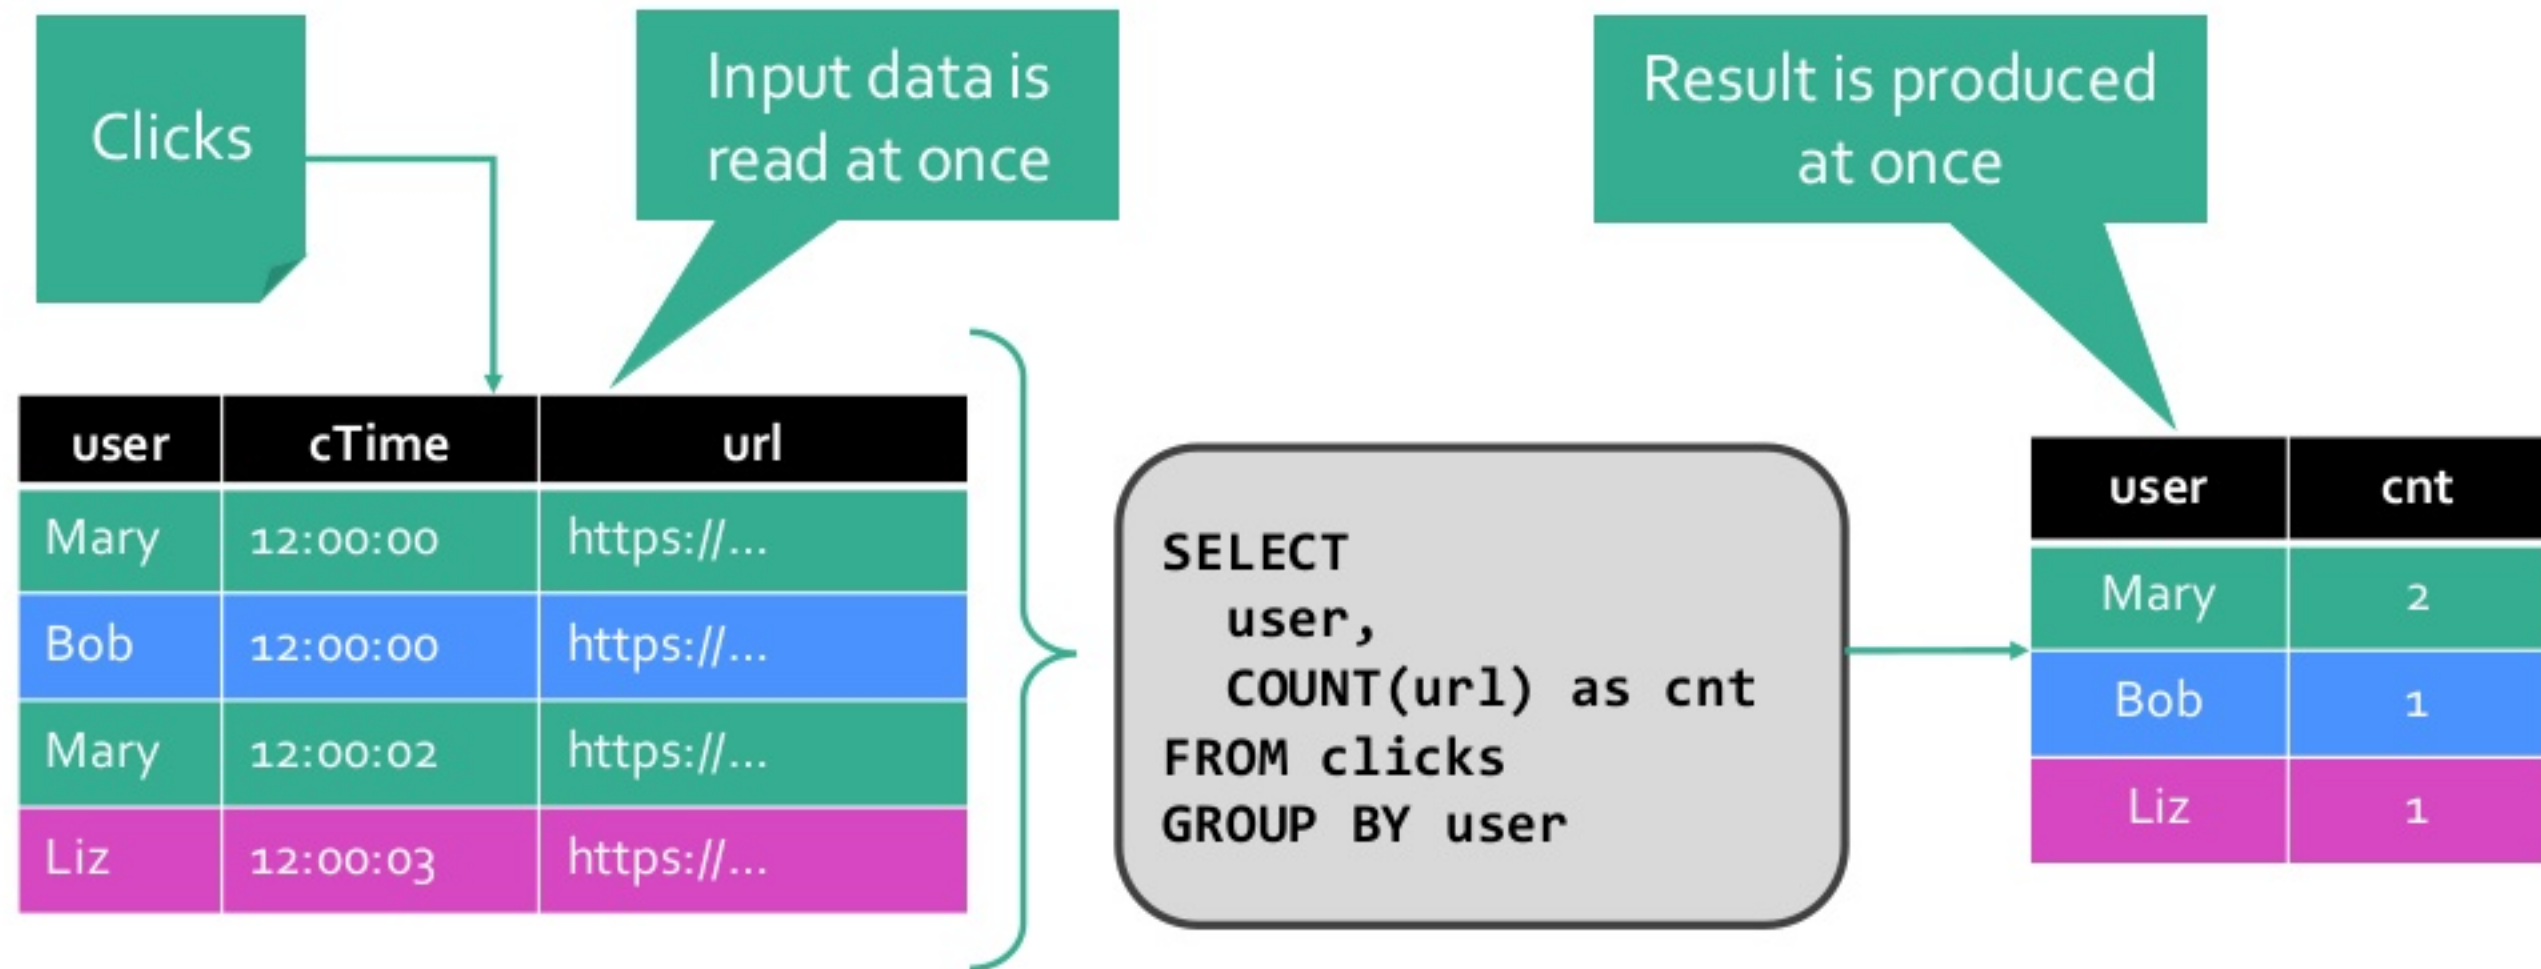
```
.groupBy('user')
```

```
.select('user', 'url.count as 'cnt')
```

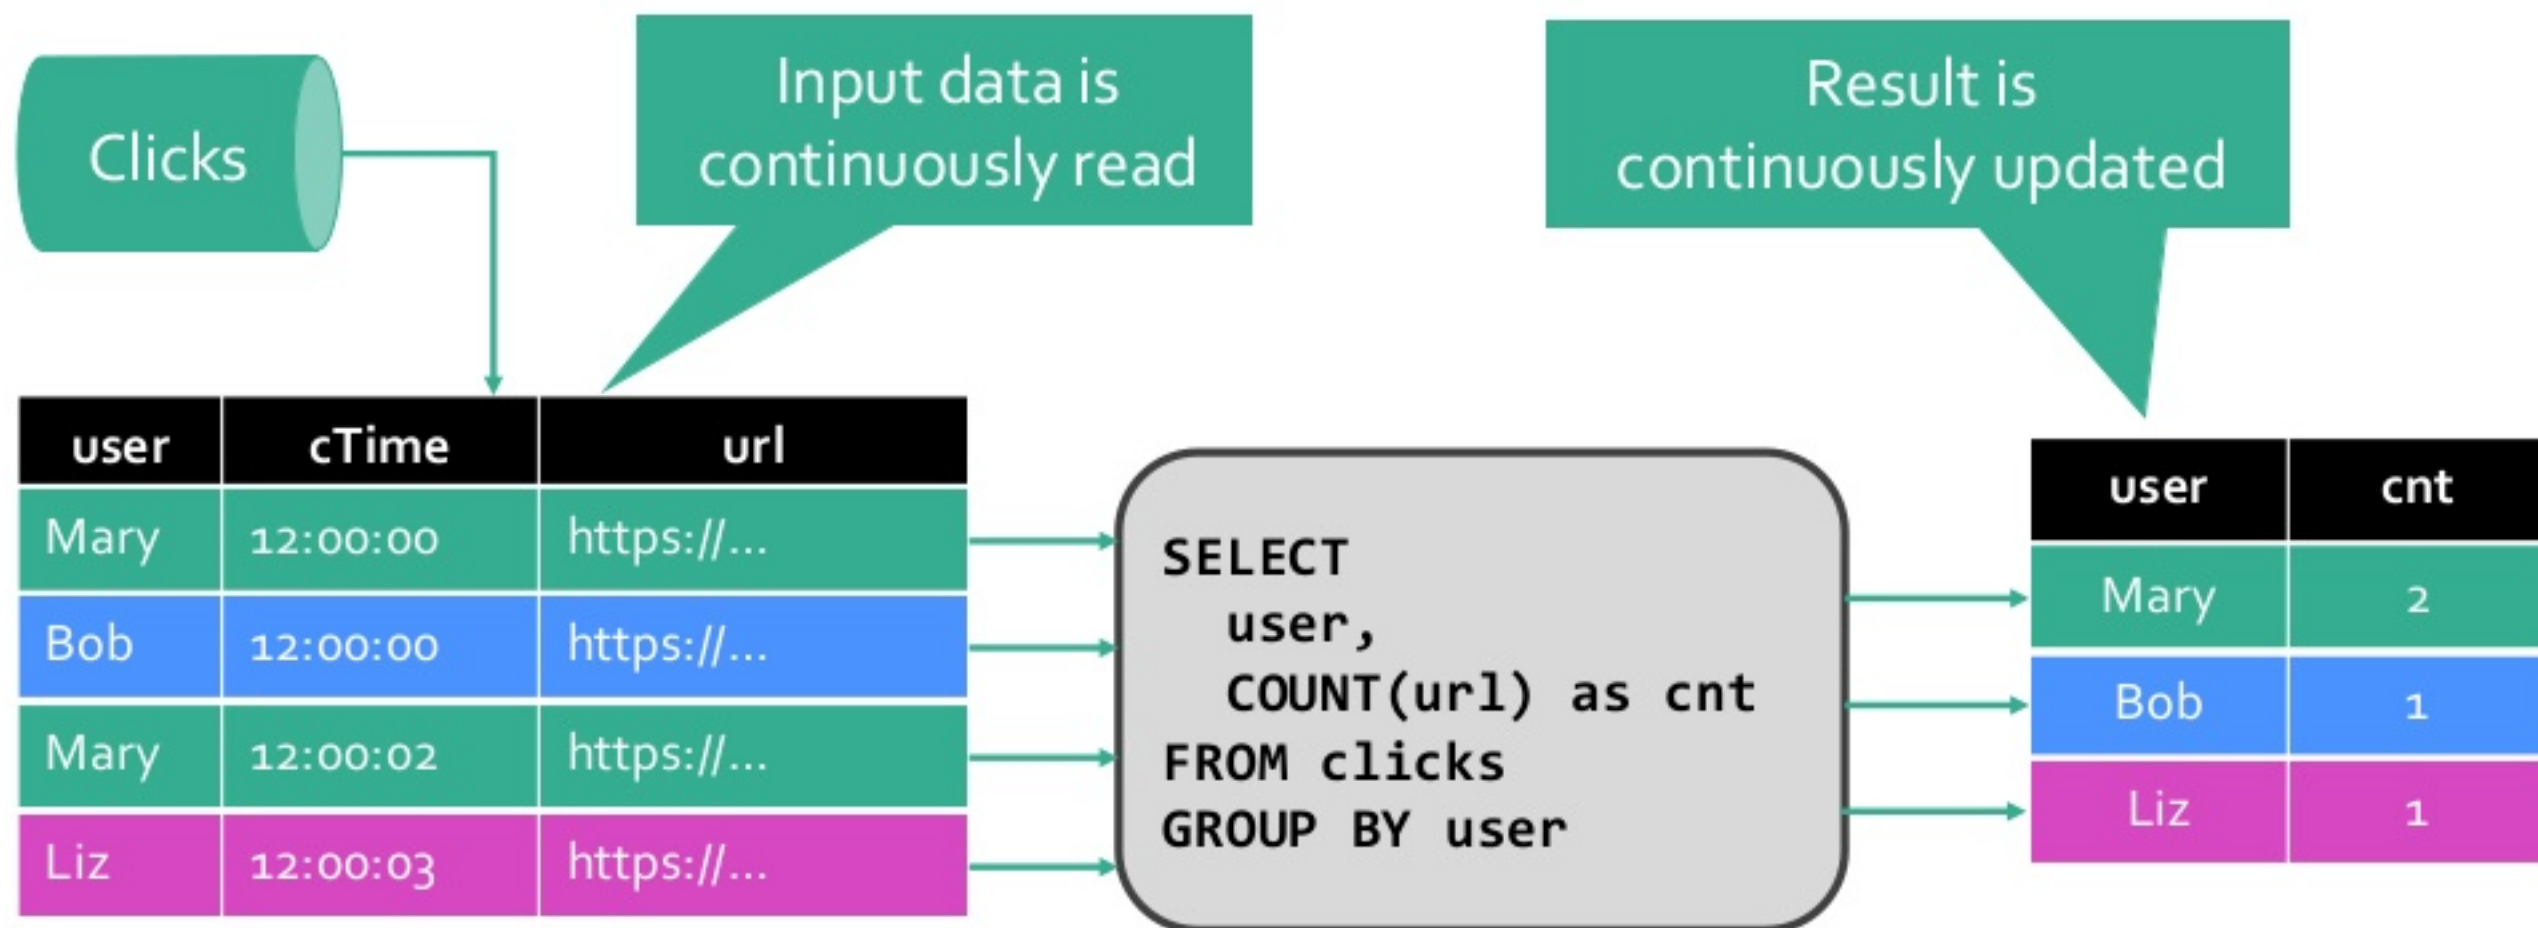
```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```



WHAT IF “CLICKS” IS A FILE?



WHAT IF “CLICKS” IS A STREAM?




The result is the same!



WHY IS STREAM-BATCH UNIFICATION IMPORTANT?

- Usability
 - ANSI SQL syntax: No custom “StreamSQL” syntax.
 - ANSI SQL semantics: No stream-specific results.
- Portability
 - Run the same query on *bounded* and *unbounded* data
 - Run the same query on *recorded* and *real-time* data



- How can we achieve SQL semantics on streams?



DATABASE SYSTEMS RUN QUERIES ON STREAMS

- Materialized views (MV) are similar to regular views, but persisted to disk or memory
 - Used to speed-up analytical queries
 - MVs need to be updated when the base tables change
- MV maintenance is very similar to SQL on streams
 - Base table updates are a stream of DML statements
 - MV definition query is evaluated on that stream
 - MV is query result and continuously updated



CONTINUOUS QUERIES IN FLINK

- Core concept is a "*Dynamic Table*"
 - Dynamic tables are changing over time
- Queries on dynamic tables
 - produce new dynamic tables (which are updated based on input)
 - do not terminate
- Stream \leftrightarrow Dynamic table conversions



STREAM ↔ DYNAMIC TABLE CONVERSIONS

- Append Conversions
 - Records are only inserted (appended)
- Upsert Conversions
 - Records are upserted/deleted
 - Records have a (composite) unique key
- Retract Conversions
 - Records are inserted/deleted

```
SELECT user, url  
FROM clicks  
WHERE url LIKE '%xyz.com'
```

```
SELECT user, COUNT(url)  
FROM clicks  
GROUP BY user
```



SQL FEATURES

SQL FEATURE SET IN FLINK 1.6.0

- SELECT FROM WHERE
- GROUP BY / HAVING
 - Non-windowed, TUMBLE, HOP, SESSION windows
- JOIN / IN
 - Windowed INNER, LEFT / RIGHT / FULL OUTER JOIN
 - Non-windowed INNER, LEFT / RIGHT / FULL OUTER JOIN
- [streaming only] OVER / WINDOW
 - UNBOUNDED / BOUNDED PRECEDING
- [batch only] UNION / INTERSECT / EXCEPT / ORDER BY



SQL FEATURE SET IN FLINK 1.6.0

- Support for POJOs, maps, arrays, and other nested types
- Large set of built-in functions (150+)
 - LIKE, EXTRACT, TIMESTAMPADD, FROM_BASE64, MD5, STDDEV_POP, AVG, ...
- Support for custom UDFs (scalar, table, aggregate)

See also:

<https://ci.apache.org/projects/flink/flink-docs-master/dev/table/functions.html>

<https://ci.apache.org/projects/flink/flink-docs-master/dev/table/udfs.html>



UPCOMING SQL FEATURES

- Streaming enrichment joins (Temporal joins) [FLINK-9712]


```
SELECT
    SUM(o.amount * r.rate) AS amount
FROM
    Orders AS o,
    LATERAL TABLE (Rates(o.rowtime)) AS r
WHERE r.currency = o.currency;
```

- Support for complex event processing (CEP) [FLINK-6935]
 - MATCH_RECOGNIZE
- More connectors and formats [FLINK-8535]



WHAT CAN I BUILD WITH THIS?

- Data Pipelines
 - Transform, aggregate, and move events in real-time
- Low-latency ETL
 - Convert and write streams to file systems, DBMS, K-V stores, indexes, ...
 - Ingest appearing files to produce streams
- Stream & Batch Analytics
 - Run analytical queries over bounded and unbounded data
 - Query and compare historic and real-time data
- Power Live Dashboards
 - Compute and update data to visualize in real-time



SQL CLIENT BETA

dataArtisans

INTRODUCTION TO SQL CLIENT

- Newest member of the Flink SQL family (since Flink 1.5)



The screenshot displays the Flink SQL Client interface. On the left, a dark blue terminal window shows the Flink SQL Client logo and a welcome message: "Welcome! Enter HELP to list all available commands. QUIT to exit." Below this, the prompt "Flink SQL>" is visible. On the right, a table of data is displayed, titled "SQL Query Results (Table)" and "Page: 1 of 1 (100)". The table has 10 columns: "id", "name", "age", "gender", "height", "weight", "blood_pressure", "sugar", "cholesterol", and "vision". The data is presented in a grid format with 10 rows and 10 columns. The bottom of the interface shows navigation controls: "Next Page", "Previous Page", "First Page", "Last Page", and "Refresh".

| id | name | age | gender | height | weight | blood_pressure | sugar | cholesterol | vision |
|----|--------------|-----|--------|--------|--------|----------------|-------|-------------|--------|
| 1 | John Doe | 35 | Male | 175 | 70 | 120/80 | 100 | 200 | 20/20 |
| 2 | Jane Smith | 28 | Female | 160 | 55 | 110/70 | 90 | 180 | 20/20 |
| 3 | Mike Johnson | 42 | Male | 180 | 85 | 130/90 | 120 | 220 | 20/20 |
| 4 | Sarah Brown | 31 | Female | 165 | 60 | 115/75 | 95 | 190 | 20/20 |
| 5 | David Wilson | 45 | Male | 170 | 75 | 125/85 | 110 | 210 | 20/20 |
| 6 | Emily Davis | 25 | Female | 155 | 50 | 105/65 | 85 | 170 | 20/20 |
| 7 | Chris Miller | 38 | Male | 172 | 68 | 118/78 | 98 | 185 | 20/20 |
| 8 | Alice Taylor | 33 | Female | 162 | 58 | 112/72 | 92 | 188 | 20/20 |
| 9 | Bob Anderson | 40 | Male | 178 | 72 | 122/82 | 105 | 205 | 20/20 |
| 10 | Grace White | 29 | Female | 158 | 52 | 108/68 | 88 | 175 | 20/20 |



INTRODUCTION TO SQL CLIENT

- Goal: Flink without a single line of code
 - only SQL and YAML
 - *"drag&drop"* SQL JAR files for connectors and formats
- Build on top of Flink's Table & SQL API
- Useful for prototyping & submission



SQL CLIENT CONFIGURATION

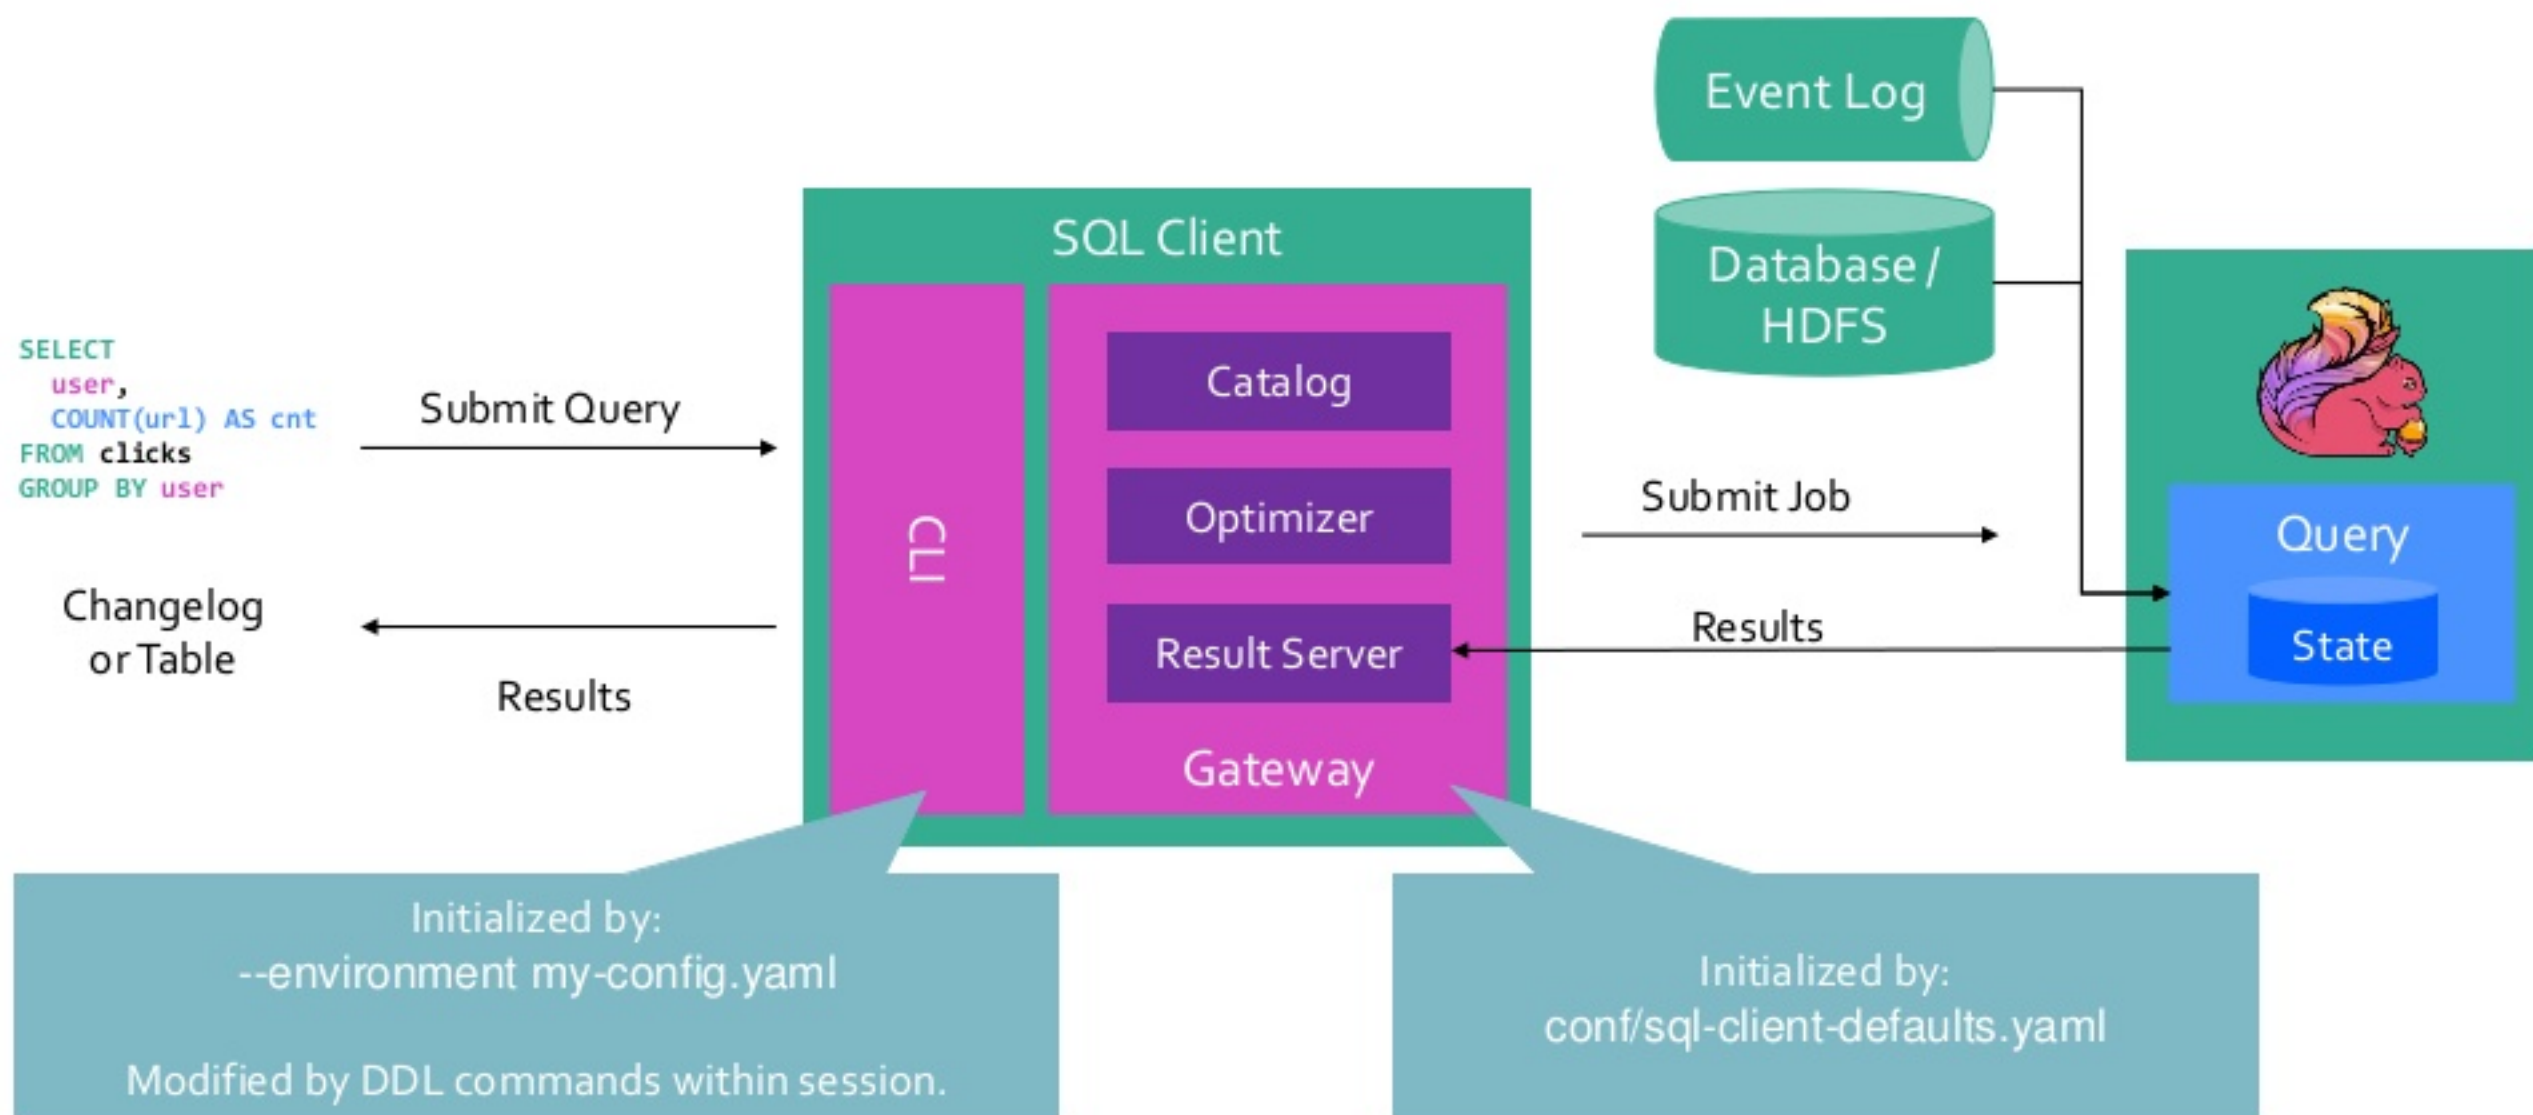
```
1 # Define table sources and sinks here.
2 tables:
3   -- name: MyTableSource
4   -- type: source
5   -- update-mode: append
6   -- connector:
7     -- type: filesystem
8     -- path: "/path/to/something.csv"
9   -- format:
10     -- type: csv
11   -- fields:
12     -- name: MyField1
13     -- type: INT
14     -- name: MyField2
15     -- type: VARCHAR
16   -- line-delimiter: "\n"
17   -- comment-prefix: "#"
18   -- schema:
19     -- name: MyField1
20     -- type: INT
21     -- name: MyField2
22     -- type: VARCHAR
23
24 # Define table views here.
25 views:
26   -- name: MyCustomView
27   -- query: "SELECT MyField2 FROM MyTableSource"
28
29 # Define user-defined functions here.
30 functions:
31   -- name: myUDF
32   -- from: class
33   -- class: foo.bar.AggregateUDF
34
35 # Execution properties allow for changing the behavior of a table program.
36 execution:
37   -- type: streaming -----# required: execution mode either 'batch' or 'streaming'
38   -- result-mode: table -----# required: either 'table' or 'changelog'
39   -- parallelism: 1 -----# optional: Flink's parallelism (1 by default)
```

See also:

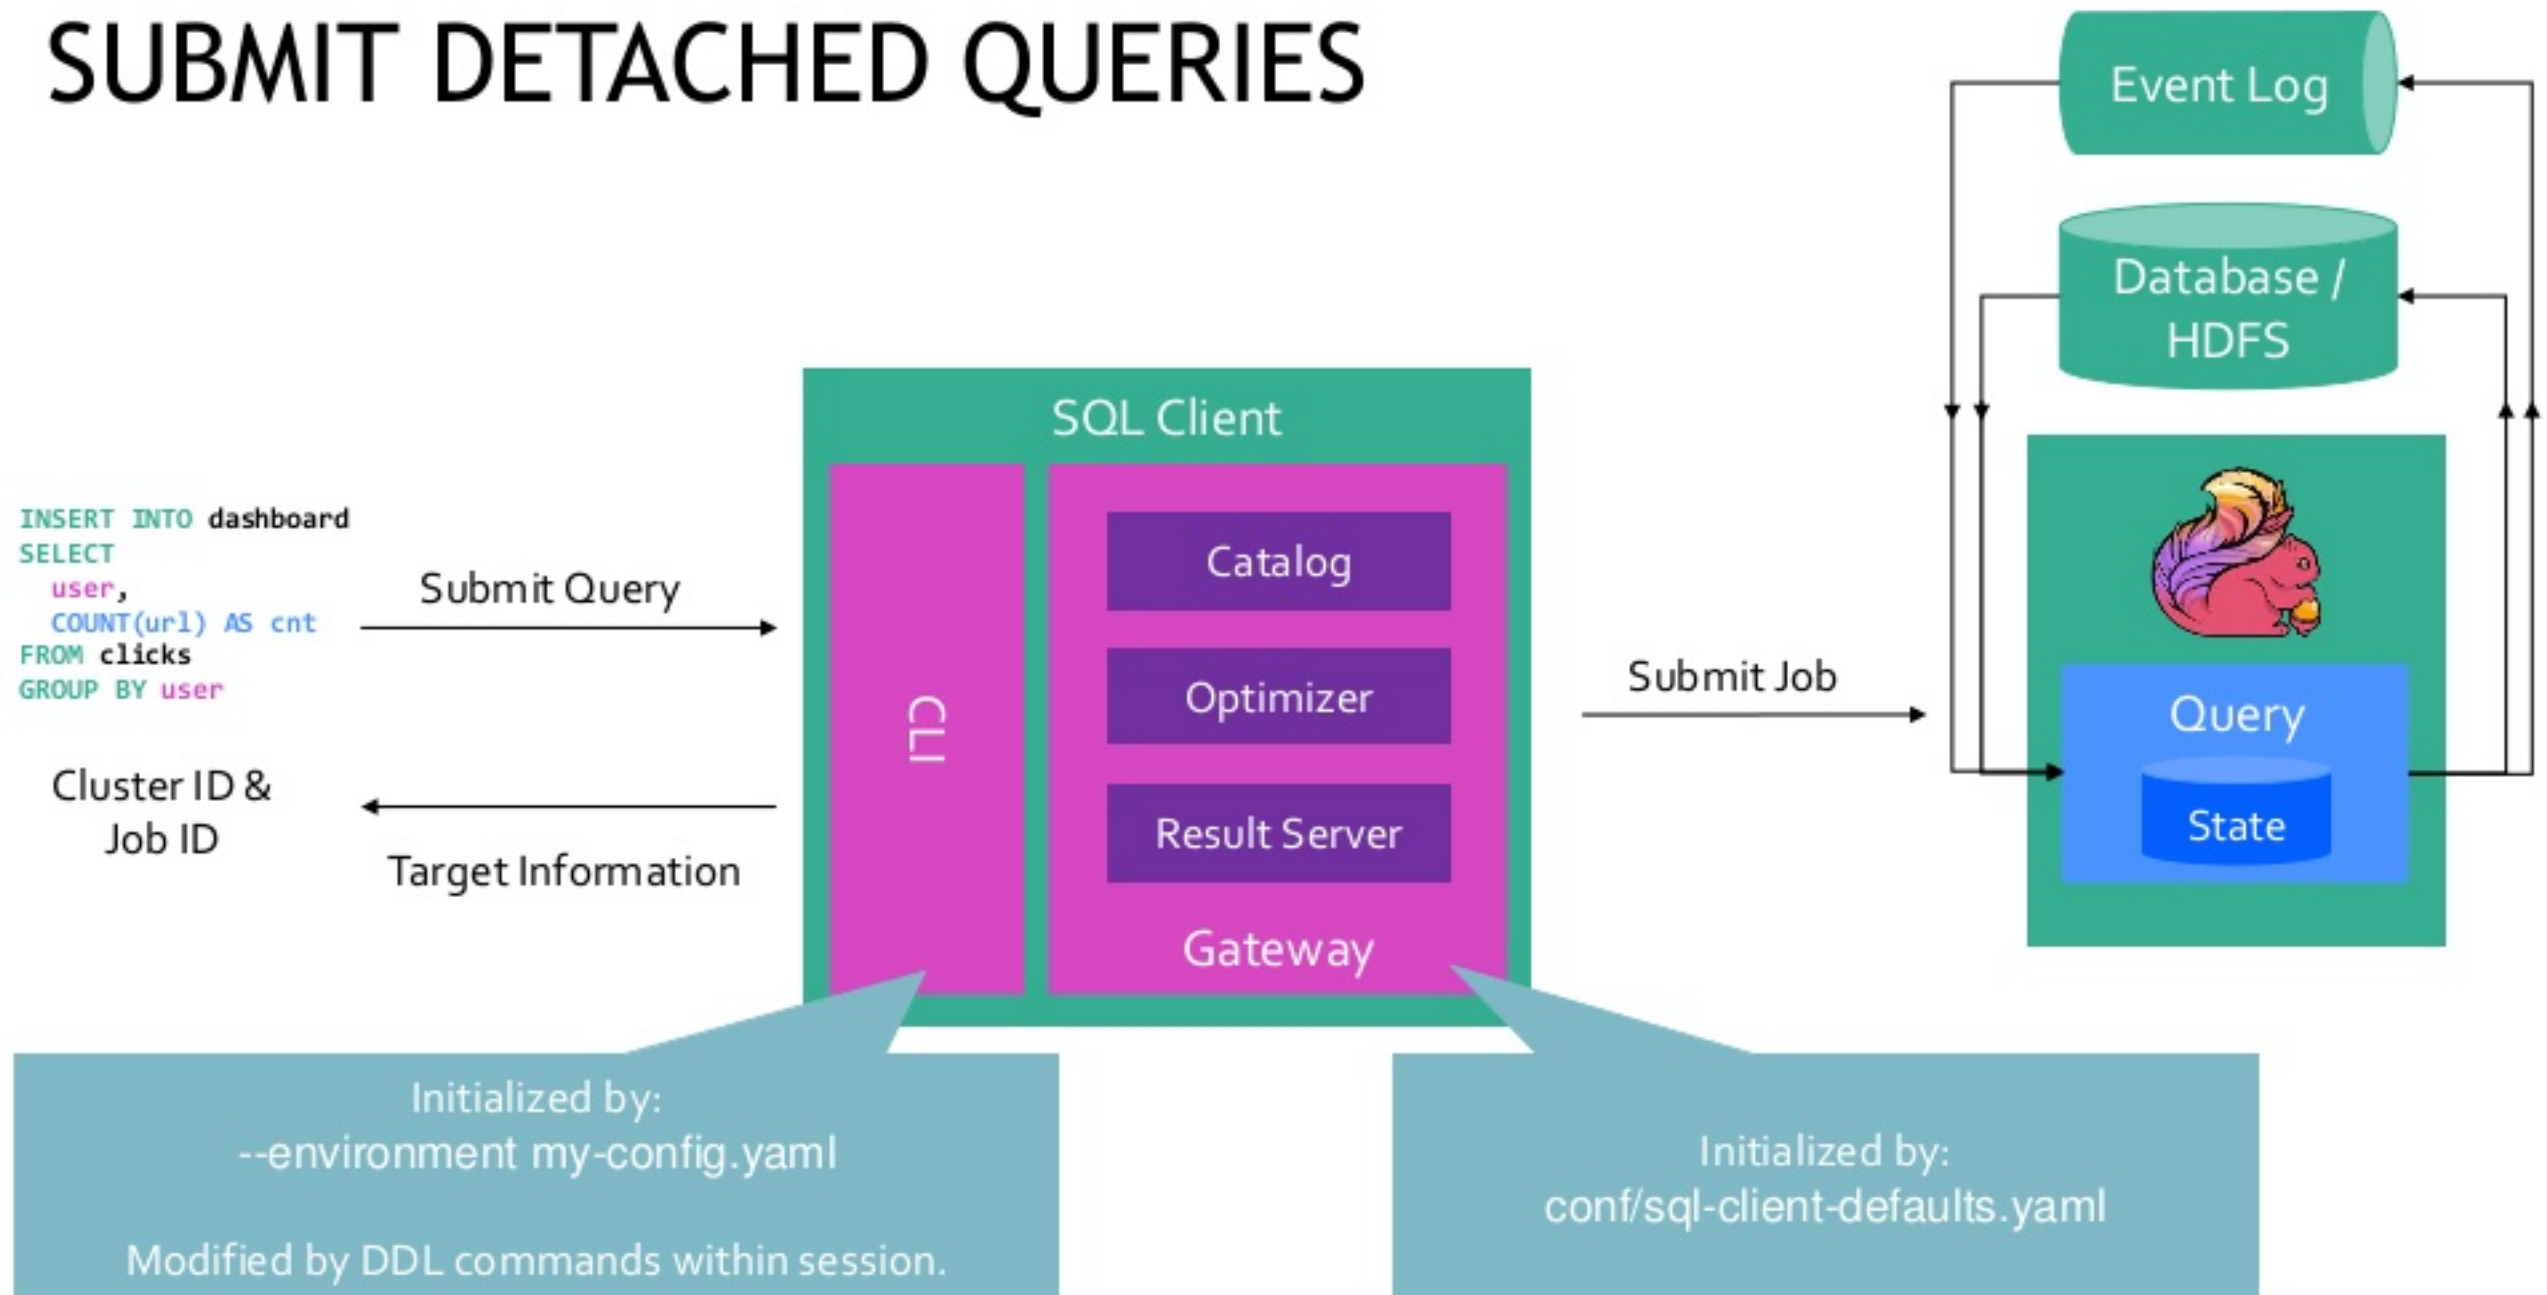
<https://ci.apache.org/projects/flink/flink-docs-master/dev/table/sqlClient.html>



PLAY AROUND WITH FLINK SQL



SUBMIT DETACHED QUERIES



ACTION TIME!

[HTTPS://GITHUB.COM/DATAARTISANS/SQL-TRAINING](https://github.com/DataArtisans/SQL-Training)

dataArtisans

SUMMARY

- Unification of stream and batch is important.
- Flink's SQL solves many streaming and batch use cases.
- Runs in production at Alibaba, Uber, and others.
- The community is working on improving user interfaces.
- Get involved, discuss, and contribute!




THANK YOU!

@twalthr

@dataArtisans

@ApacheFlink



Available on O'Reilly Early Release!

WE ARE HIRING

data-artisans.com/careers

dataArtisans