

Lessons learnt from Migrating to a Stateful Streaming Framework

Wei-Che(Tony) Wei

Software Engineer , Appier

Appier

Outline

- **Use case in Appier**
- **Moving from micro batch to true streaming**
- **Tips to conquer obstacles during the migration**
- **Challenge from design in streaming way**

Outline

- **Use case in Appier**
- Moving from micro batch to true streaming
- Tips to conquer obstacles during the migration
- Challenge from design in streaming way

About Appier

Appier is a technology company which aims to provide artificial intelligence platforms to help enterprises solve their most challenging business problems.

<http://www.appier.com/en/index.html>






Use case

- Retargeting
- Fraud detection
- Real time recommendation
- ...

Outline

- Use case in Appier
- **Moving from micro batch to true streaming**
 - Disadvantage from micro batch
 - New requirements
 - Our solution by using Flink
- Tips to conquer obstacles during the migration
- Challenge from design in streaming way

Dynamic Rule Service: Previous design




Disadvantage from micro batch

- High latency
- Struggle with back pressure
- Hard to maintain states

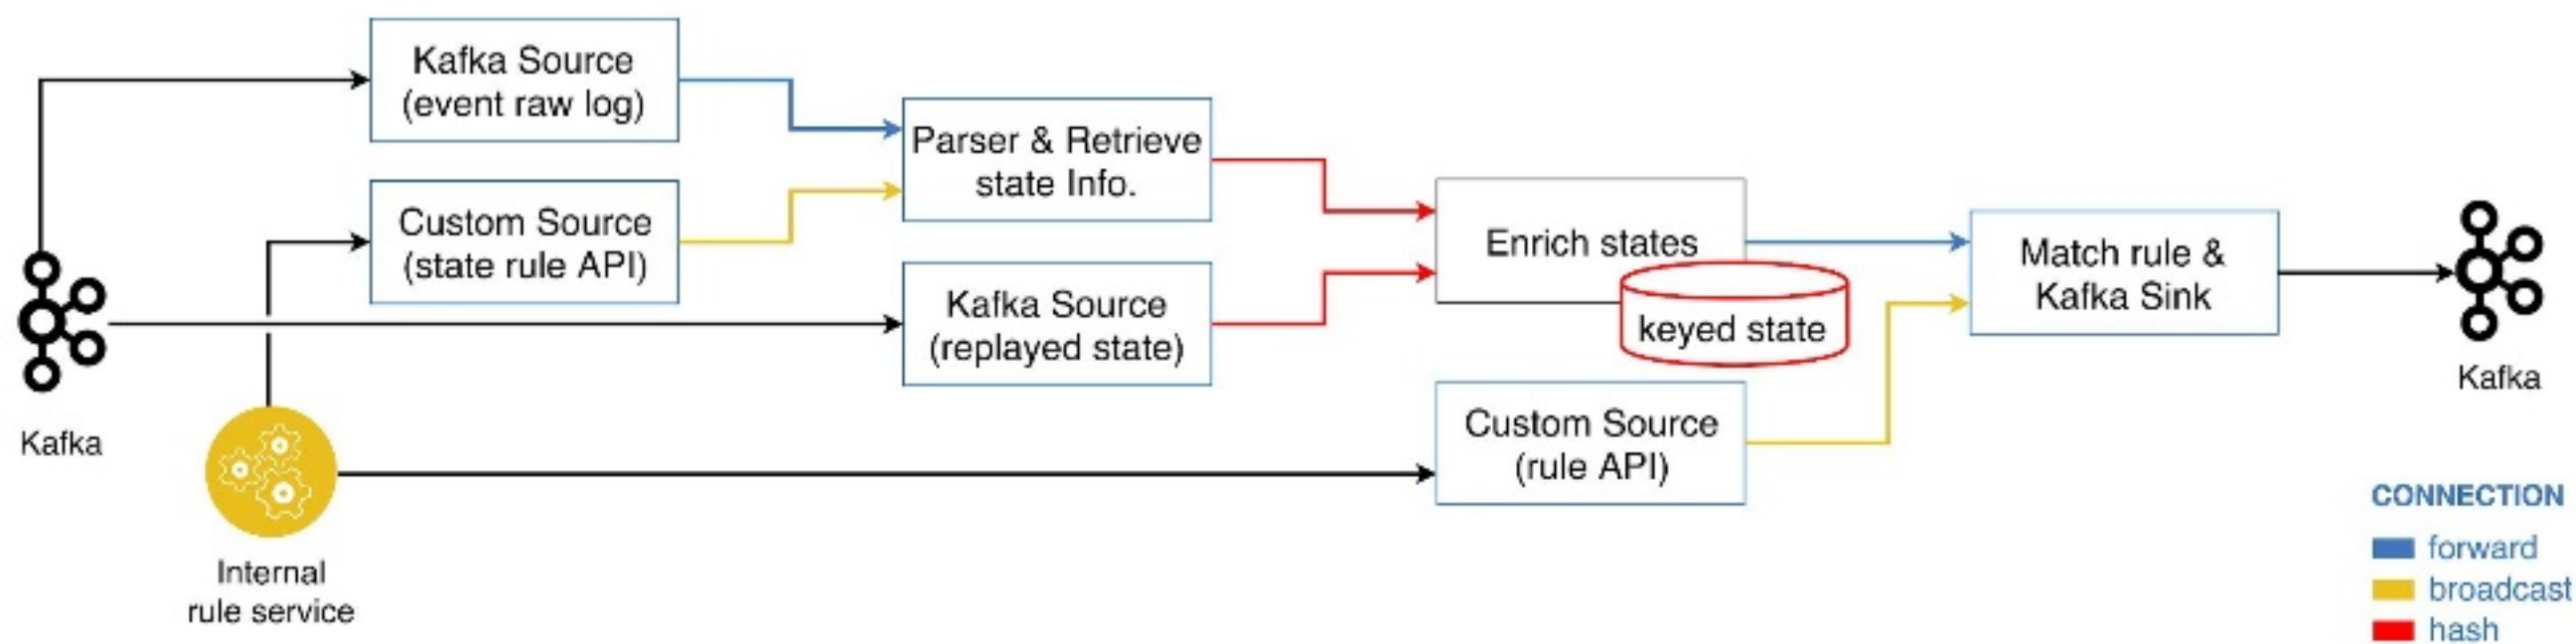
New requirements

- “Who visited this site twice in the past week”
- “Who viewed this product in the site but didn’t put into cart”
- “Who visited the certain pattern of some pages and bought this product”

Dynamic Rule Service with Apache Flink



Detail of JobGraph (Apache Flink 1.4.0 release)



Improvements

- **Support more needs for stateful rules**
- **Flexible architecture**
- **Performance efficiency**
- **Cost efficiency**

Outline


- Use case in Appier
- Moving from micro batch to true streaming
- **Tips to conquer obstacles during the migration**
- Challenge from design in streaming way

Tip 1:

Document and Mailing List are your best friends

Case Study:
taskmanager.exit-on-fatal-akka-error: true

<https://ci.apache.org/projects/flink/flink-docs-release-1.4/ops/config.html#jobmanager-taskmanager>



Kafka Consumer Throughput



buffers input Queue Length



buffers output Queue Length







Tip 2:
Monitor and alert
are important



Tip 2:
Monitor and alert
are important



Tip 2:
Monitor and alert
are important



Tip 2:
Monitor and alert
are important

FlinkCheckpointAlignmentTimeOver30Secs (0 active)

```

alert: FlinkCheckpointAlignmentTimeOver30Secs
expr: max(flink_taskmanager_job_task_checkpointAlignmentTime)
      BY (job_name, job_id, task_attempt_num, task_name) / 1000 / 1000 / 1000 > 30
for: 100m
labels:
  receiver: rt-alert
annotations:
  description: 'Flink Job: {{labels.job_name}}{{labels.job_id}} checkpoint alignment
    time takes more than 30 seconds on task {{labels.task_name}}'

```

FlinkCheckpointDuration10Mins (0 active)

FlinkCheckpointFailed (0 active)

FlinkProcessTimeOver30Secs (0 active)

```

alert: FlinkProcessTimeOver30Secs
expr: sum(increase(flink_taskmanager_job_task_operator_process_time_latency{operator_name="match_rule"}[2m]))
      BY (job_name, job_id, task_attempt_num) / sum(increase(flink_taskmanager_job_task_operator_numRecordsIn{operator_name="match_rule"}[2m]))
      BY (job_name, job_id, task_attempt_num) > 30000
for: 10m
labels:
  receiver: rt-flink-alert
annotations:
  description: 'Flink Job: {{labels.job_name}}{{labels.job_id}} process time takes
    more than 30 seconds for 10 minutes.'

```

FlinkTaskManagerHighGC (0 active)

```


alert: FlinkTaskManagerHighGC
expr: max(increase(flink_taskmanager_Status_JVM_GarbageCollector_G1_Old_Generation_Time[10m])
      + increase(flink_taskmanager_Status_JVM_GarbageCollector_G1_Young_Generation_Time[10m]))
      WITHOUT (tm_id) / 10 / 60 / 1000 > 0.03
for: 30m
labels:
  receiver: rt-flink-alert
annotations:
  description: Flink Taskmanager instance {{labels.instance}} high GC ratio

```

Tip 2:

Monitor and alert

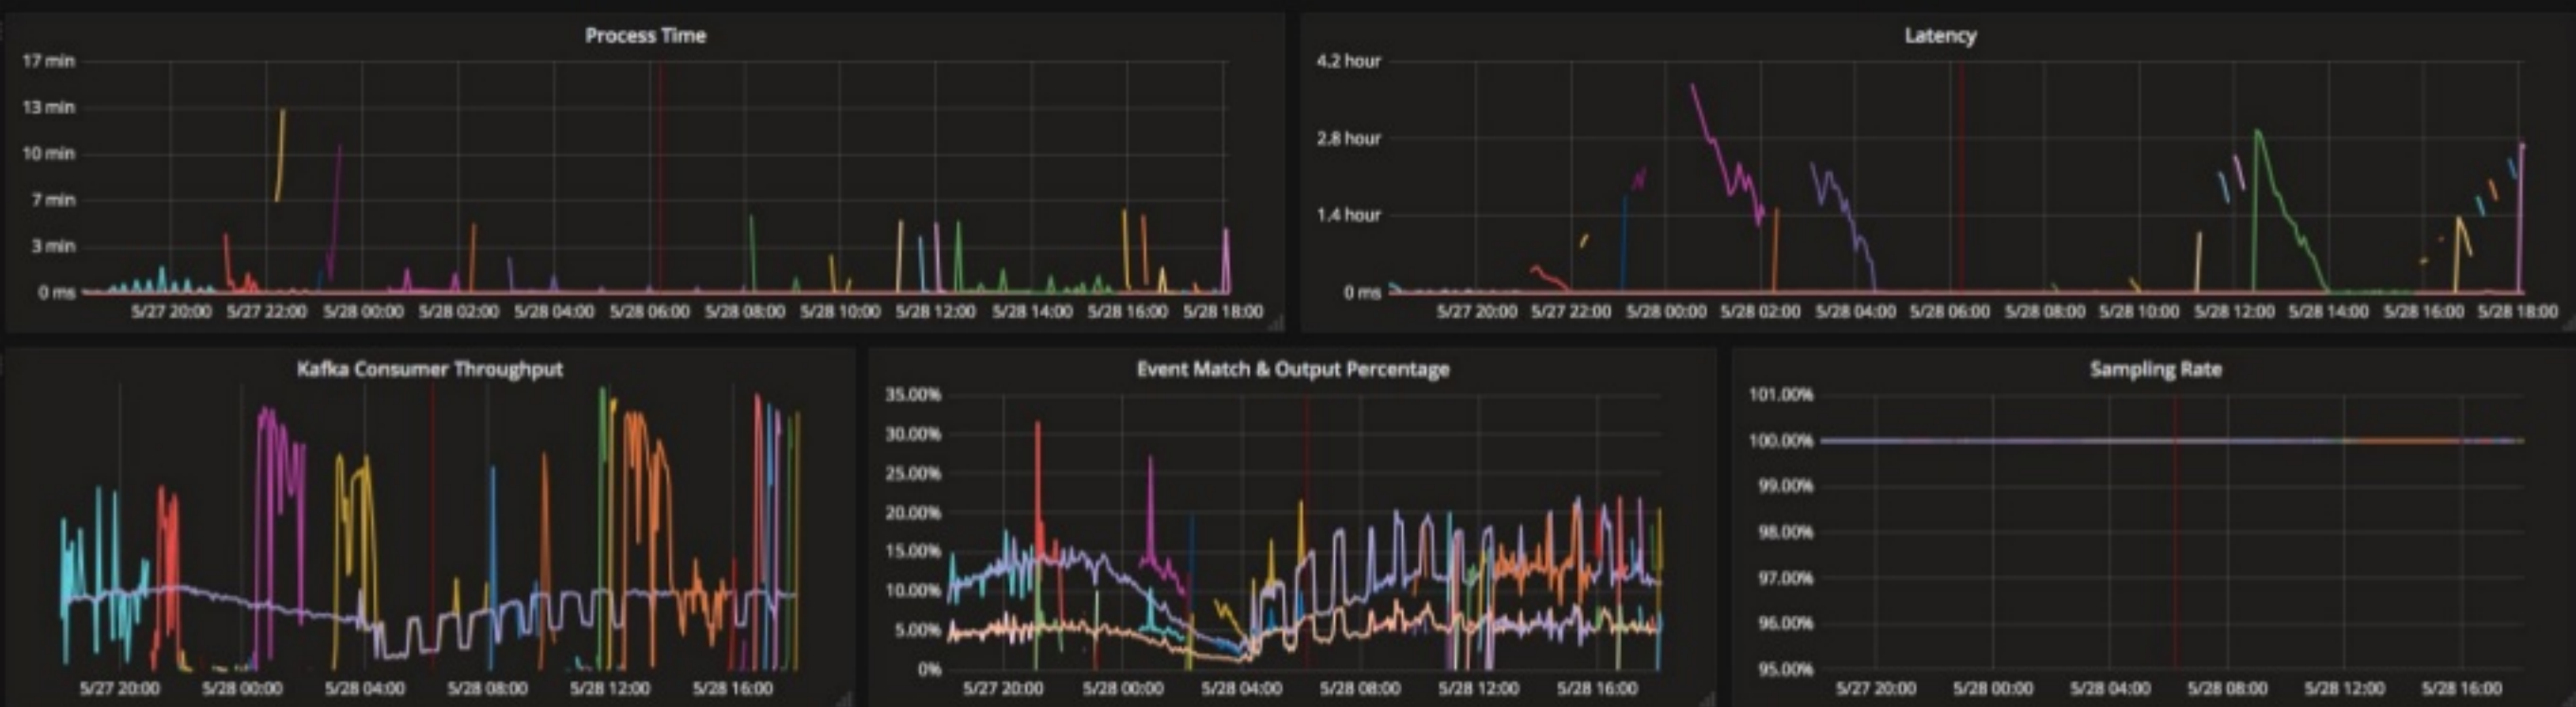
are important



Tip 3:
**Be familiar with
your environment
and your job**

Case Study:

[checkpoint stuck with rocksdb statebackend and s3 filesystem](#)



Tip 3:
**Be familiar with
your environment
and your job**

Case Study: Verify bottleneck of our streaming job

- ✗ network performance issue
- ✓ memory bound job met the resource limitation

Outline

- Use case in Appier
- Moving from micro batch to true streaming
- Tips to conquer obstacles during the migration
- **Challenge from design in streaming way**

Challenge from design in streaming way

- **Replay expired data or rebuild states are complex**
- **How to expose user-defined metrics well**
 - Current matched count for each rule
 - Too expensive to use querible states
- **End-to-end verification is a tough job**
 - How to verify those replay data is prepared

TL;DR

- Documents and information from mailing list help a lot.
- Monitor and alert let you respond to problem and diagnose it quickly.
- Stateful streaming is environment sensitive. Be careful of it.
- Community is always behind you.

A photograph of a modern office interior, overlaid with a semi-transparent blue filter. The office features long white desks, ergonomic chairs, and several people working at their computers. Large windows on the right side offer a view of a city skyline. In the foreground, a large, angular white desk is visible, with the word 'APPIER' in large, blue, 3D block letters on its side. A small potted plant sits on the desk. The ceiling has exposed ductwork and modern pendant lights.

Thank you