Alibaba Group

Unified Engine for Data Processing and AI

Xiaowei Jiang

Sept, 2018

About Me

Alibaba Group

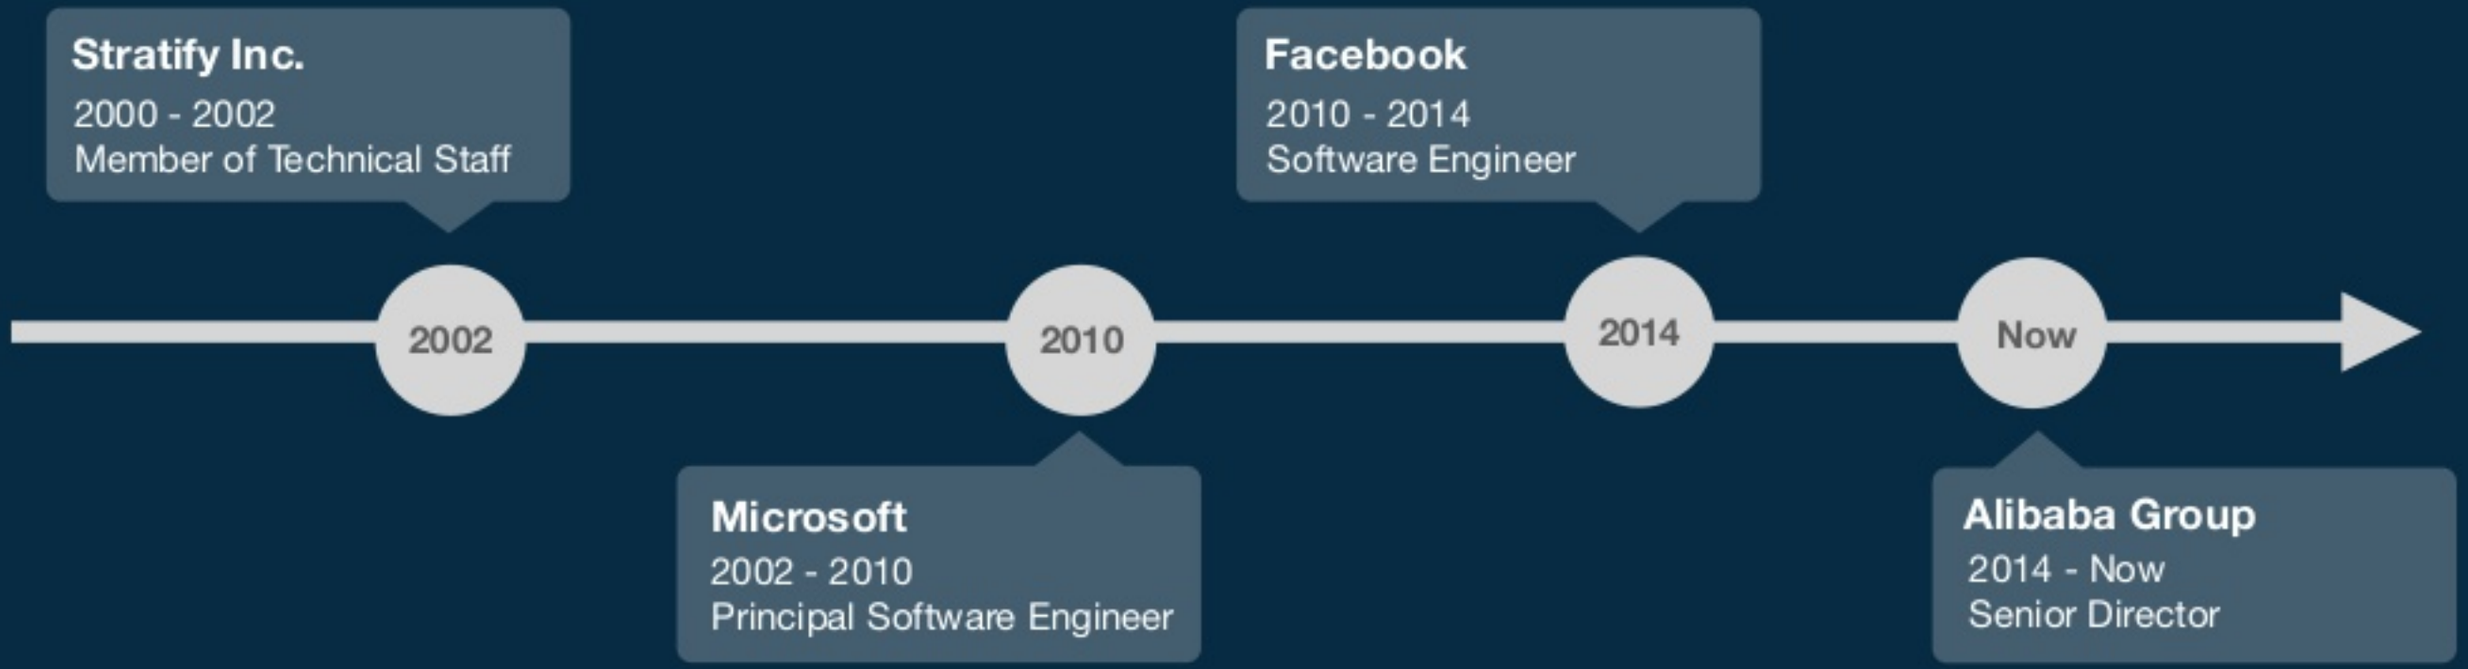

About Alibaba

Alibaba Group

聚划算 淘宝网 Taobao.com 天猫 阿里云 Alibaba Group 支付宝 ALIPAY CAINIAO 菜鸟 AliExpress

EB Total

PB Everyday

1T Event/Day

472M Events/sec

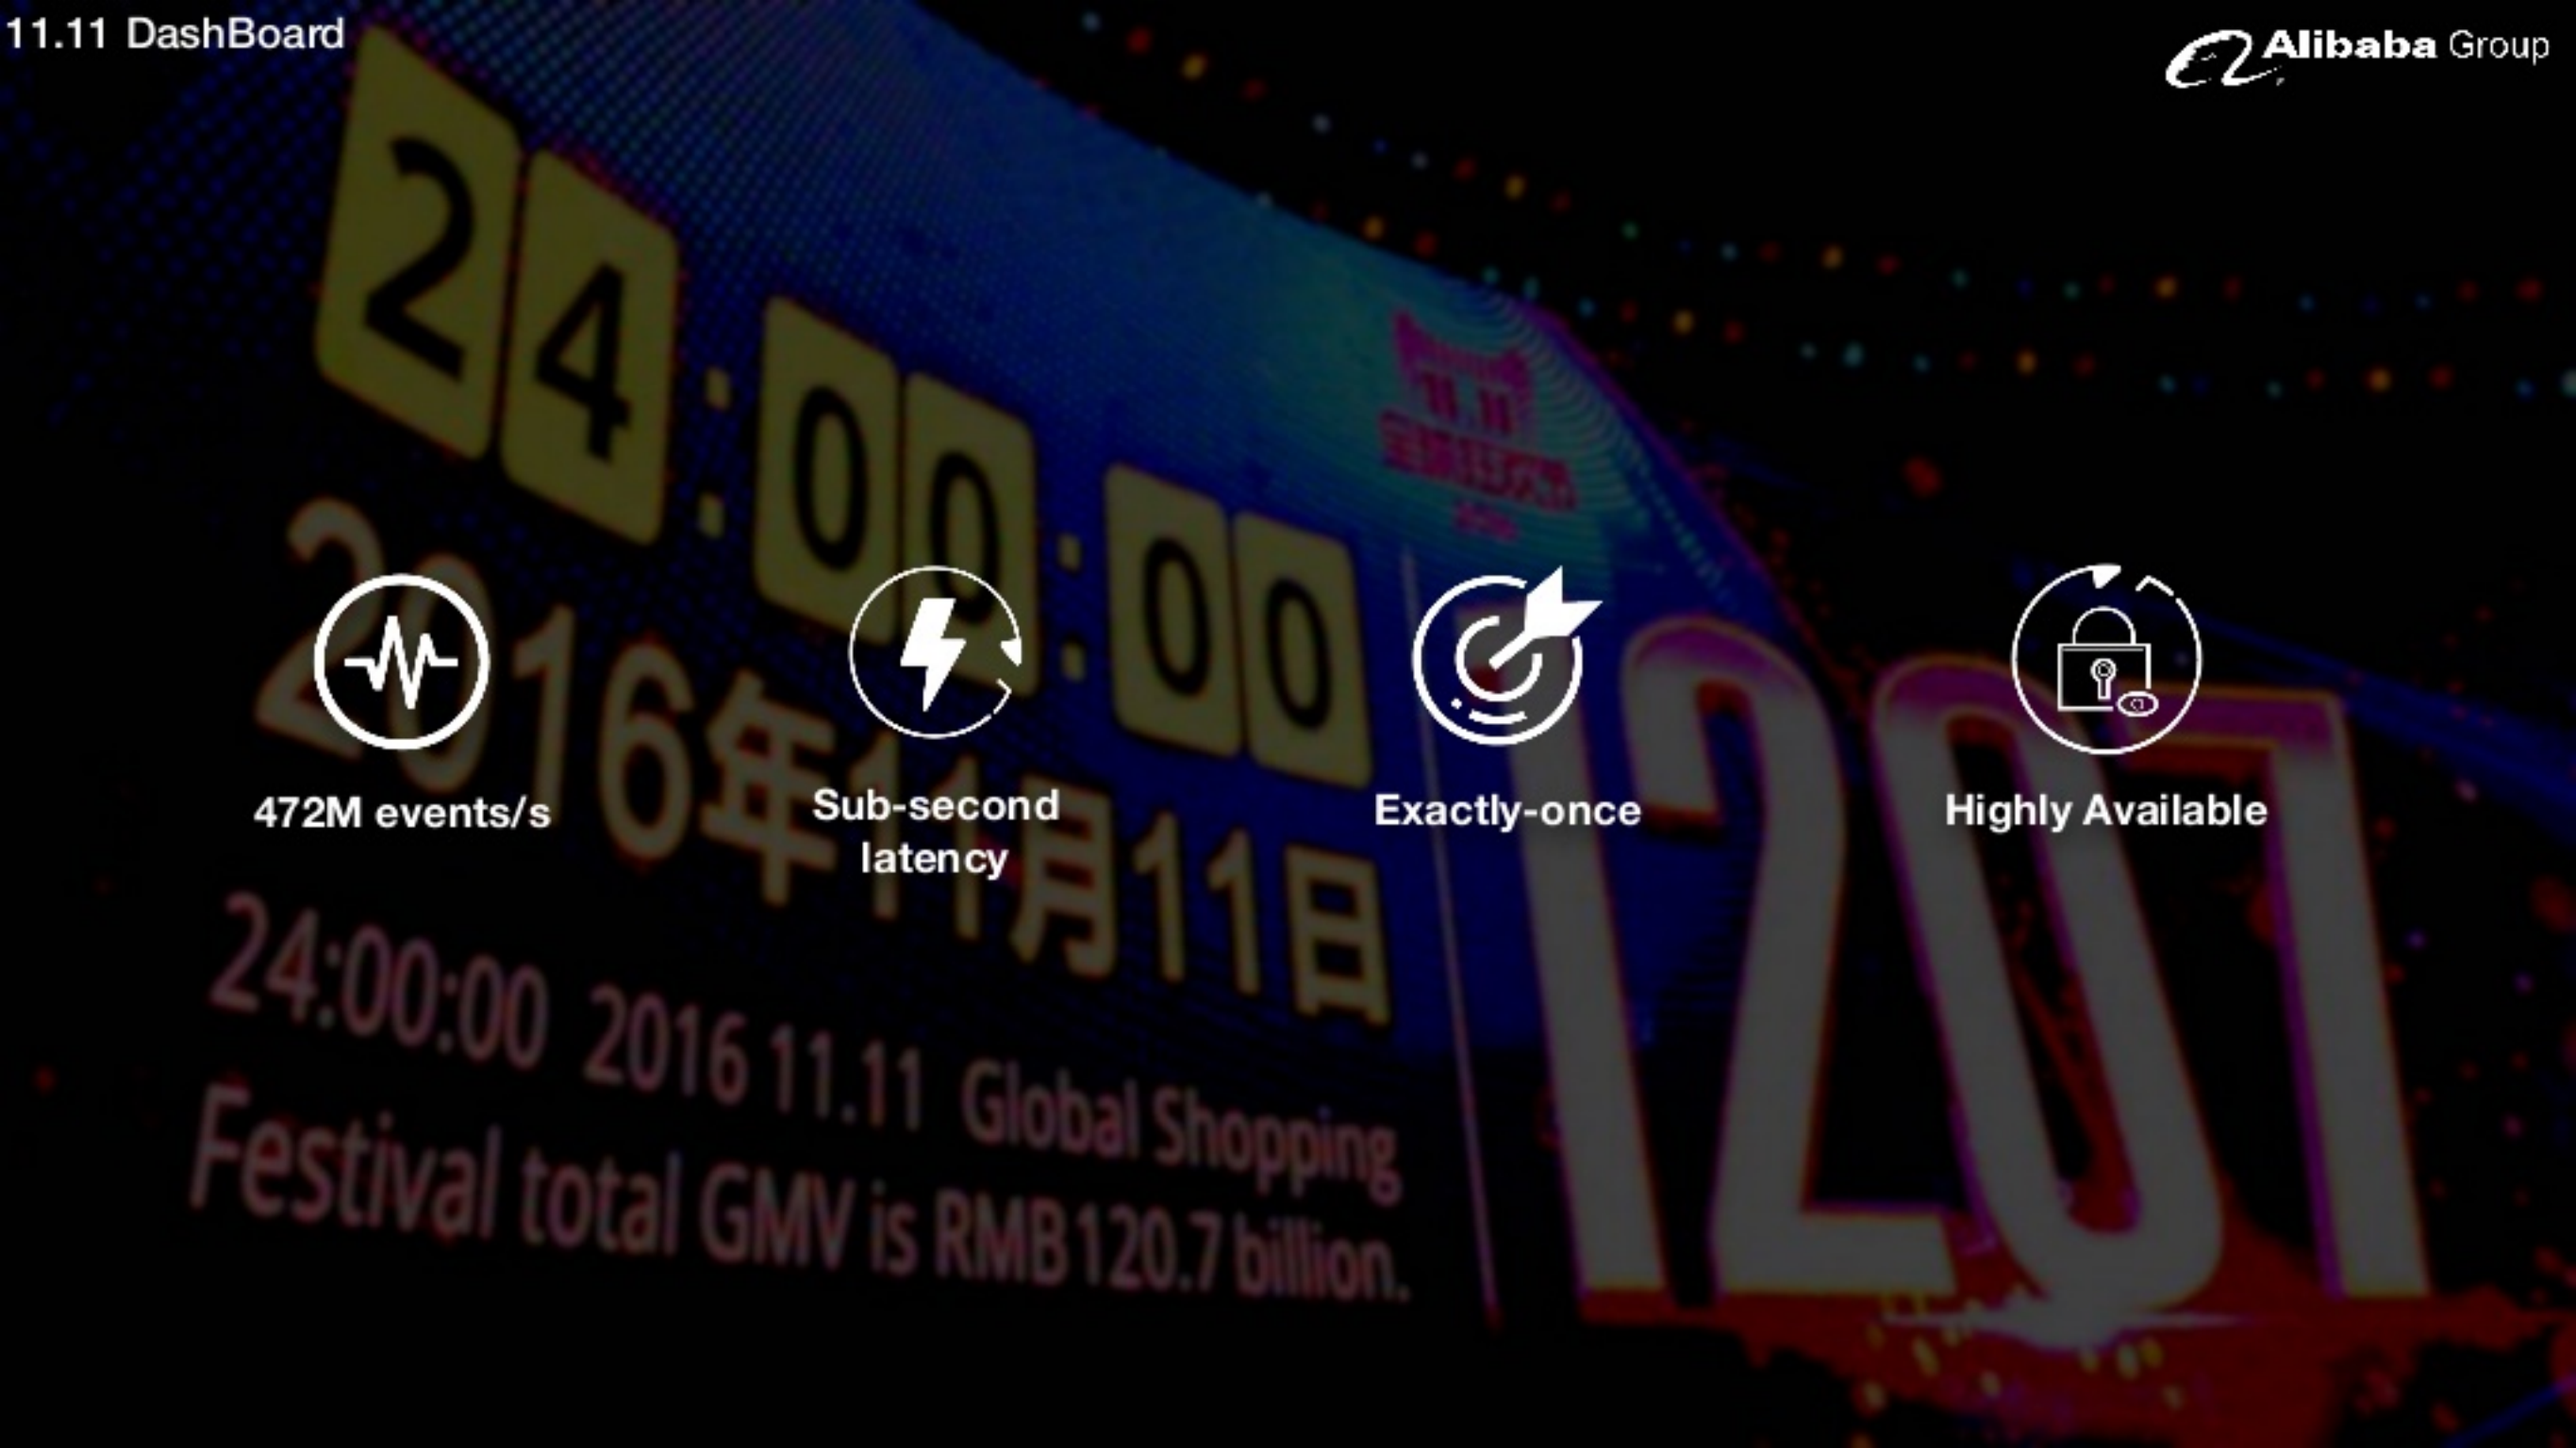
11.11 DashBoard
Alibaba Group
24:00:00
2016年11月11日
24:00:00 2016 11.11 Global Shopping
Festival total GMV is RMB 120.7 billion.
1207
472M events/s
Sub-second latency
Exactly-once
Highly Available

Scenarios for Data Processing

Alibaba Group

Low Latency
Fixed Query

Stream Processing

Periodic/Continuous
Batch Jobs

Progressive Processing

High Latency
Flexible Query

Batch Analytics

Streaming as the **core abstraction**, Batch as a **special case** of Streaming

Flink Architecture

Alibaba Group

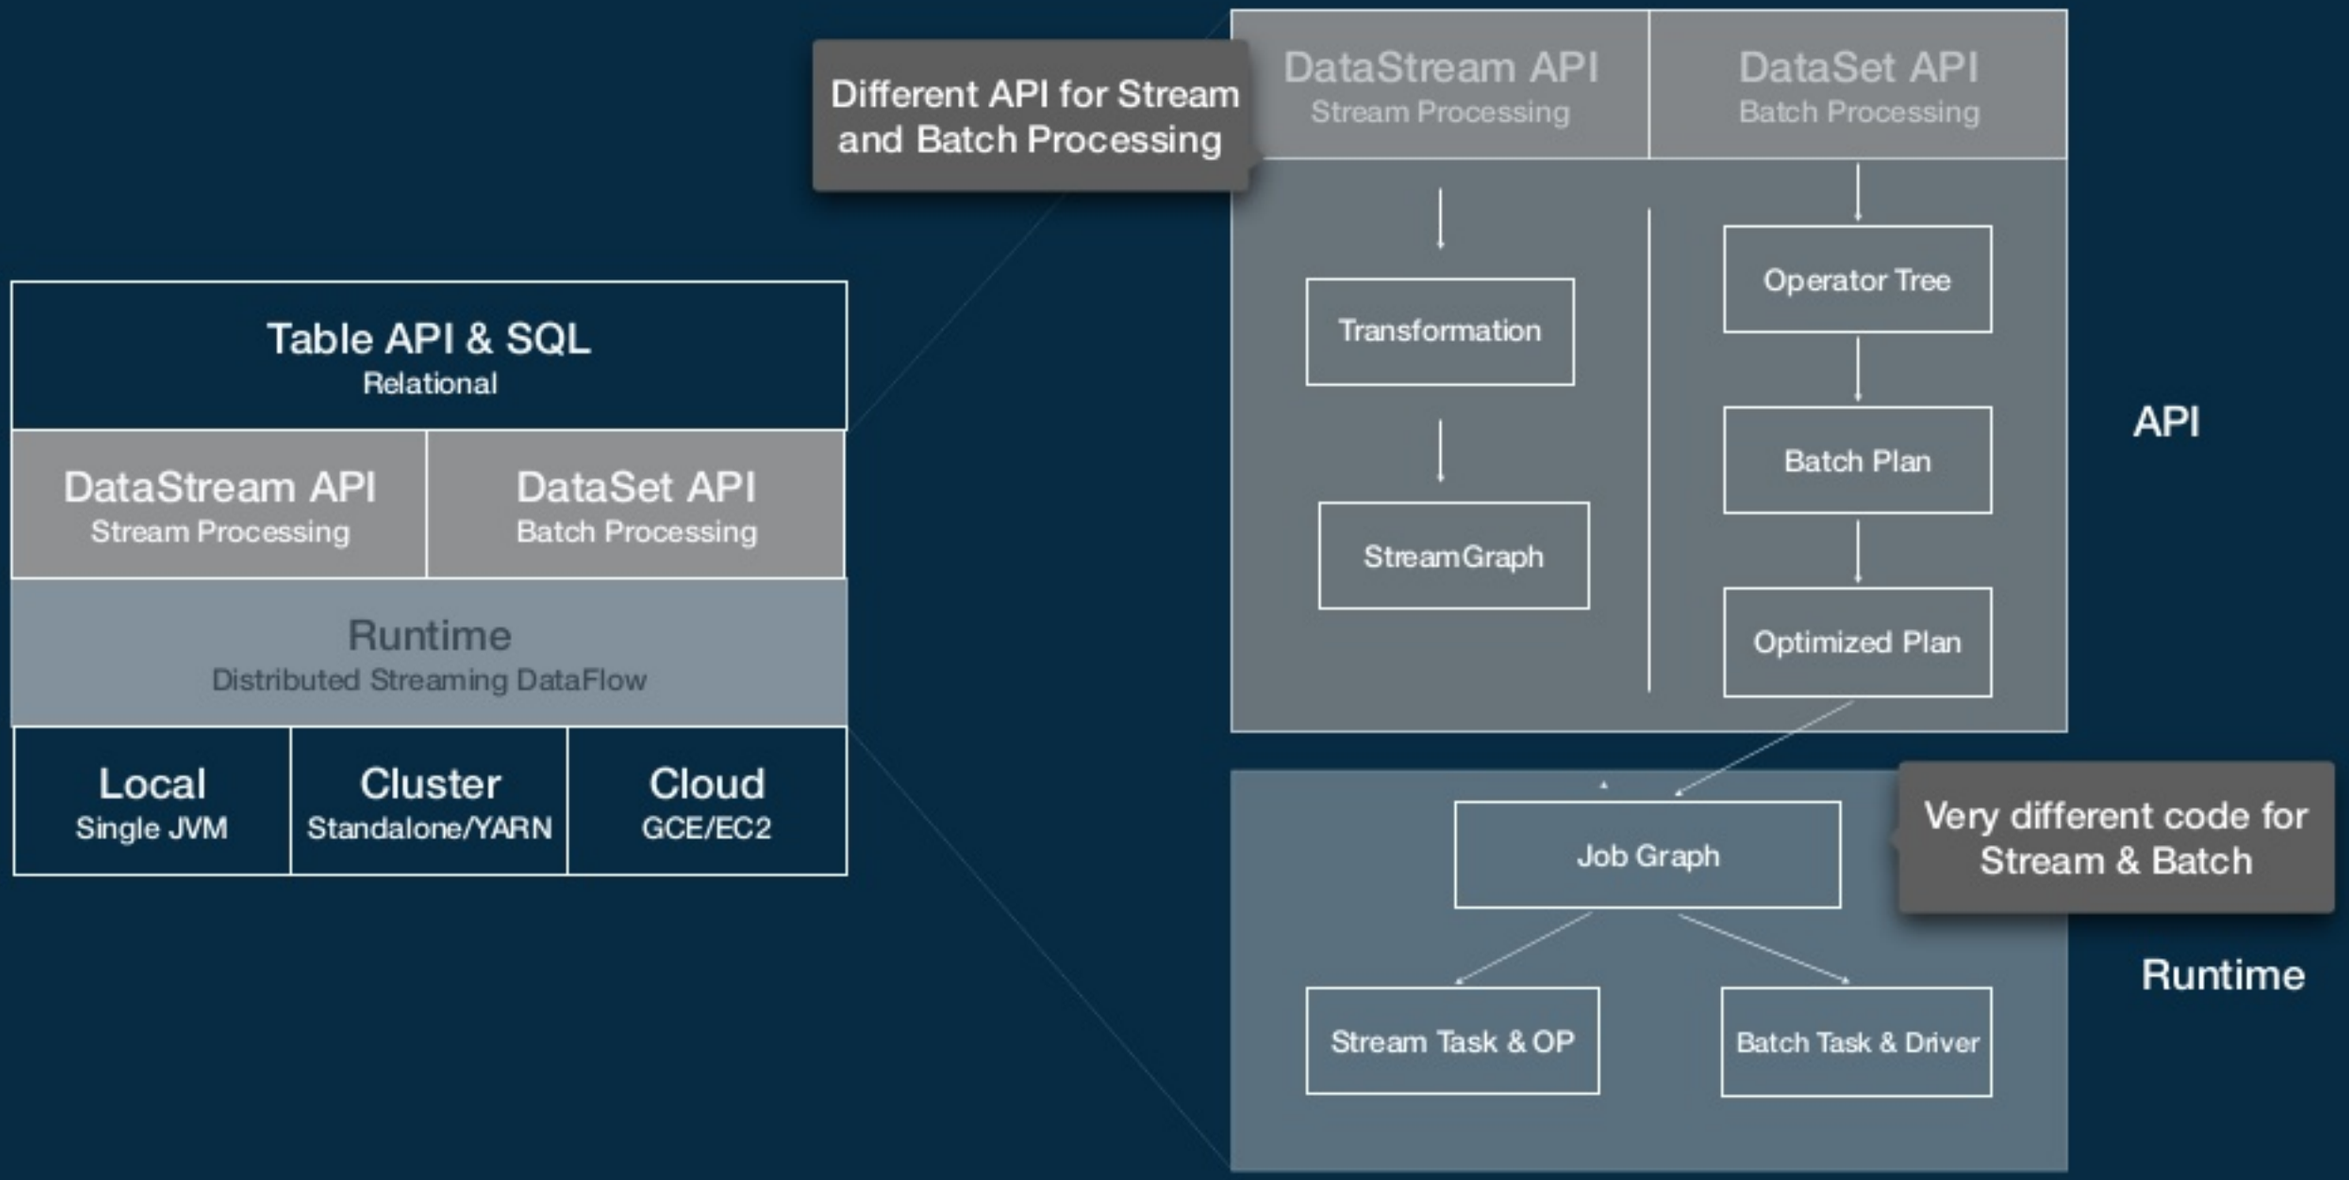

Alibaba Group

Stream Processing Engines

Apache Flink is the most sophisticated open-source Stream Processor

Alibaba Group

Batch Processing Engines

Can Apache Flink become the most sophisticated open-source batch processor?

Batch Micro-benchmark

Alibaba Group

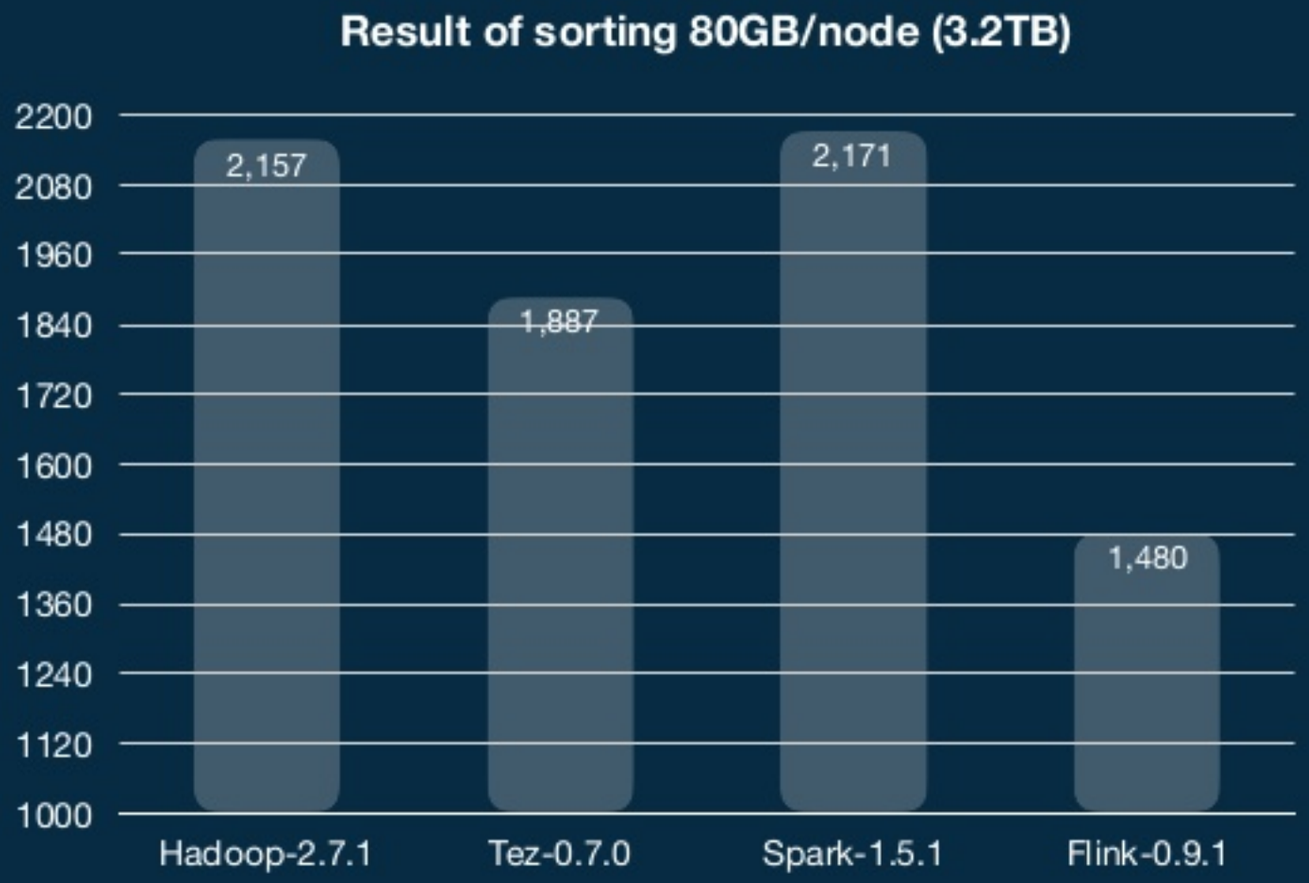

Flink is the fastest due to its pipelined execution

Tez and Spark do not overlap 1st and 2nd stages
MapReduce is slow despite overlapping stages

A Comparative Performance Evaluation of Flink, Dongwon Kim, POSTECH, Flink Forward 2015

Unified Engine

Functionality

Performance

Reliability

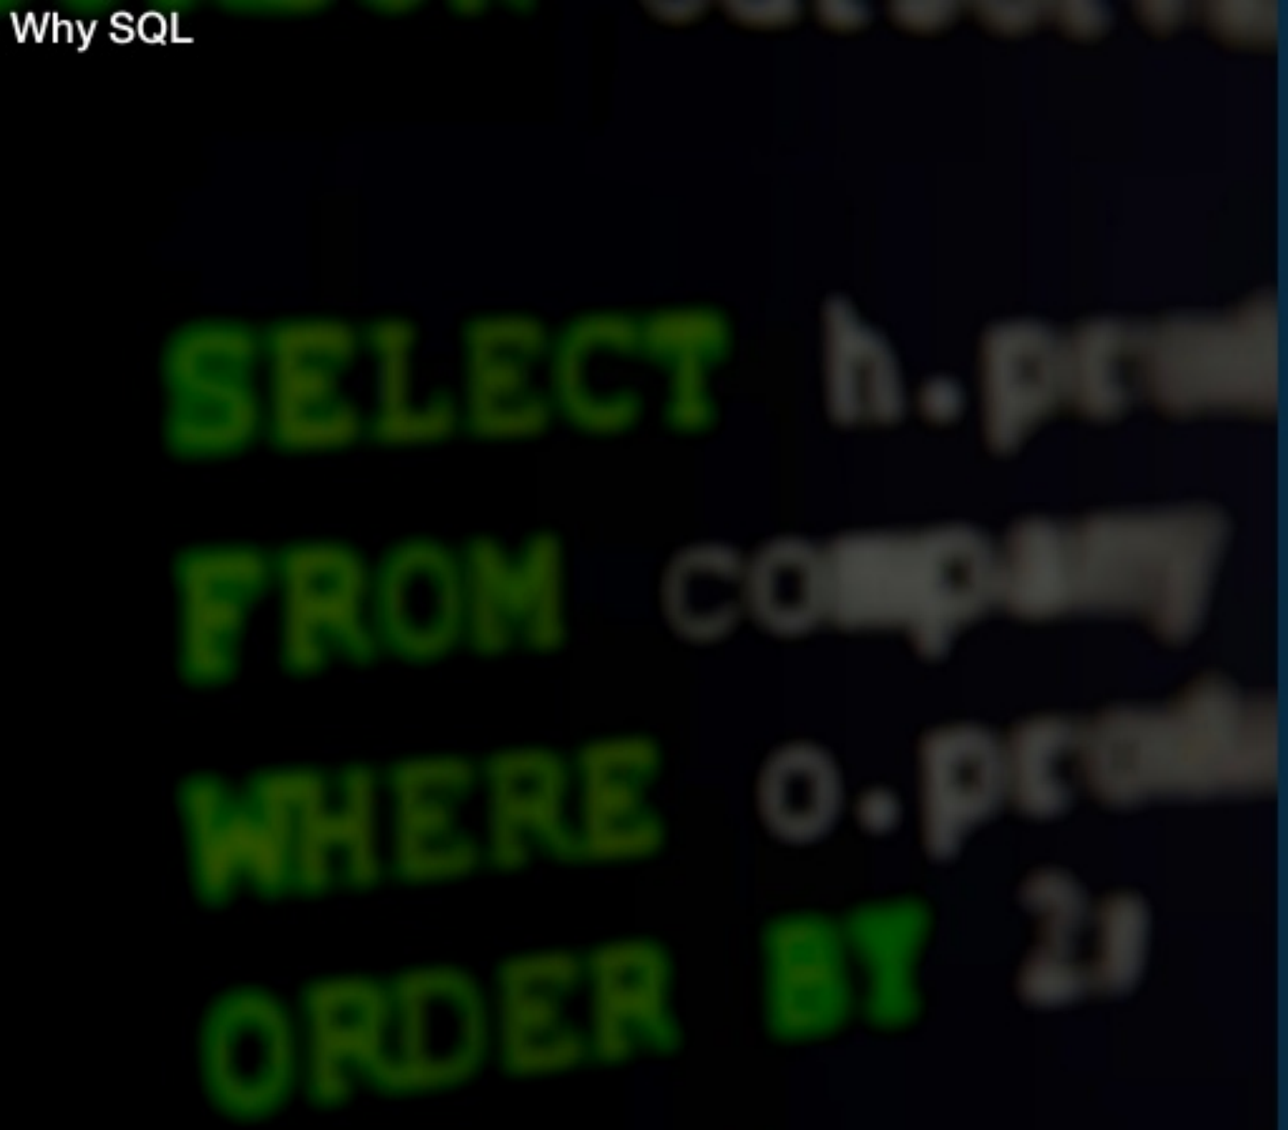

Alibaba Group

- Declarative
- Optimizable
- Understandable
- Stable
- Unified

Dynamic Table

Alibaba Group

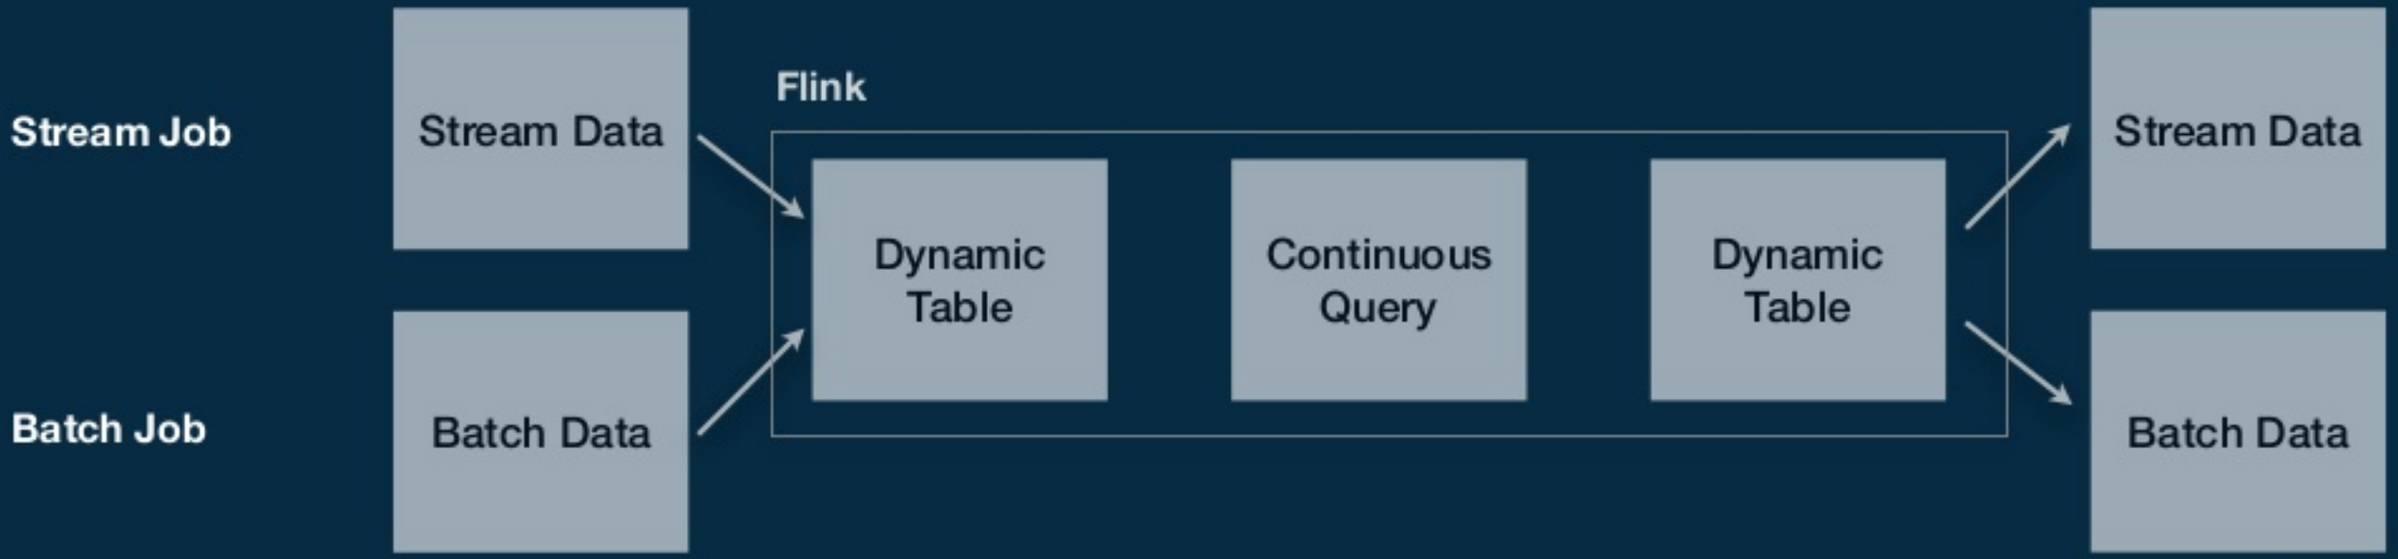

Streaming SQL Functionalities

Alibaba Group


Stream Join


Agg w/ Retraction


Window

UDX


DDL Support

Connector

Alibaba Group

TPCH & TPCDS

2K+ CASE & SQL

Existing Execution Model

Alibaba Group

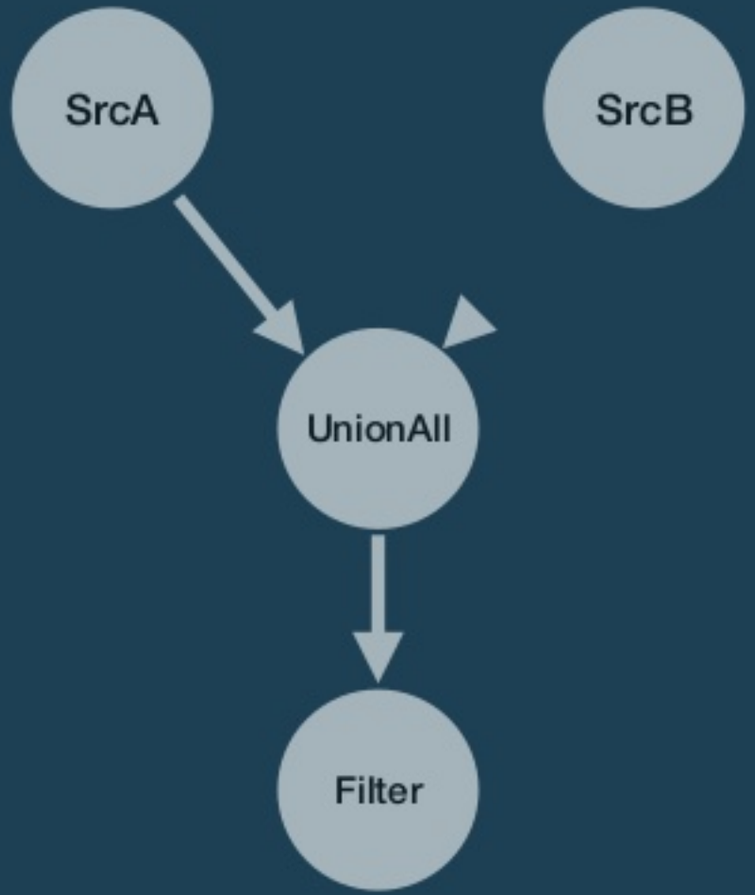

Stream Model

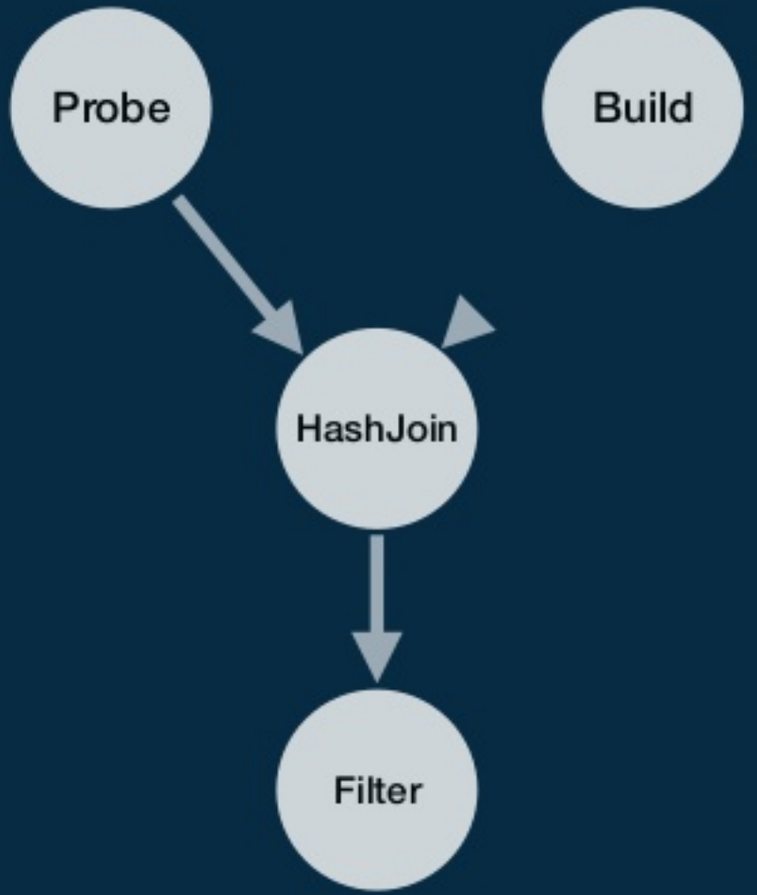

Batch Model

Unified Operator Framework

Alibaba Group

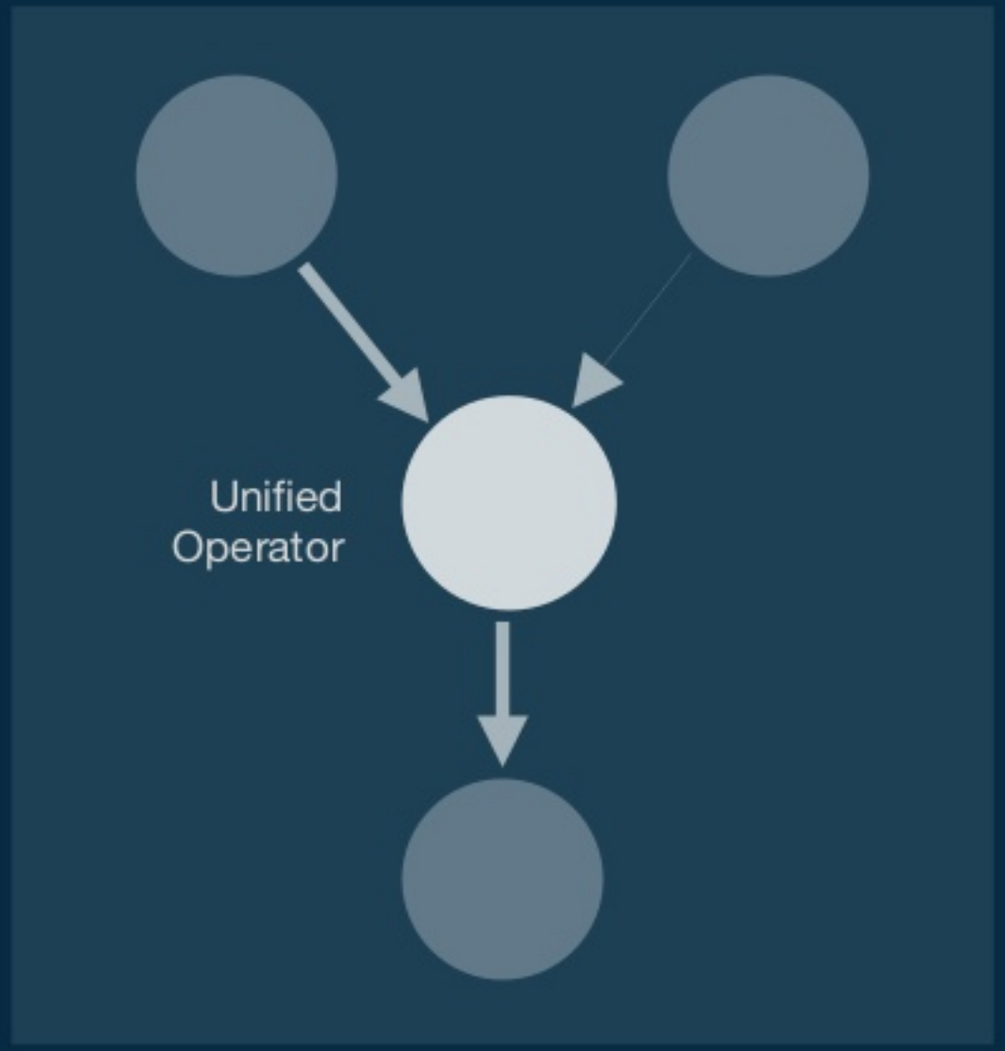

Unified Operator Abstraction

- *Operators can choose inputs*
- *Operators can be chained easily*
- *Helps batch as well as streaming*

Flink Architecture: New Design

Alibaba Group

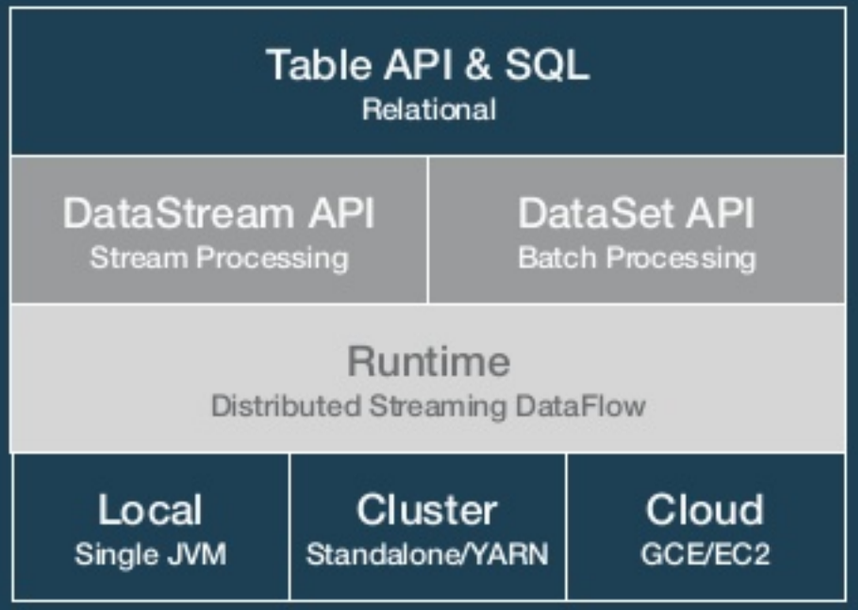

Old Design

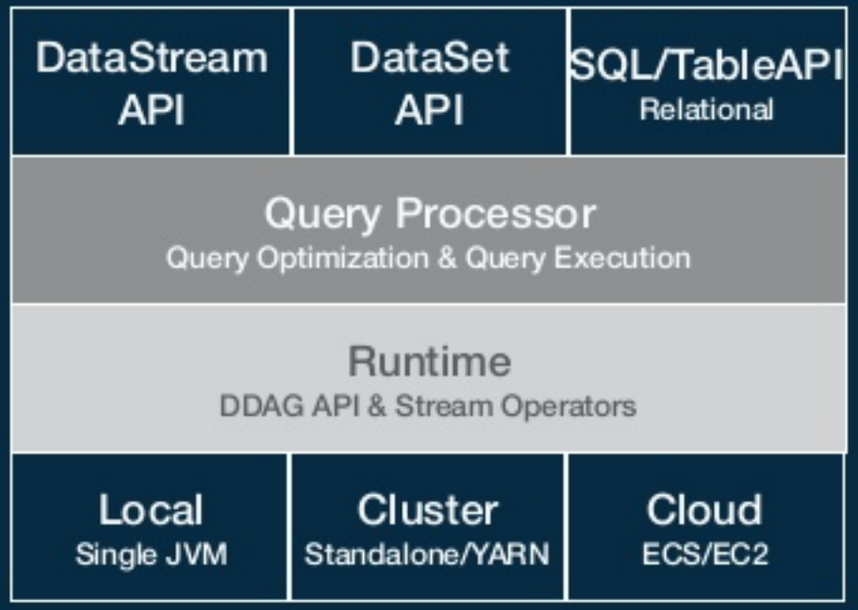

New Design

Query Processing

Alibaba Group

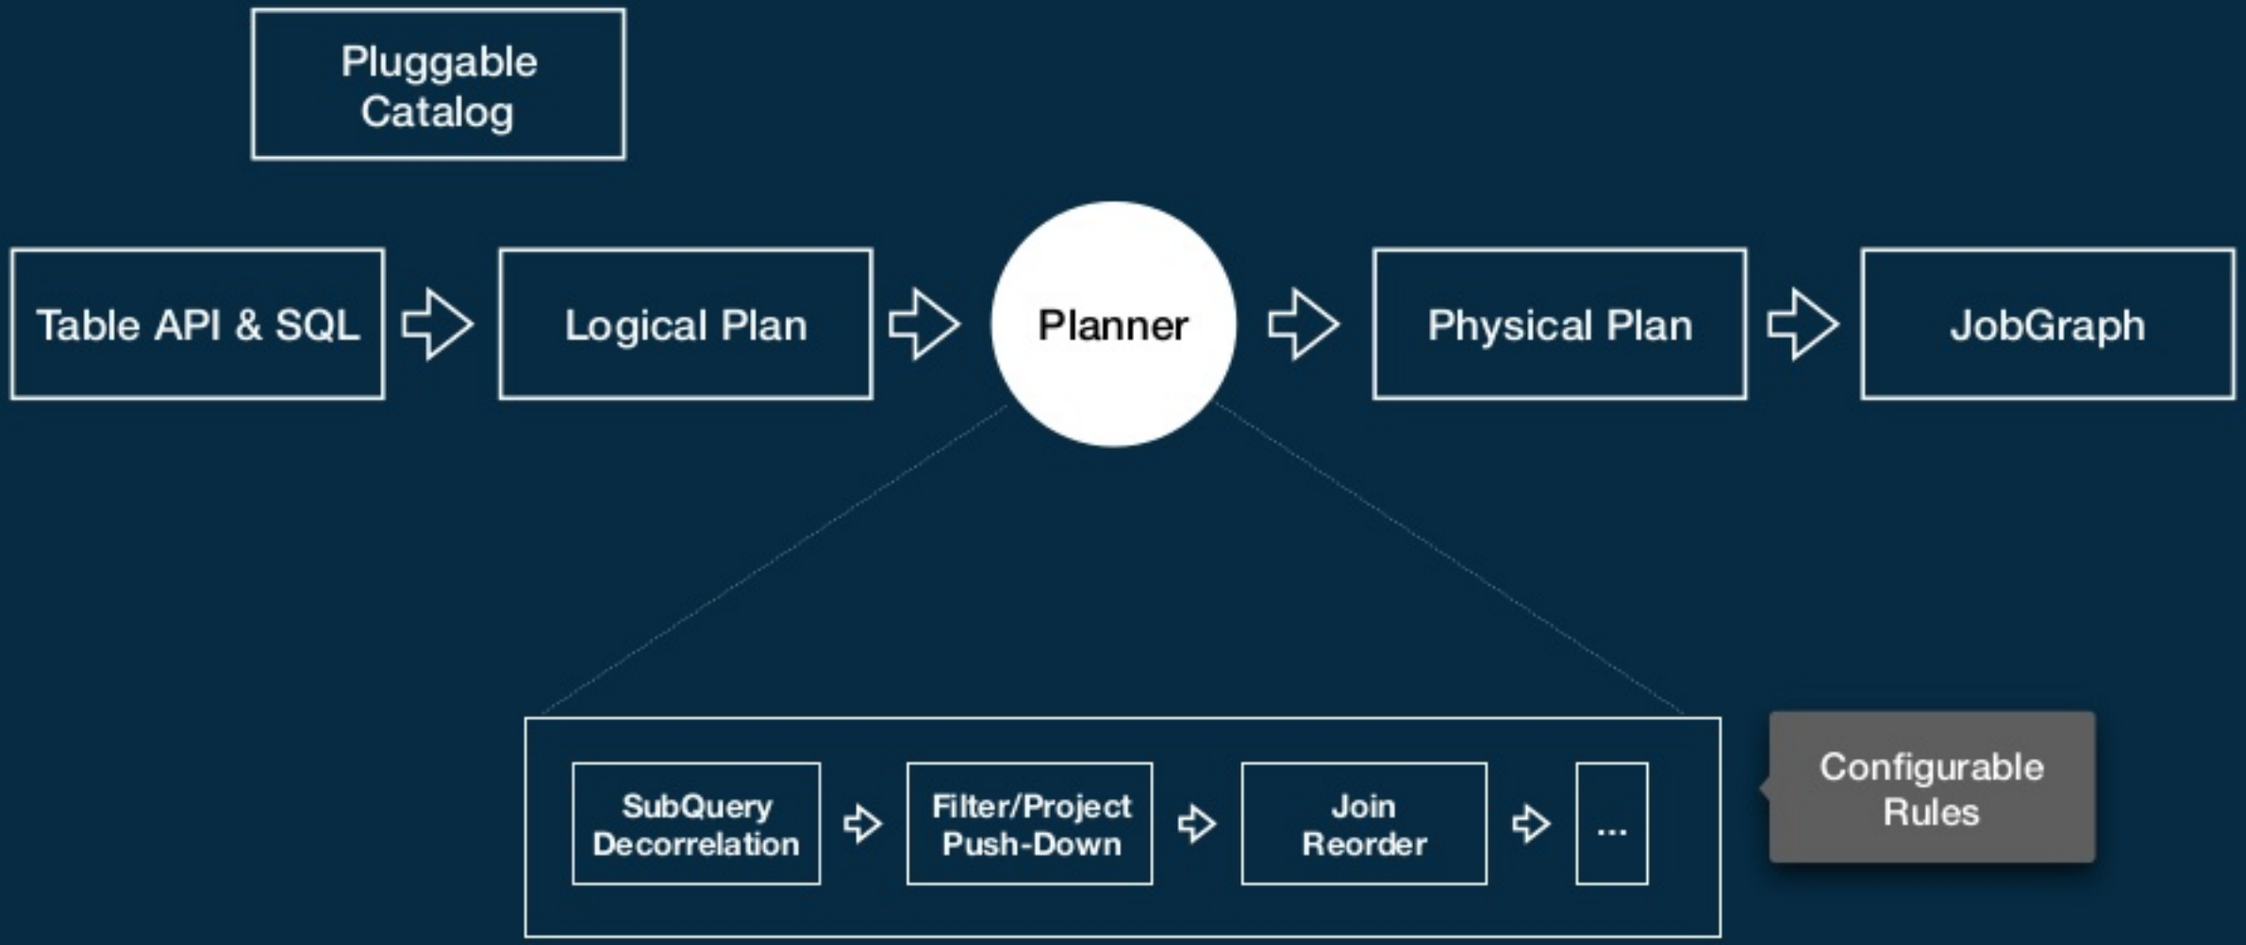

Unified Engine

Functionality

Performance

Reliability

Batch Performance

Alibaba Group

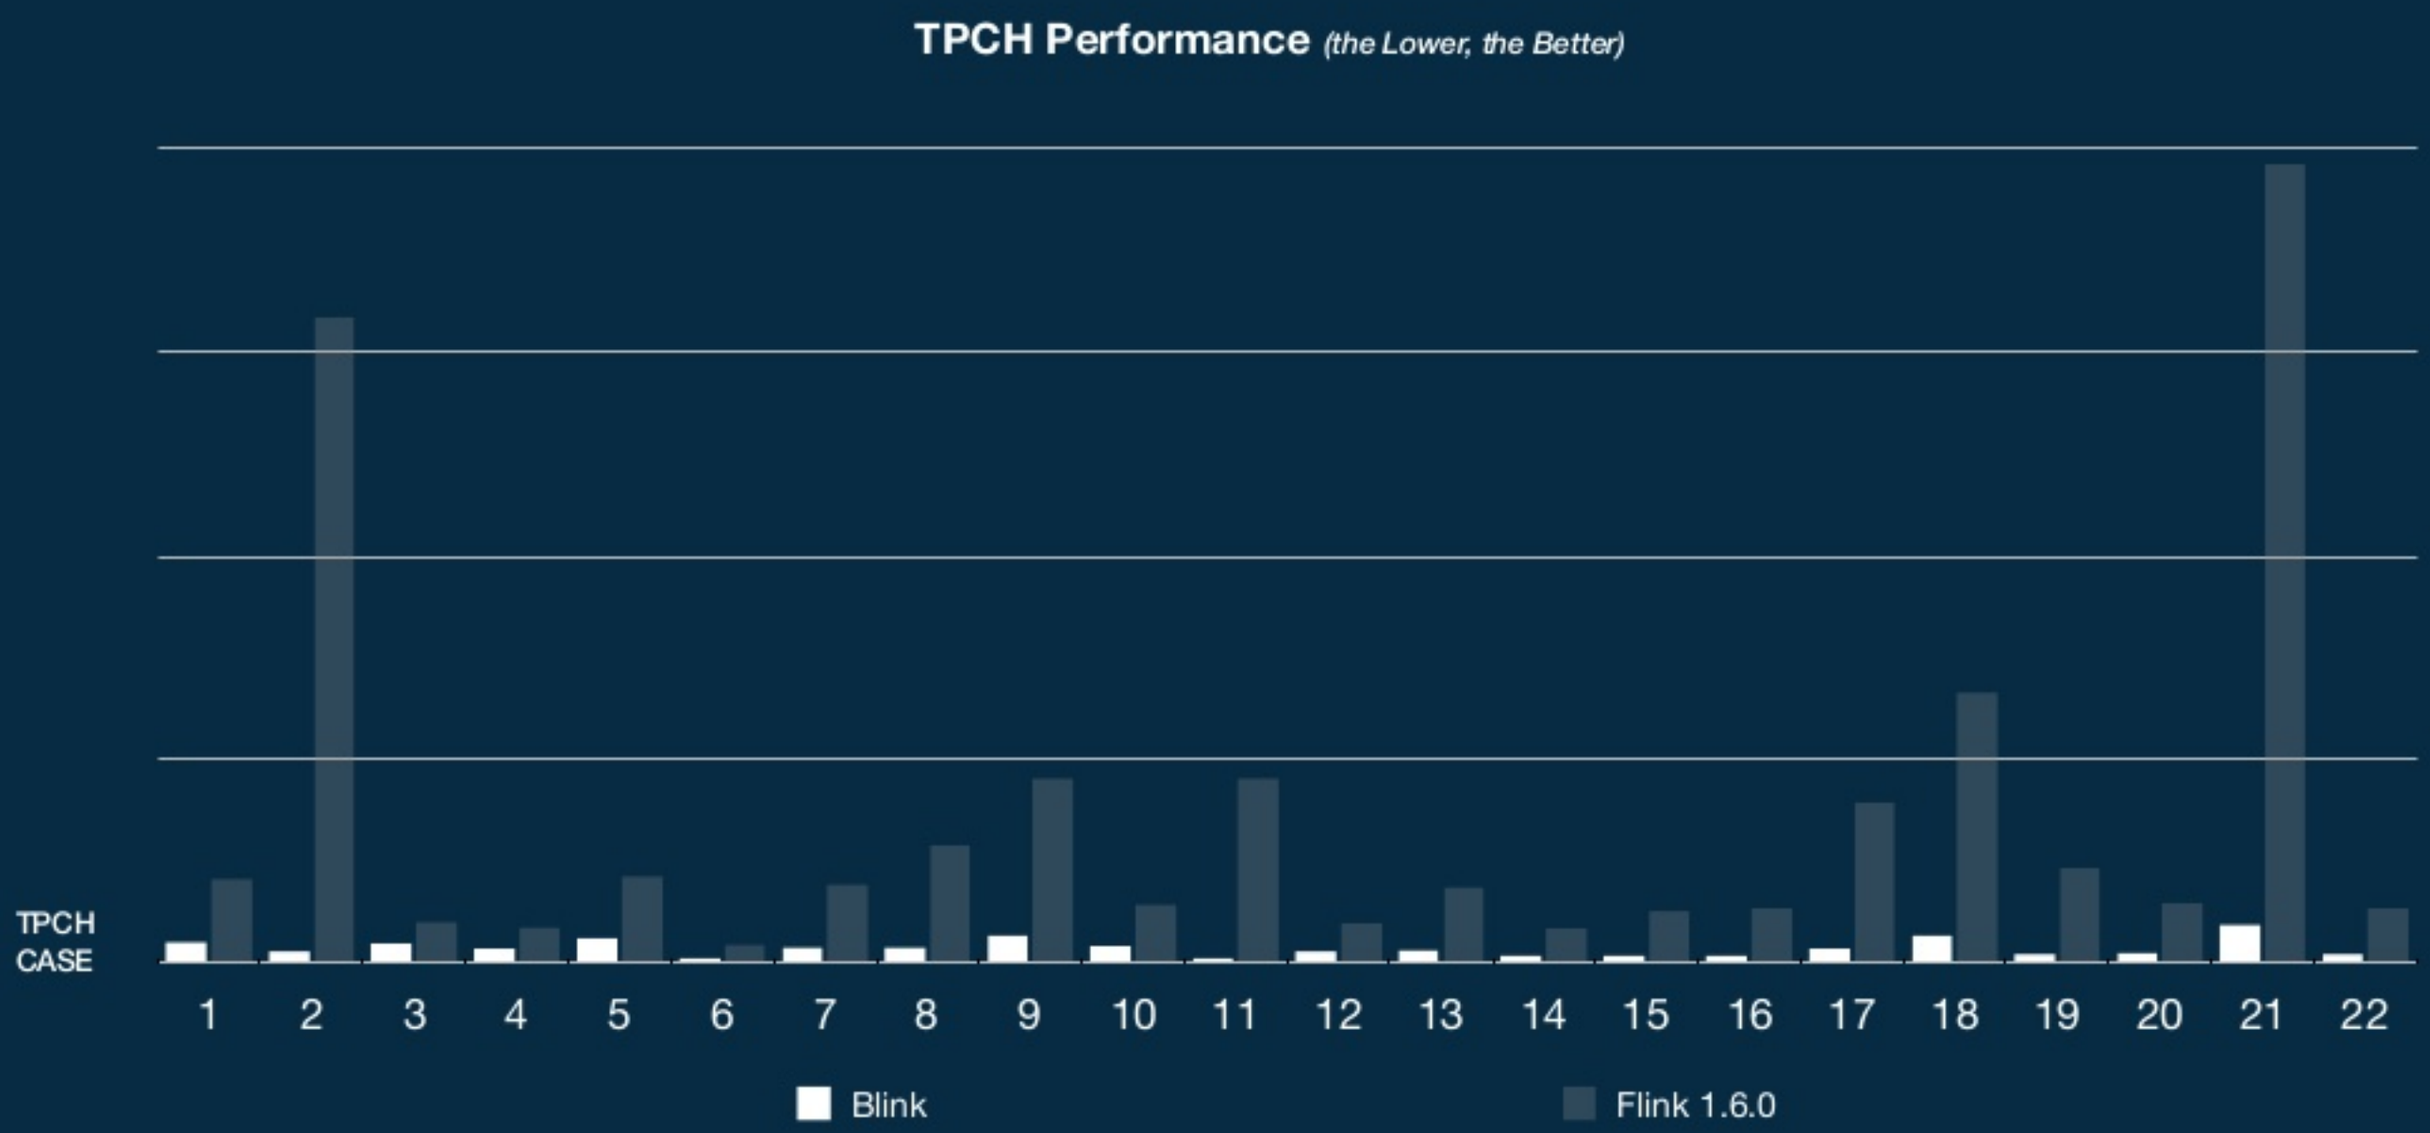

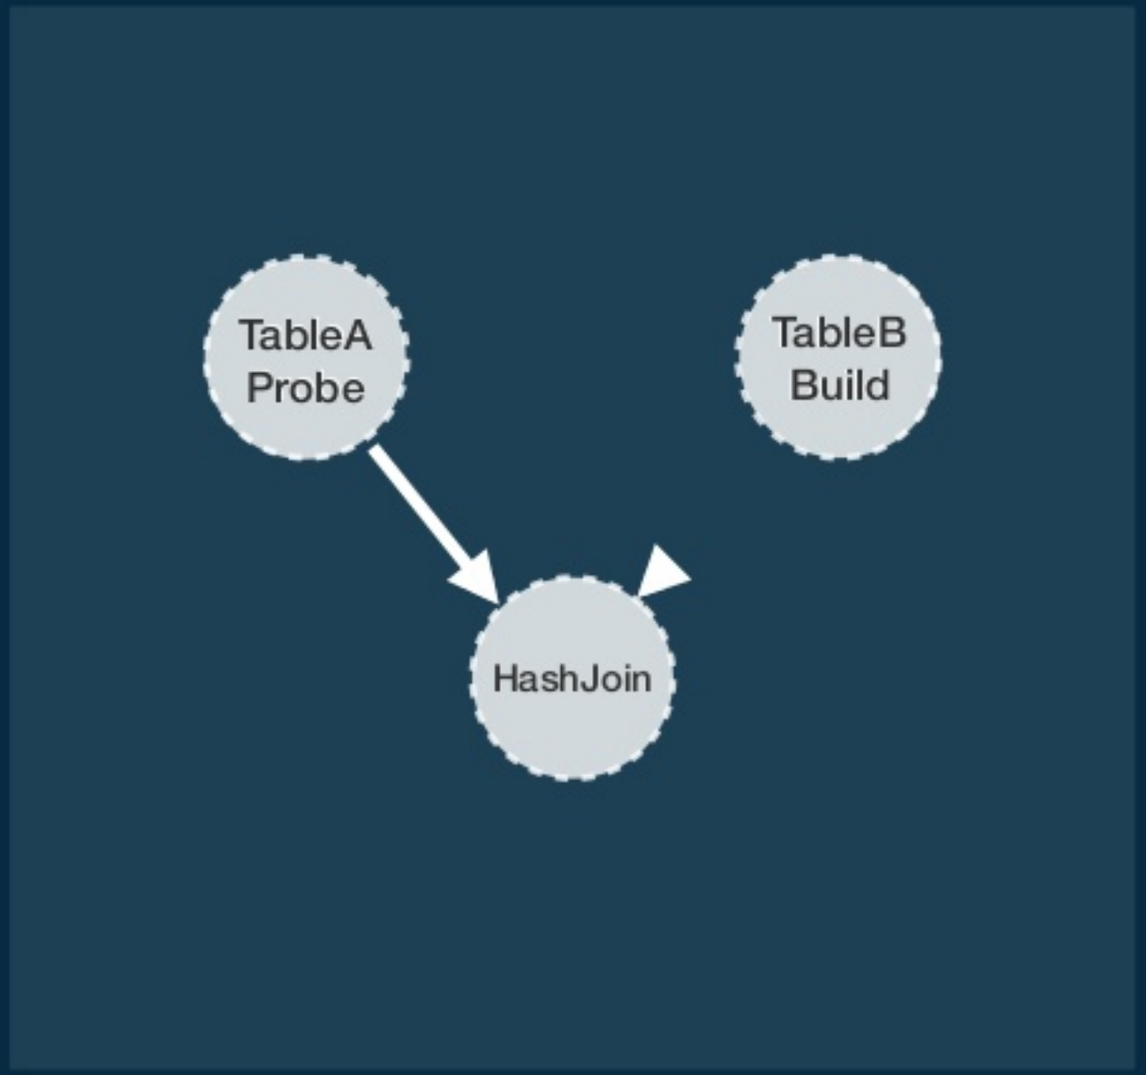

Customizable Scheduling

- *Flexible control over when tasks get scheduled*
- *Much better resource usage achieved*

Flexible Chaining

Alibaba Group

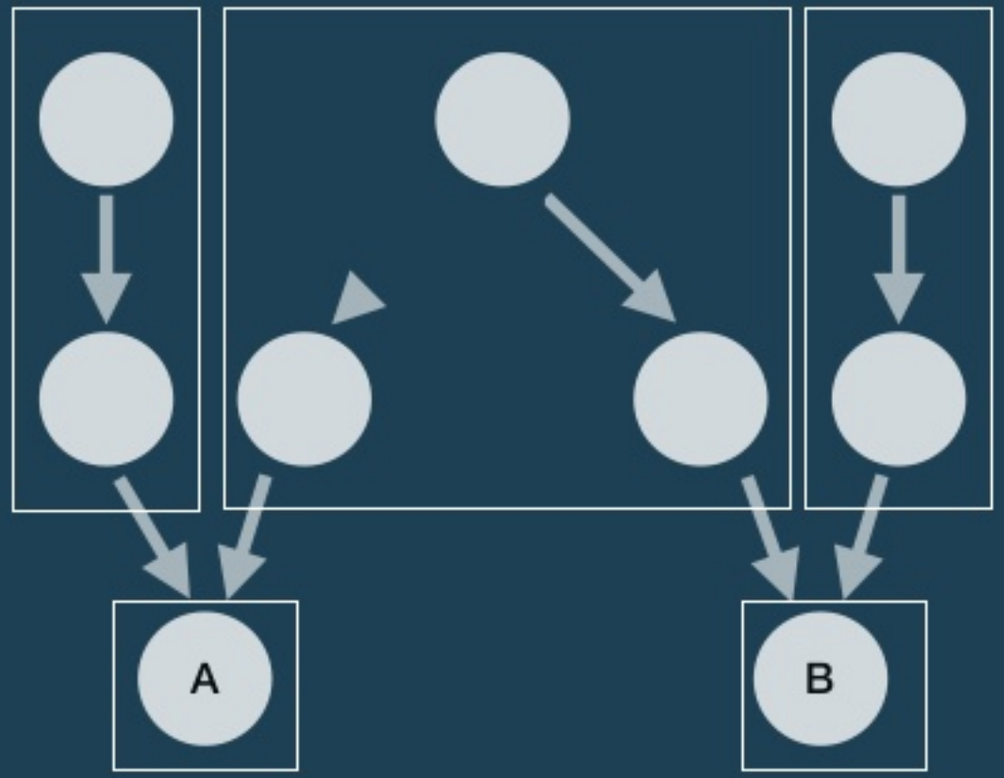

Only One-Input Operators can be chained

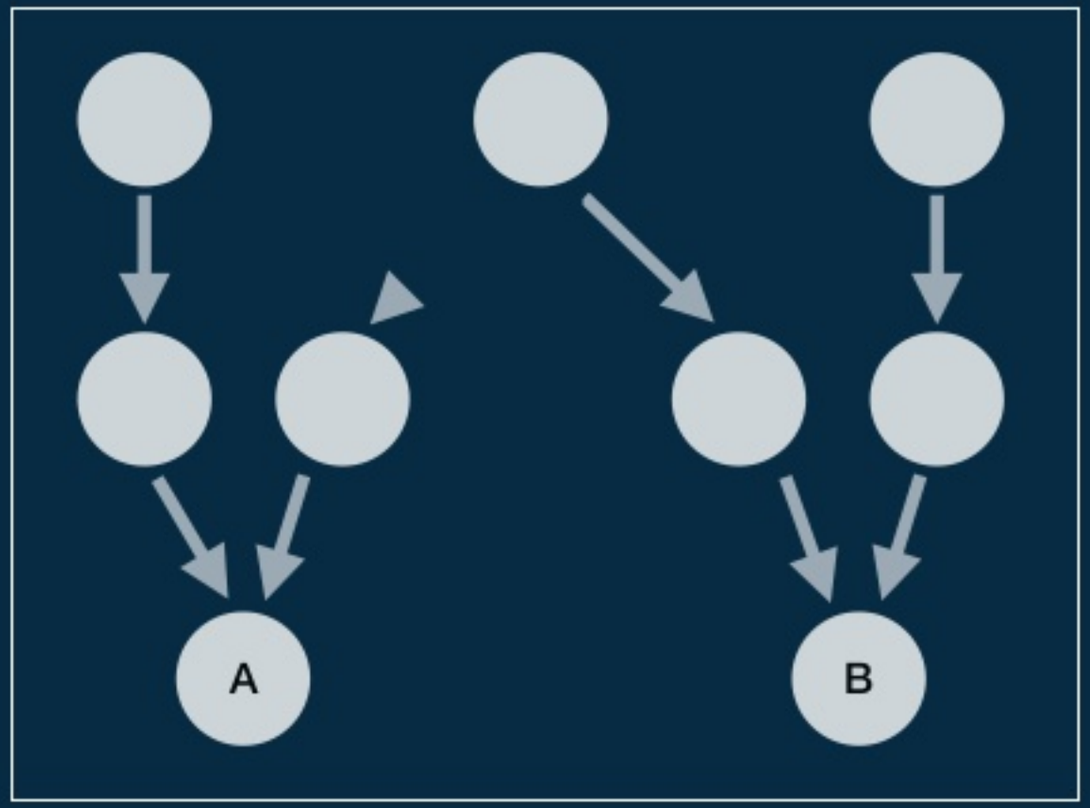

Multi-Input Operators can also be chained

Record Format

- Introduced new row format: BinaryRow
- Tight integration with memory management
- Avoid deserialization cost

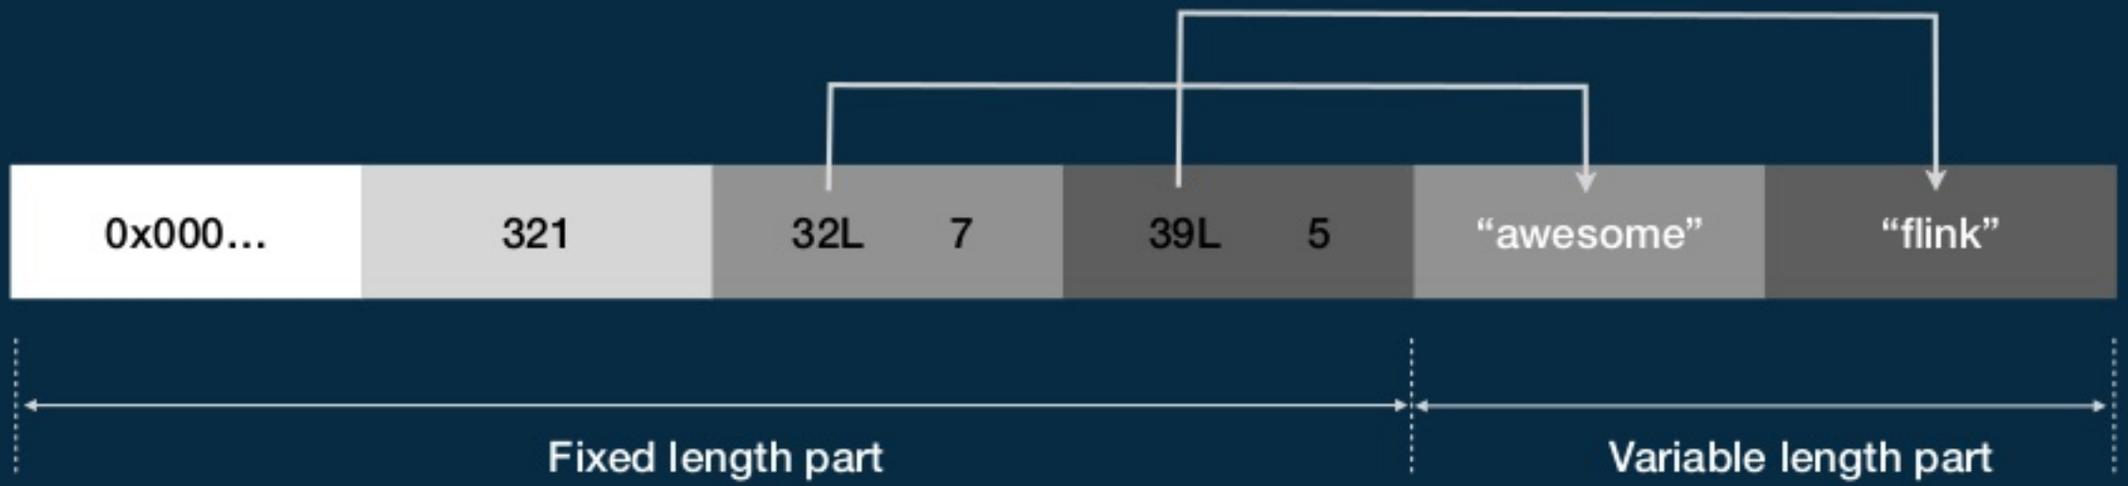

Query Execution

Alibaba Group

Expression Optimizations

Operate directly on binary data
JVM intrinsics
Hot method codegen

Performant Operators

Operator codegen
HashAgg
Improved HashJoin
Semi/Anti join

Resource Optimizations

Stats based estimation
Dynamic memory allocation

Alibaba Group

Cost Based

- Join order
- Join type
- Agg strategy
- ……

Advanced Rules

- Subplan reuse
- Join condition expansion
- Shuffle removal
- Distinct Agg rewrite
- ….

Rich Stats

- NDV
- NULL count
- Avg length
- Max length
- Min
- Max

Performance Improvements - Recap

Alibaba Group

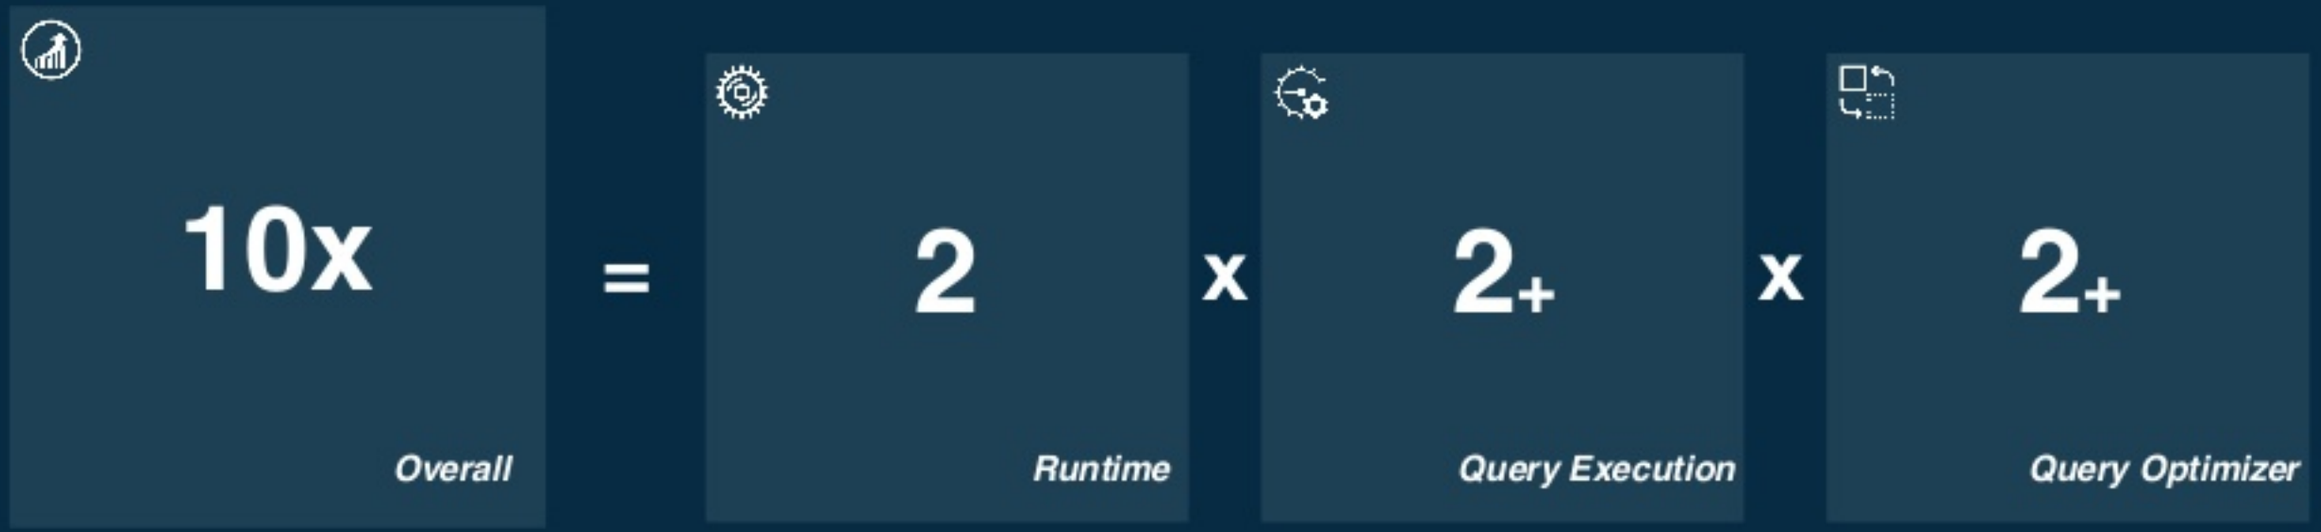

Unified Engine

Functionality

Performance

Reliability

Reliability Improvements

Alibaba Group

Failover

- Region Based Failover
- JM Failover
- Blacklist
- ....

Shuffle

- Decoupled from TM
- Yarn Shuffle Service
- Async mode

More Details?

Sept 4th, 2018
5:10 PM - 5:50 PM
Maschinenhaus
Feng Wang, Alibaba

Runtime Improvements for Flink Batch Processing

Next Steps

Grand Unification of Data Processing

Switch between batch processing and streaming seamlessly

Flink Machine Learning/AI

PyFlink, TableAPI, DL Integration, Flink ML Improvements

Takeaway

Alibaba Group

Already the best stream processor

Becoming the best batch processor

Unified approach benefits both batch & stream processing

Working on seamless experience for AI

Flink Forward China

Dec 20th-21st @ Beijing National Conference Center

First Flink Forward Conference in Asia, 3000+ participants expected

Flink Community

Joint efforts by all major players in Flink community from China

Call For Talks & Sponsors

Submit your talks to flink-forward-china@list.alibaba-inc.com